

A06: Practical Considerations for moving to a more Agile Integration Architecture

Using App Connect with Containers

Andy Garratt, IBM UK

andyg@uk.ibm.com



IBM Cloud

IBM

Define goals, targets and success factors

Agile Integration Architecture is great!

So is App Connect V11!.

So are Clouds, Containers, Kubernetes, Service Meshes, CI/CD etc

They can help you do a lot – but what do you ACTUALLY want to do?

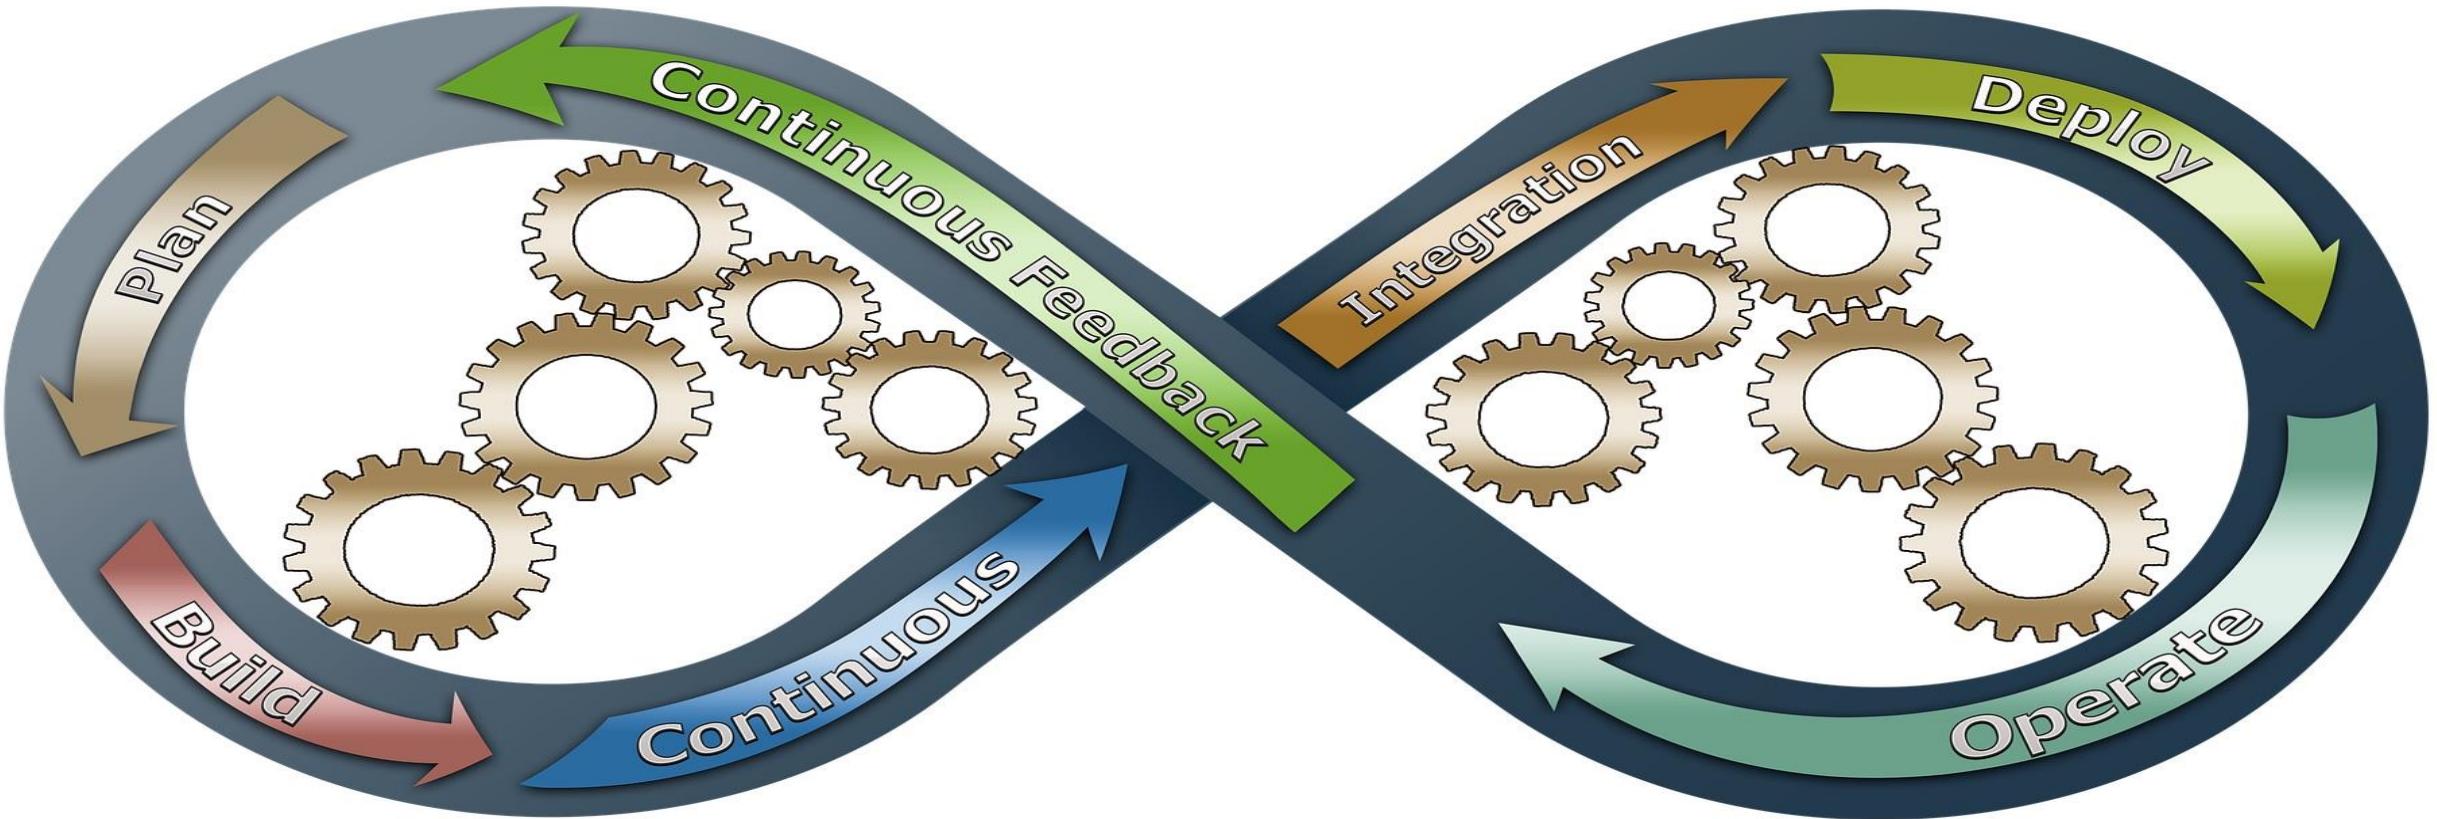


Defining “Agile”

- Individuals and interactions over processes and tools?
- Working software over comprehensive documentation?
- Customer collaboration over contract negotiation?
- Responding to change over following a plan?



“Our highest priority is to satisfy the customer through early and *continuous delivery* of valuable software”



“Deliver working software frequently, from a couple of ~~weeks~~ hours to a couple of months,
with a preference to the shorter timescale.”

Risk Free Deployment

"We want to be able to upgrade and/or patch without months of retro-testing"

"We want to deploy or patch an API or a service independently of all other services"

"We want to deploy and patch 'live' without impacting our users"

Cloud / Location Independence and portability

"We want to be able to run our integrations wherever makes most sense. We also want to move them around"

"Minimum/Zero dependencies – embrace the container concept"

"Just another container. Deploy integrations as microservices"

Embrace DevOps

"Everything is built using a fully-automated pipeline"
"Each integration has independent pipelines"

"What gets built in DEV is what's deployed to PROD"

"Developer testing is on a multi-instance cluster. No more 'it works on my laptop'"

Escape the "backlog trap"

"If they want to build integrations with their teams, that's fine"

"We want to be more responsive to our API consumers and be able to build faster"

"Re-use by copy – look at NPM / git copy+fork approach"

Agile Integration Architecture

Modernizing integration
to enable
business agility

Fine grained deployment

Architecture & design

Decentralized Ownership

People & Process

Cloud native infrastructure

Infrastructure & Technology



Improve build independence and production velocity
(deployment agility)

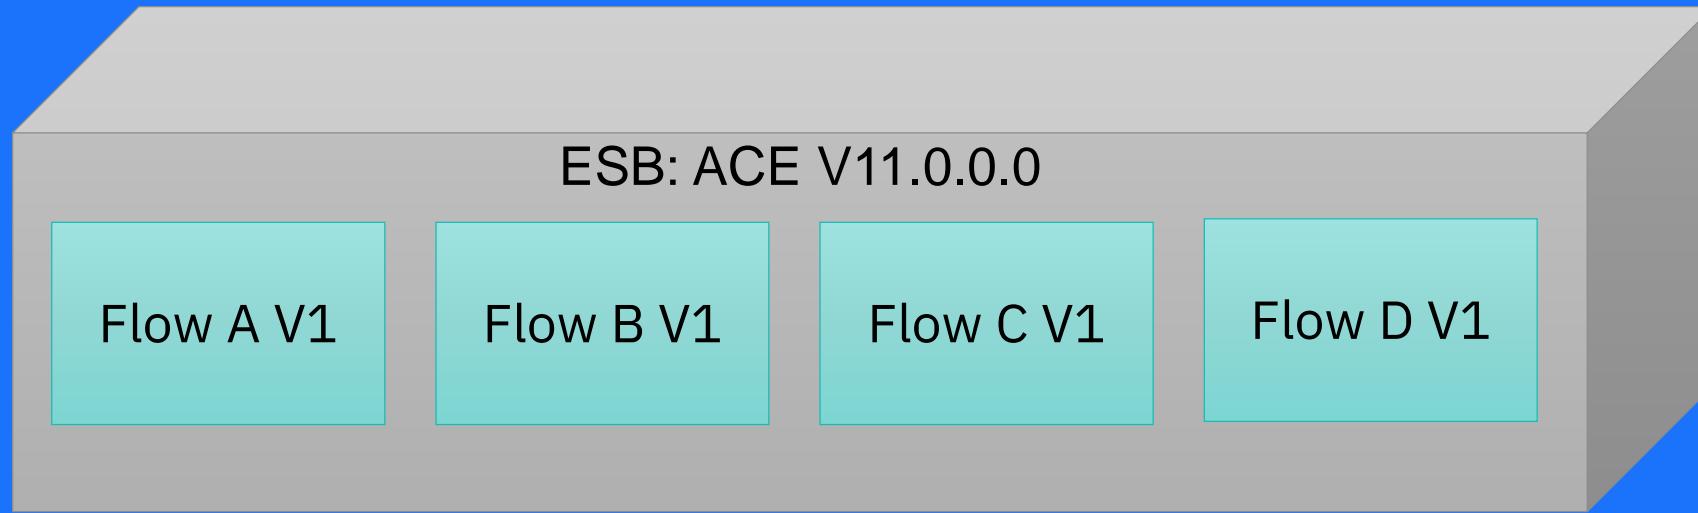


Accelerate agility and innovation
(development agility)

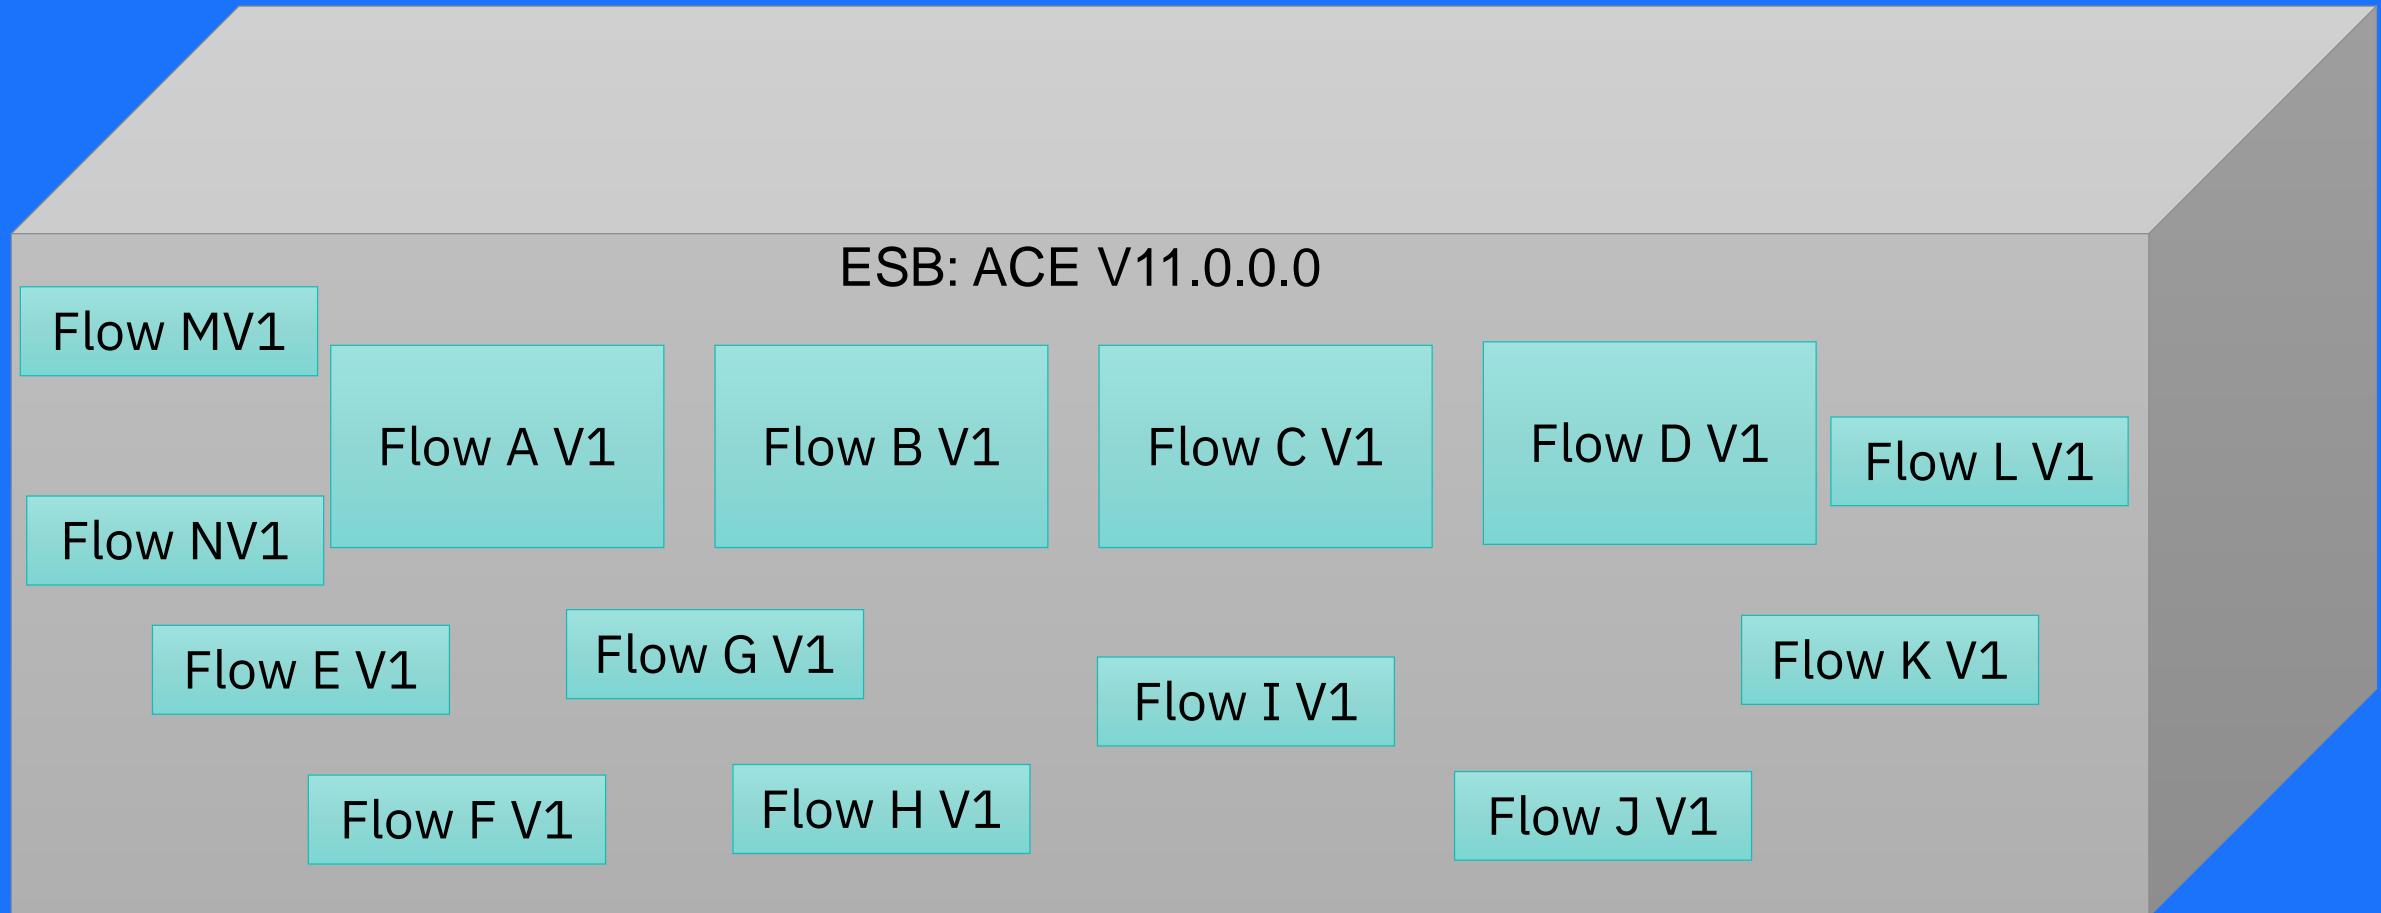


Dynamic scalability and inherent resilience
(operational agility)

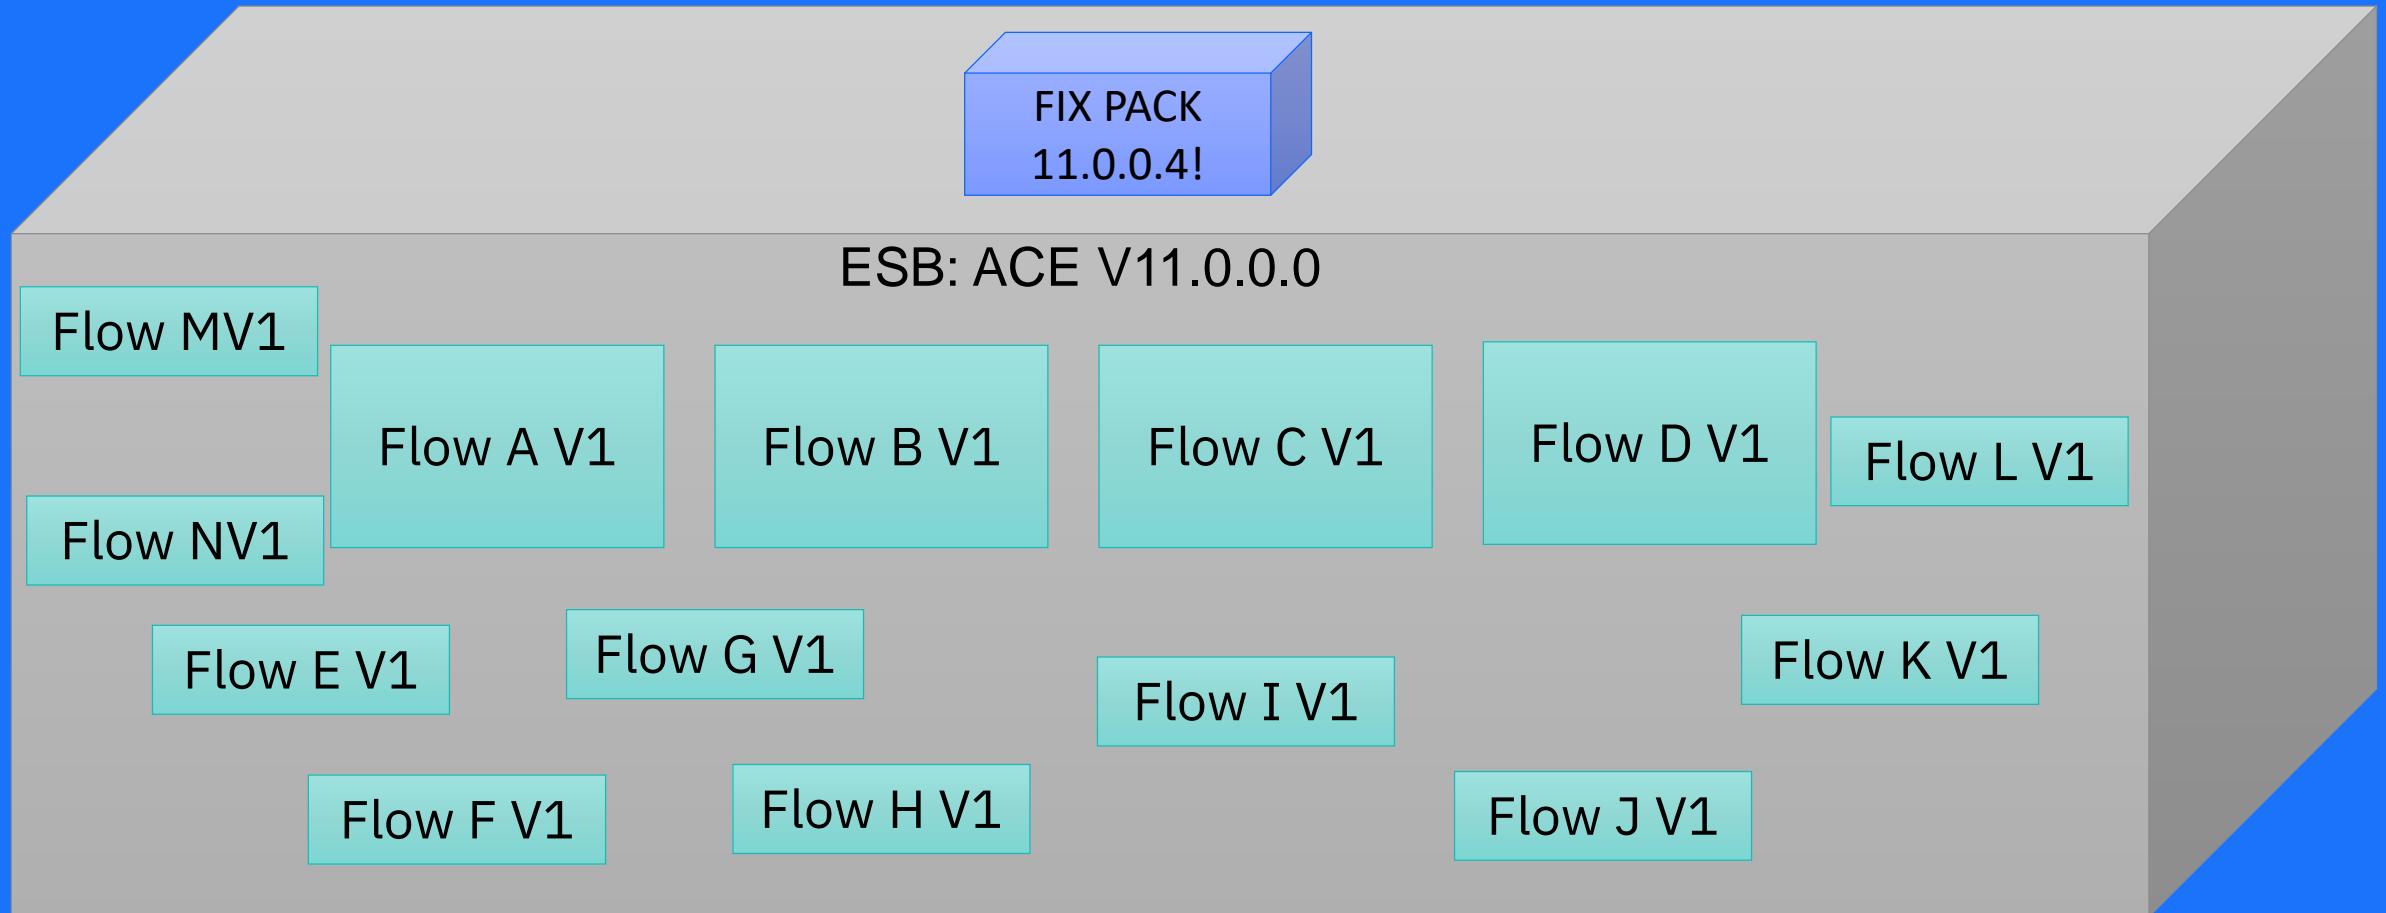
Importance of Independence



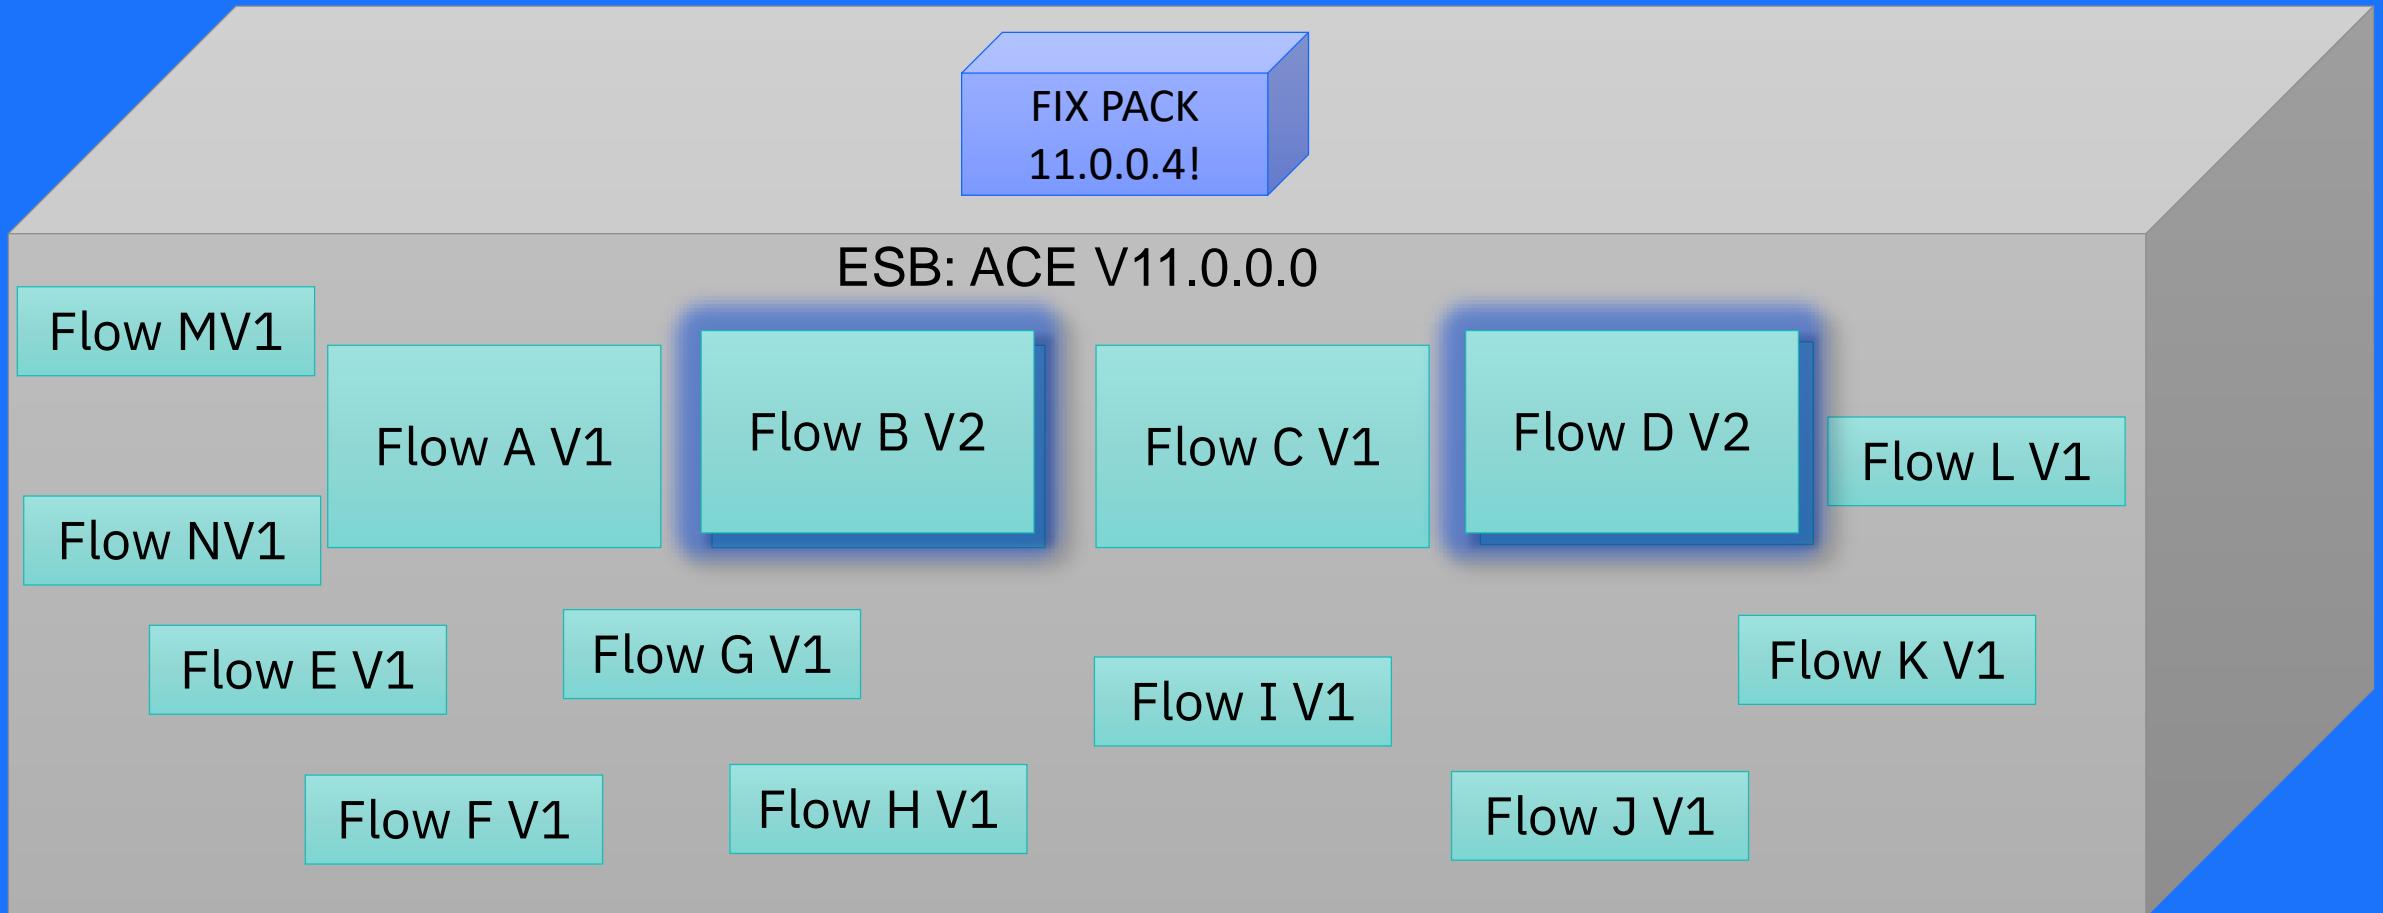
Importance of Independence



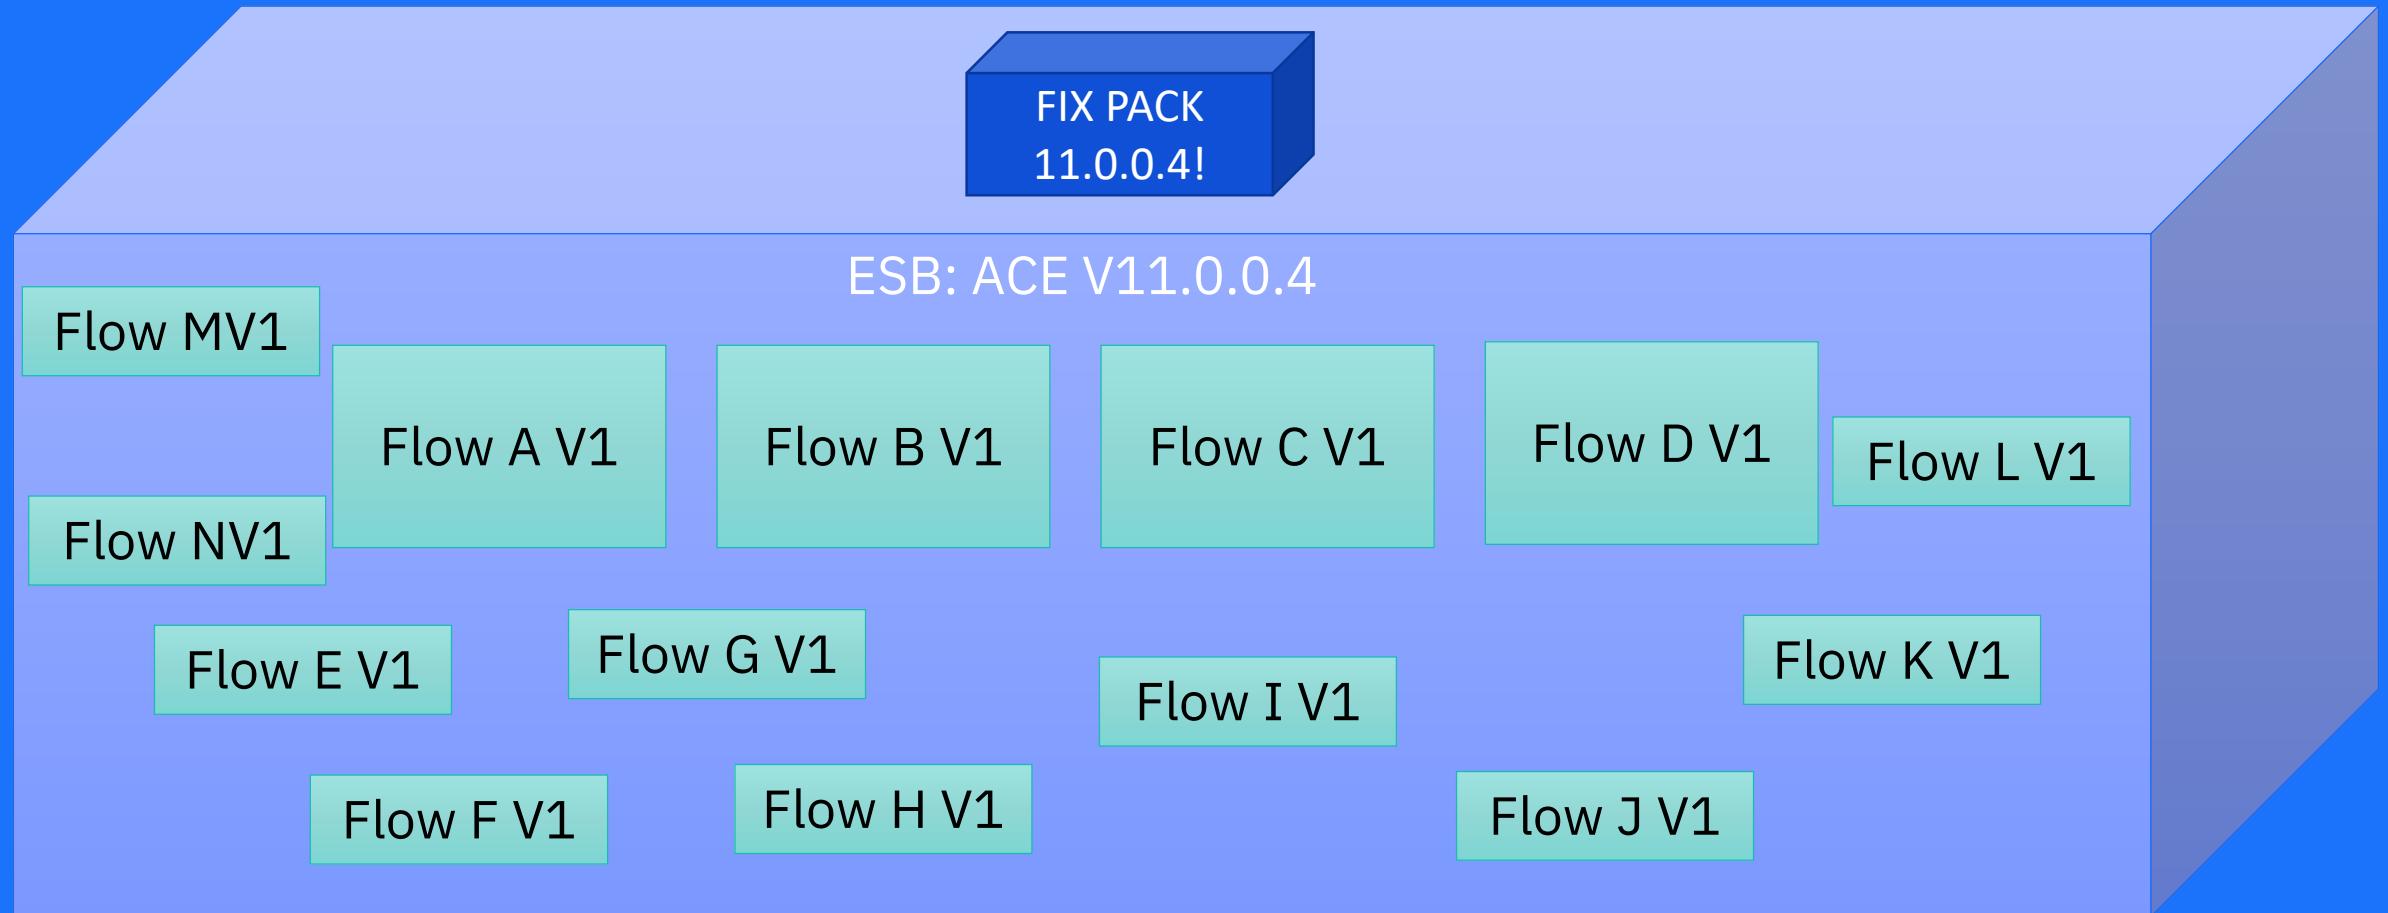
Importance of Independence



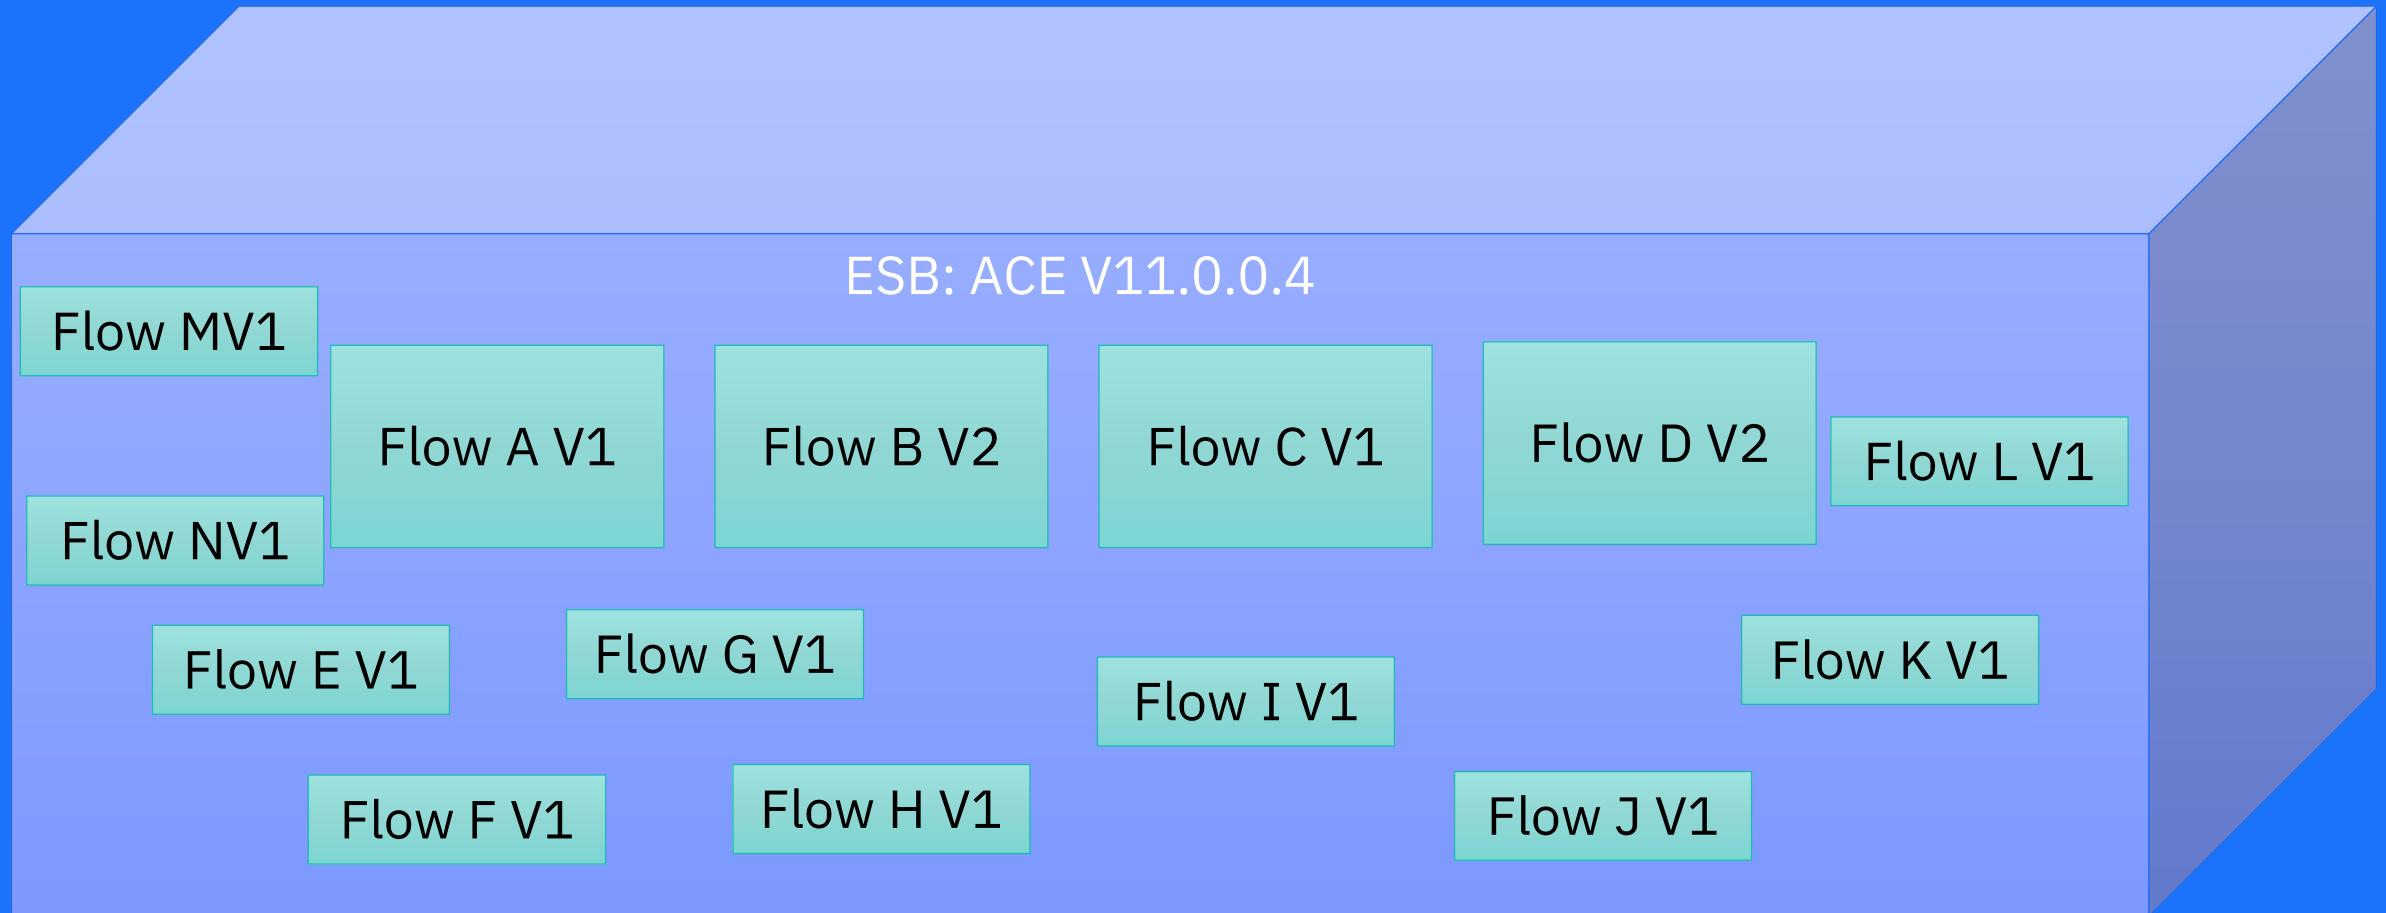
Importance of Independence



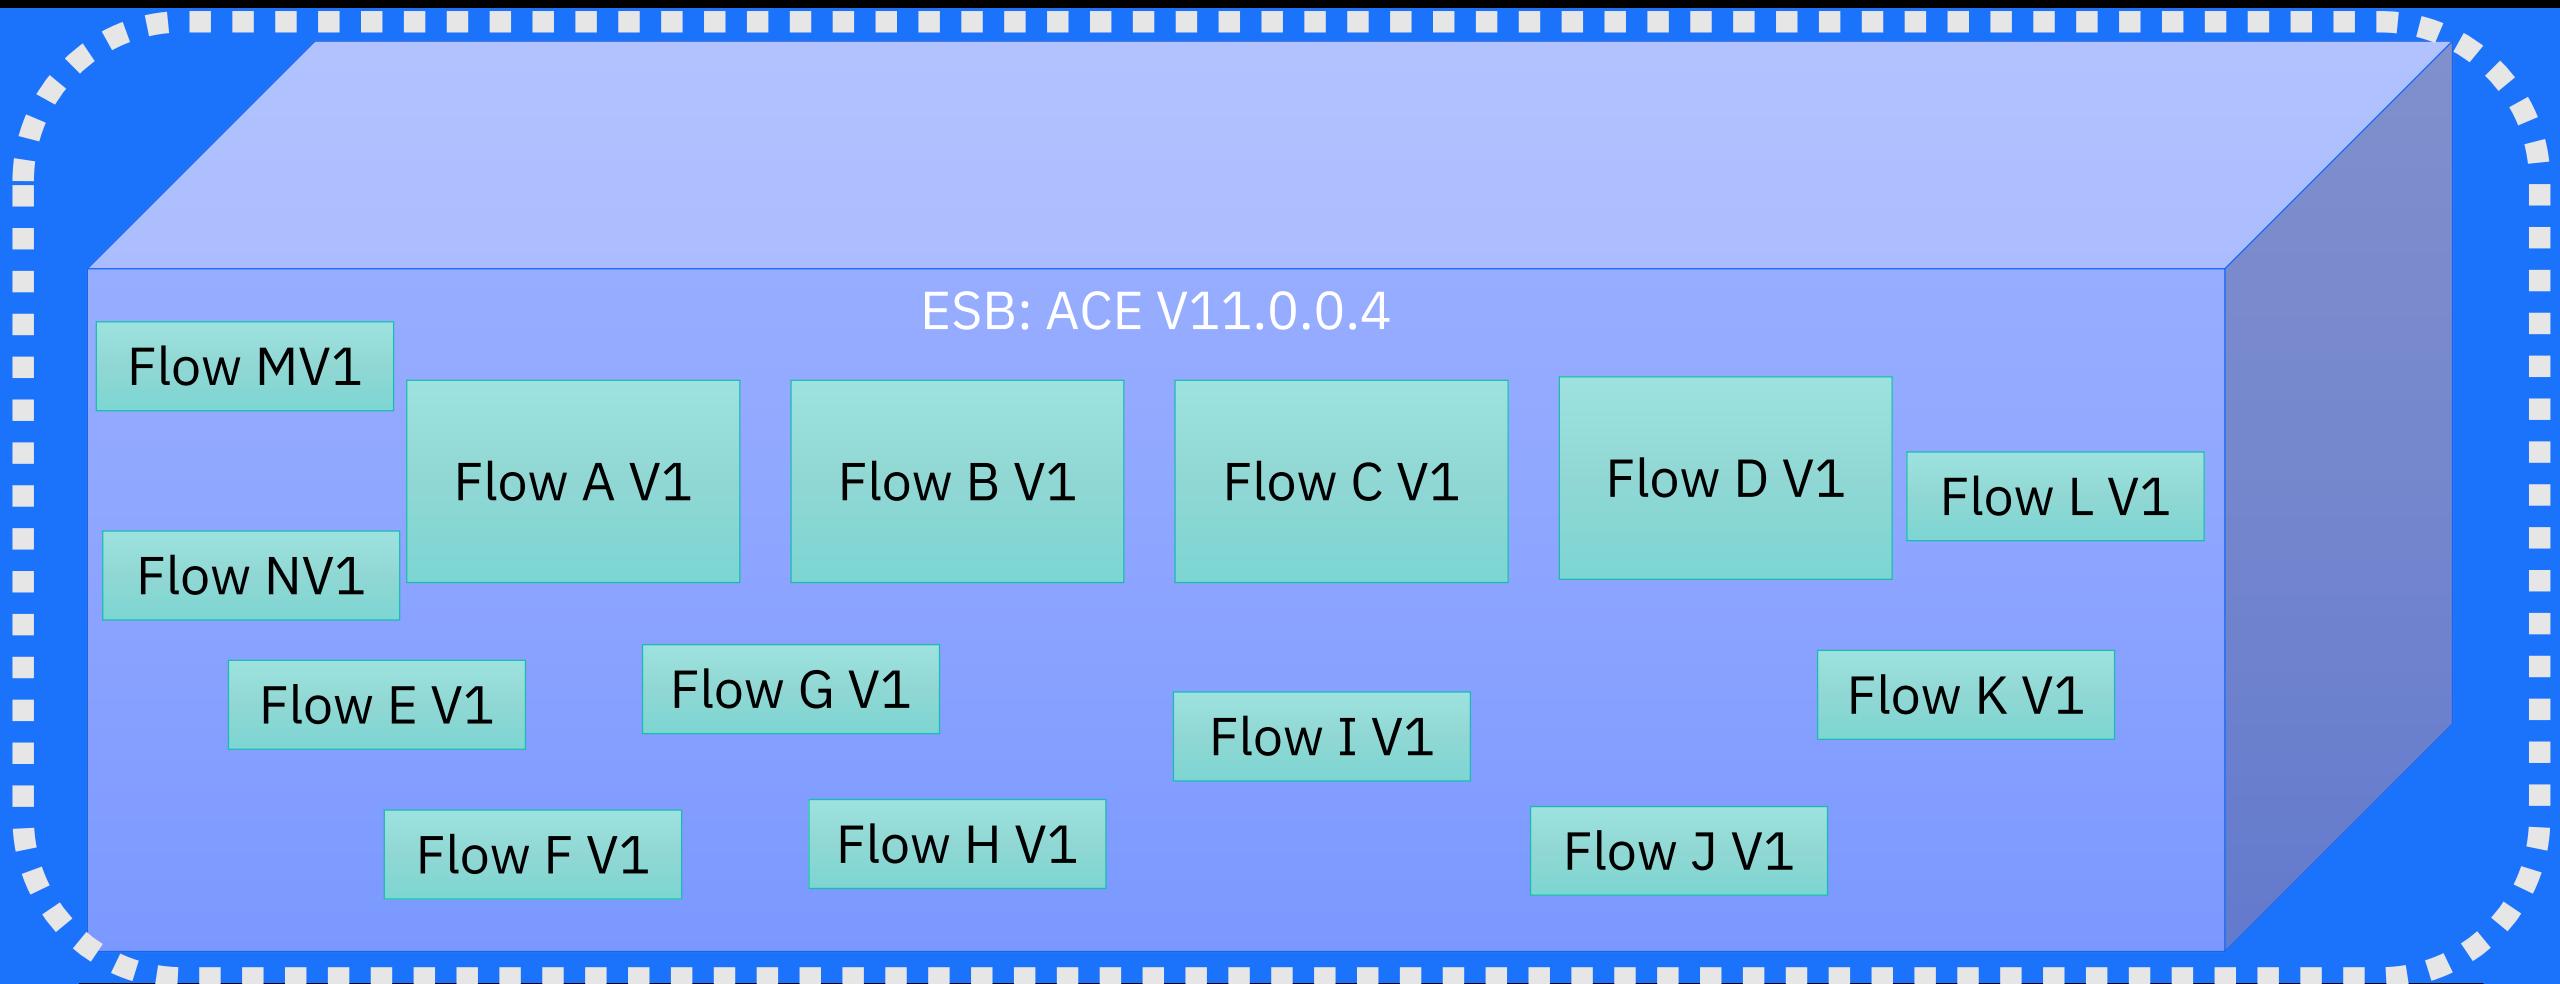
Importance of Independence



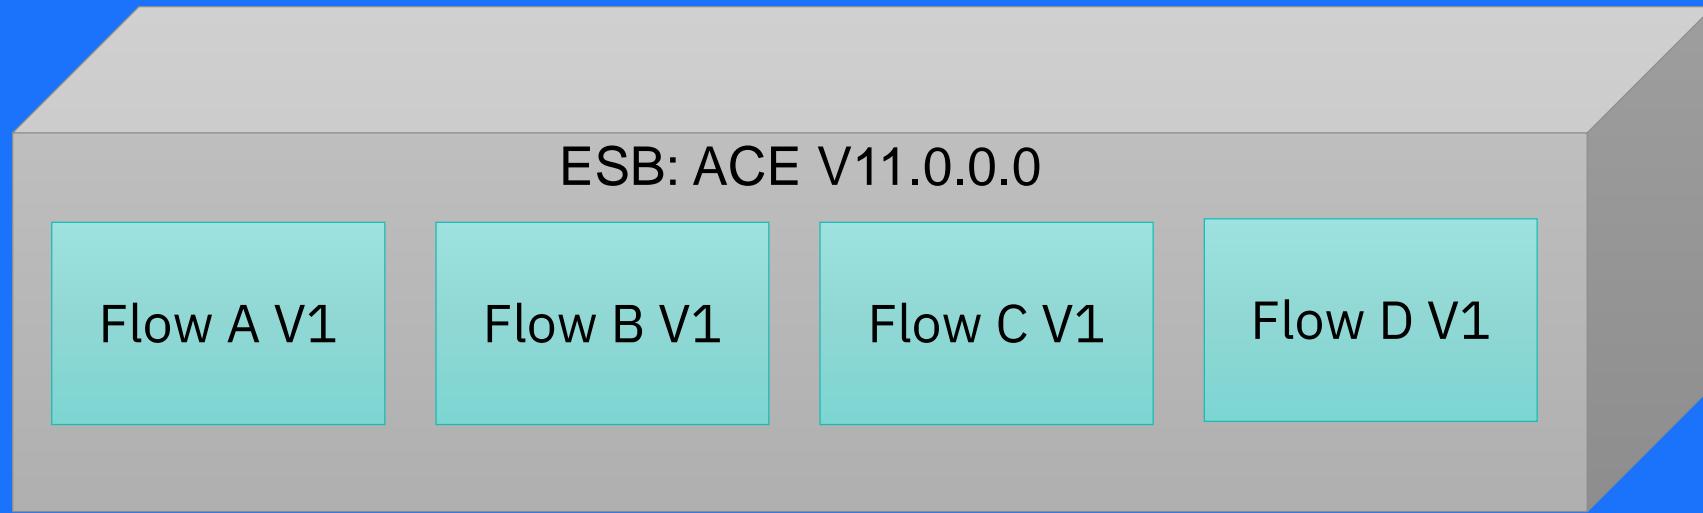
Importance of Independence



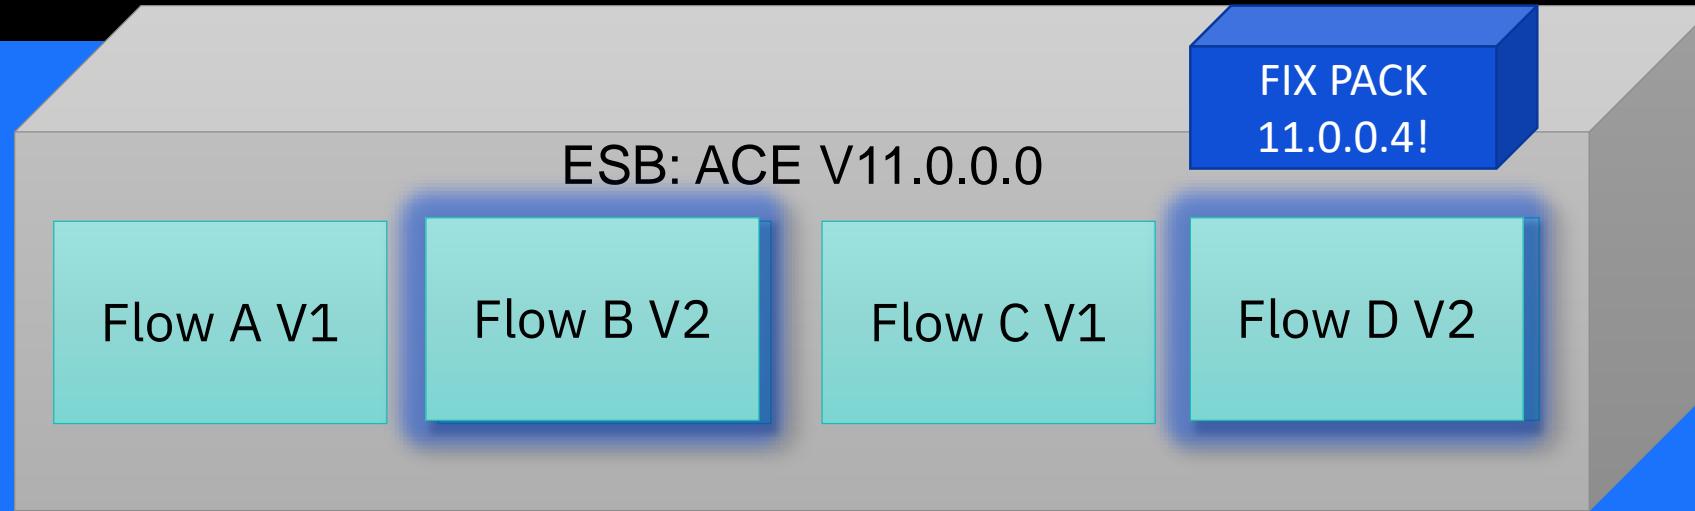
Importance of Independence



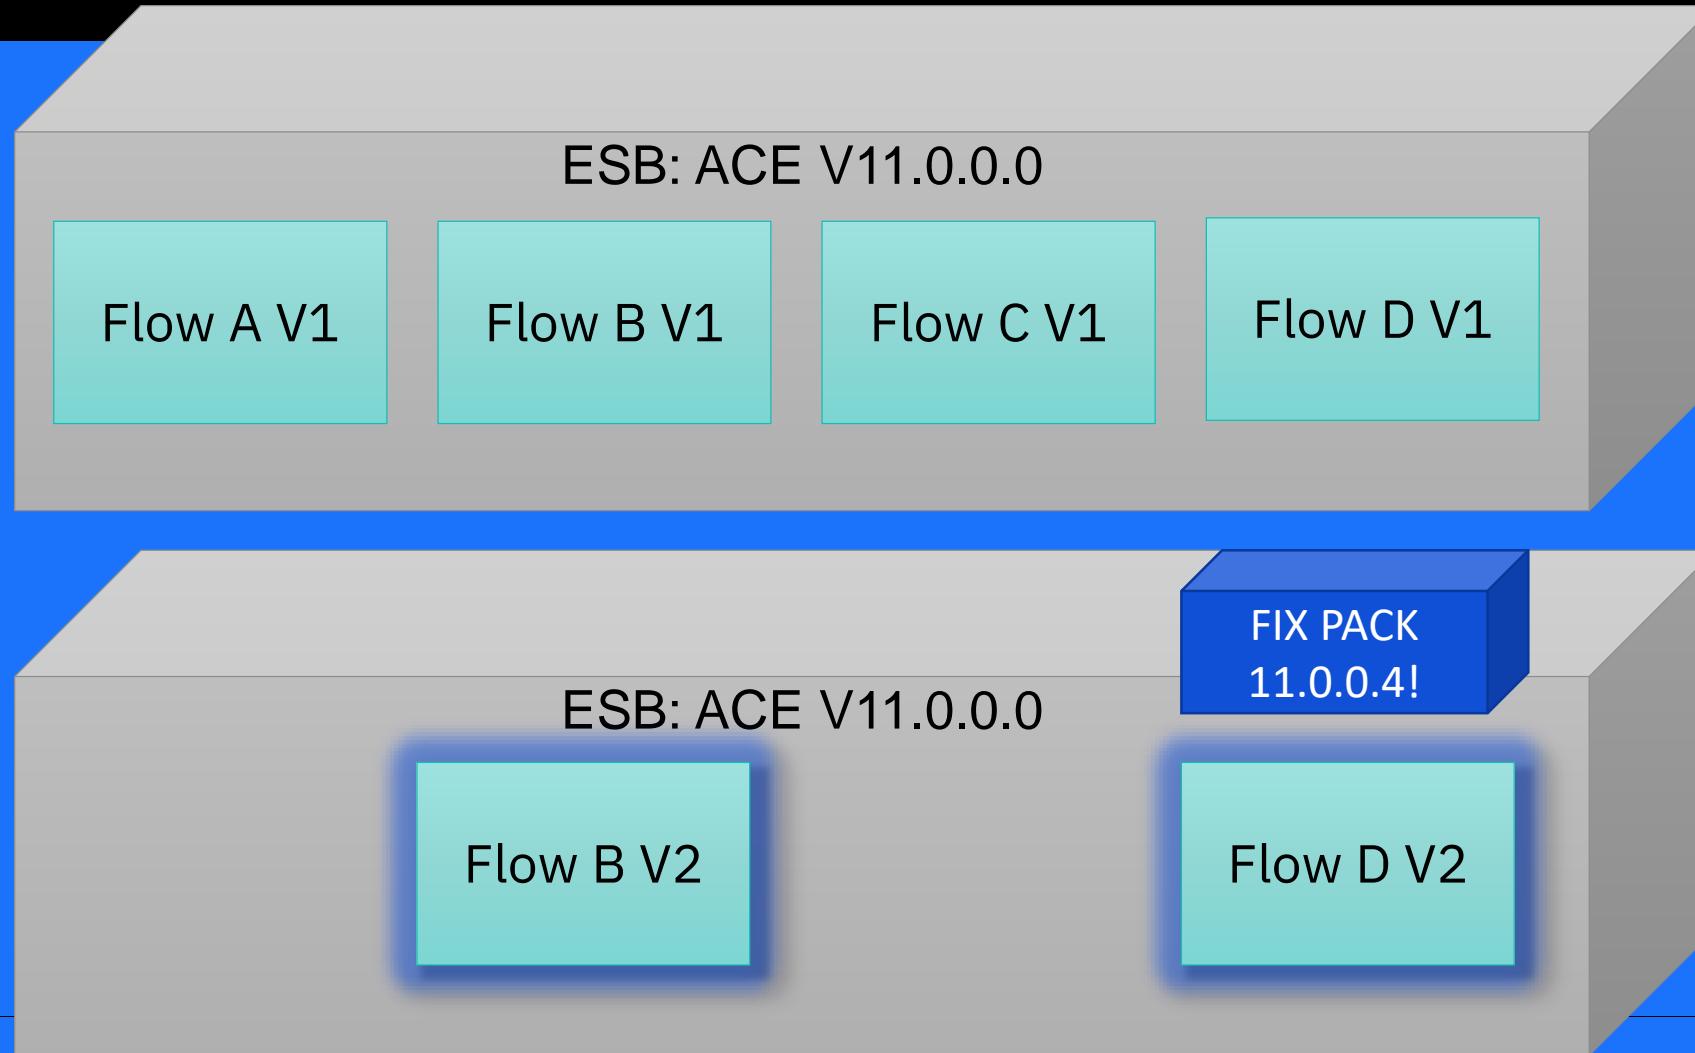
Importance of Independence



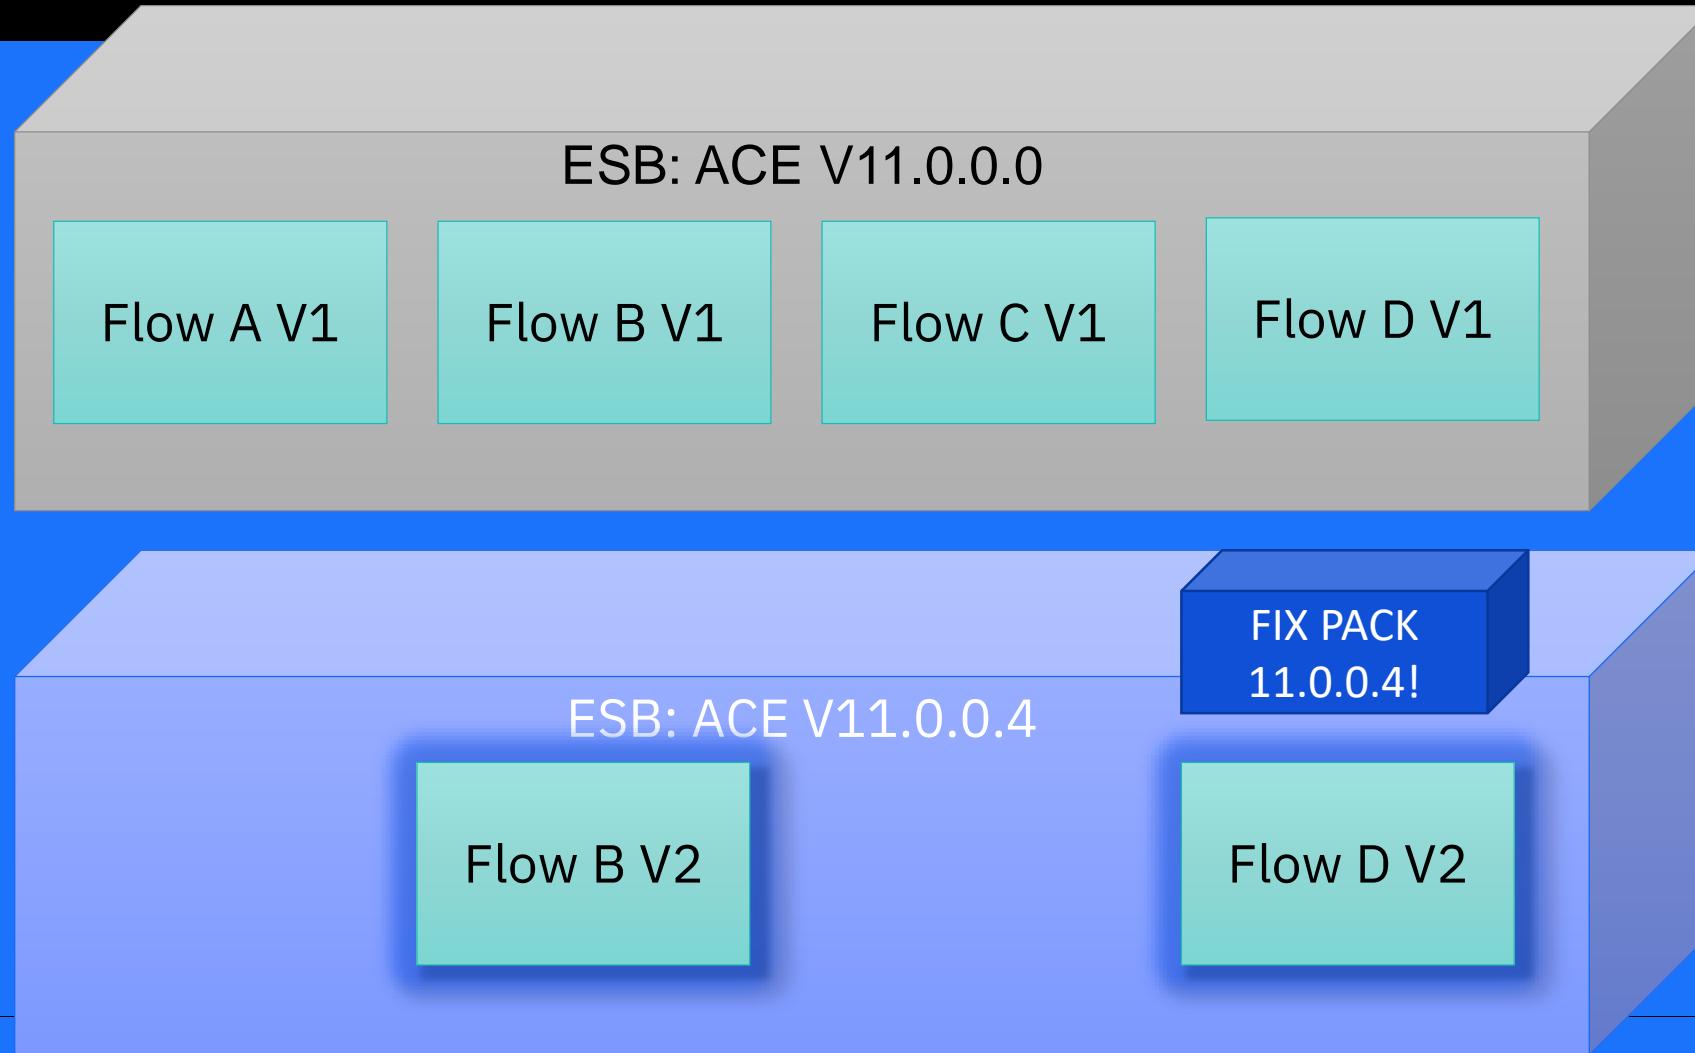
Importance of Independence



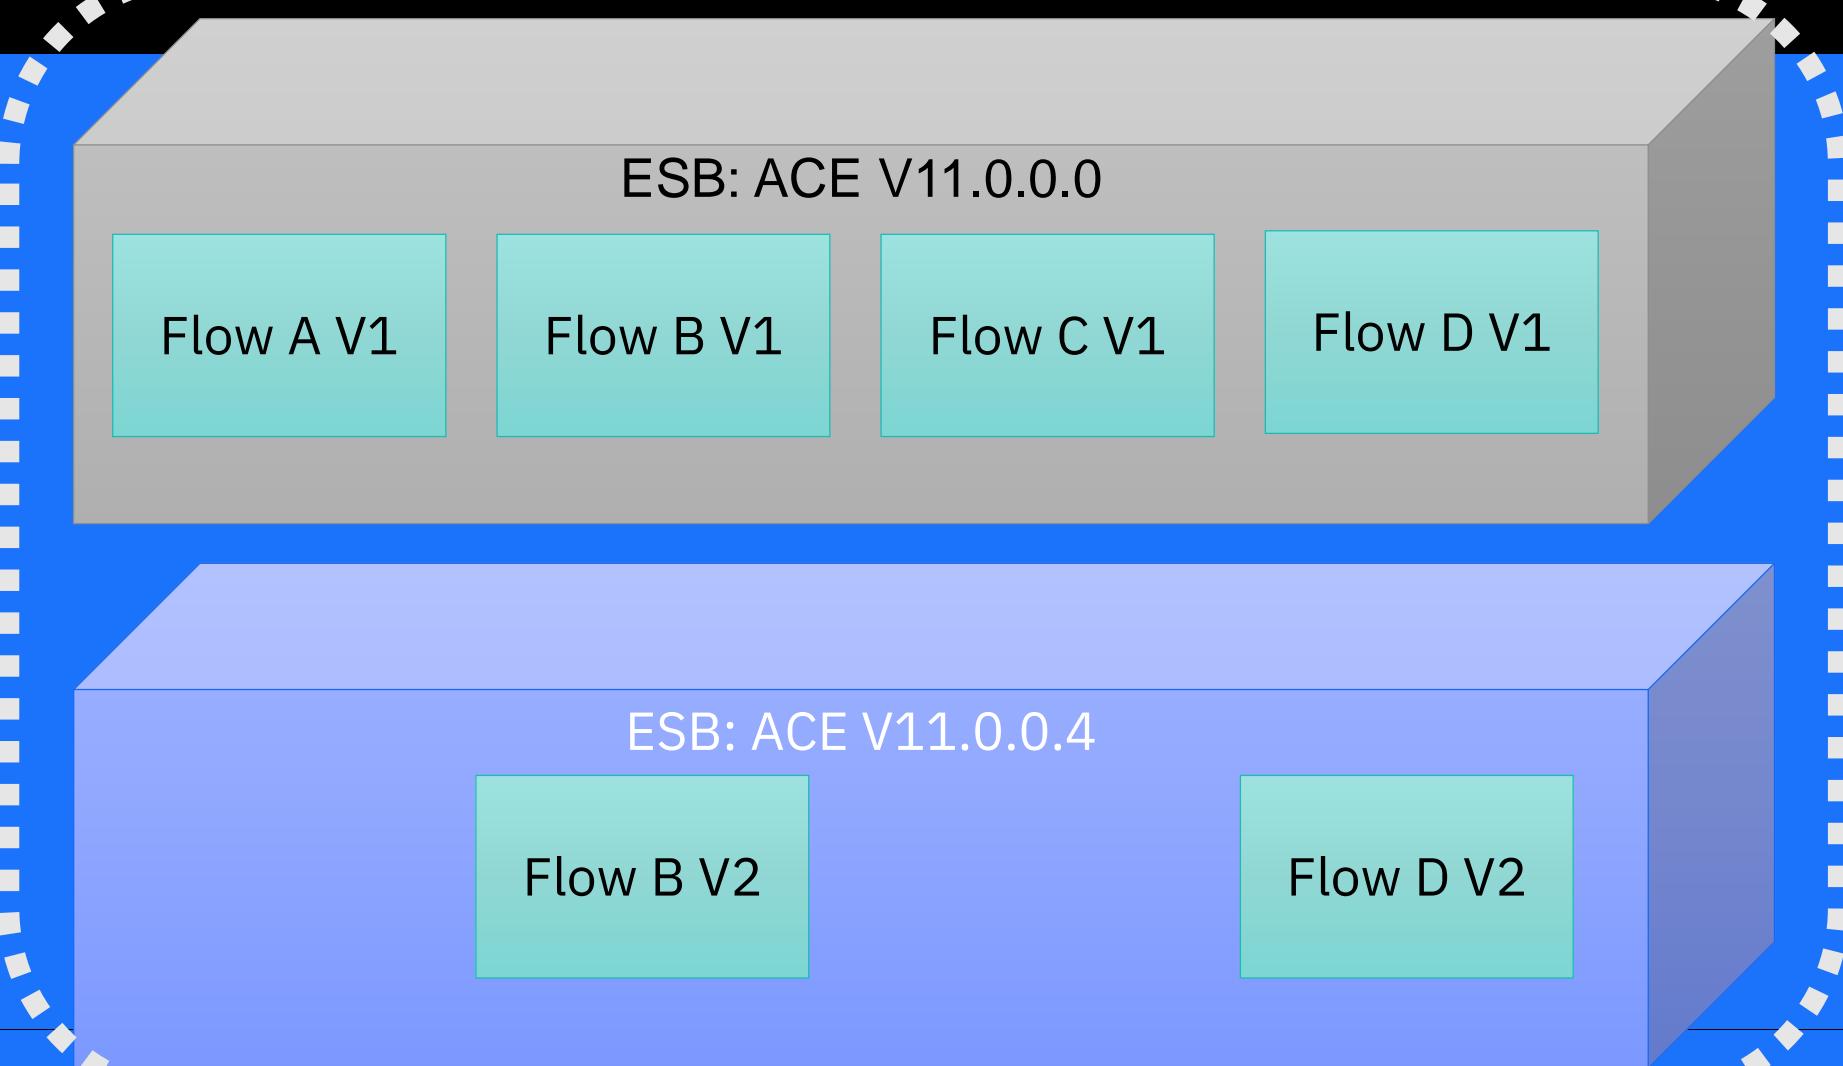
Importance of Independence



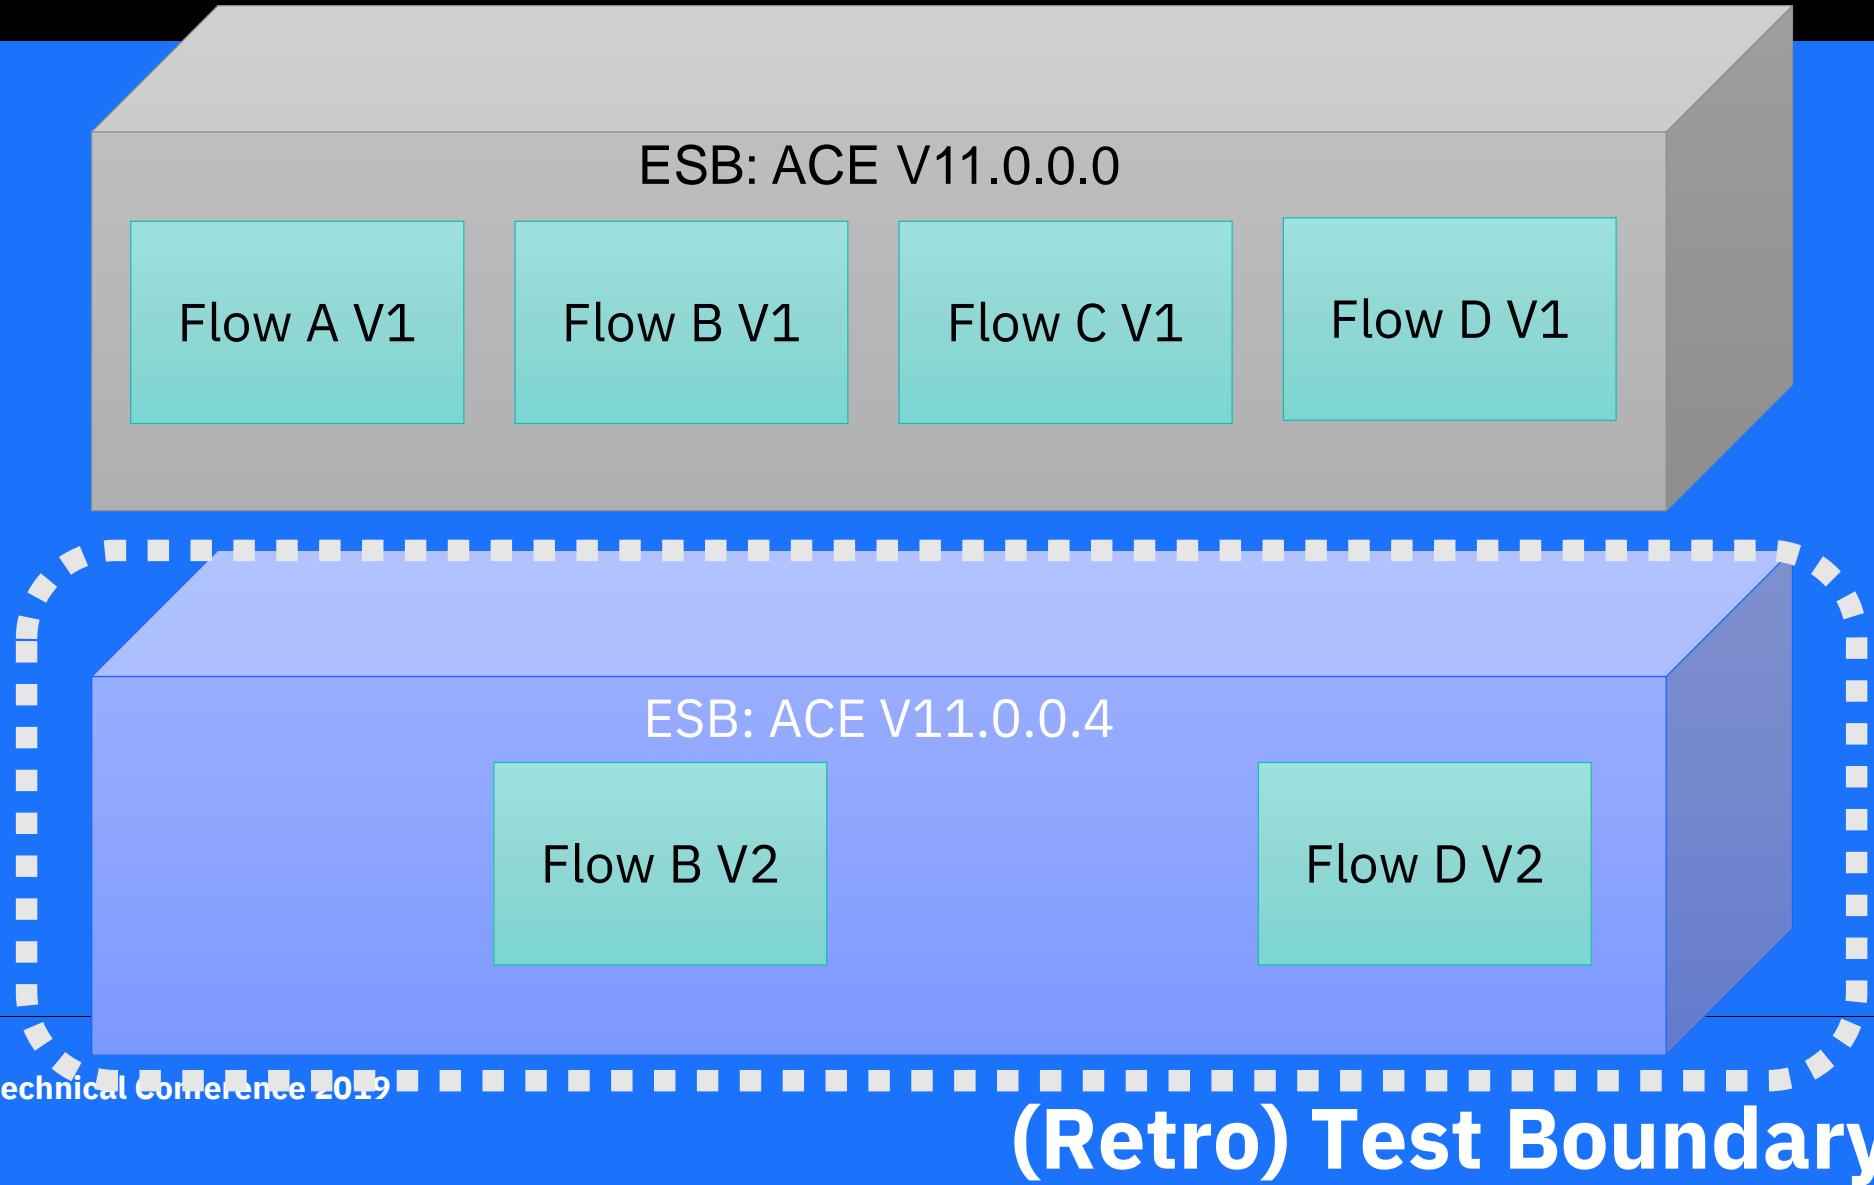
Importance of Independence



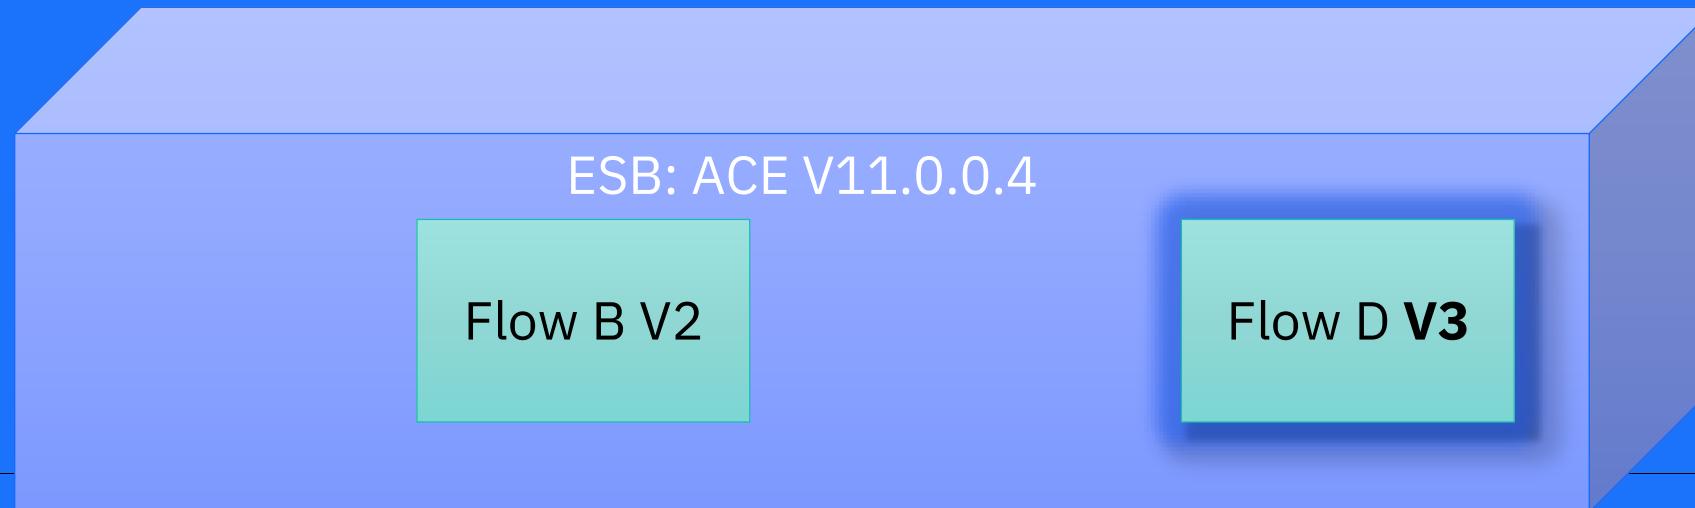
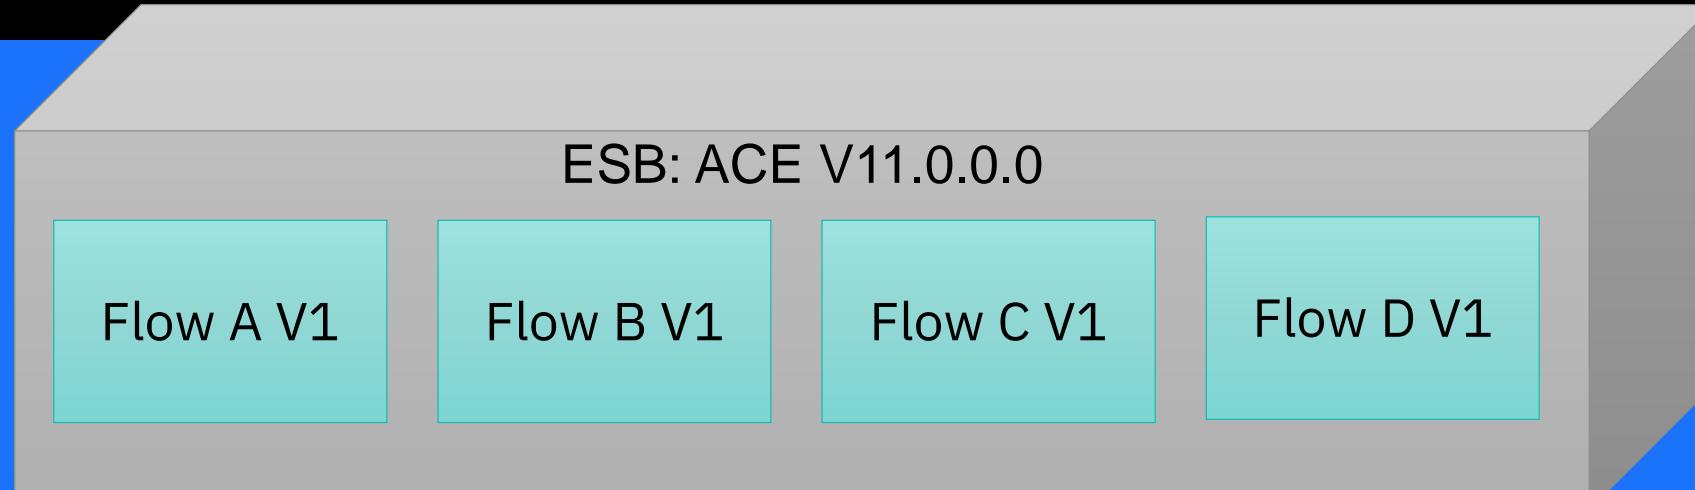
Importance of Independence



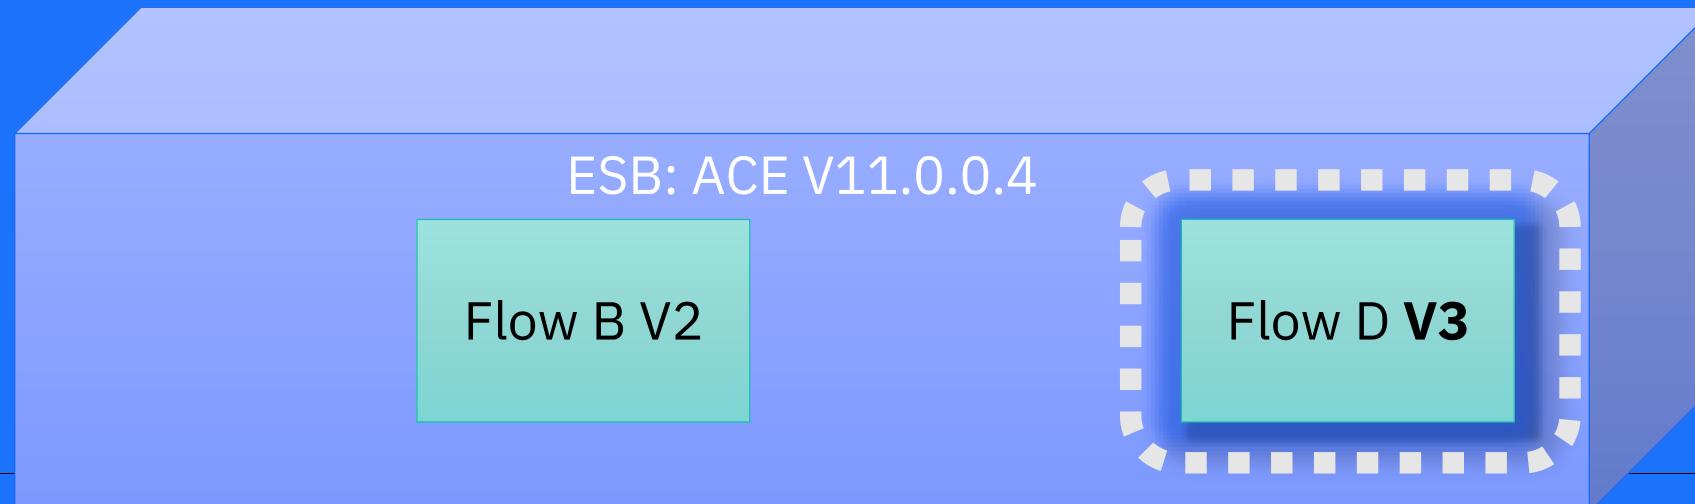
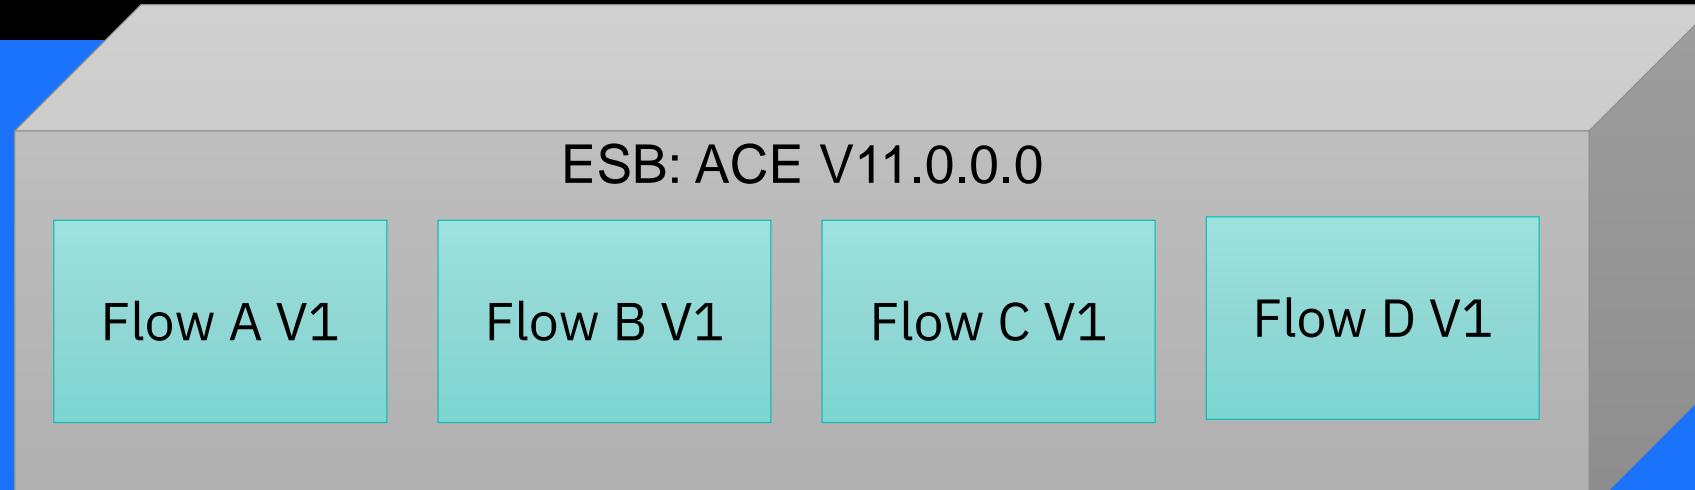
Importance of Independence



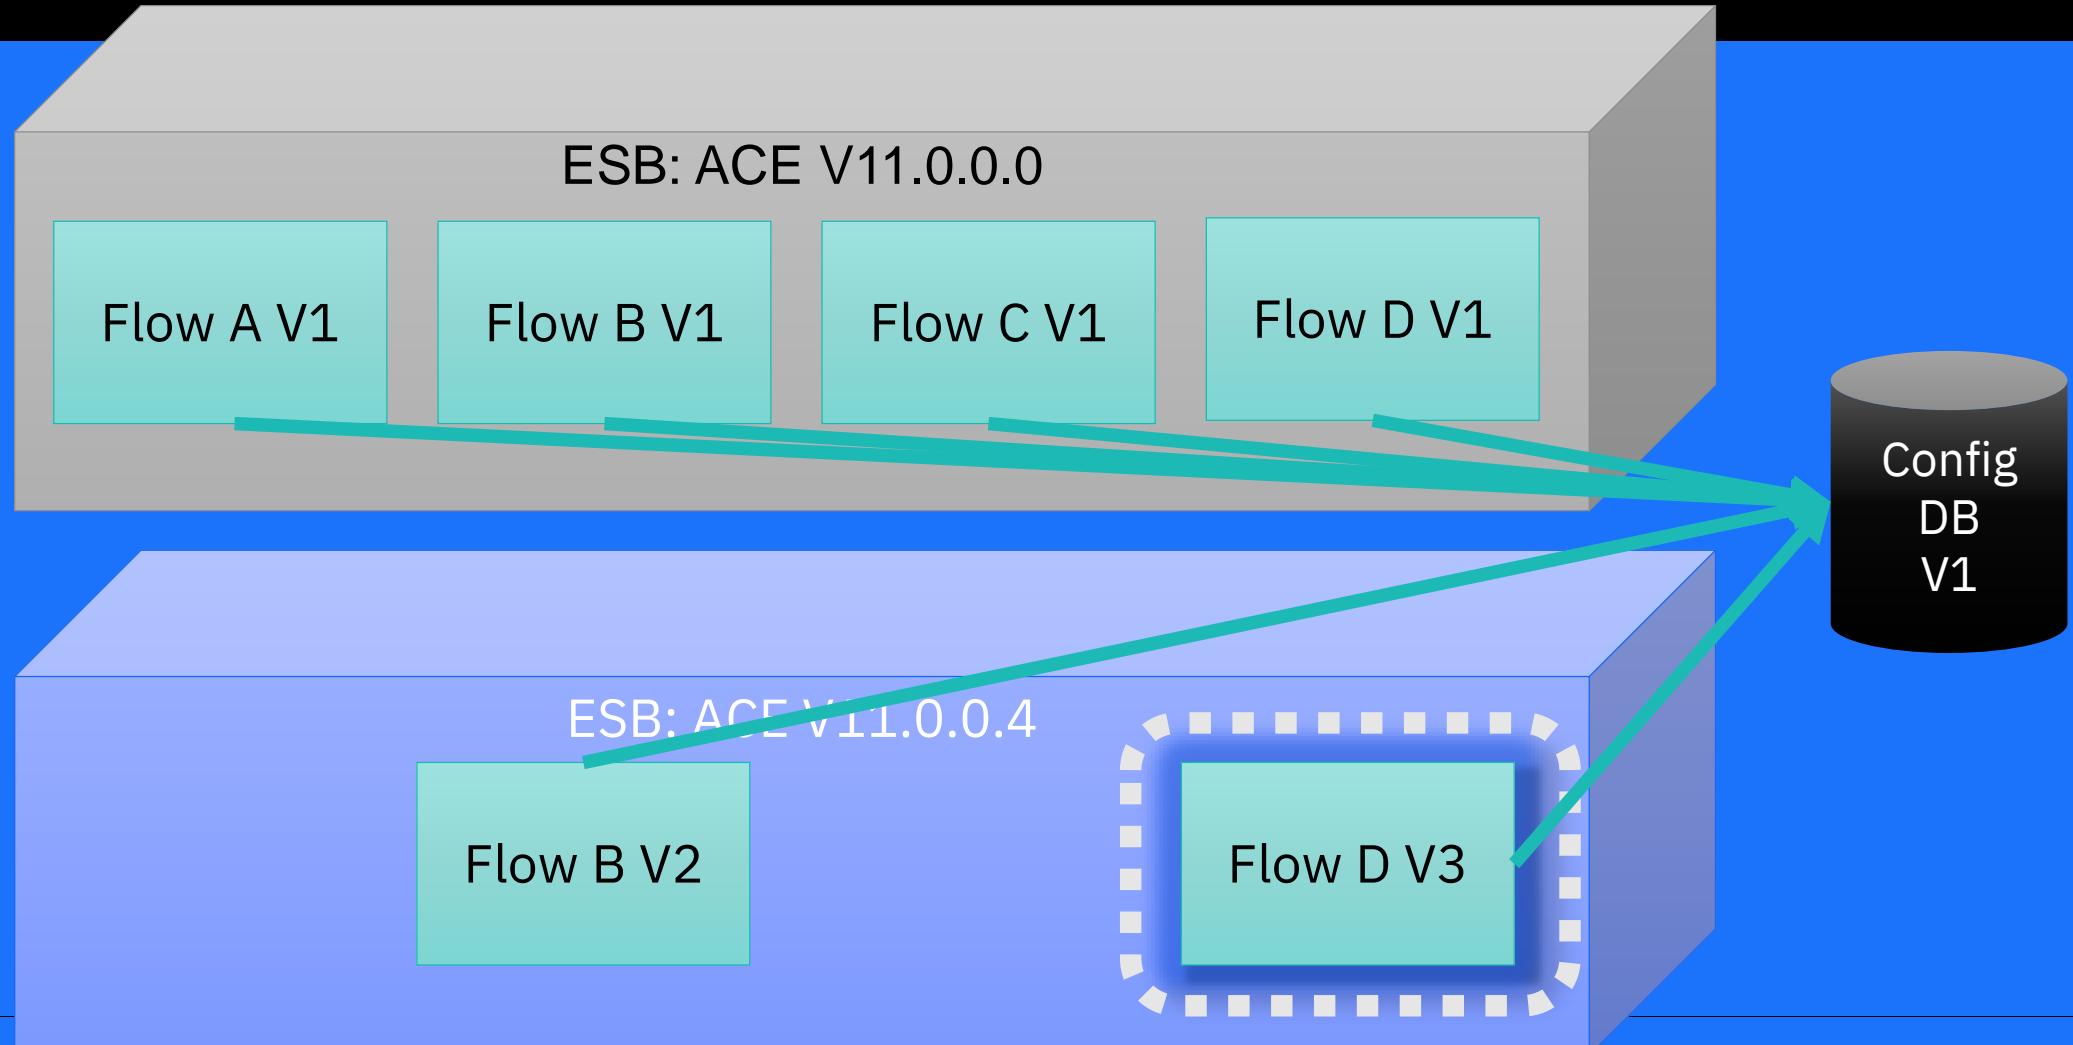
Importance of Independence



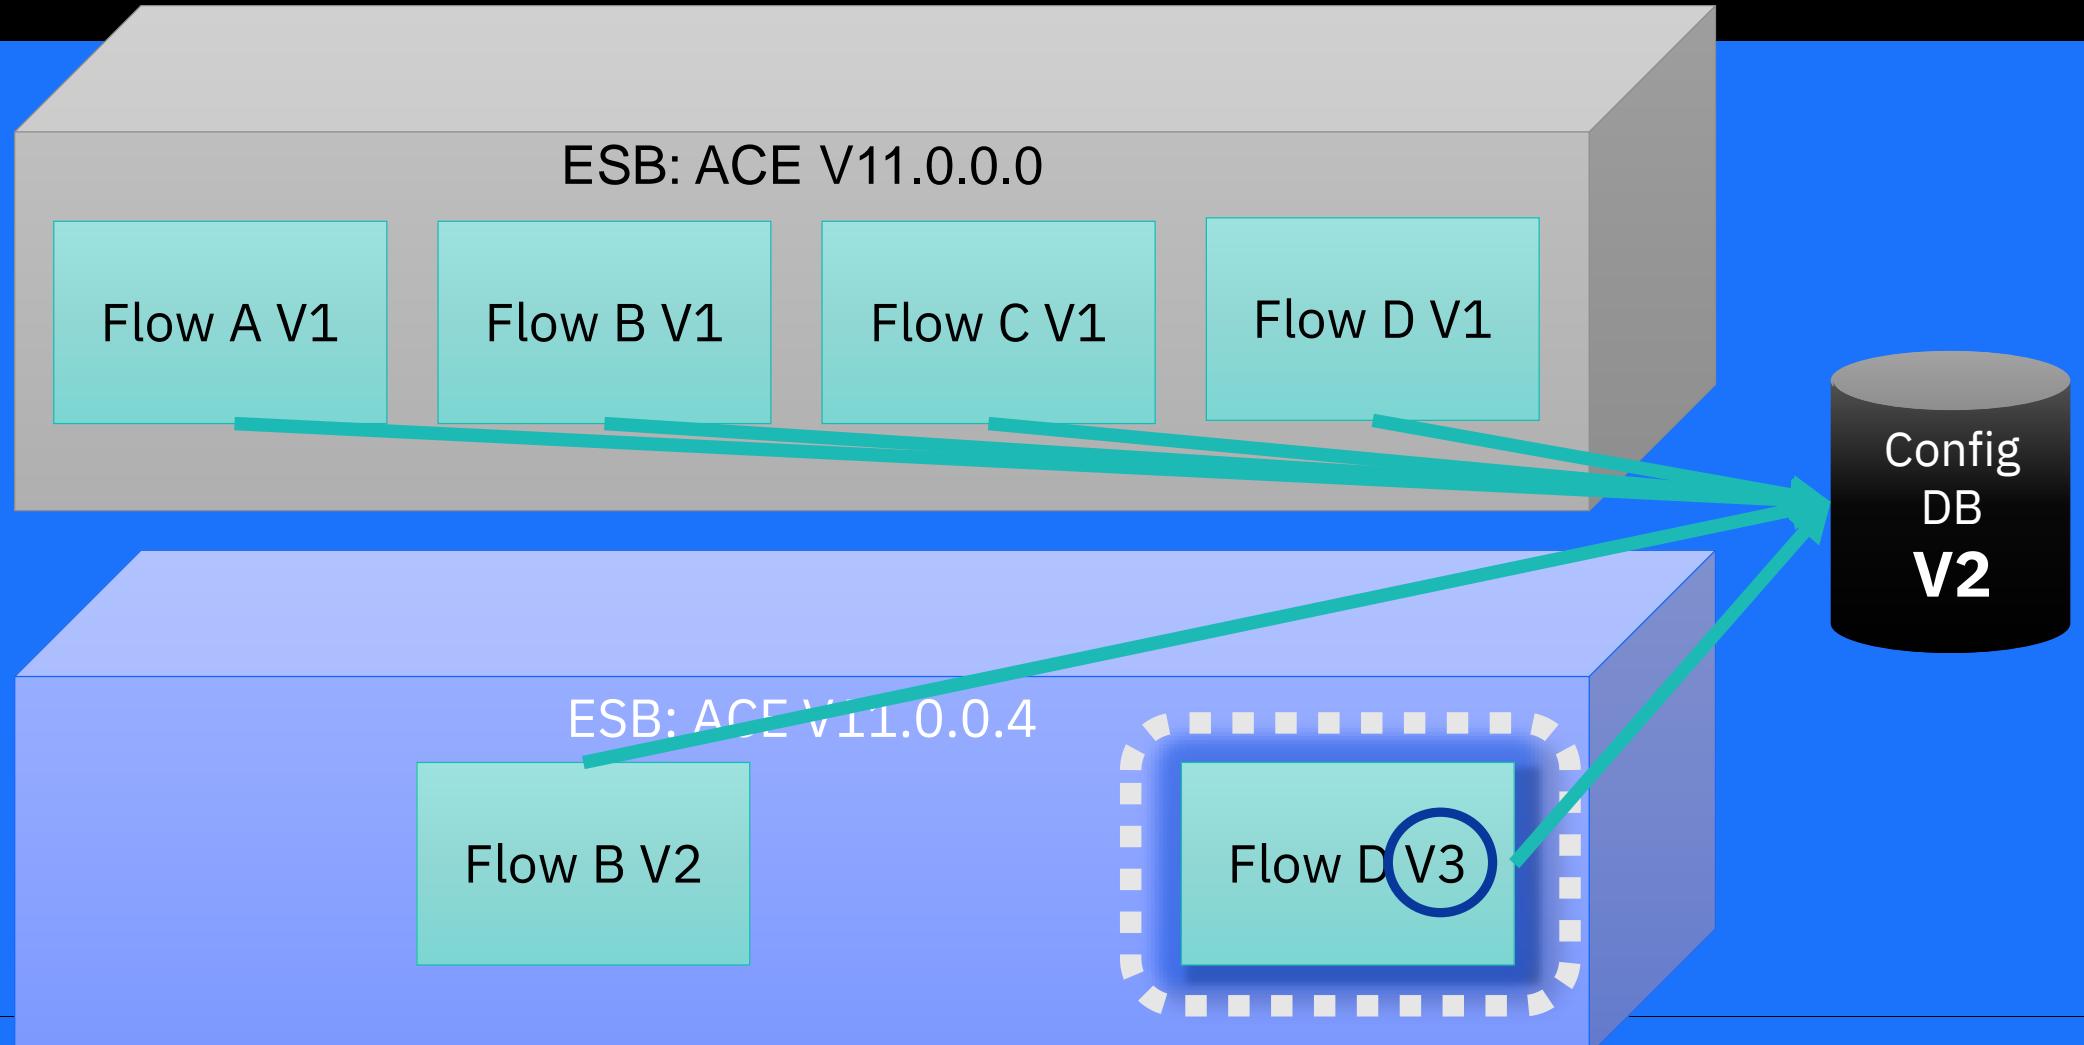
Importance of Independence



Importance of Independence

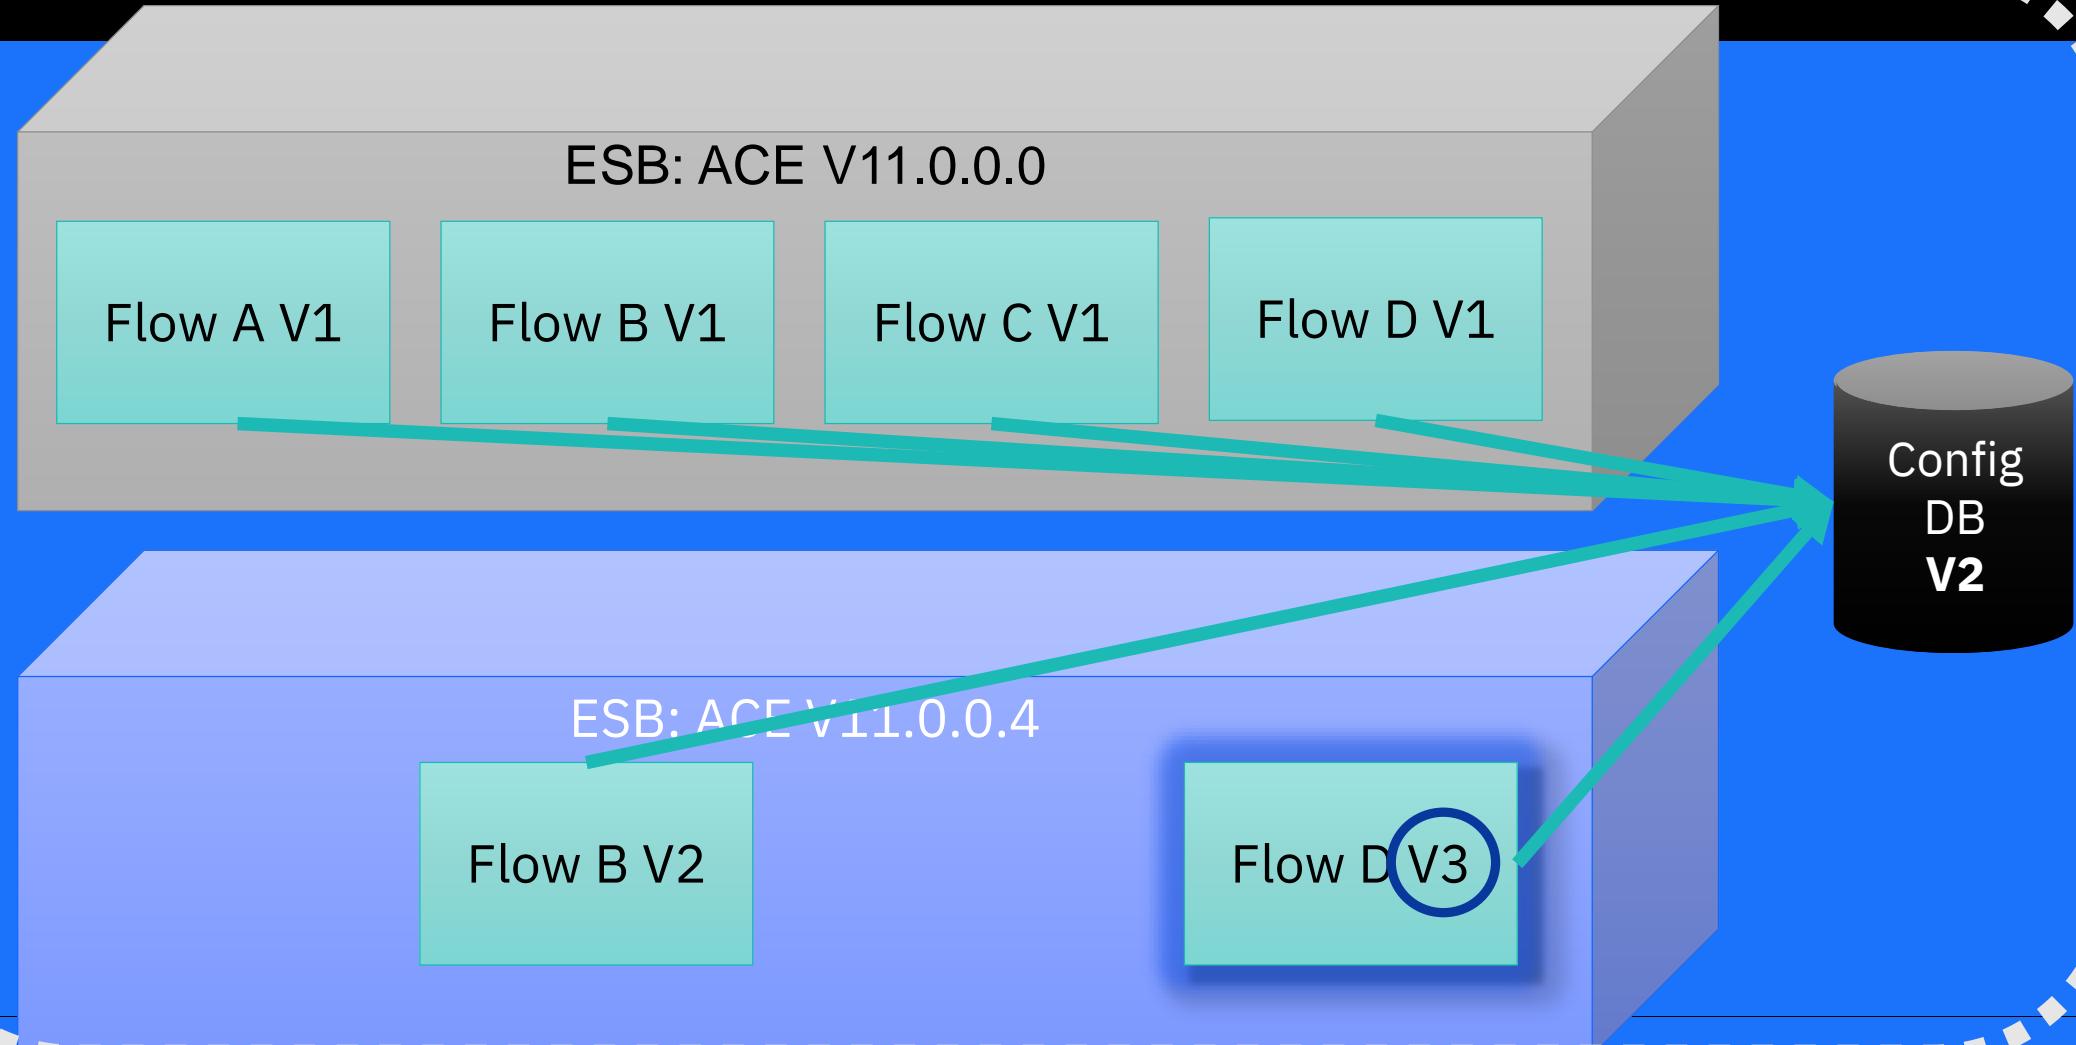


Importance of Independence



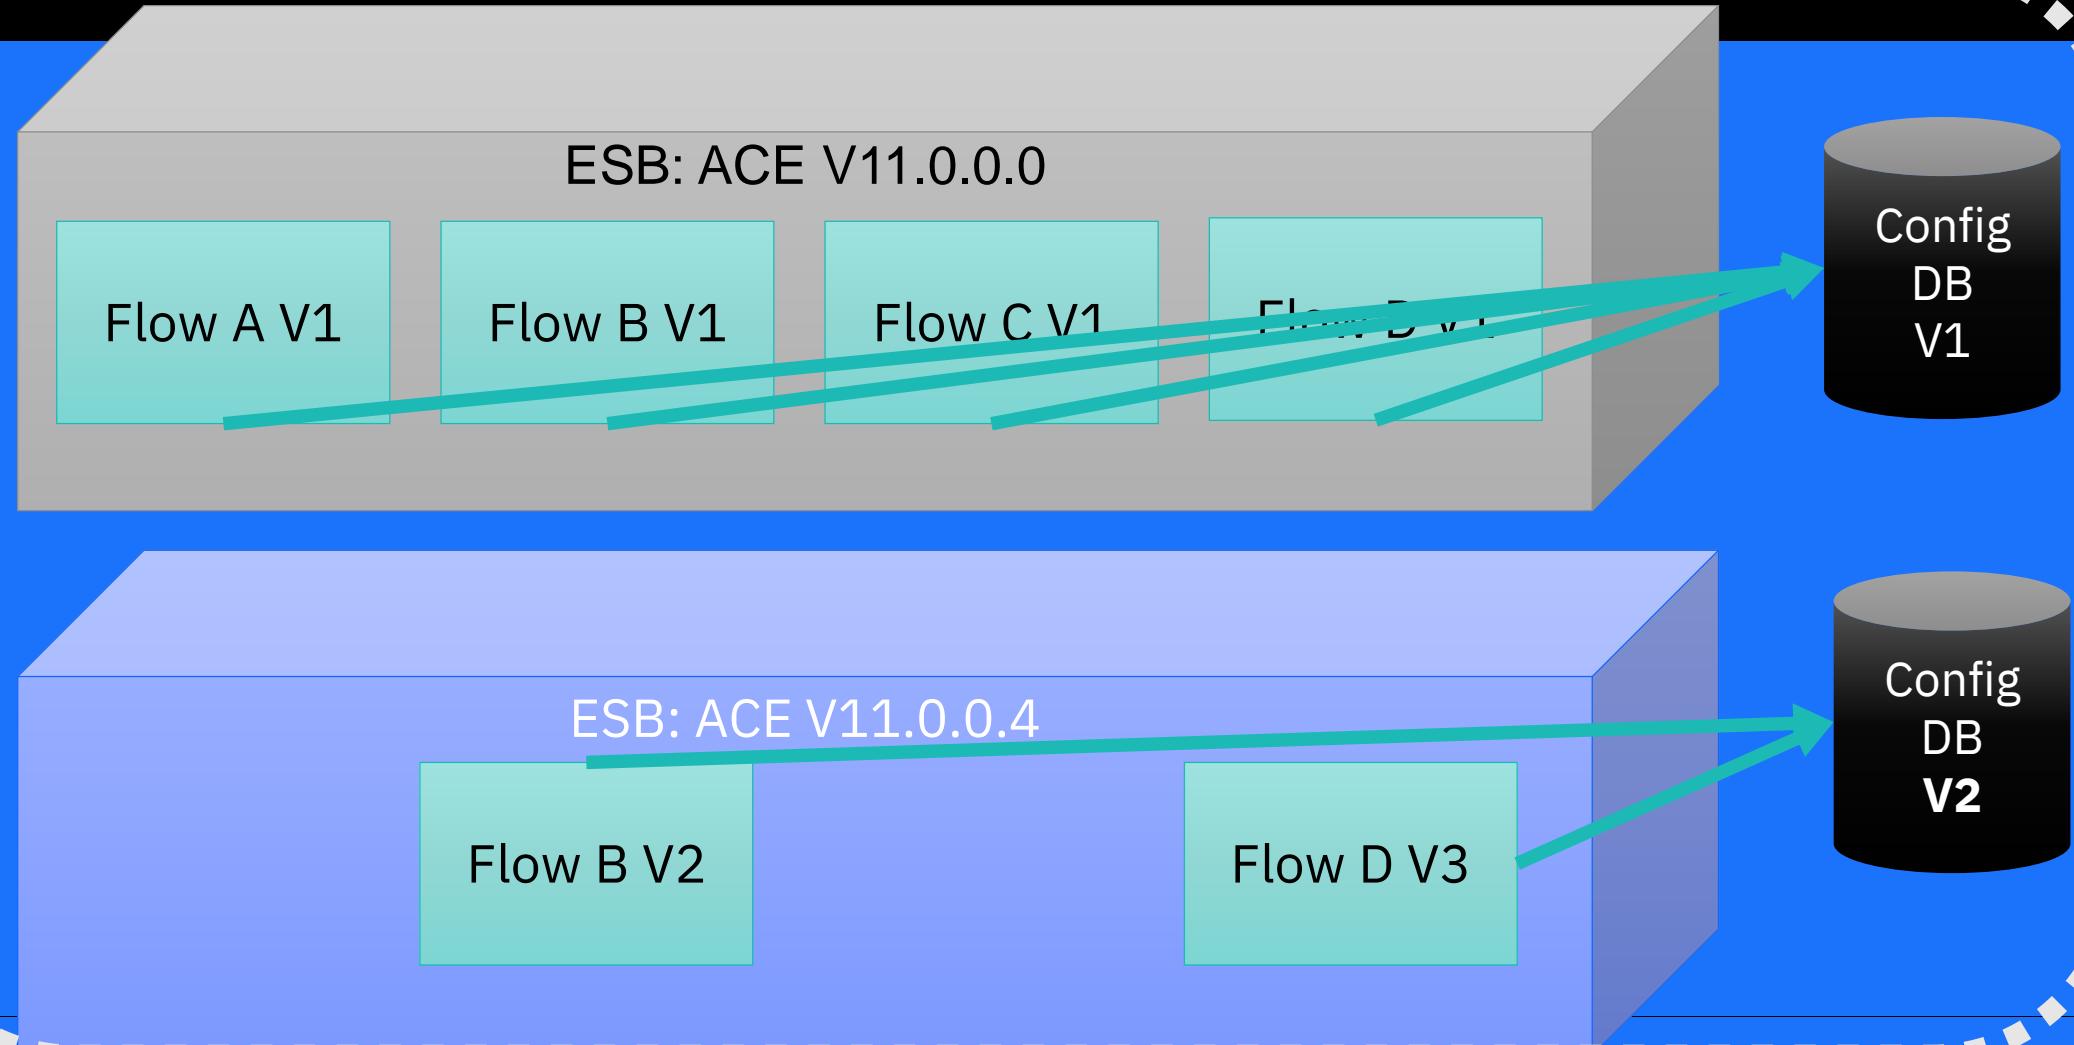
Importance of Independence

IBM

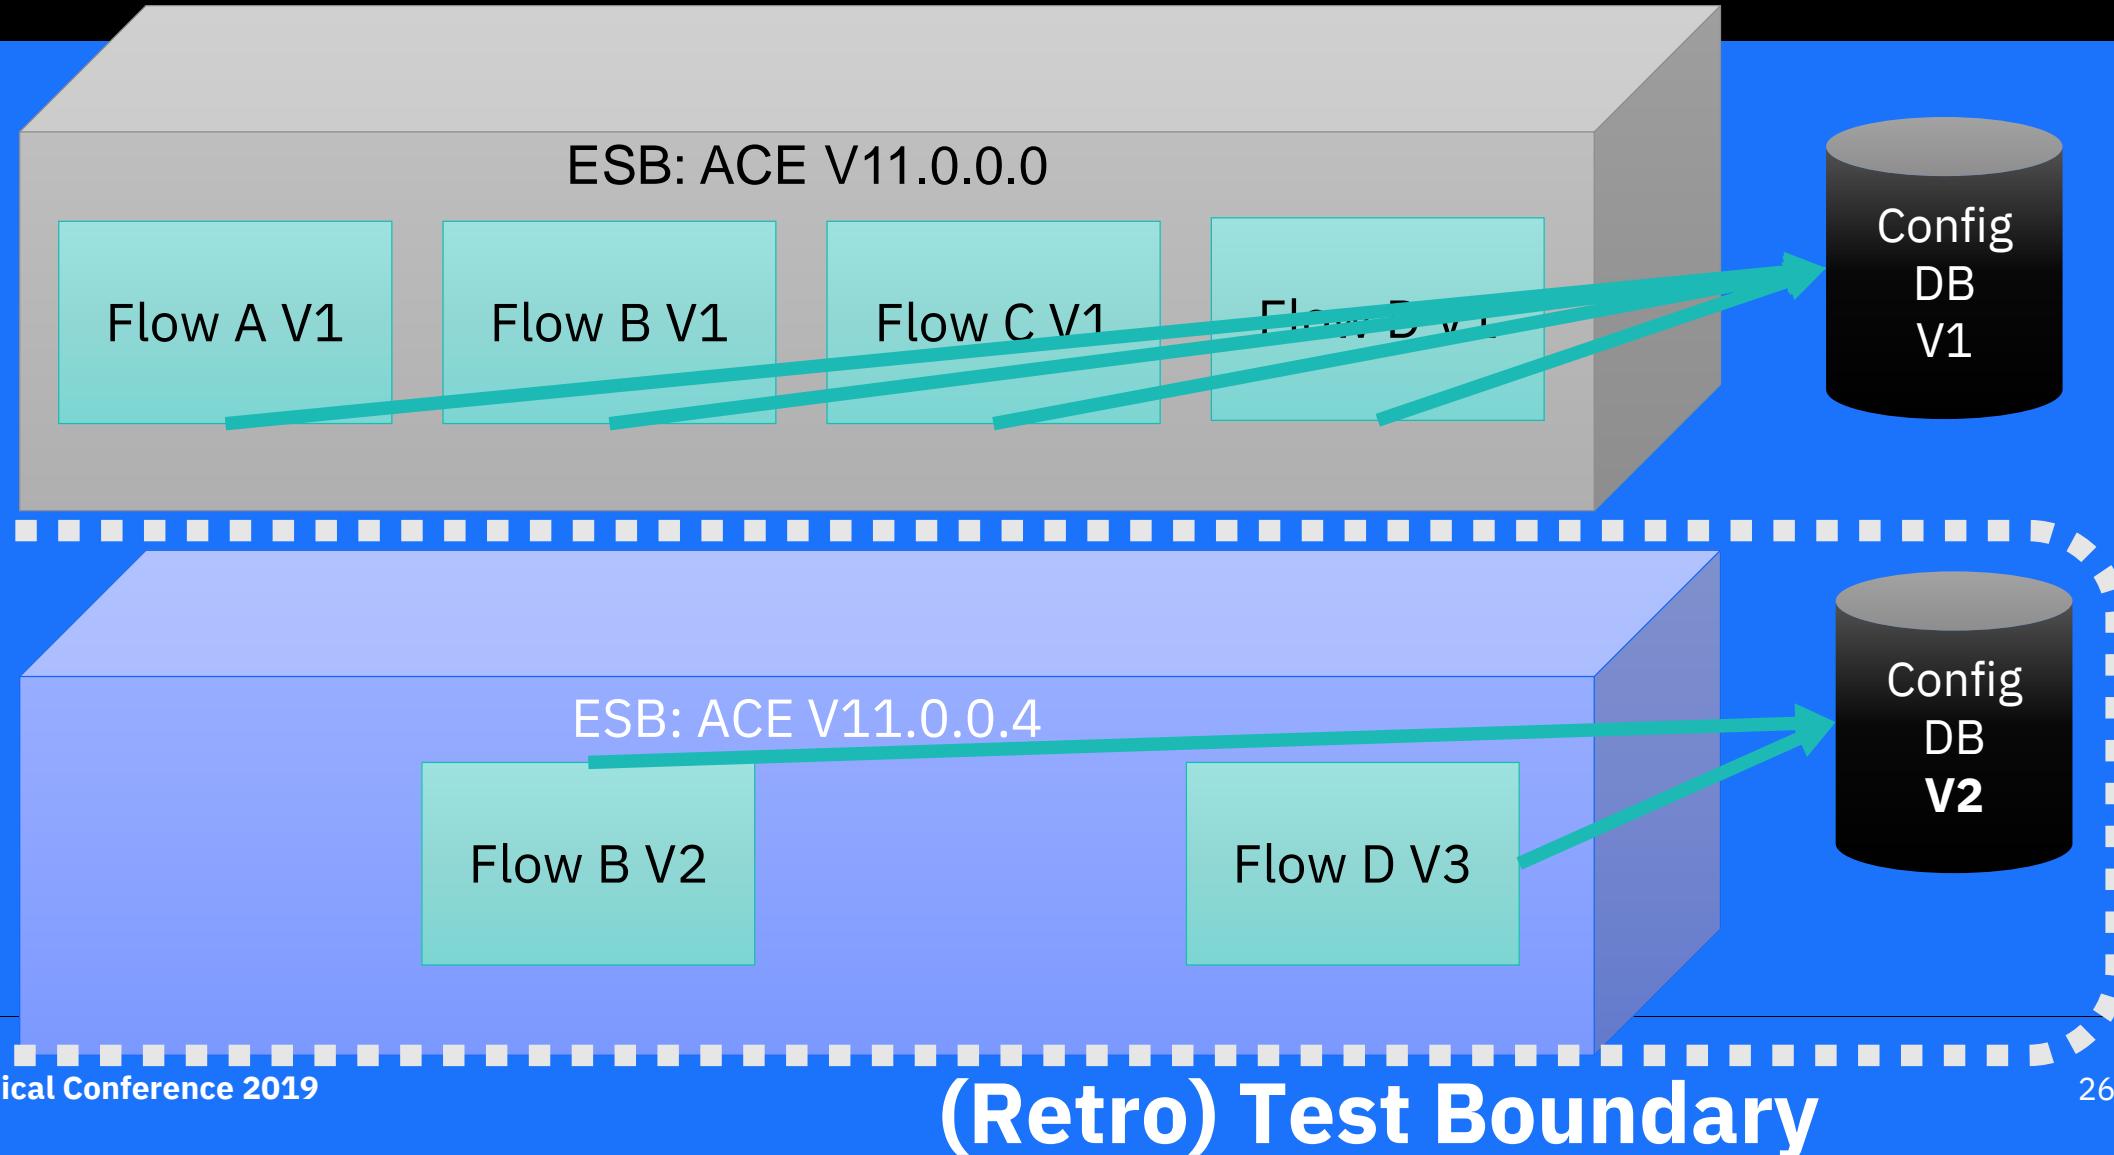


Importance of Independence

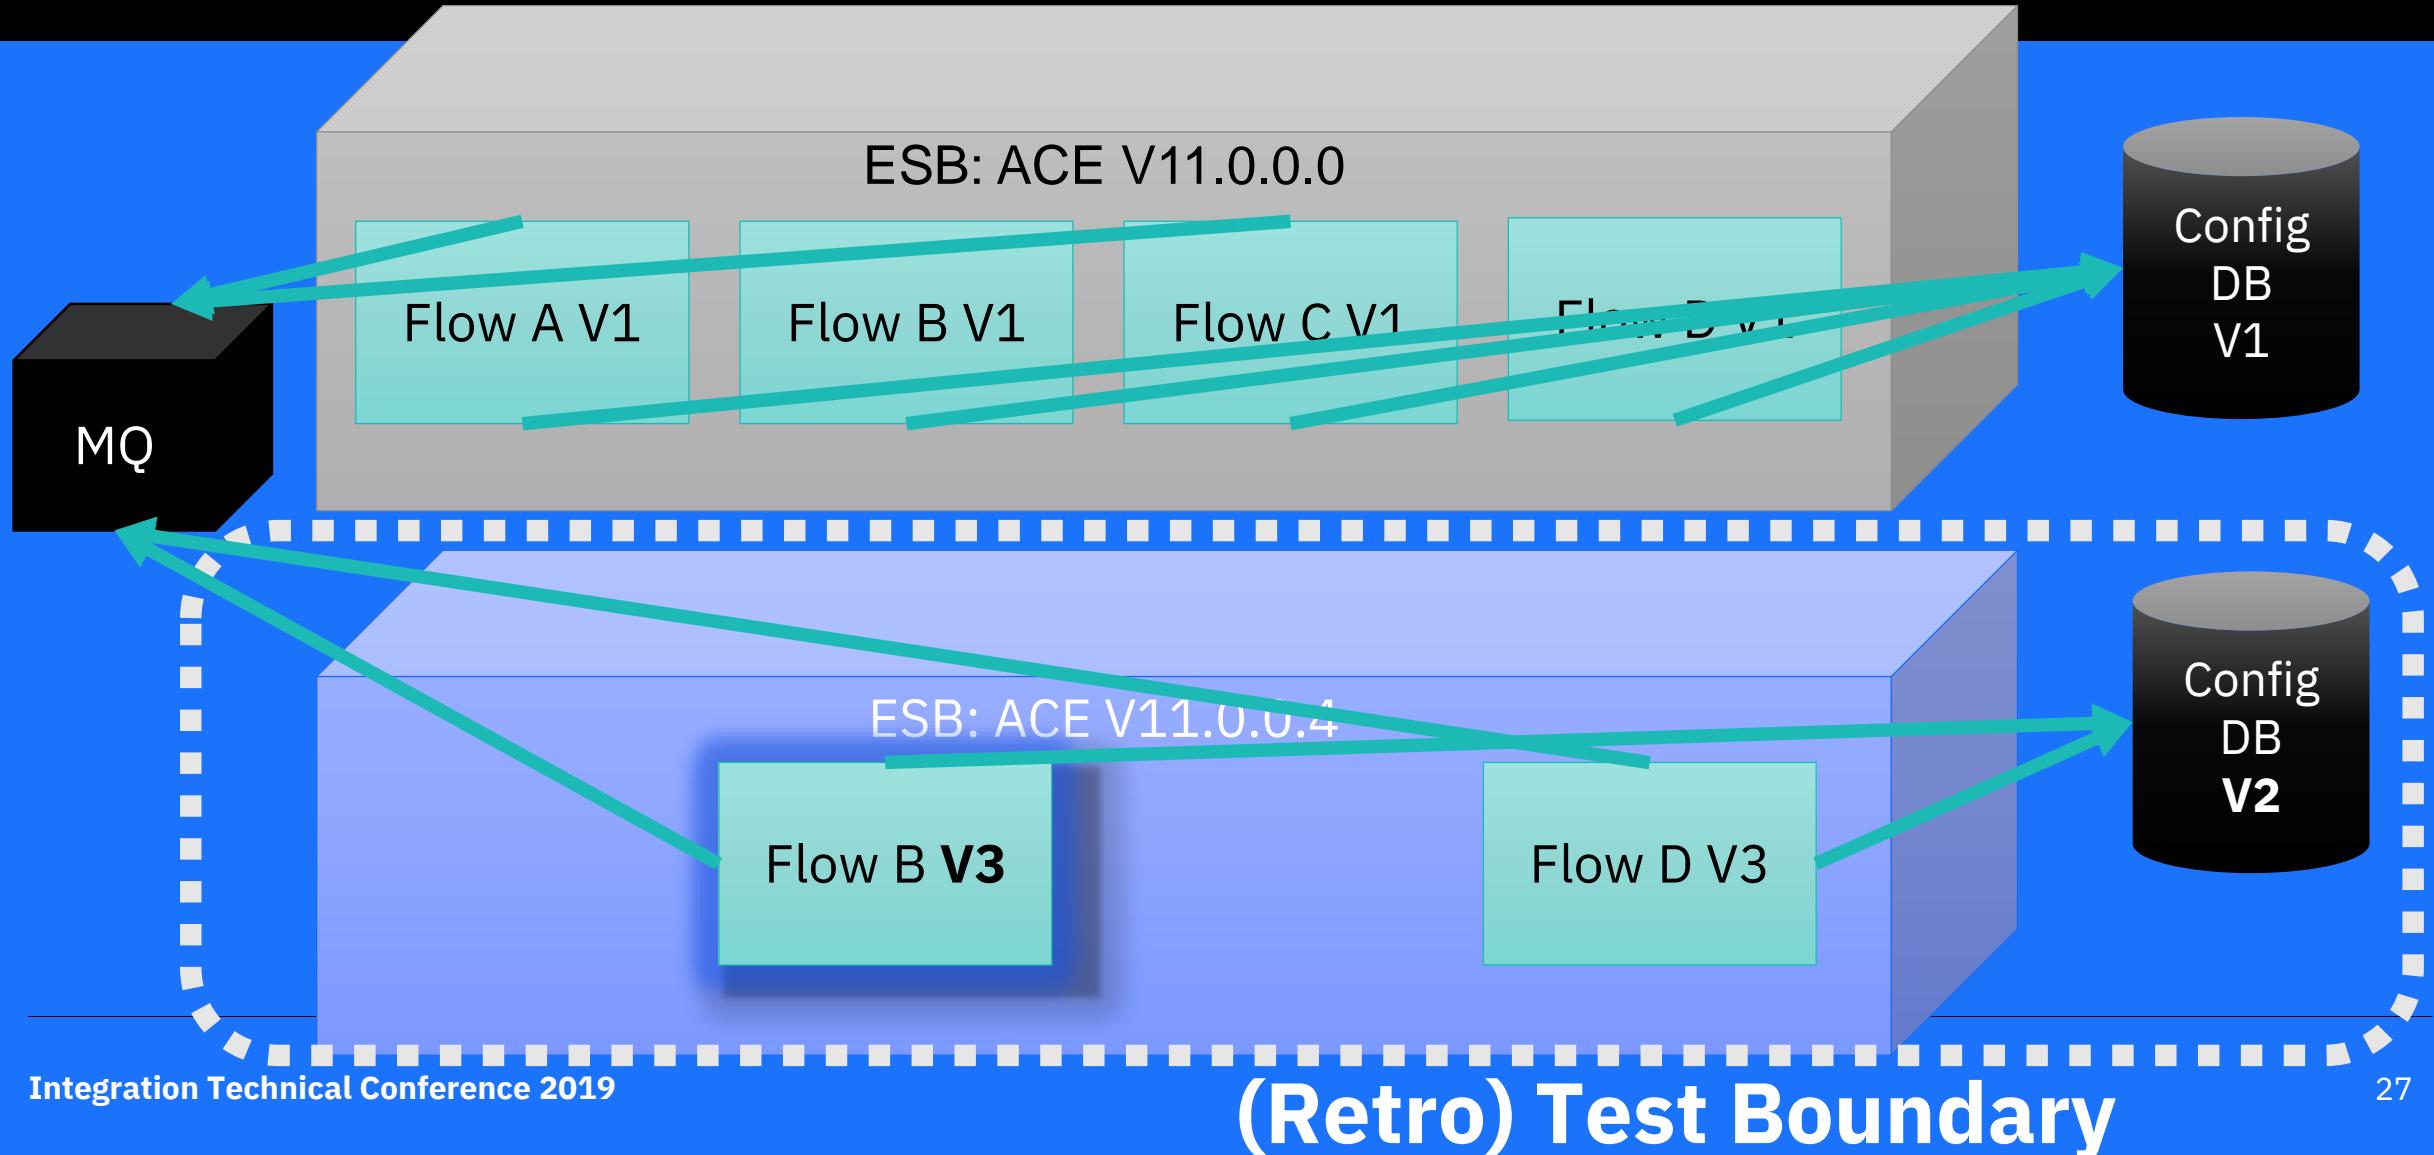
IBM



Importance of Independence

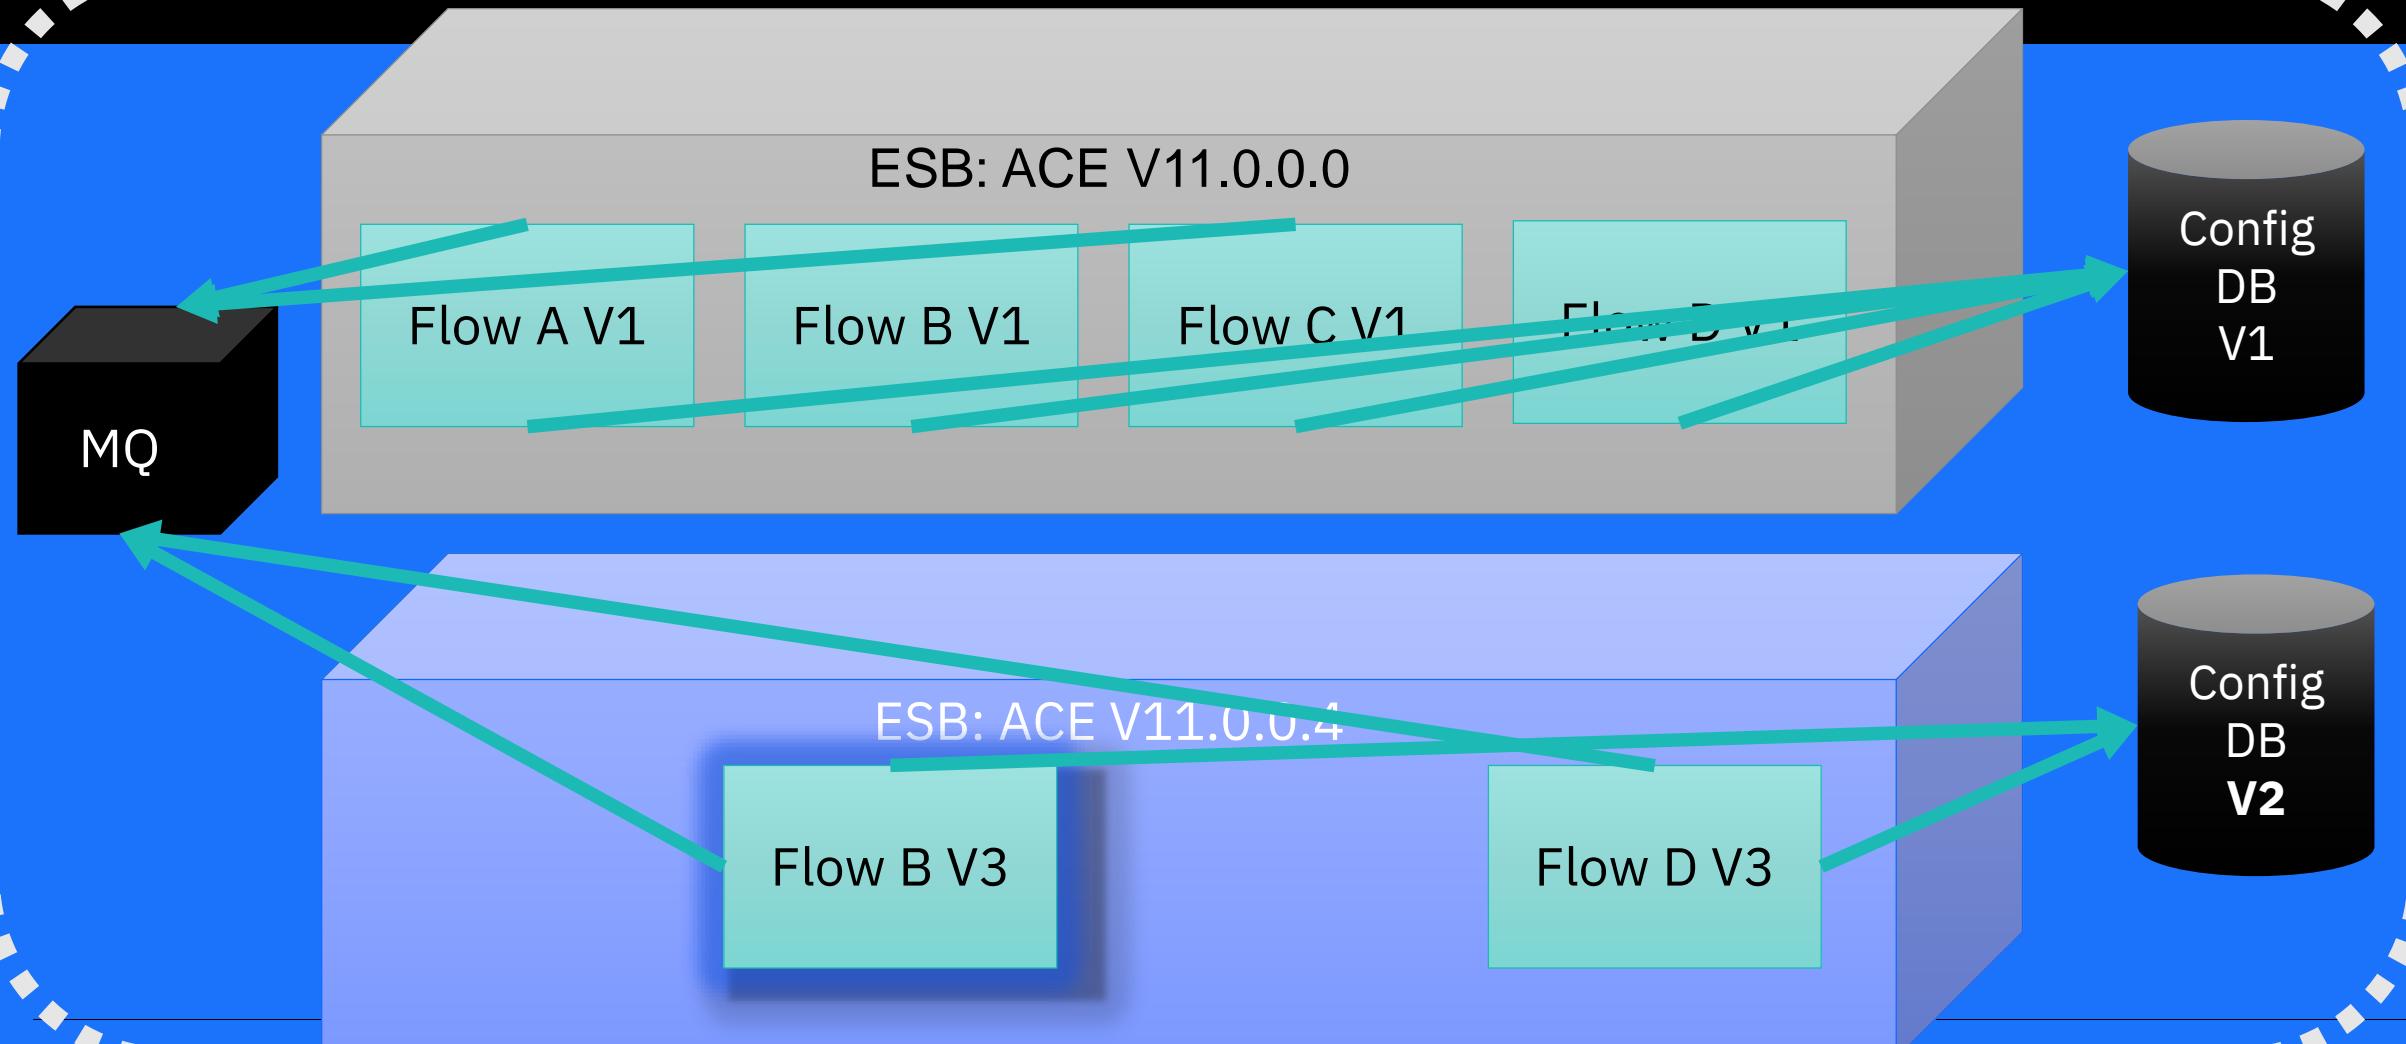


Importance of Independence



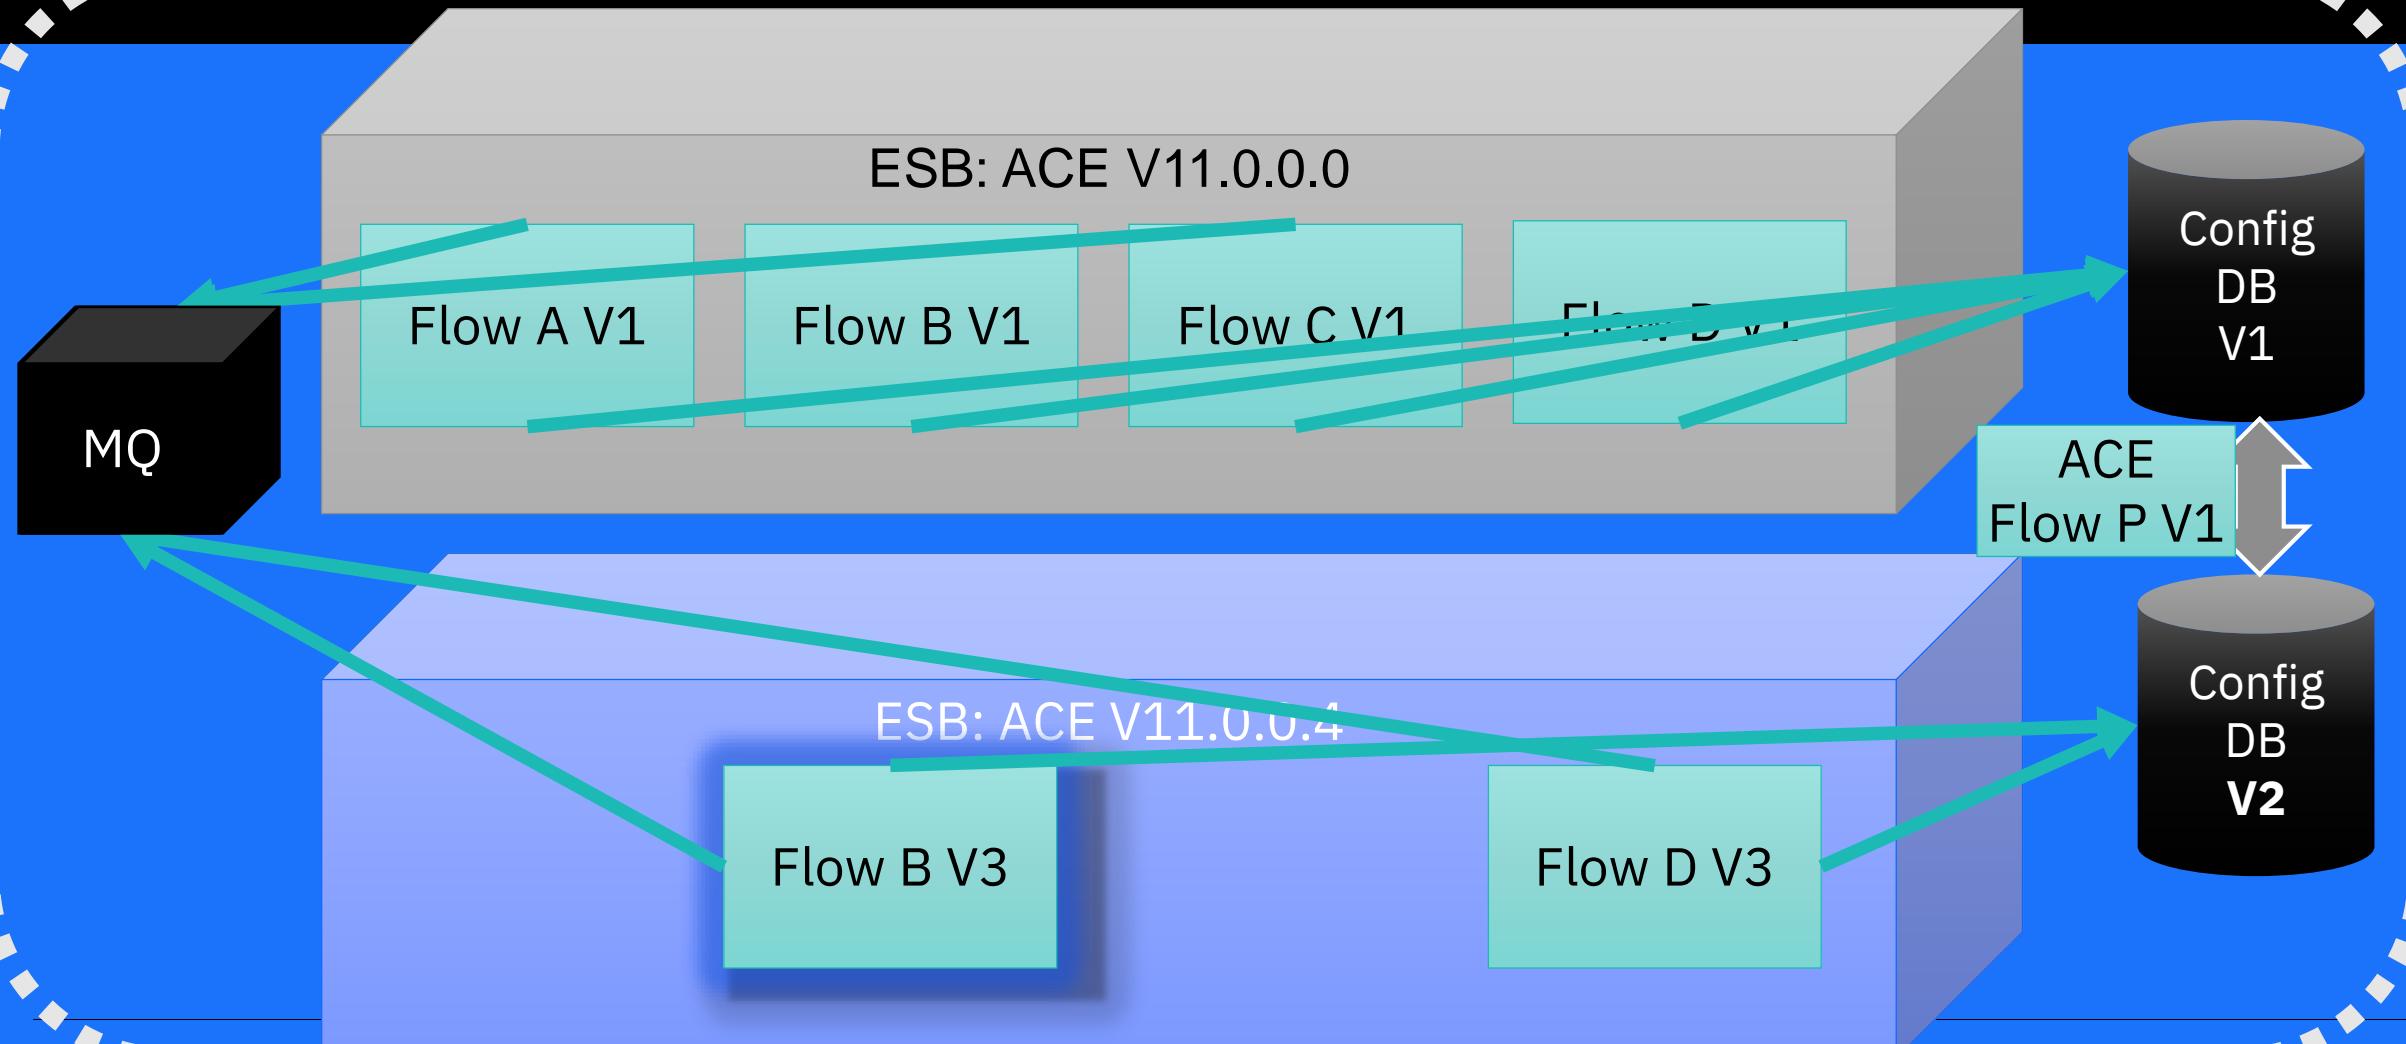
Importance of Independence

IBM



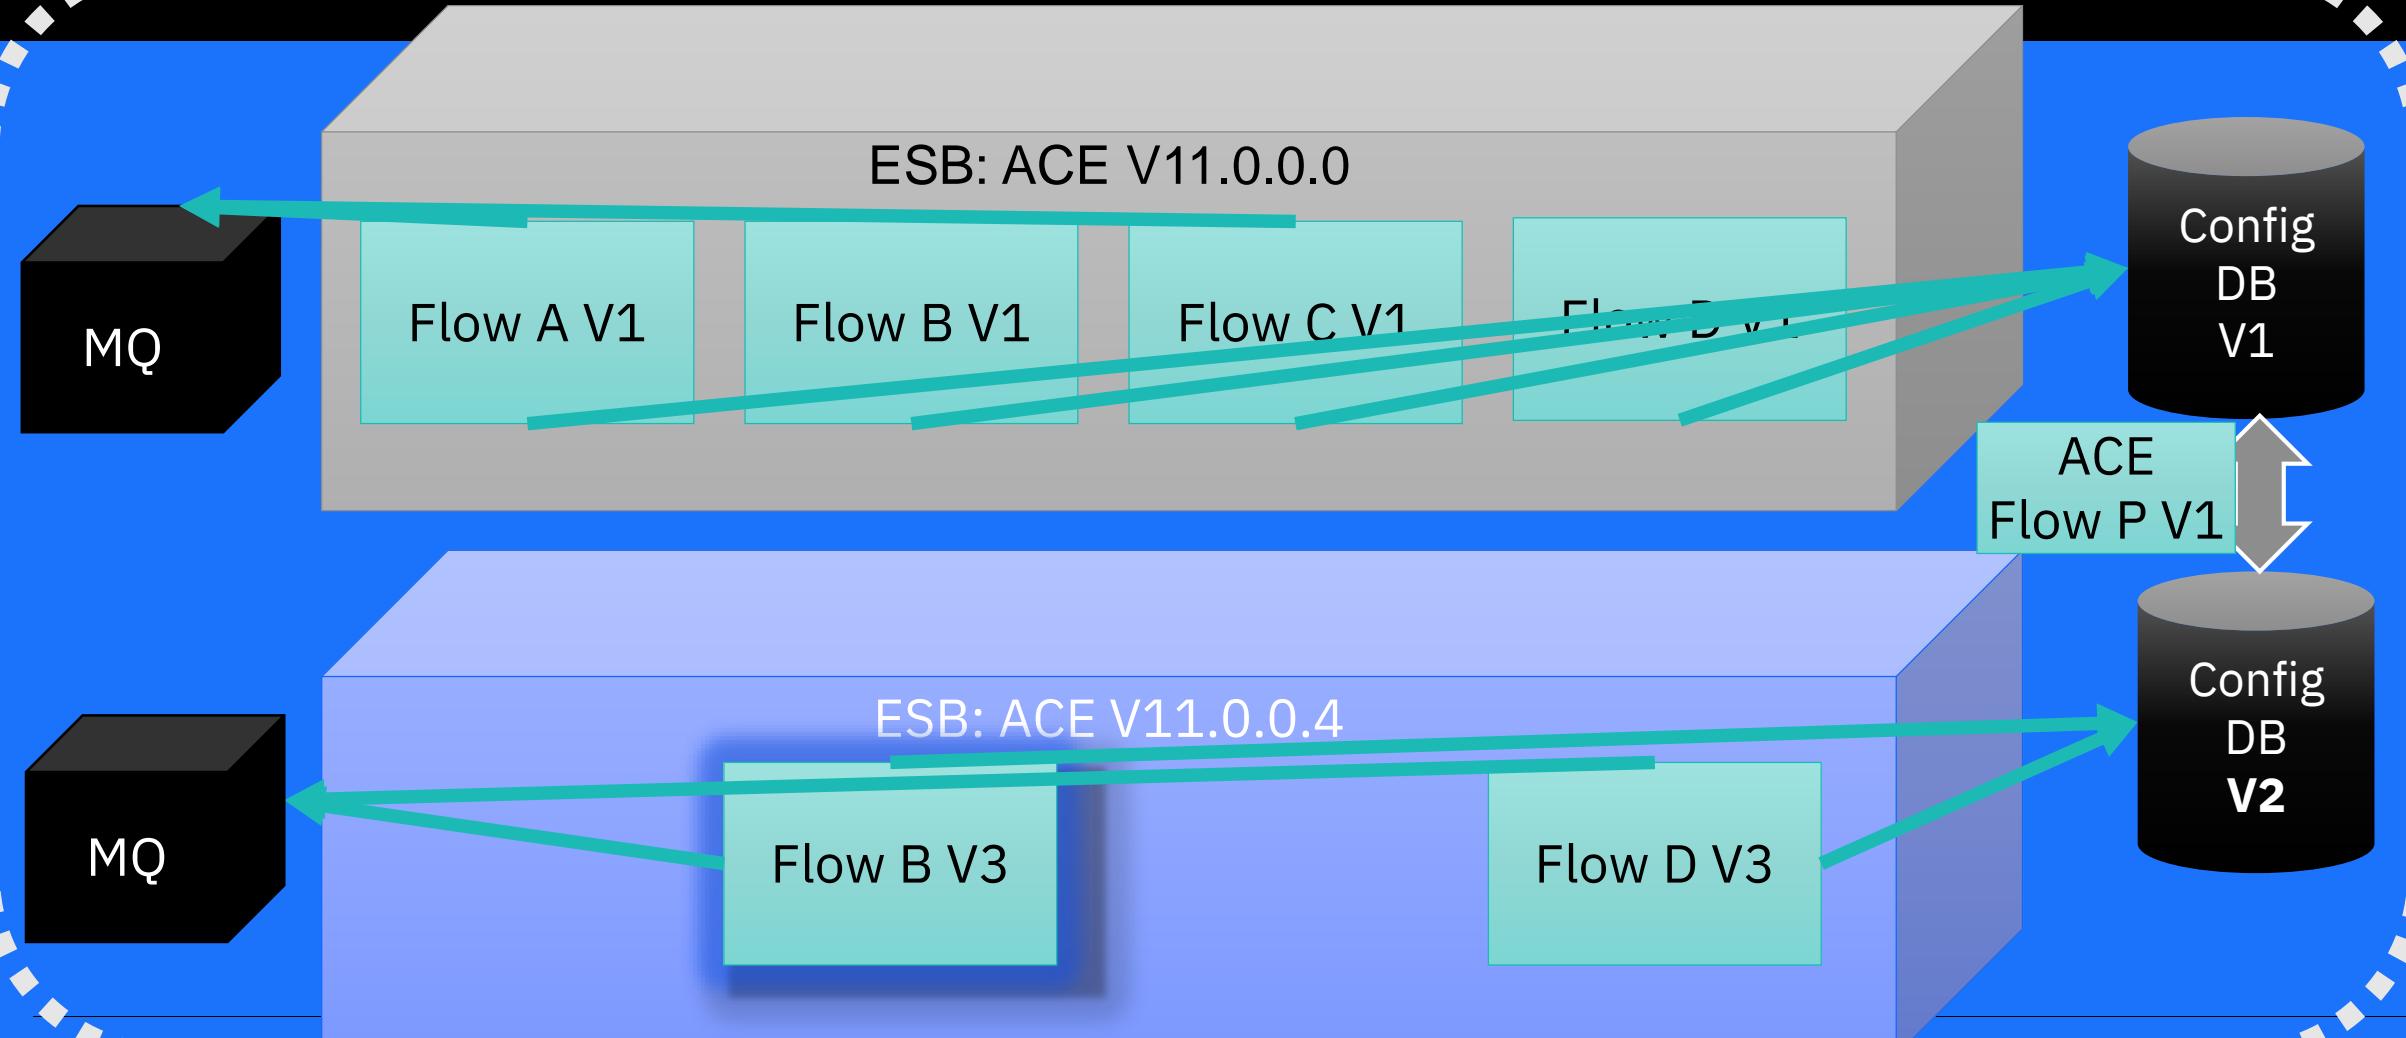
Importance of Independence

IBM

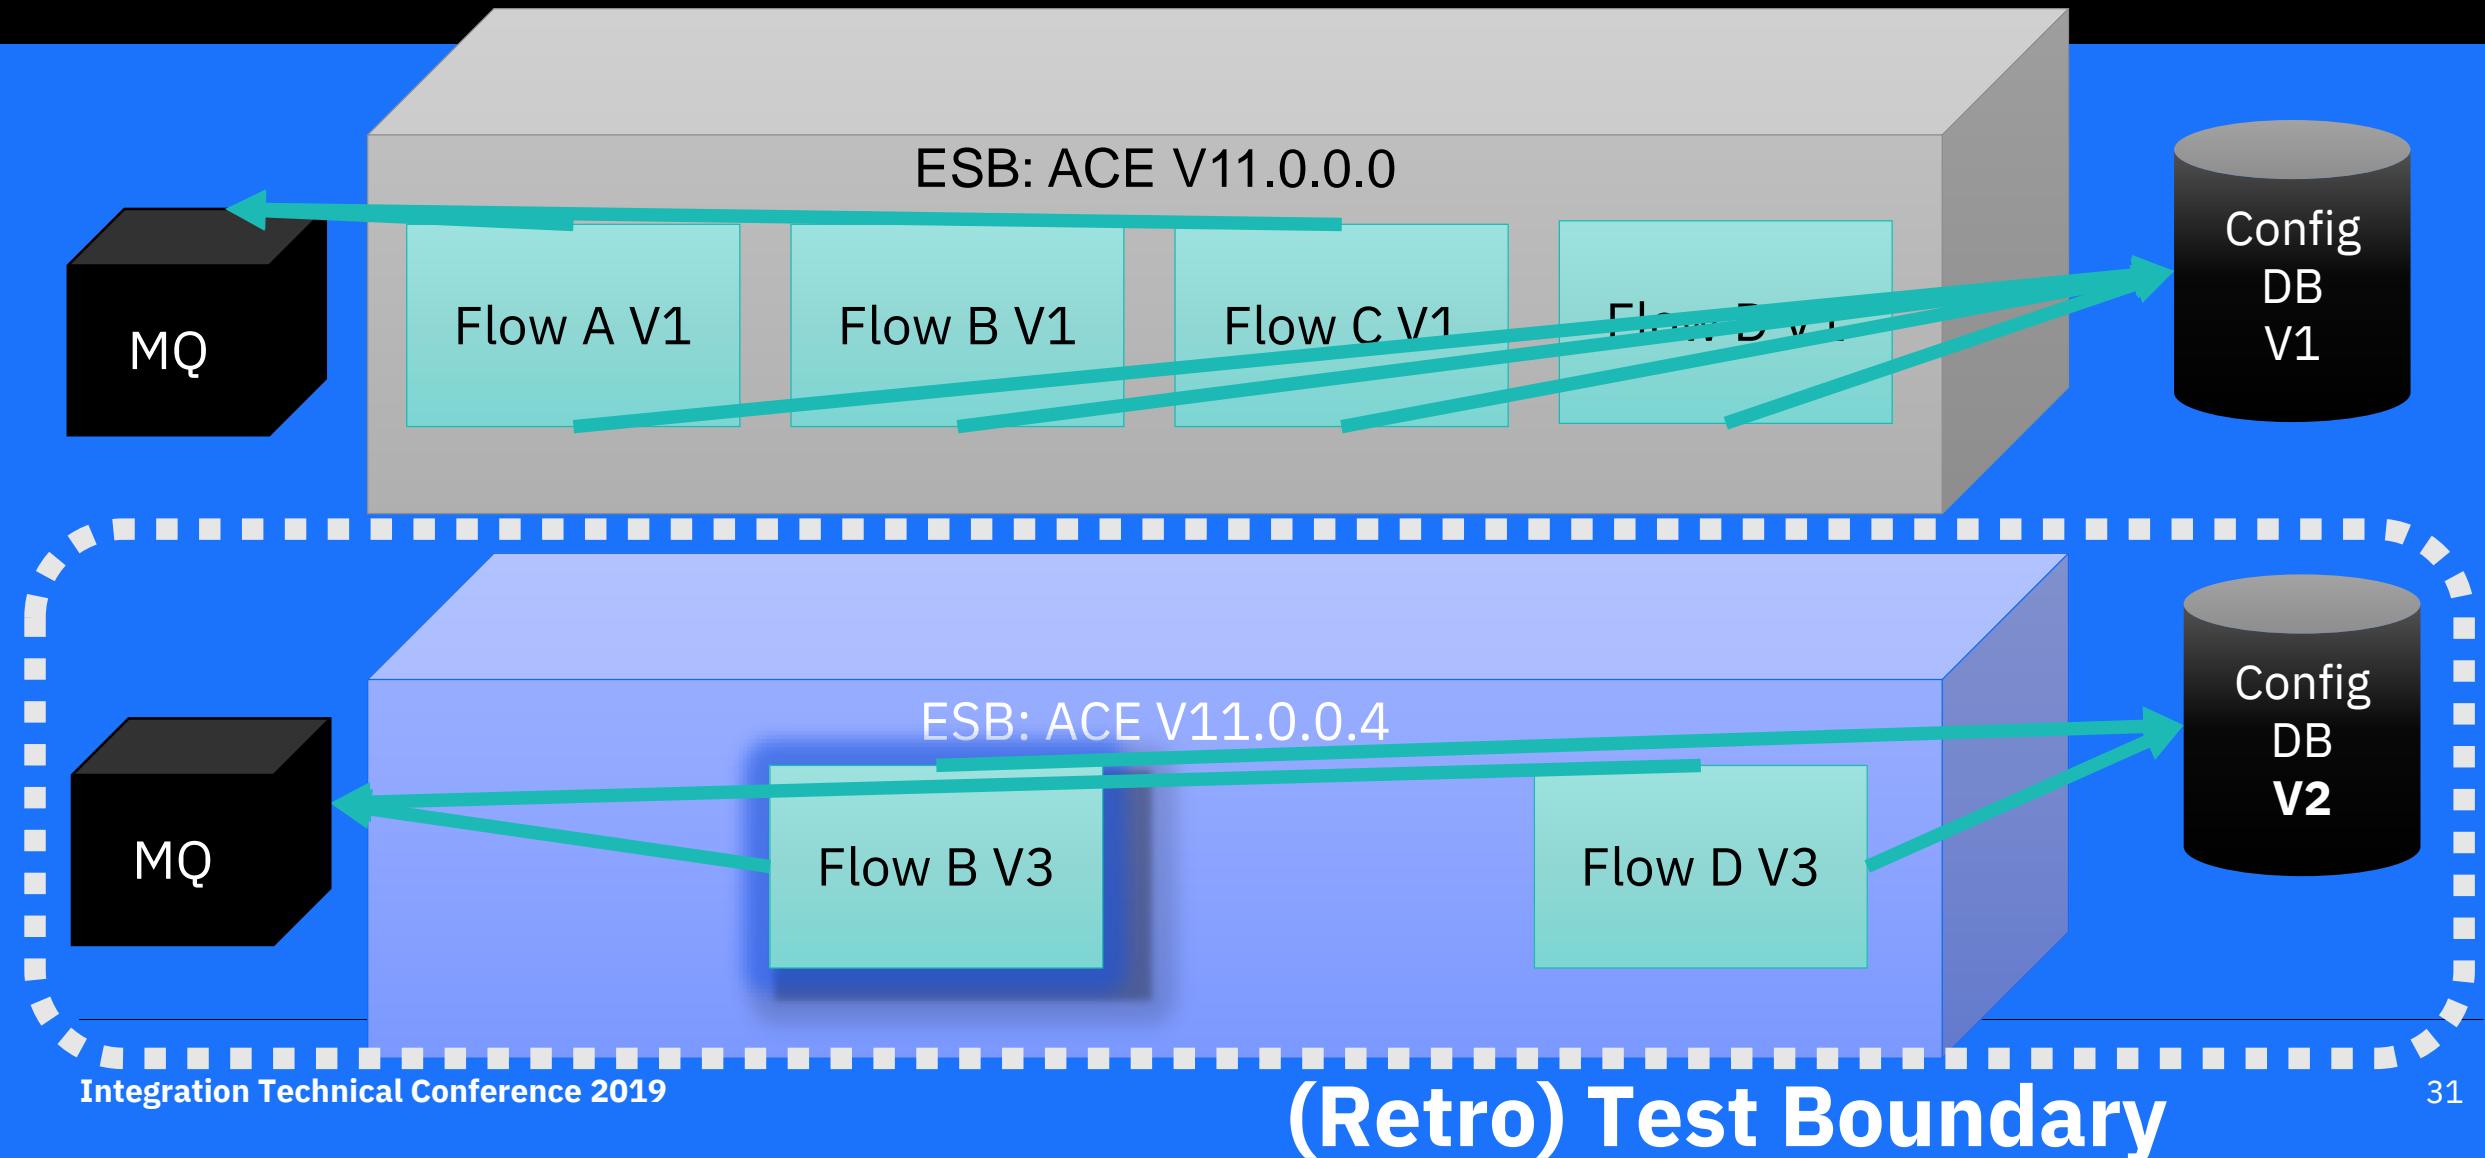


Importance of Independence

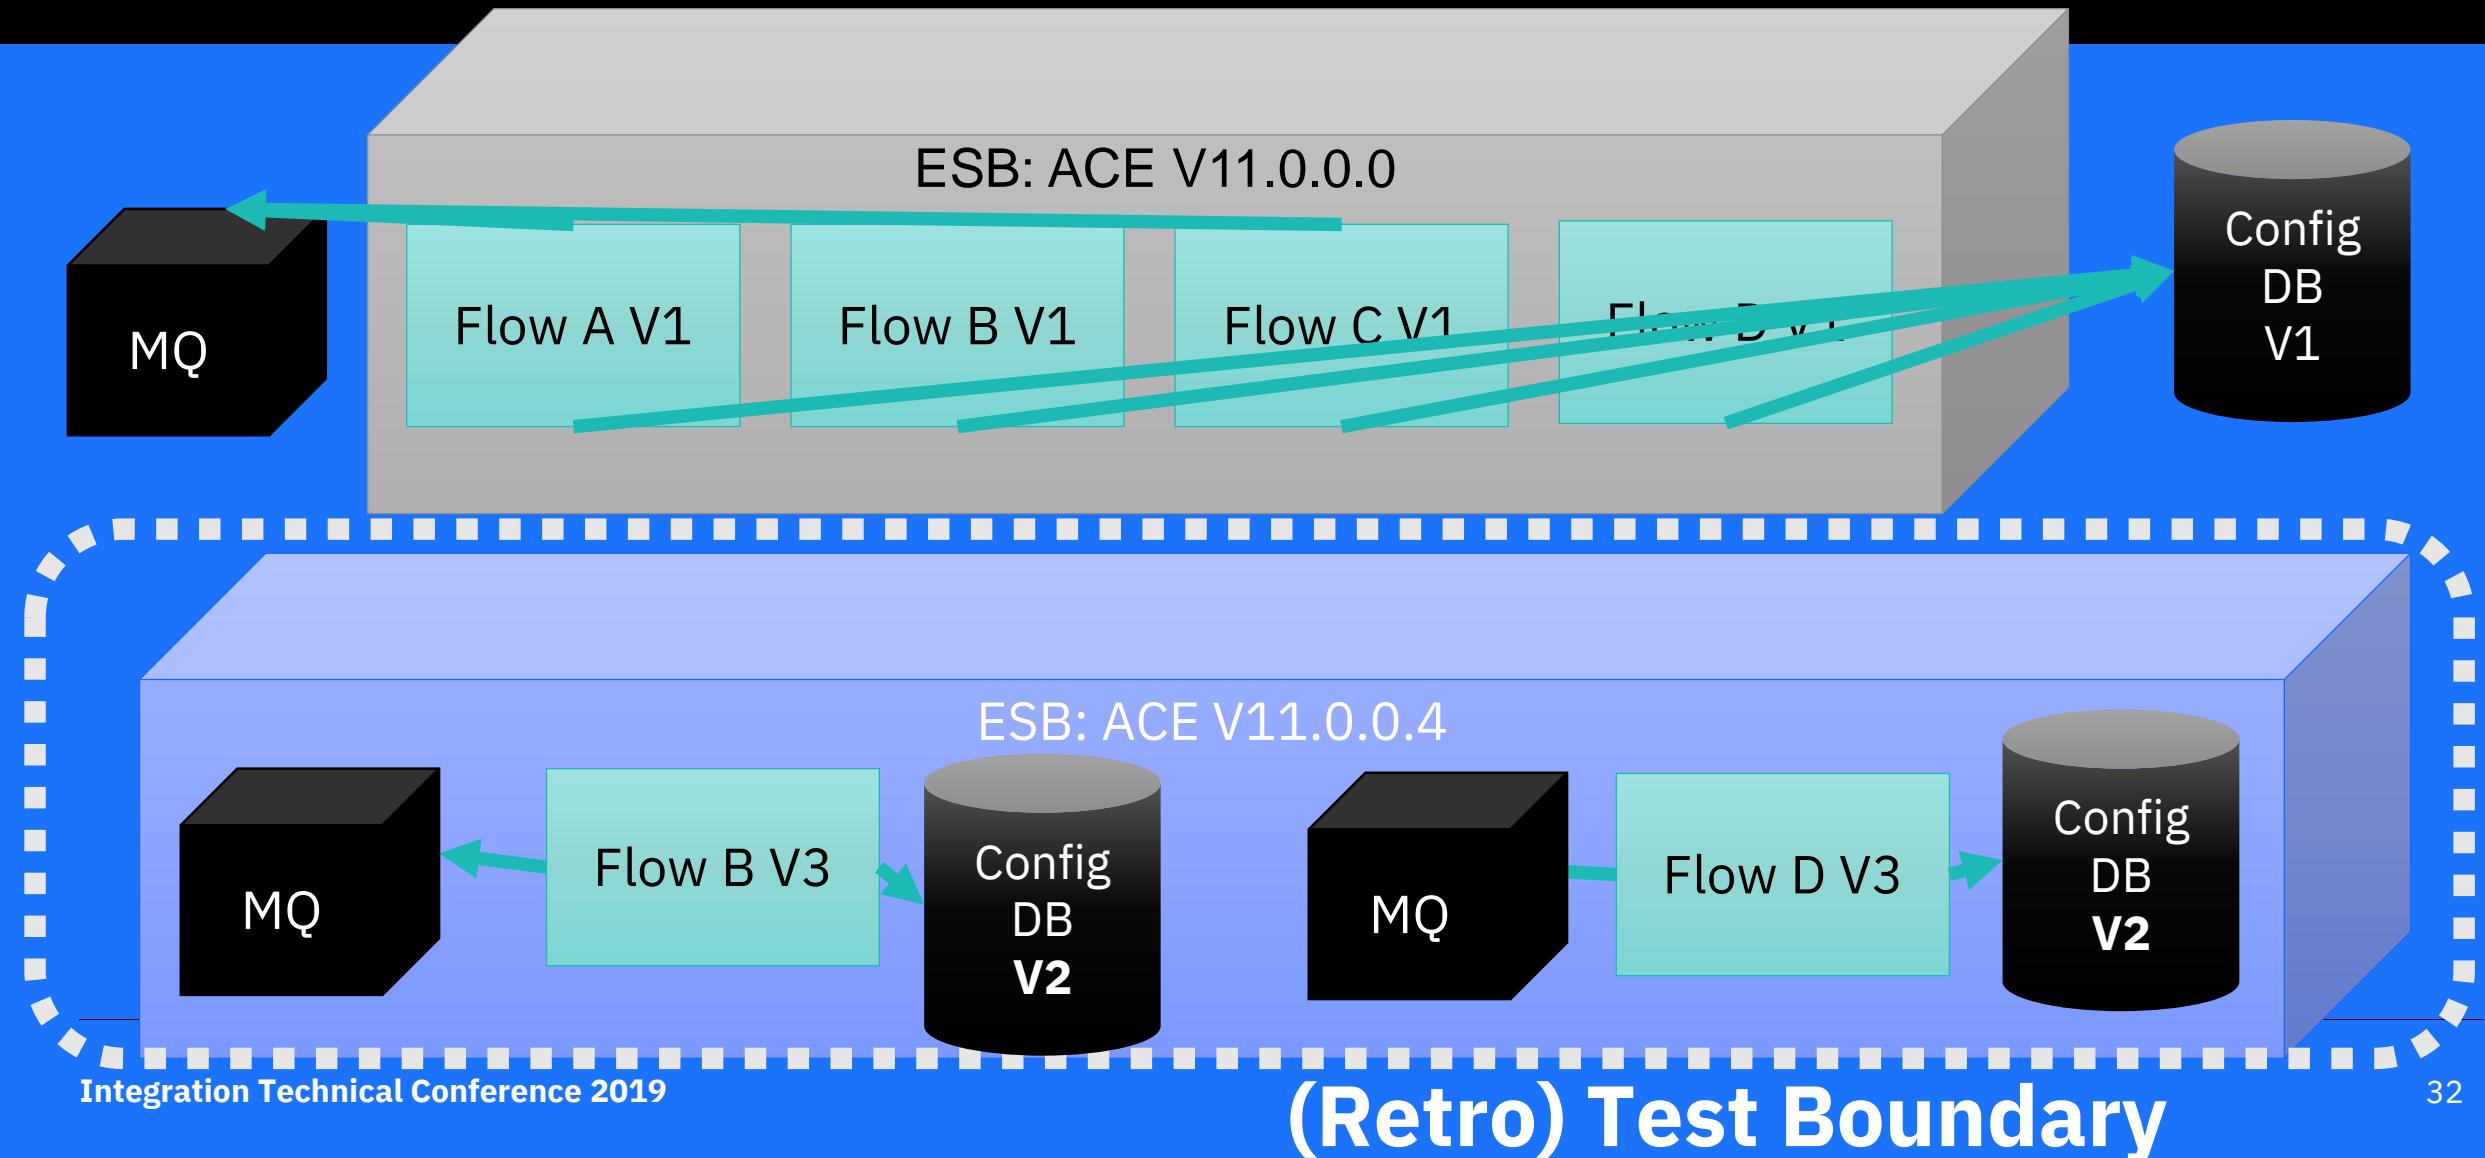
IBM



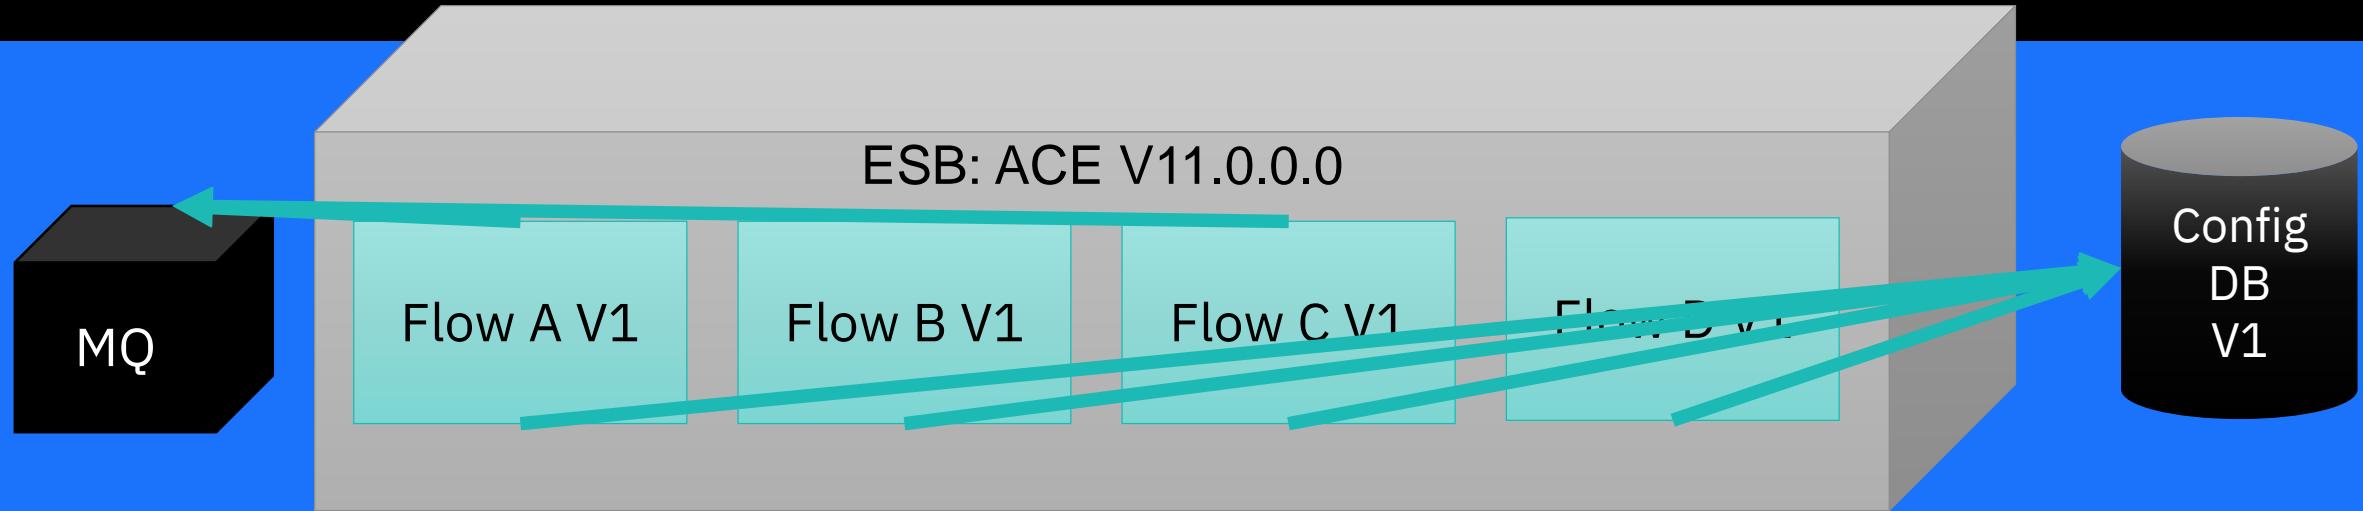
Importance of Independence



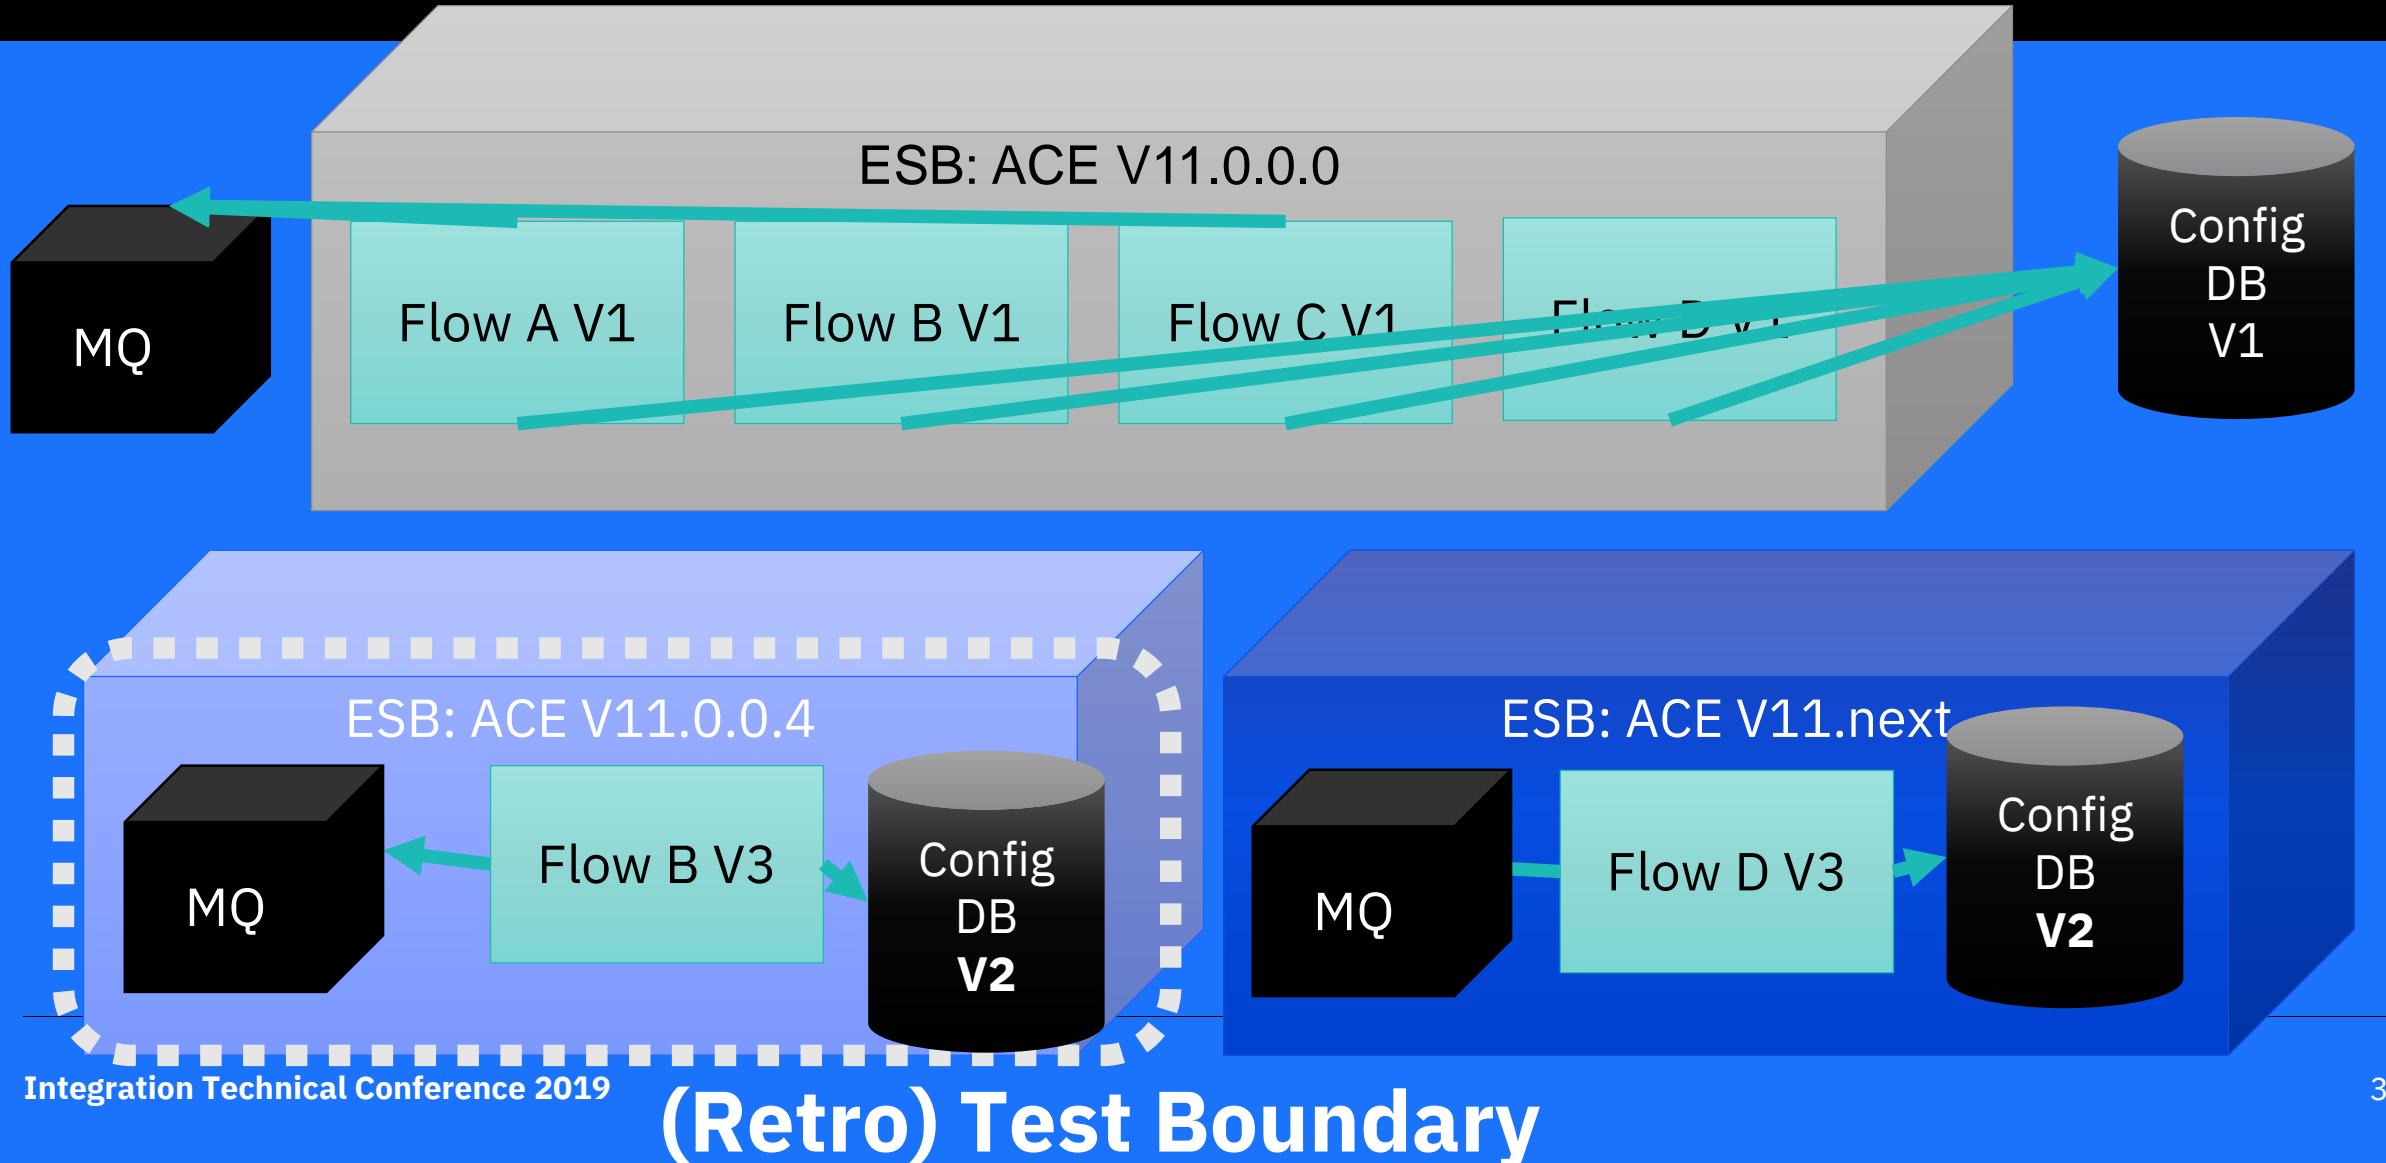
Importance of Independence



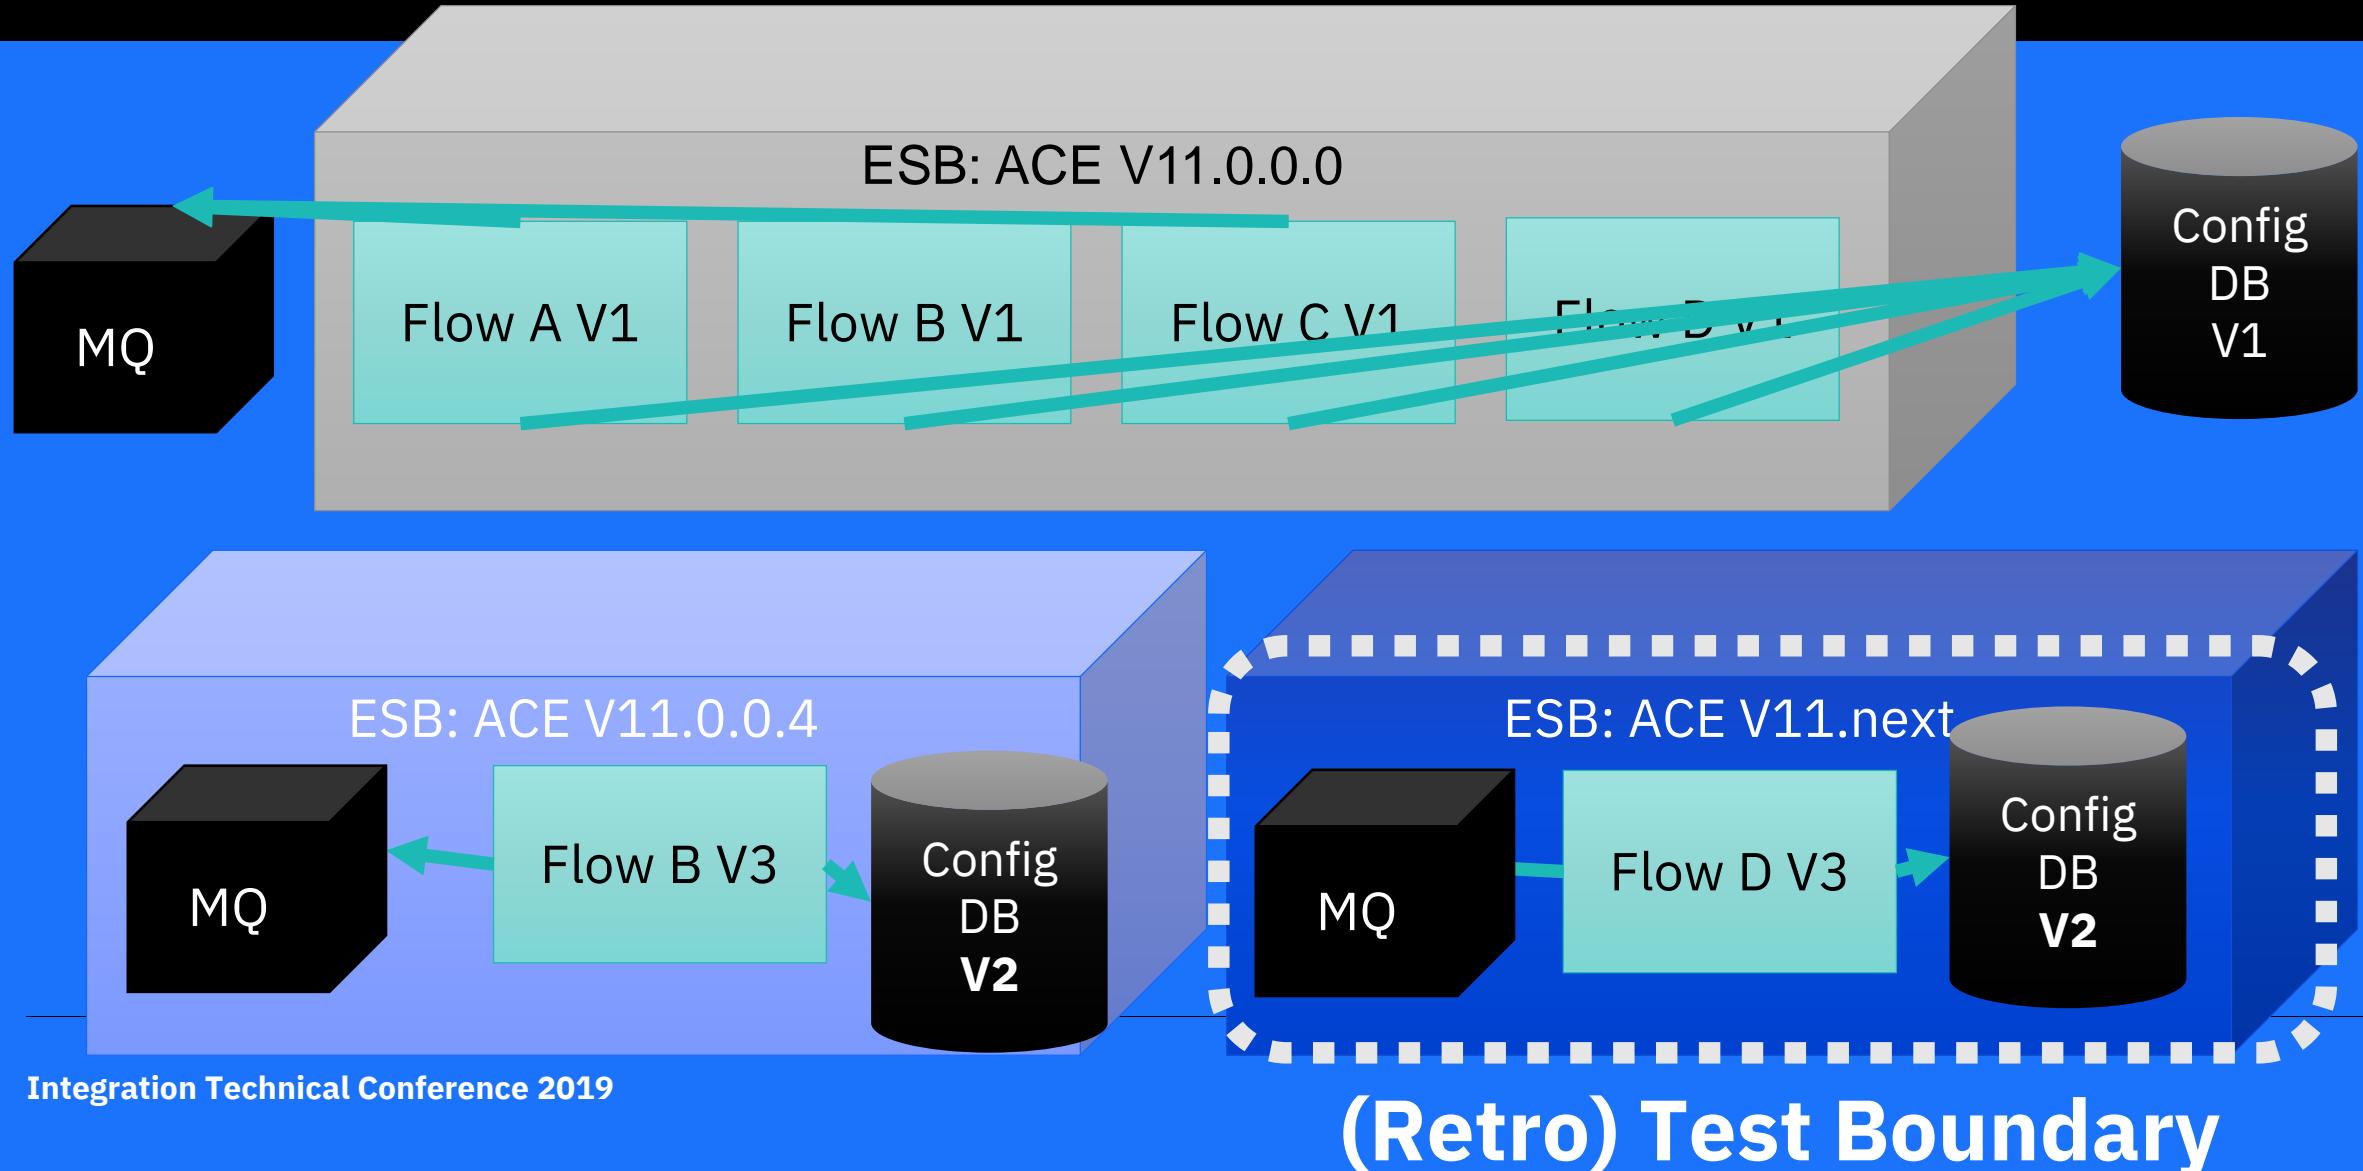
Importance of Independence



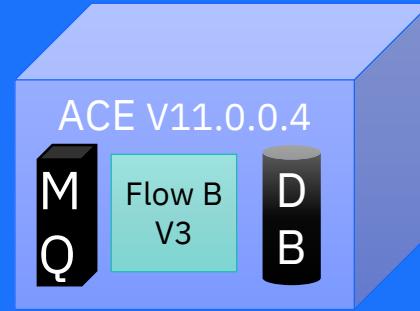
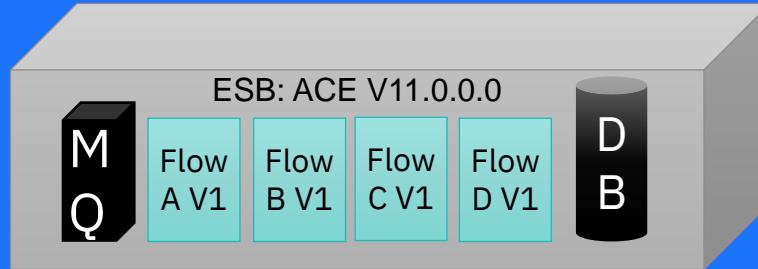
Importance of Independence



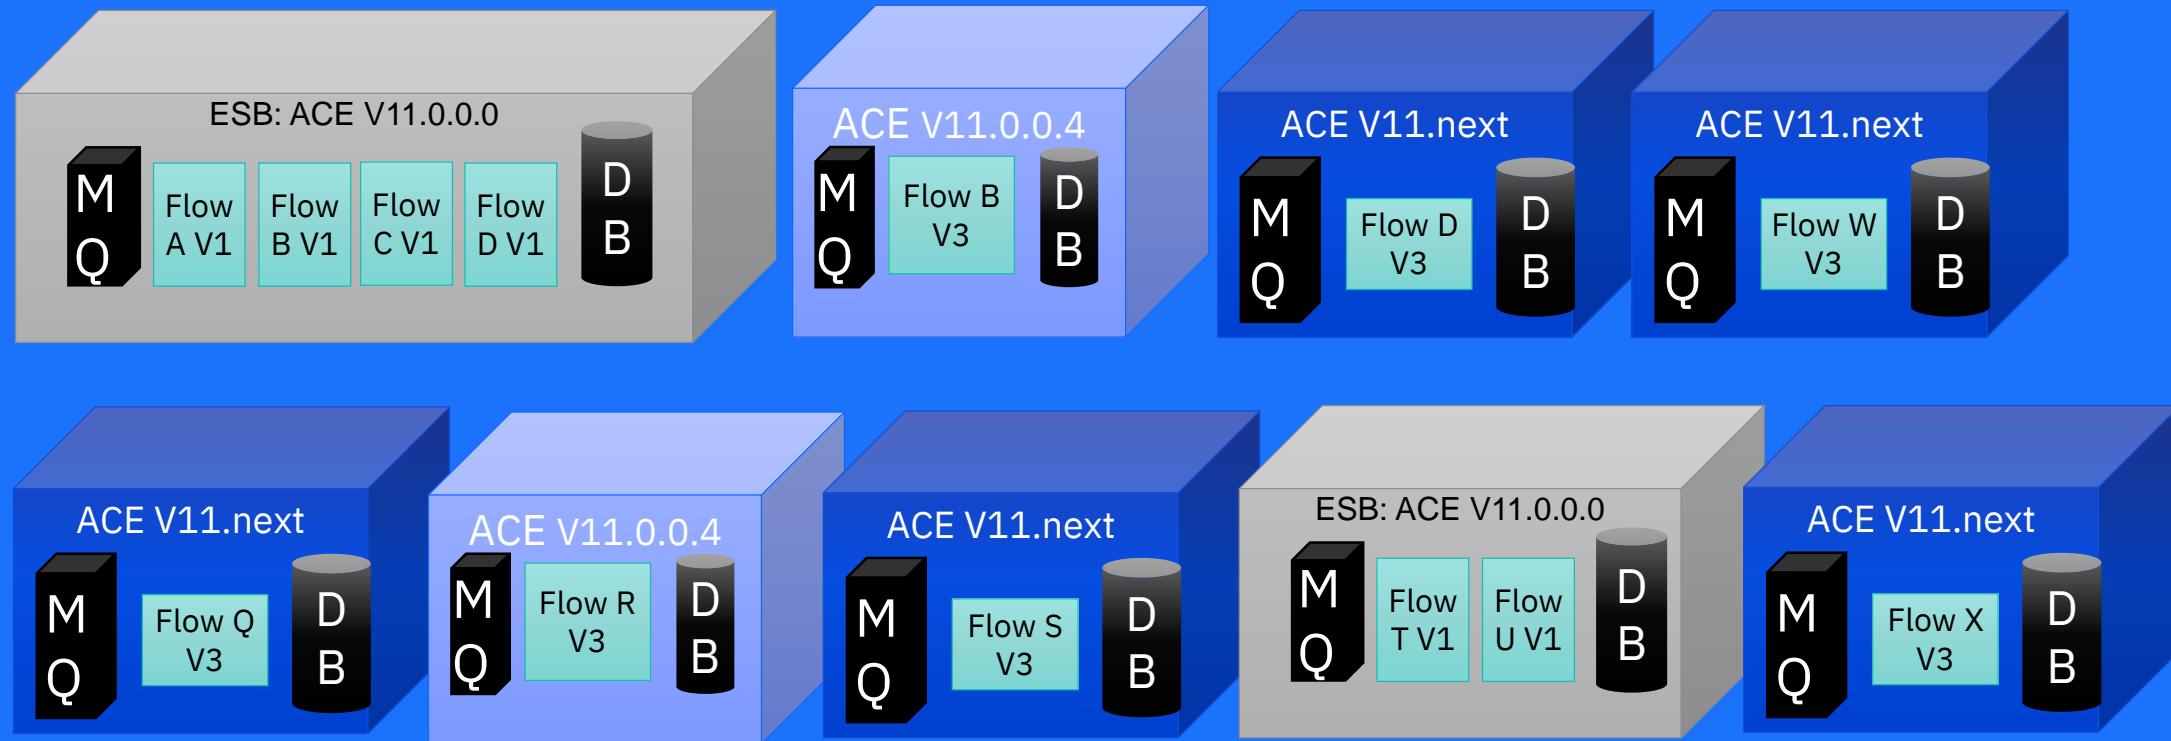
Importance of Independence



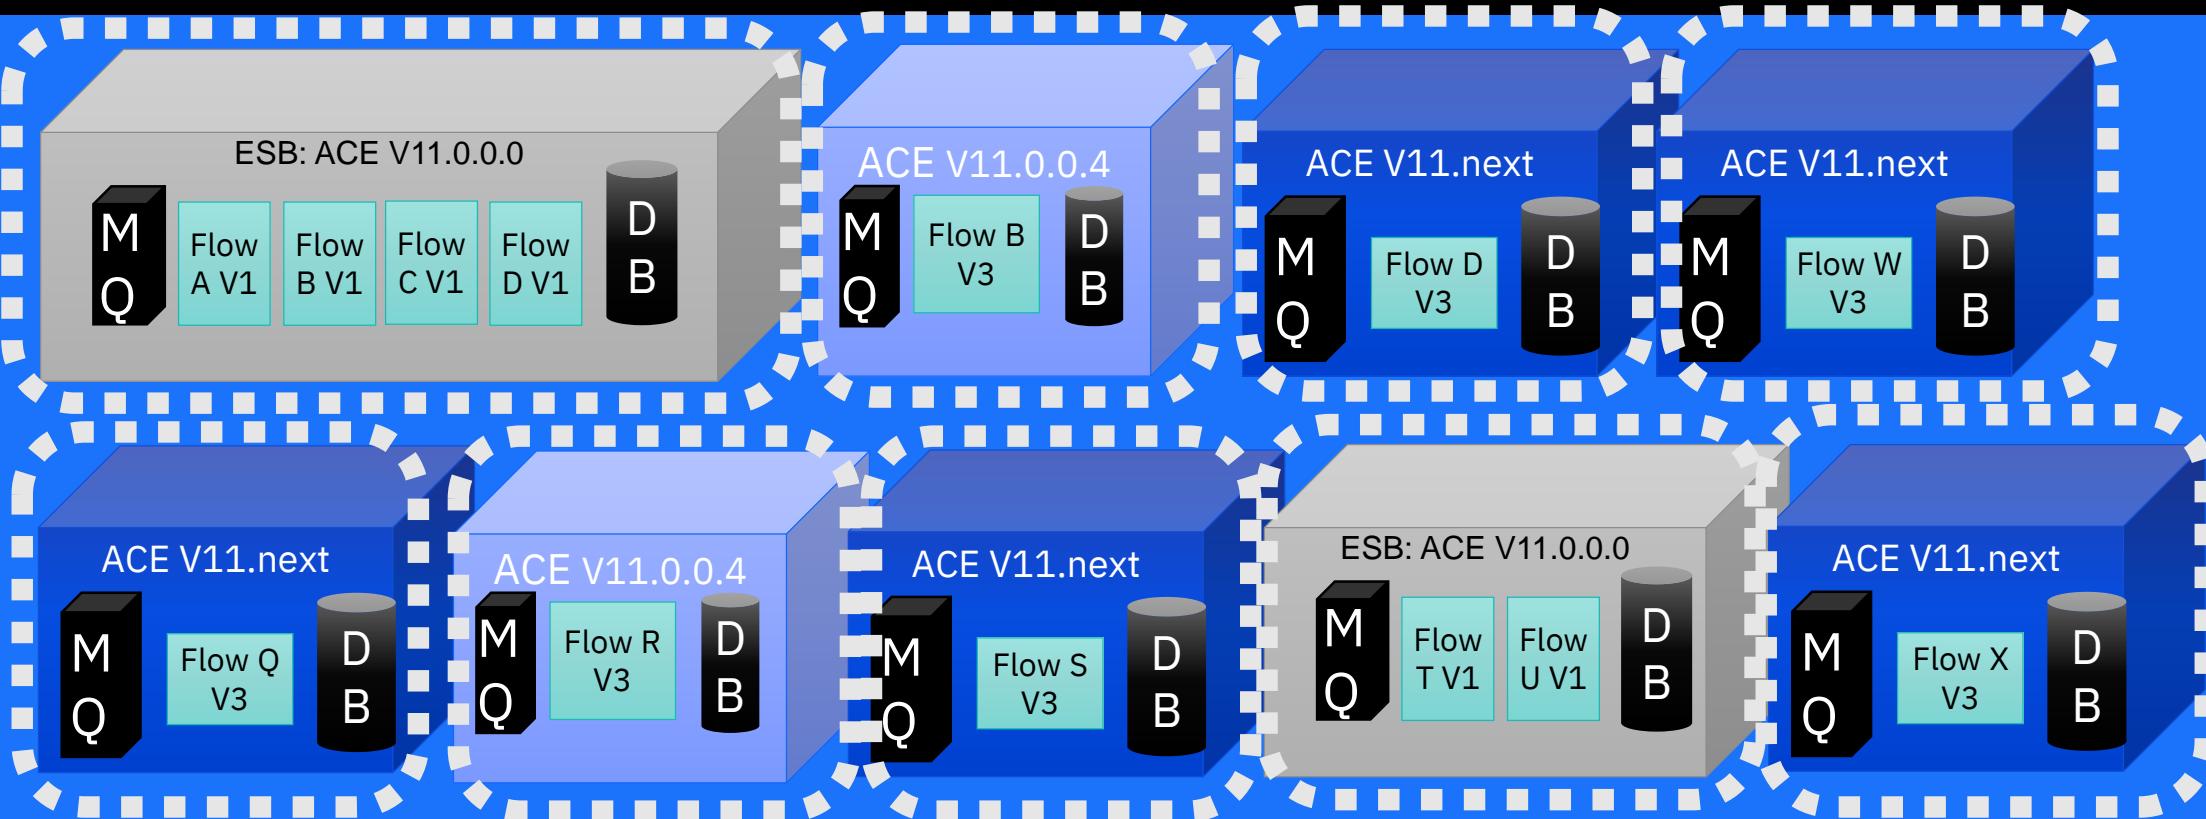
What about importance of size?



What about importance of size?



What about importance of size?



(Retro) Test Boundary

Granularity

What is your deployment unit?

One Operation?

One service / API e.g. Swagger / WSDL?

A group of Integrations co-located on a server?

What will change (Deploy)?

Where is your testing boundary?

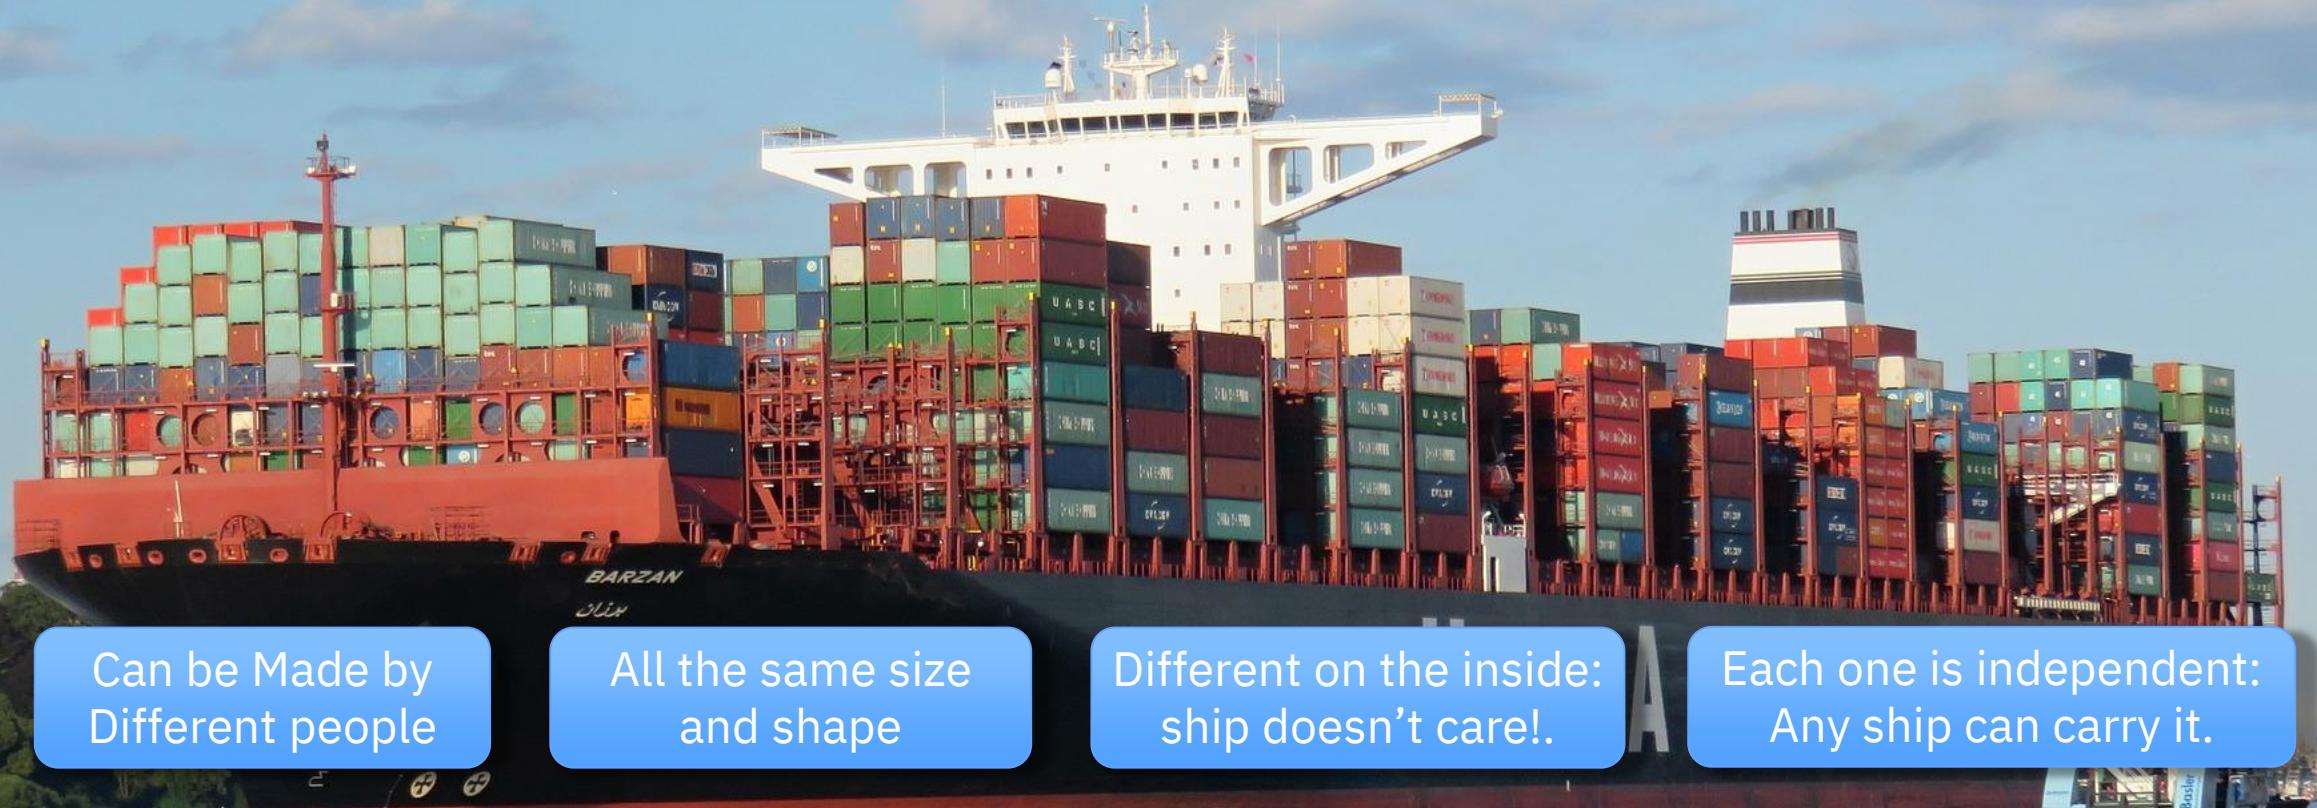
What takes resources (Scale)?



Containers: History and Philosophy



To summarise some key container points:



Can be Made by
Different people

All the same size
and shape

Different on the inside:
ship doesn't care!.

Each one is independent:
Any ship can carry it.

Small ships and big ones carry containers.
Big ones just carry more

Need to ship more/do more integrations?
Deploy more containers

Heavy containers take more fuel to ship – only put in what you need for each container.

Agile Integration Architecture

Modernizing integration
to enable
business agility

Fine grained deployment

Architecture & design

Decentralized Ownership

People & Process

Cloud native infrastructure

Infrastructure & Technology



Improve build independence and production velocity
(deployment agility)



Accelerate agility and innovation
(development agility)



Dynamic scalability and inherent resilience
(operational agility)

Anatomy of a container:

Container Image

Anything else

Helpers and Utilities

Static Configuration

Integrations (BARs)

App Connect

Operating System



Anatomy of a container:

Container Image

Anything else

Helpers and Utilities

Static Configuration

Integrations (BARs)

MQ client or QM?

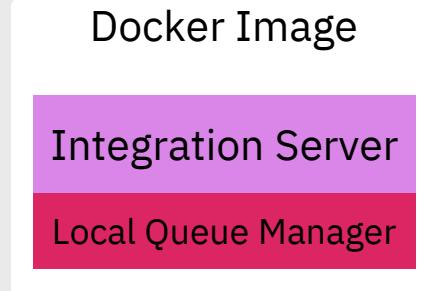
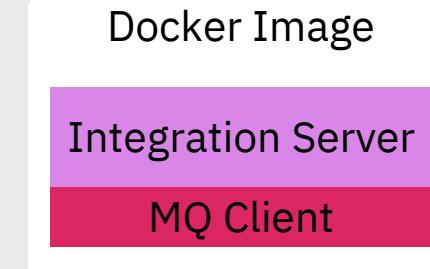
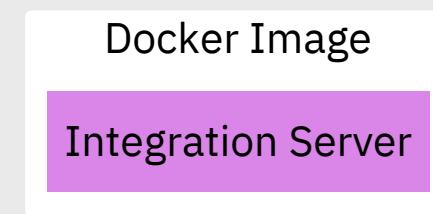
App Connect

Operating System



Image content tailored to each integration.

Different on the inside:
ship doesn't care!.

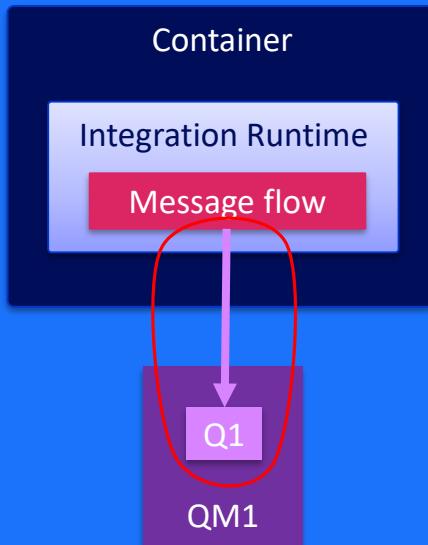


MQ connectivity	Yes (HTTP API)	Yes (client binding)	Yes (server binding)
1PC for MQ	No	Yes	Yes
EDA nodes	No	No	Yes
2PC	No	No	Yes
Horizontally scalable	Yes	Yes	Yes (with loss of sequencing)
Persistent volume	Not required	Not required	Required (if durability desired)
Start up	Fast	Fast	Slower
Disk space	Smallest	Medium	Largest

Why? Example: Single vs two phase commit

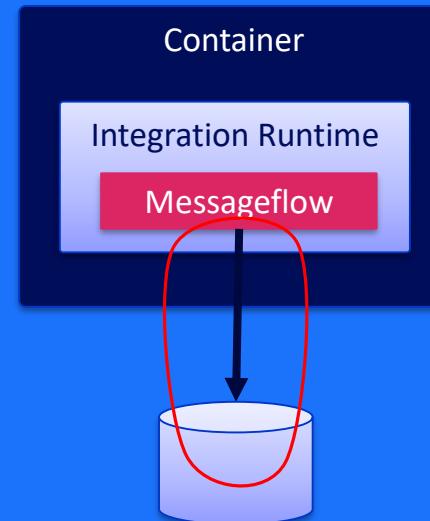


a) Single transaction with an MQ queue



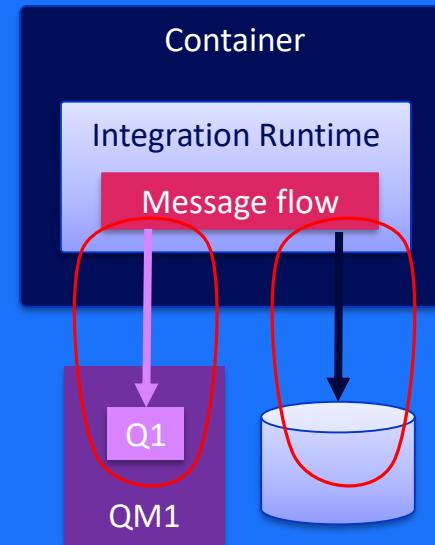
- One-phase commit sufficient
- No local queue manager required

b) Single transaction with a database



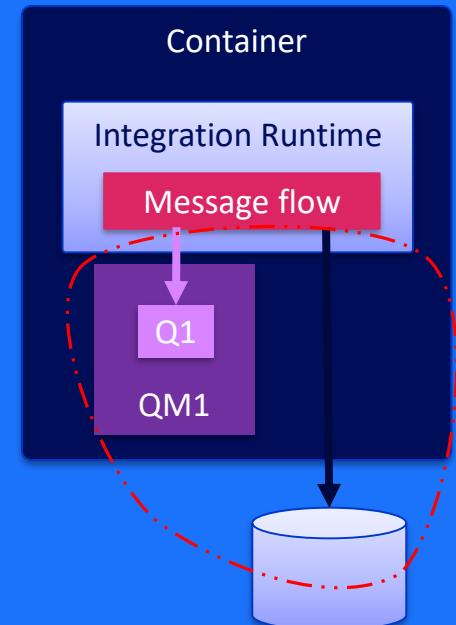
- One-phase commit sufficient
- No local queue manager required

c) Separate transactions to a database and a queue



- One-phase commit sufficient
- No local queue manager required

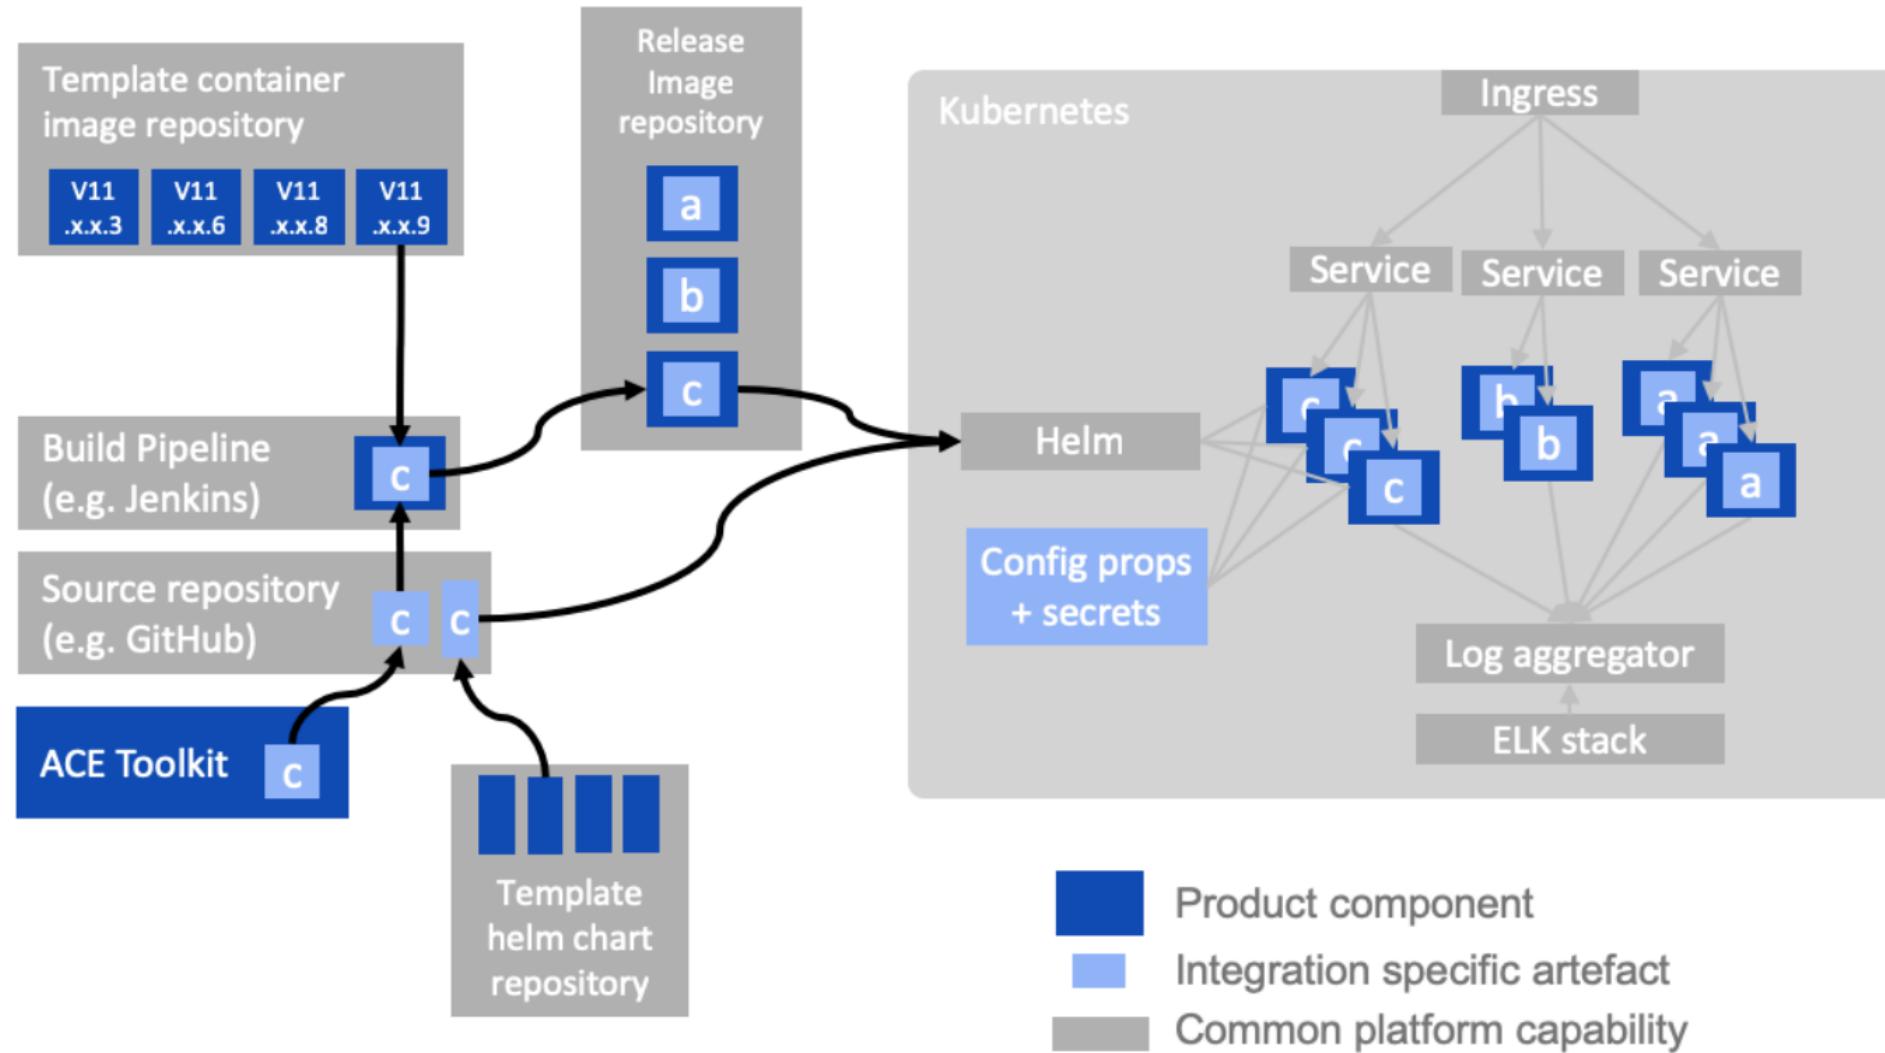
d) Combined transaction to a database and a queue



- Two-phase commit required
- Local queue manager required to co-ordinate.

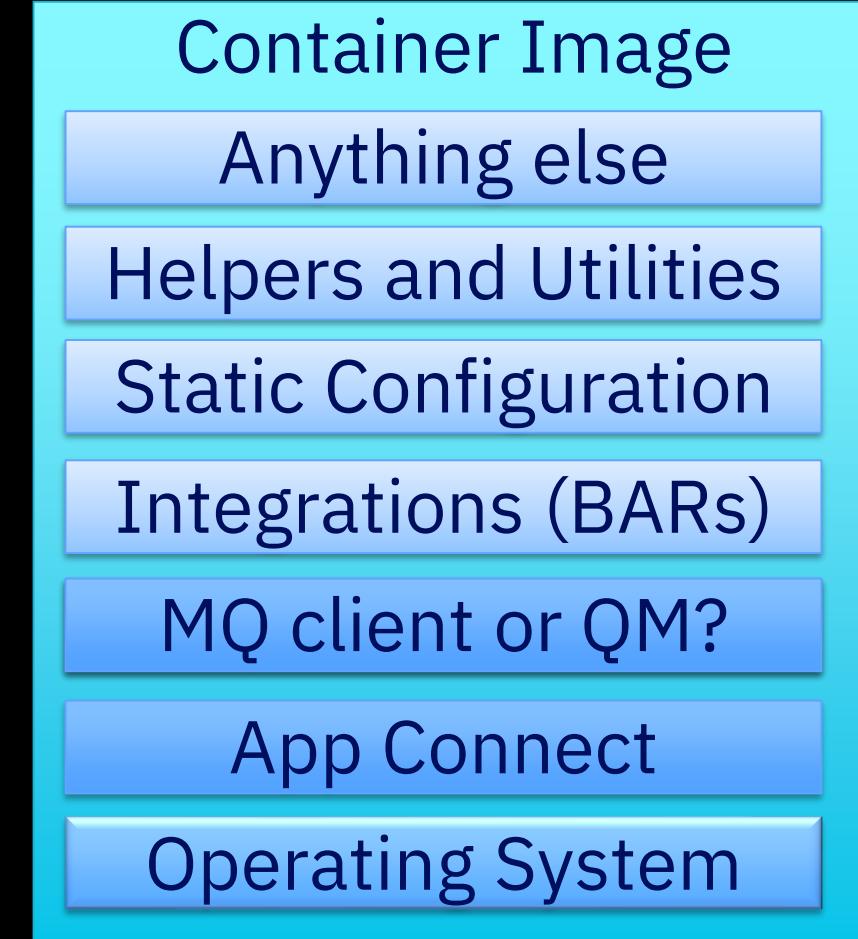
Image + Config + Orchestration = Runtime

IBM



Step 1: Create your image(s)

- Docker images are ‘layered’ – each component is built on top of the previous one
- Use pre-built layers e.g. o/s, App Connect pre-built images or build your own.
- The dockerfile defines what is in your image: We have samples that we use to build our images.
- Create repositories of images at various stages of build for re-use
- Define ‘static’ config: Non/slow mover config can go in the build.



IBM pre-built examples: Search for `ibmcom/ace` on hub.docker.com



Screenshot of the Docker Hub search results for `ibmcom/ace`.

The search bar shows `ibmcom/ace`. The results are filtered by `Containers`. The sorting is set to `Most Popular`, showing 1 - 25 of 2,819 results.

Filters:

- Docker Certified (i)
 - Docker Certified
- Images
 - Verified Publisher (i)
Docker Certified And Verified Publisher Content
 - Official Images (i)
Official Images Published By Docker

Results:

Image Details	Downloads	Stars
 ibmcom/ace By ibmcom • Updated 2 months ago Official IBM App Connect Enterprise for Developers image <small>Container Linux x86-64</small>	100K+	14

Instructions, dockerfile and code/scripts that we used for ours – so you can use/copy/change them to build your own



- 11.0.0.3 , latest (Dockerfile)

Building a container image

Download a copy of App Connect Enterprise (ie. ace-11.0.0.2.tar.gz) and place it in the deps folder. When building the image use build-arg to specify the name of the file: --build-arg ACE_INSTALL=ace-11.0.0.2.tar.gz

- **Important:** Only ACE version **11.0.0.2 or greater** is supported.

Choose if you want to have an image with just App Connect Enterprise or an image with both App Connect Enterprise and IBM MQ Advanced.

Using App Connect Enterprise for Developers

Get ACE for Developers edition. Then place it in the deps folder as mentioned above.

Using MQ production image

When building an image with both ACE and MQ, the docker file uses the [MQ Advanced for Developers image in docker registry](#) as base image by default on Ubuntu. On RedHat Enterprise Linux the image can be built using the [MQ instructions](#) or downloaded from [IBM Passport Advantage] (<https://www.ibm.com/software/passportadvantage/>).

When building a production image with MQ, follow the [MQ instructions](#) to build your own production MQ image. Then, when building the ACE with MQ image use build-arg to set the BASE_IMAGE to your production MQ image. More details below.

Build an image with App Connect Enterprise and MQ

Info on how to get the Developers or production image for MQ

ot4i / ace-docker

Branch: master → ace-docker / ubuntu / Dockerfile.aceonly

IBMRob RHEL & Hostname/Port Override Update

1 contributor

75 lines (56 sloc) | 2.58 KB

```
1 ARG BASE_IMAGE=ubuntu:16.04
2
3 FROM golang:1.10.3 as builder
4 WORKDIR /go/src/github.com/ot4i/ace-docker/
5 ARG IMAGE_REVISION="Not specified"
6 ARG IMAGE_SOURCE="Not specified"
7 COPY cmd/ ./cmd
8 COPY internal/ ./internal
9 COPY vendor/ ./vendor
10 RUN go build -ldflags "-X \"main.ImageCreated=$(date --iso-8601=seconds)\" "
11 RUN go build ./cmd/chkaceready/
12 RUN go build ./cmd/chkacehealthy/
13 # Run all unit tests
14 RUN go test -v ./cmd/runaceserver/
15 RUN go test -v ./internal/...
16 RUN go vet ./cmd/... ./internal/...
17
18 FROM ubuntu:16.04 as aceinstall
19 ARG ACE_INSTALL=ace-11.0.0.2.tar.gz
```

ot4i / ace-docker

No description, website, or topics provided.

26 commits | 2 branches | 4 releases | 4 contributors

Branch: master → New pull request | Create new file

IBMRob Fix file location

cmd RHEL & Hostname/Port Override Update

deps 11.0.0.2 update

internal 11.0.0.2 update

rhel MQ Client & fixes

sample RHEL & Hostname/Port Override Update

ubuntu Fix file location

vendor 11.0.0.2 update

README.md MQ Client & fixes

ace_compile_bars.sh 11.0.0.3 update

ace_config_bars.sh 11.0.0.3 update

ace_config_keystore.sh 11.0.0.3 update

ace_config_logging.sh 11.0.0.2 update

ace_config_odbcini.sh 11.0.0.3 update

ace_config_policy.sh 11.0.0.3 update

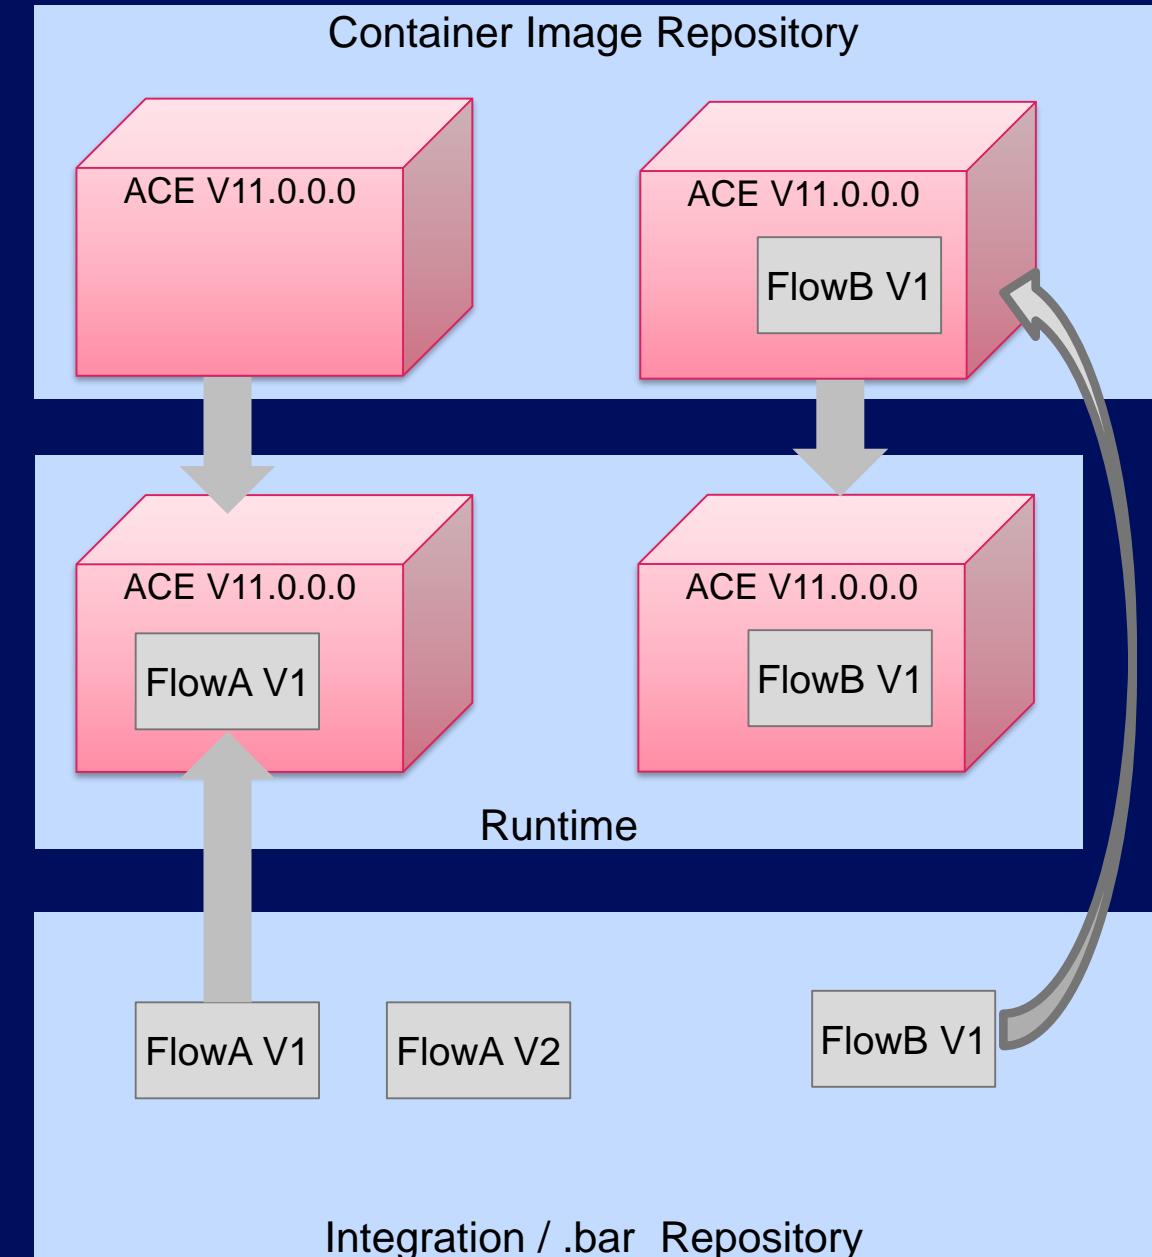
Immutability==Consistency

Do you deploy an App Connect container and then deploy flows at runtime?

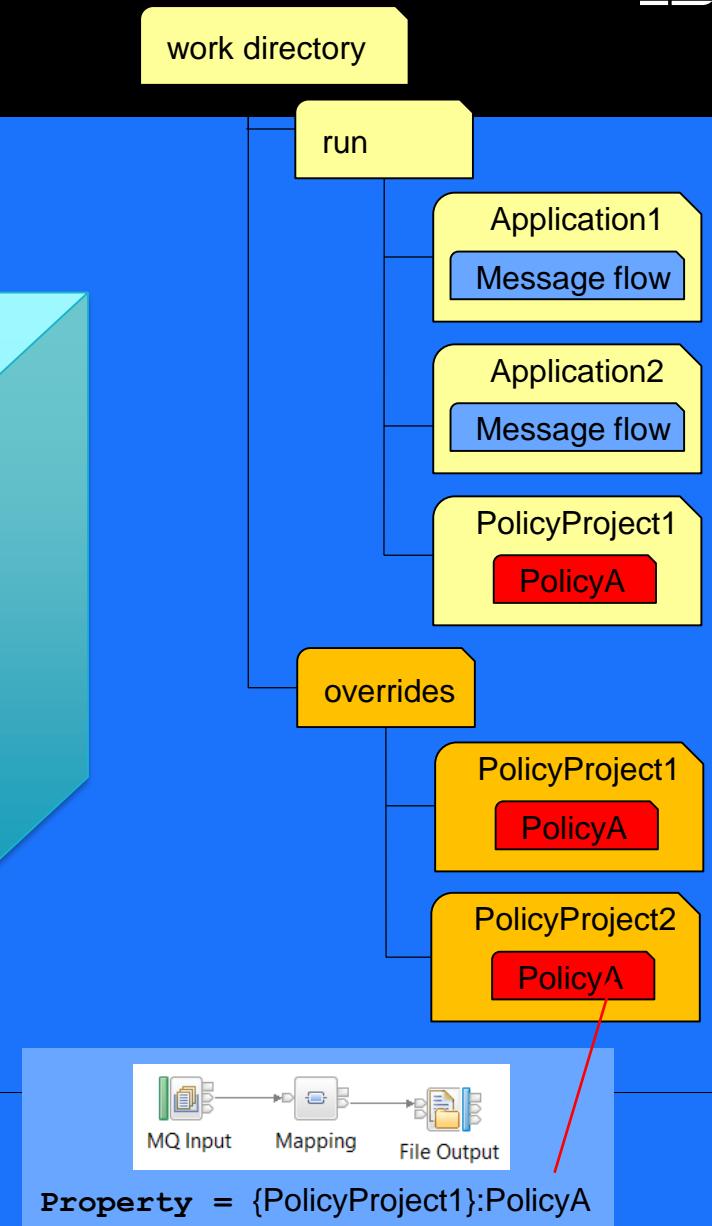
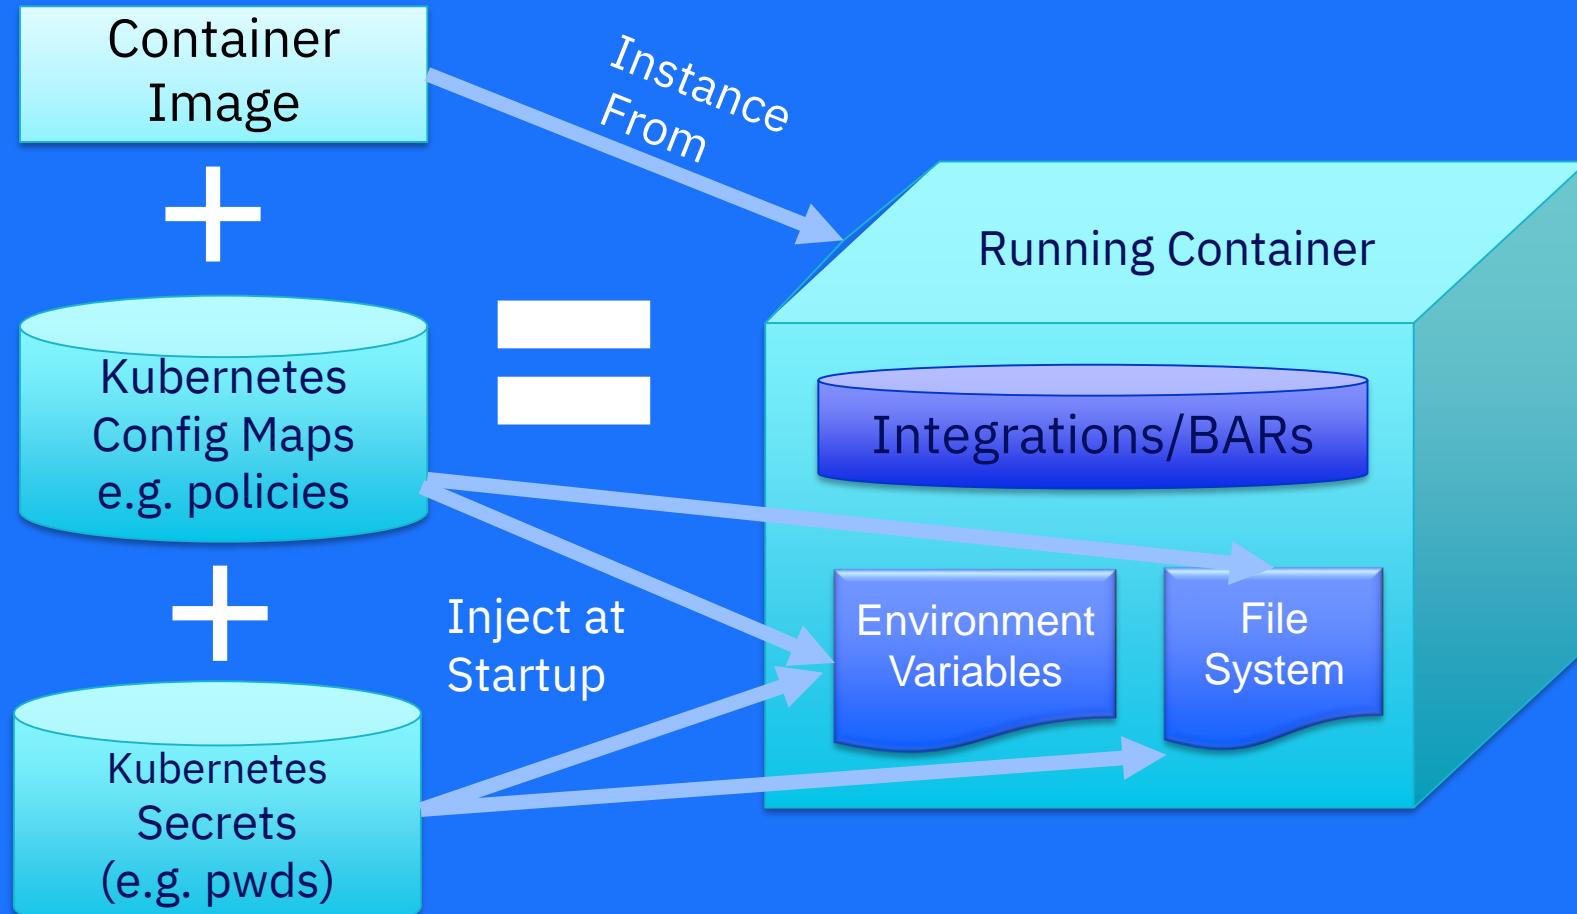
Or do you build a container with all of the flows ready-built into it?

Which makes DEV look more like PROD?

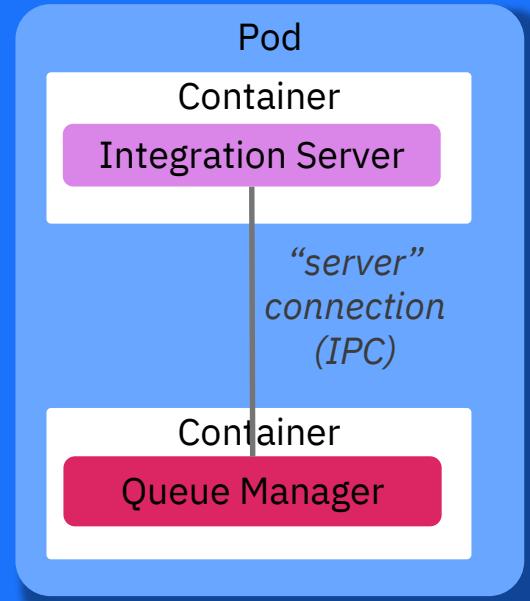
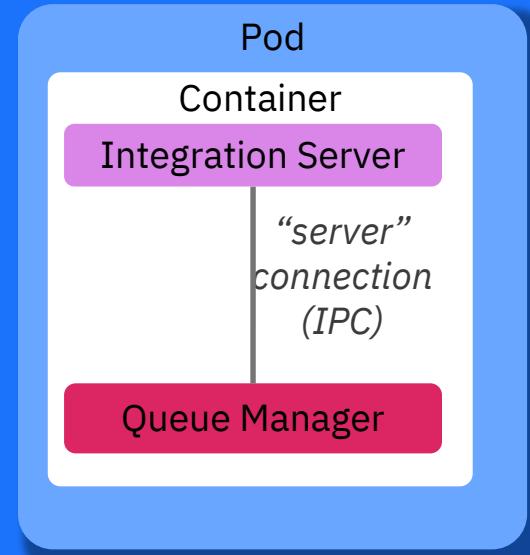
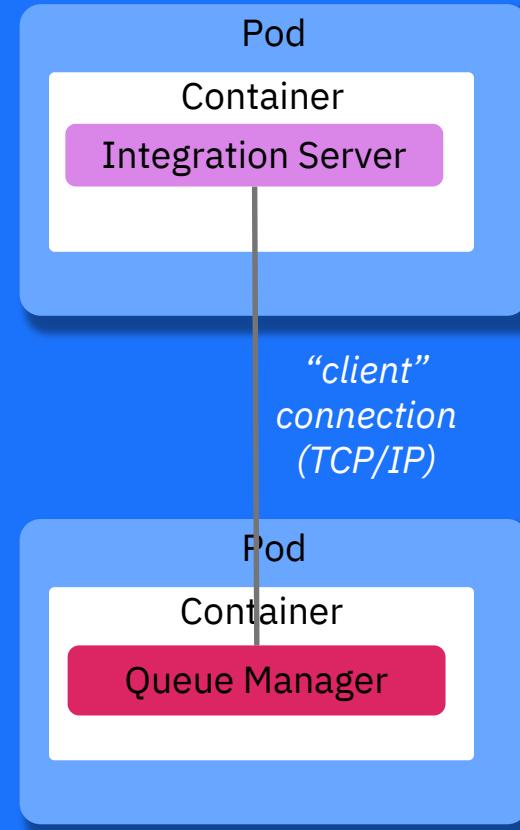
Which is simpler to deploy?



Step 2: Define your configuration points



Step 3: Create your pod definitions



A pod is more than just a group of containers

A pod shares namespace, filesystem and ip address (localhost)

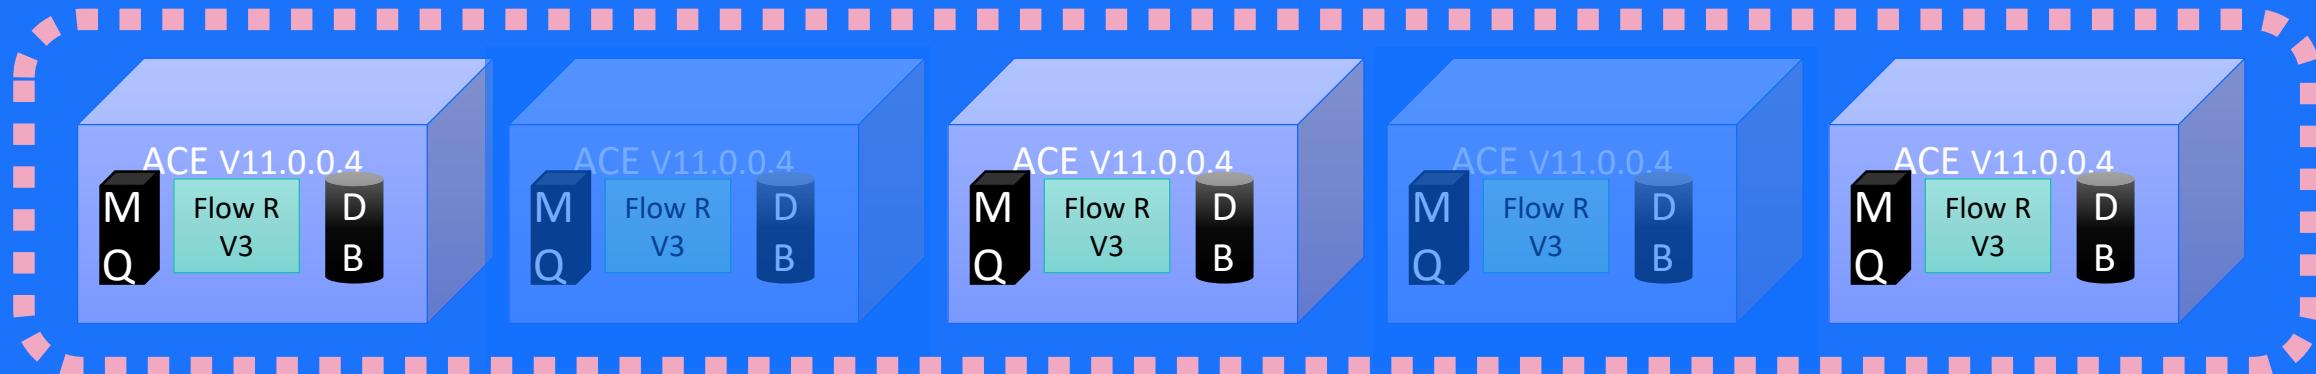
Multiple containers in a pod allows independence of container builds / versions / updates etc. with a consistent deployable unit

Step 4: Create your deployment charts**



'For this integration, Make it like this' – **Kubernetes is declarative**

- E.g. Use a Pod with:
 - Flow R, Version 3 on ACE 11.0.0.4
 - MQ version AAA
 - Database version BBB
- Make it a set of replicas of the same image
- Minimum of 3 running at all times
- Maximum of 5 running – K8s: you scale it
- Expose ports x, y and z



Did we mention we have sample helm charts?



The screenshot shows a GitHub repository page for `ot4i/ace-helm`. The repository has 1 issue and 0 pull requests. It contains files for versions 11.0.0.3 and 11.0.0.2, including `Chart.yaml`, `README.md`, and `values.yaml`. The `values.yaml` file is displayed on the right side of the screen, showing configuration for a Helm chart. The code includes template logic for service accounts and container ports, along with environment variables for license and QMGR usage.

```
72      - {{ .Values.arch }}
73      serviceAccountName: {{ include "fullname" . }}
74      {{- if .Values.image.pullSecret }}
75      imagePullSecrets:
76        - name: {{ .Values.image.pullSecret }}
77      {{- end }}
78      hostNetwork: false
79      hostPID: false
80      hostIPC: false
81      securityContext:
82        runAsUser: 1000
83        runAsNonRoot: true
84      fsGroup: {{ default 1001 .Values.fsGroupGid }}
85      containers:
86        - name: {{ $deploymentName }}
87          image: "{{ .Values.image.repository.aceonly }}:{{ .Values.image.tag }}"
88          imagePullPolicy: {{ .Values.image.pullPolicy }}
89          ports:
90            - containerPort: {{ .Values.service.webuiPort }}
91              name: webui
92            - containerPort: {{ .Values.service.serverlistenerPort }}
93              name: ace-http
94            - containerPort: {{ .Values.service.serverlistenerTLSPort }}
95              name: ace-https
96          env:
97            - name: LICENSE
98              value: {{ .Values.license | quote }}
99            - name: USE_QMGR
100             value: "false"
```

Container Philosophy Part 2:

Containers are plentiful and cheap

If we lose one (crashes), just spin up another**

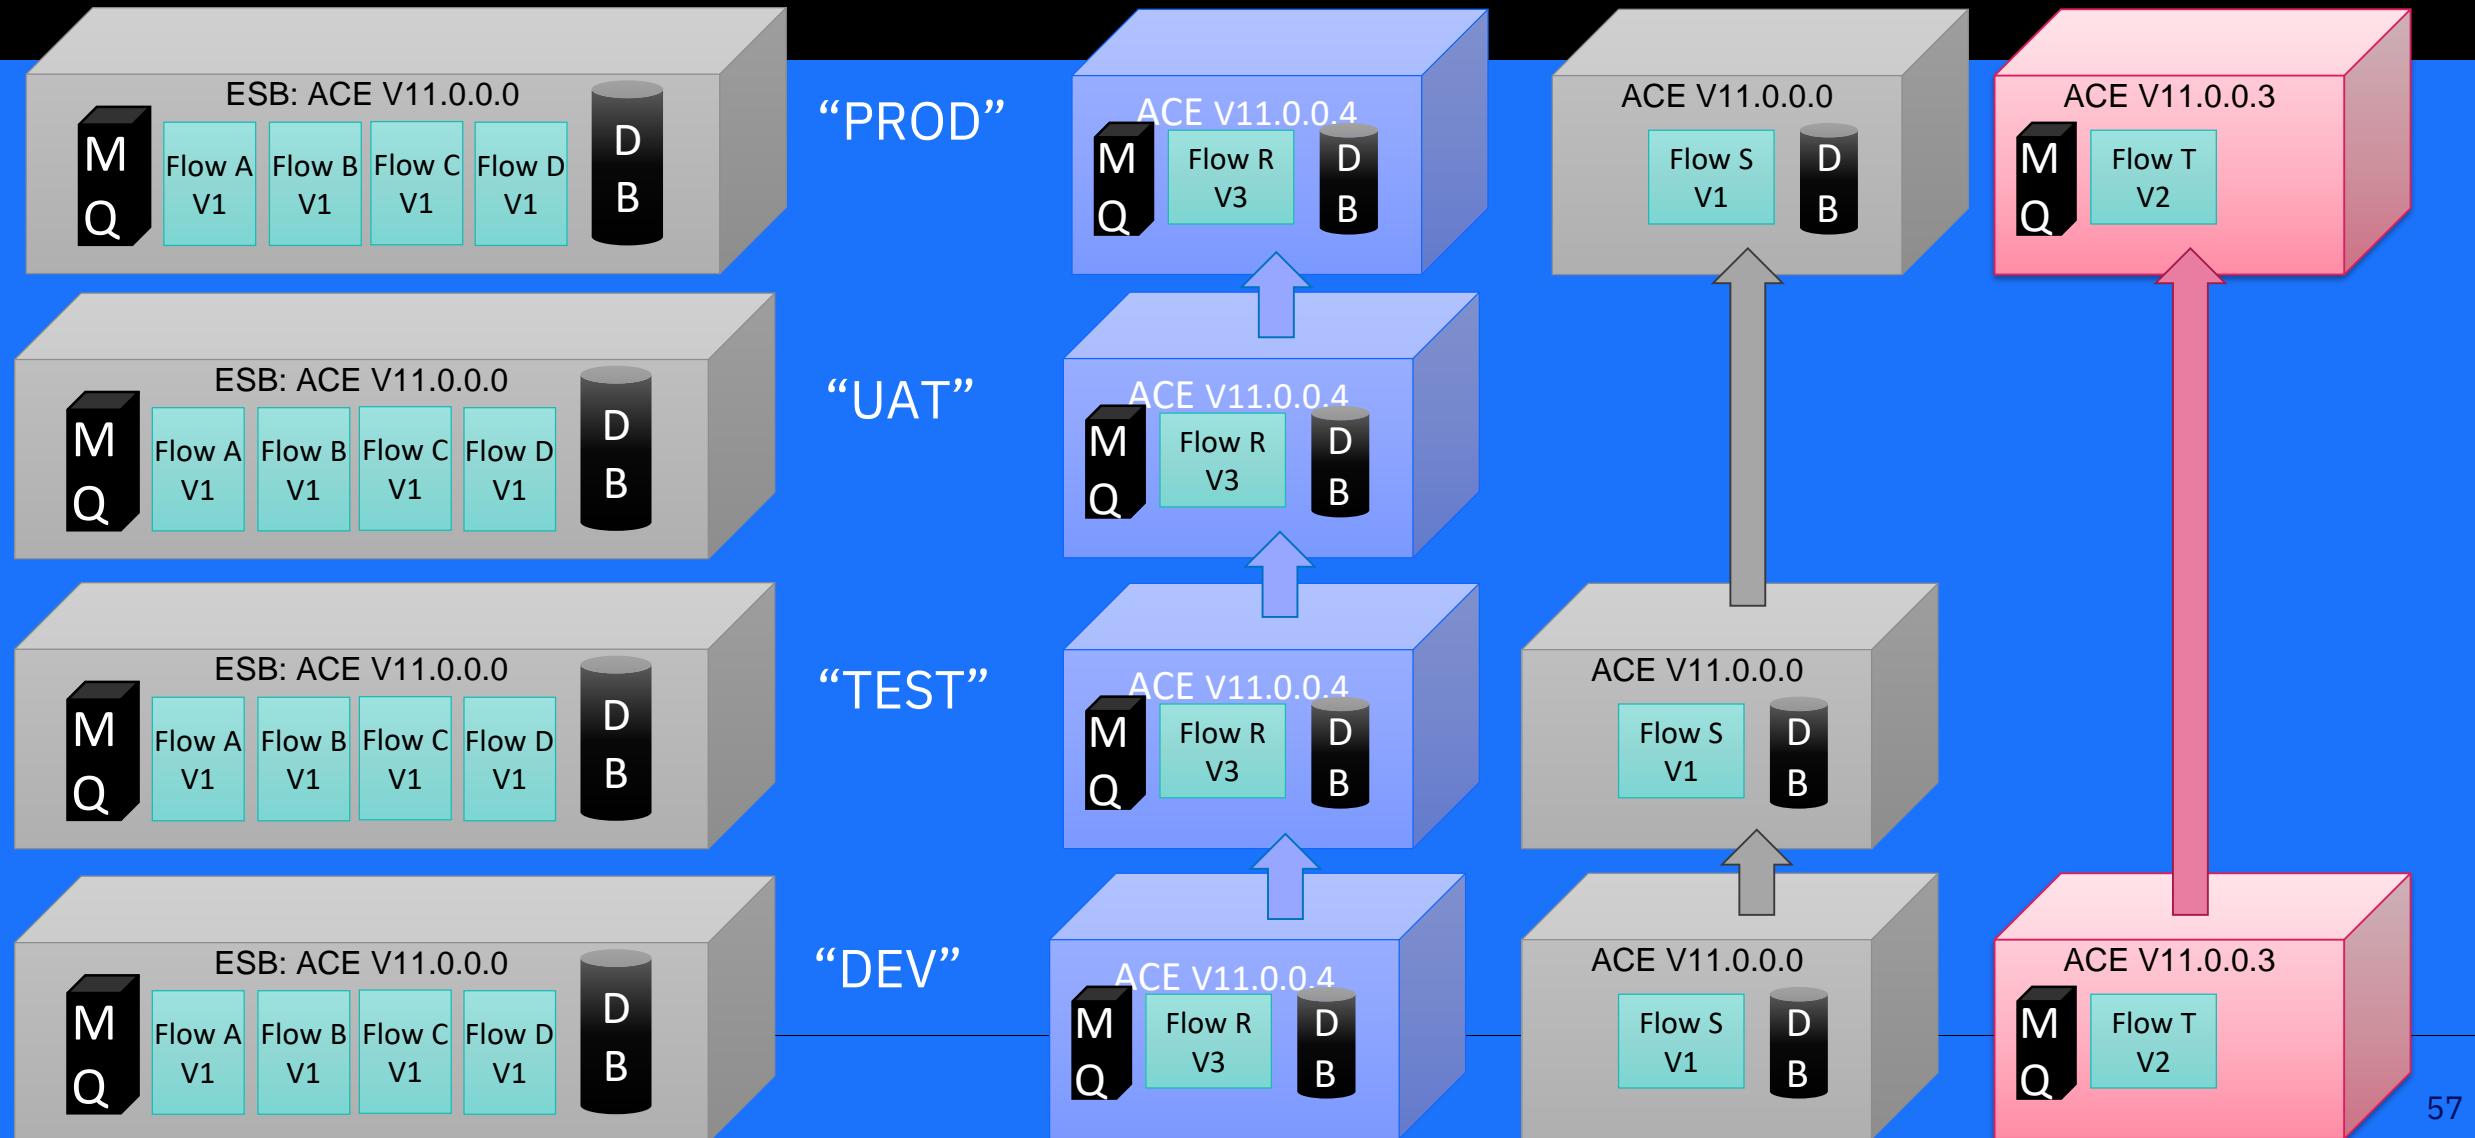
“Just Another Container” – not “an ACE container”?

** If you have state, make sure you know where it is

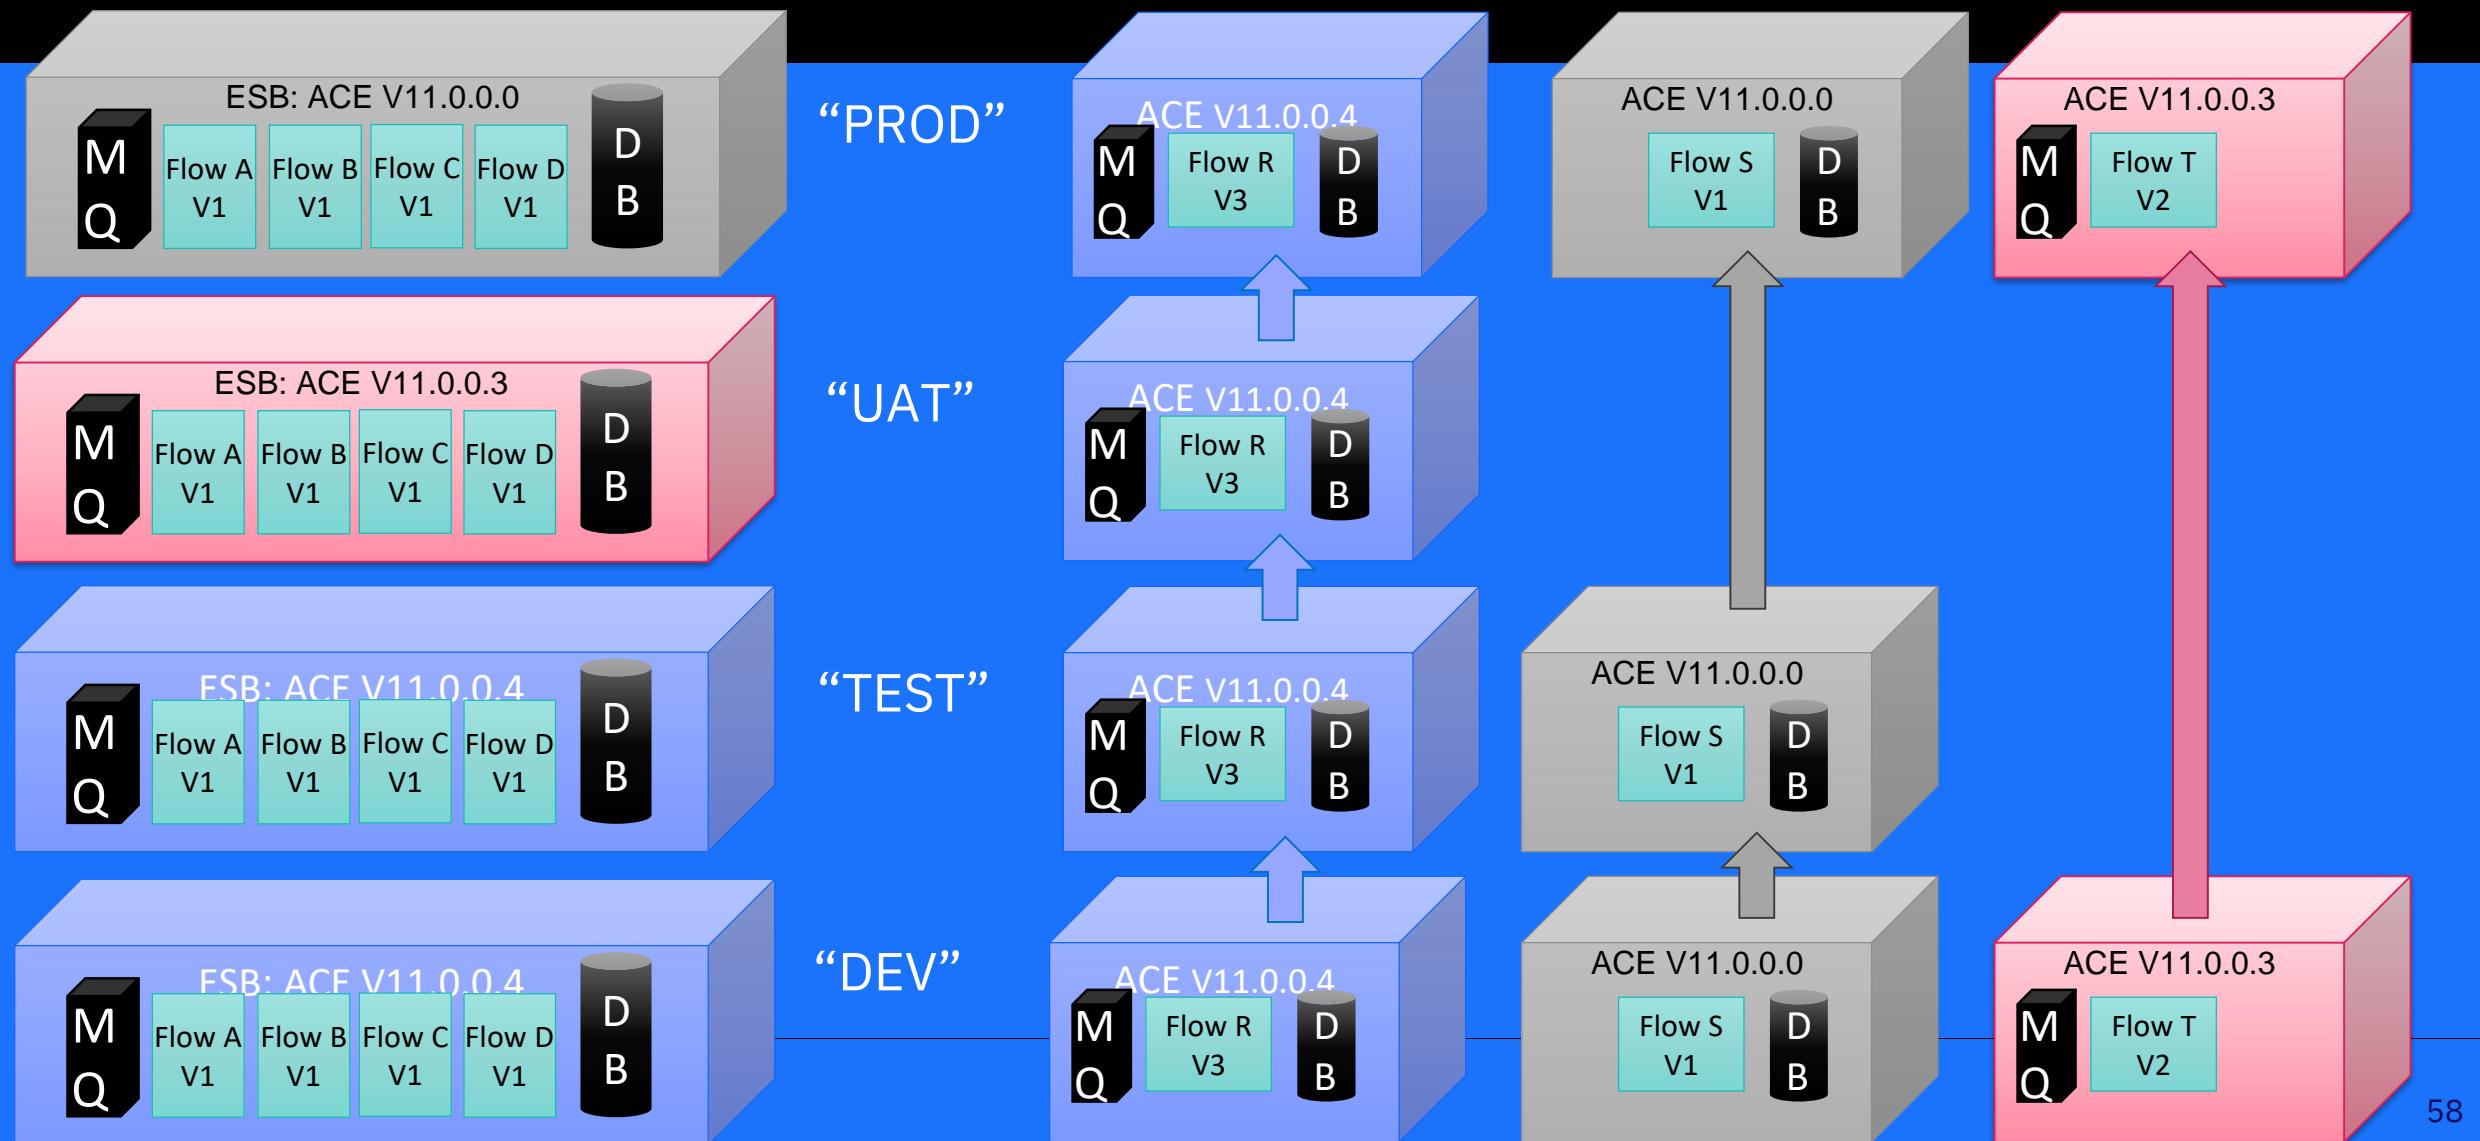


Infrastructure as code – no more “environments”?

IBM

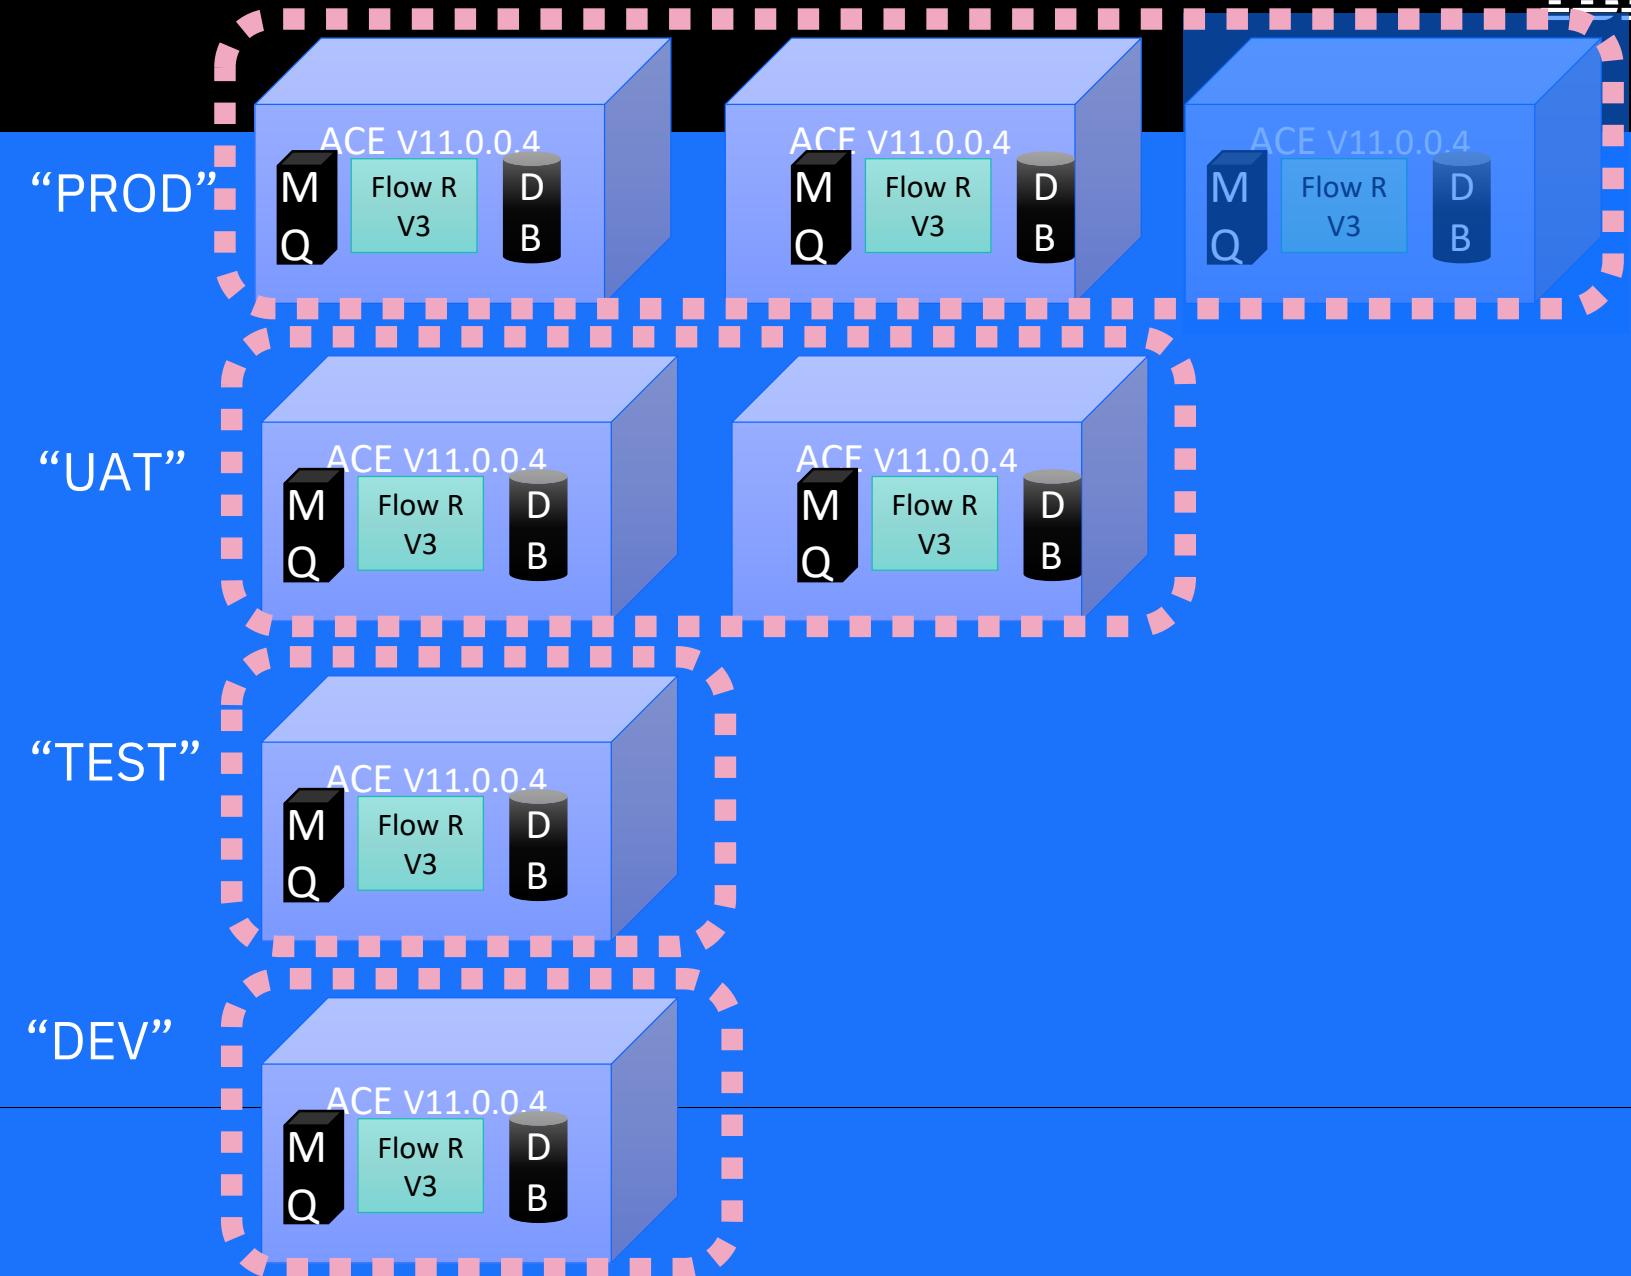
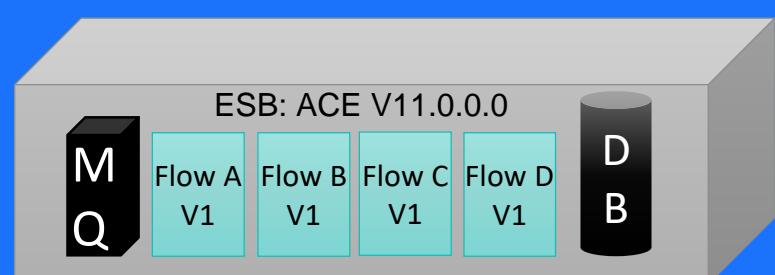
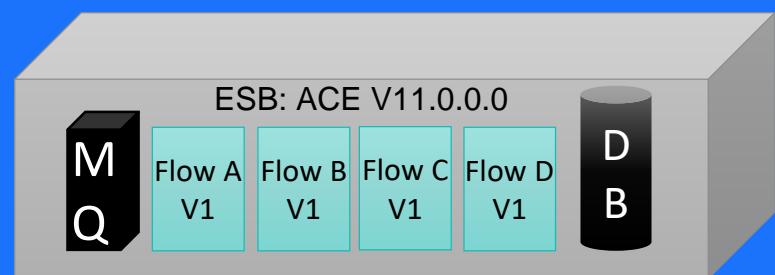
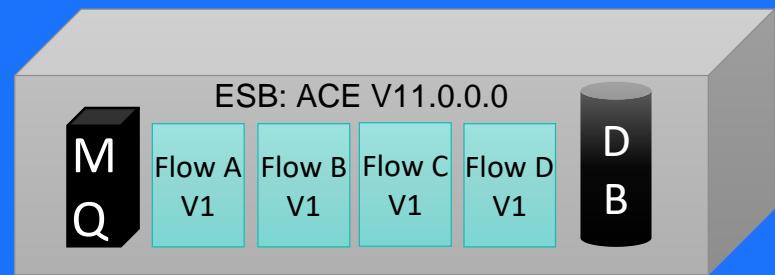
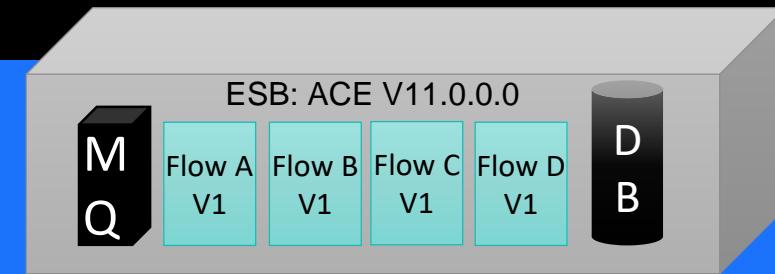


No more “Environment Drift?”



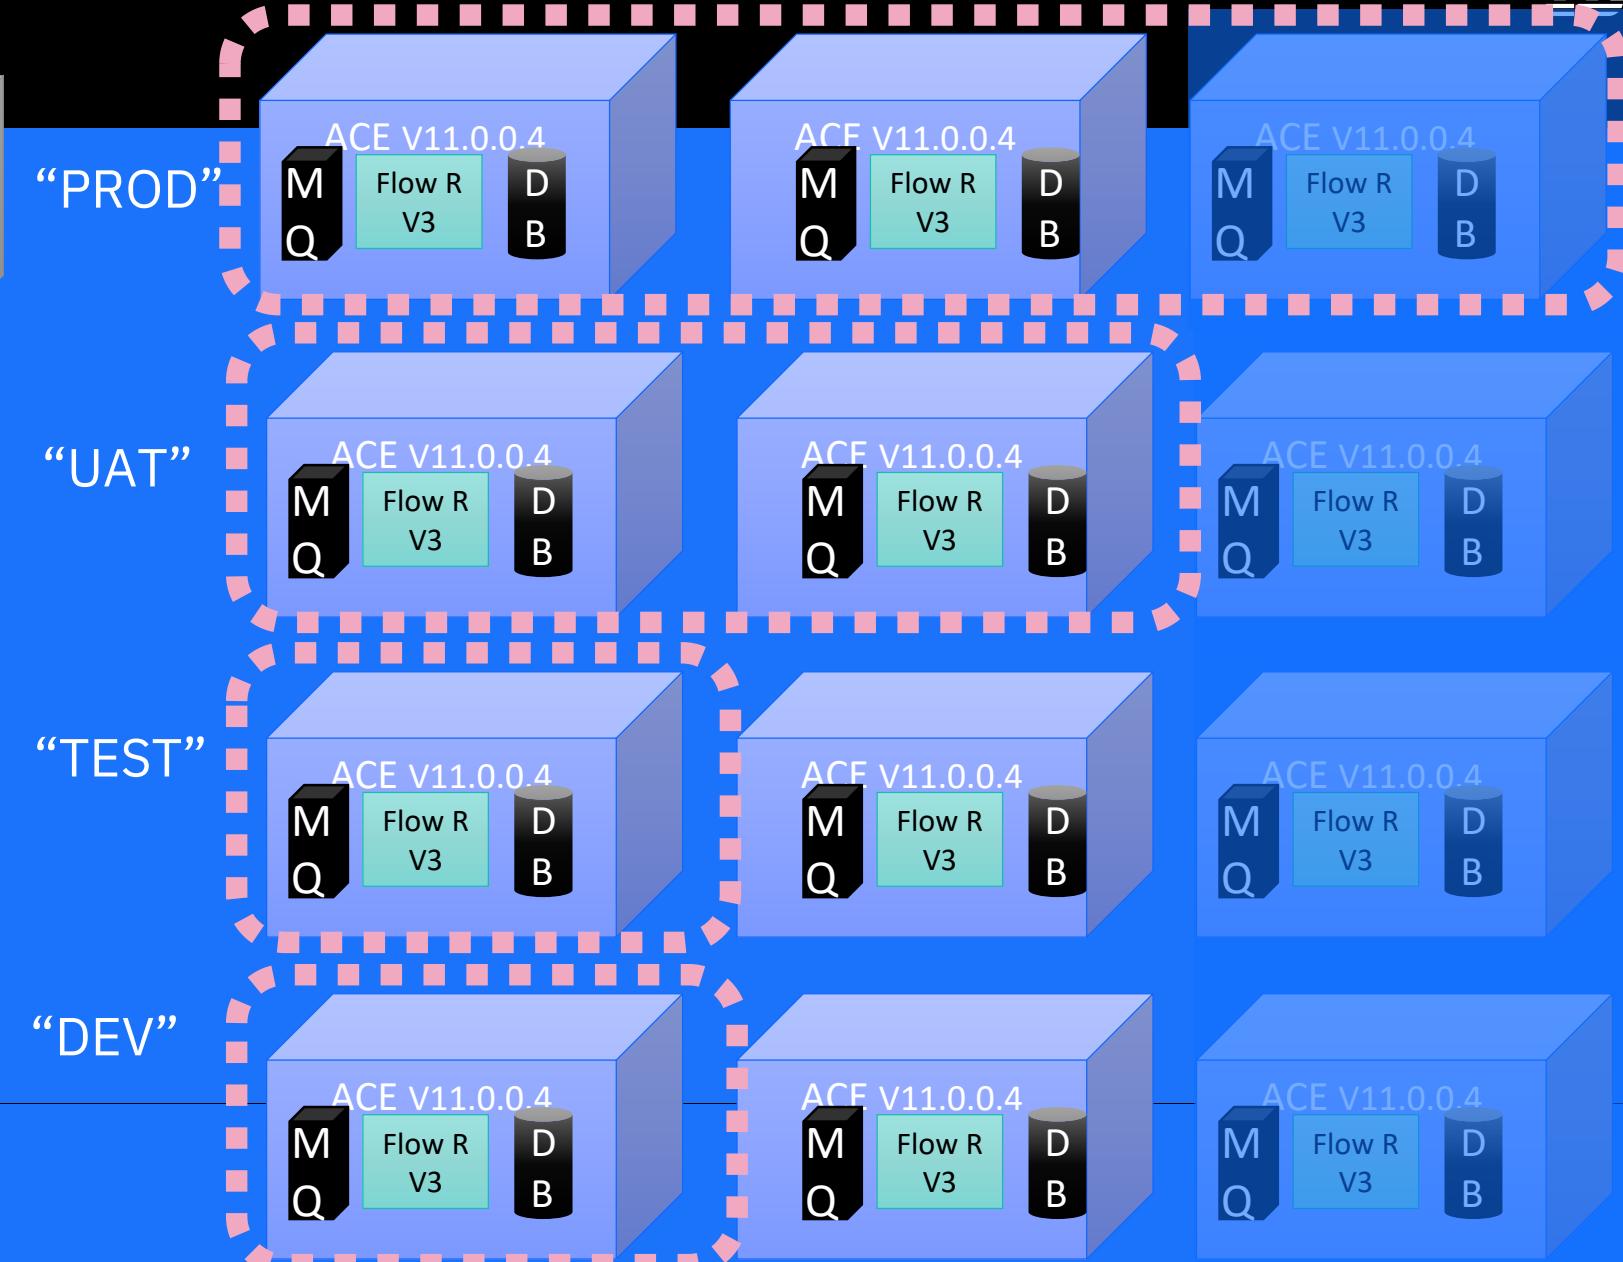
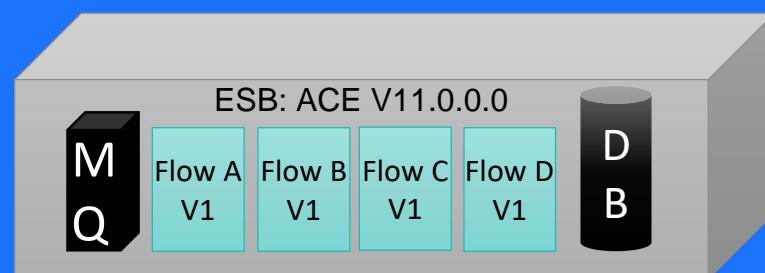
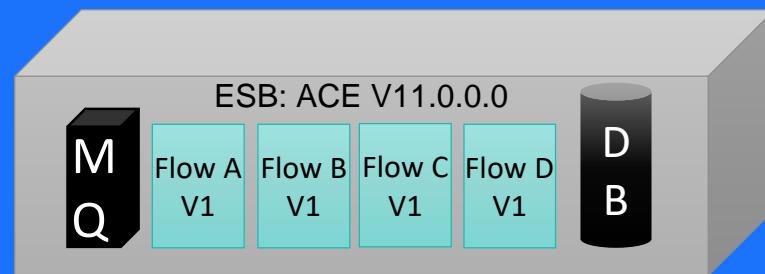
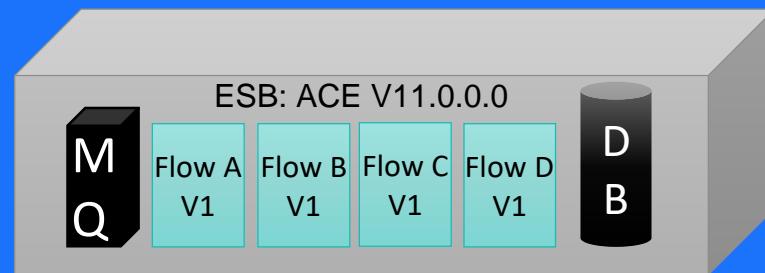
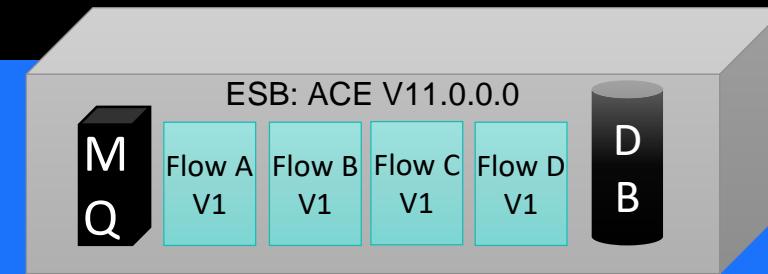
Shift DevOps “Left”: Performance and HA/DR in DEV

IBM



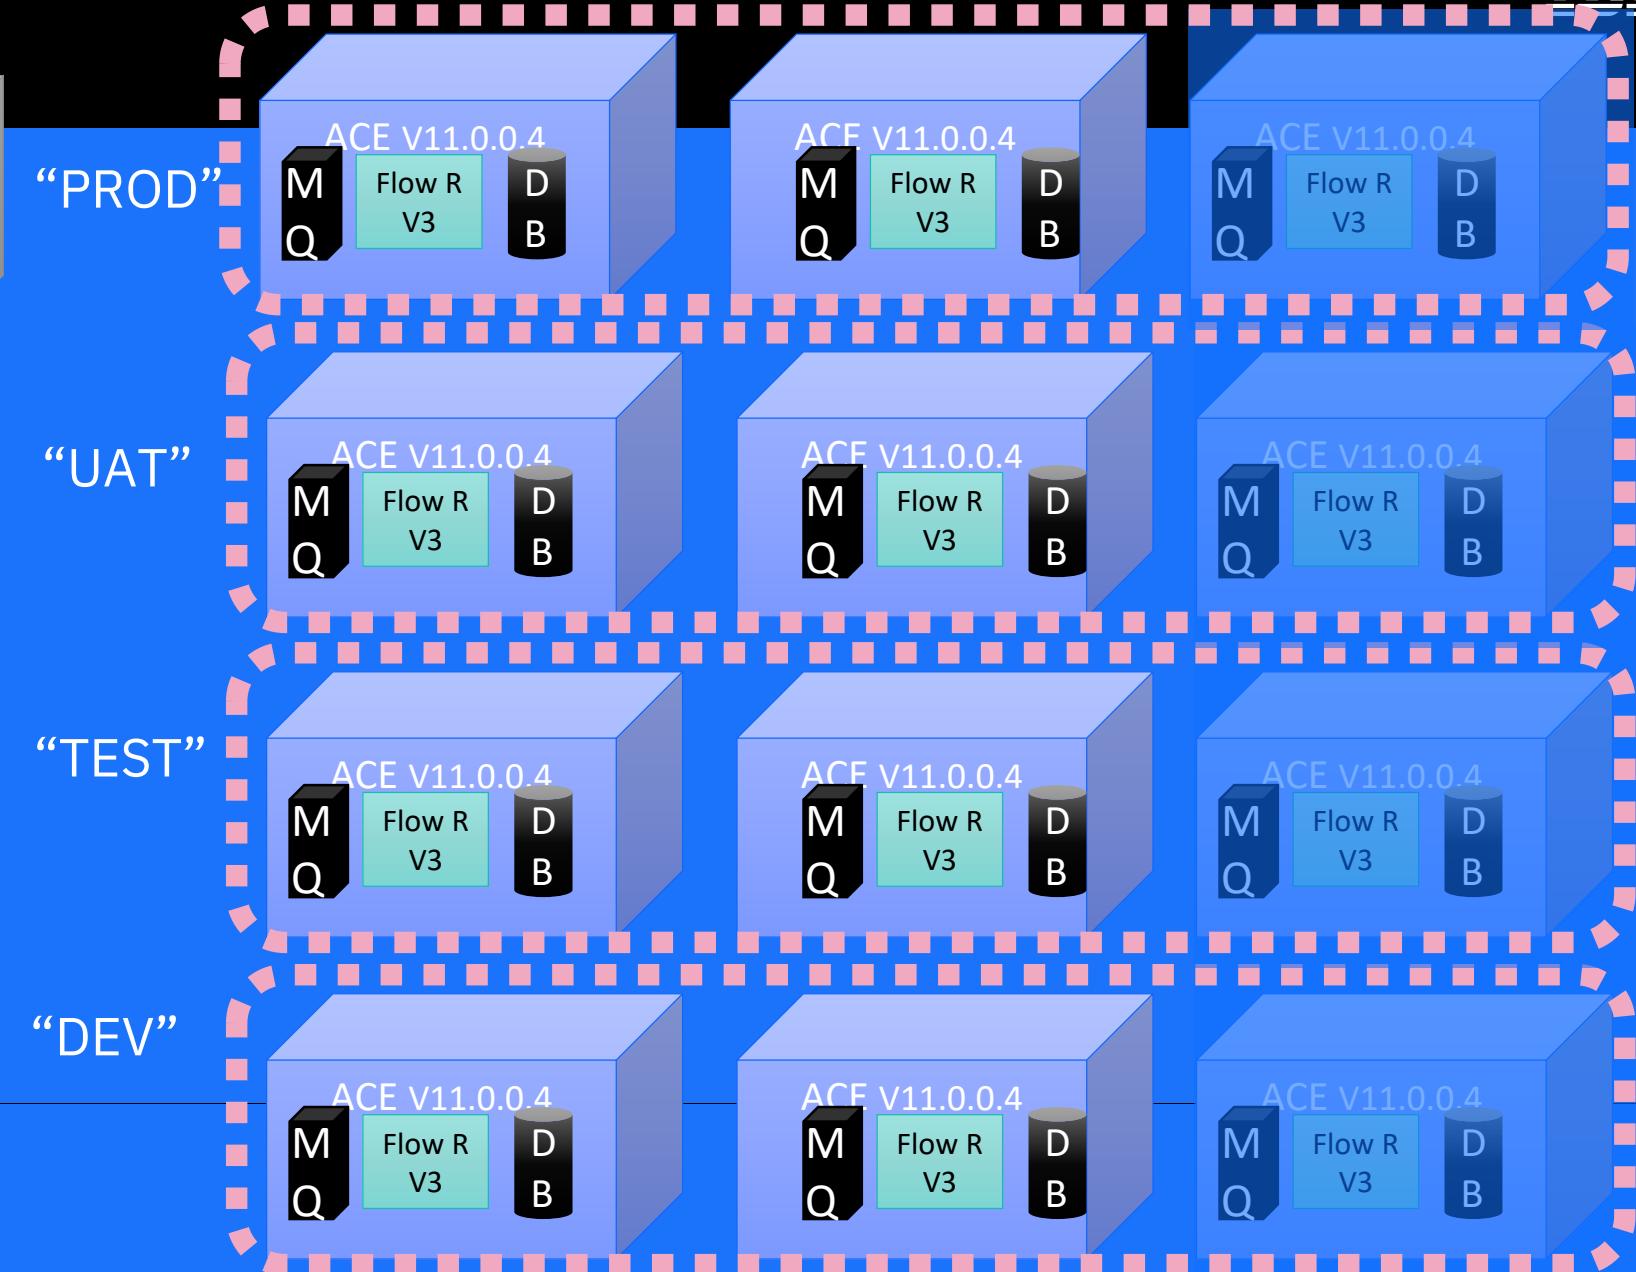
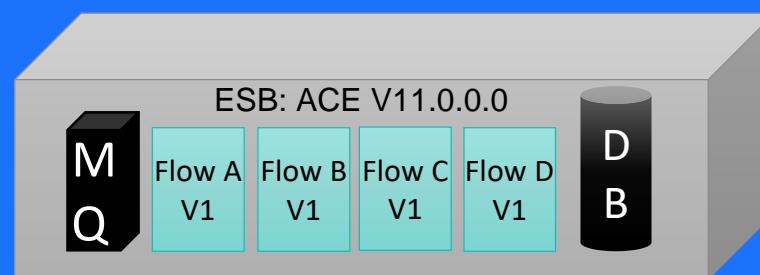
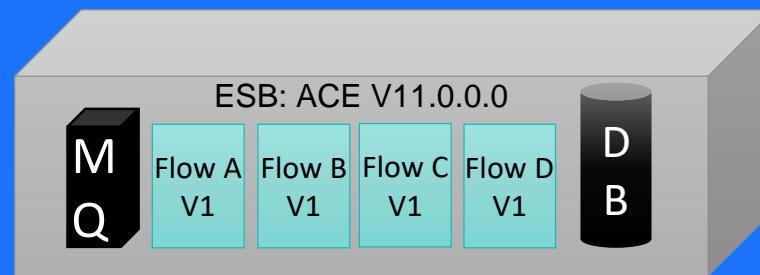
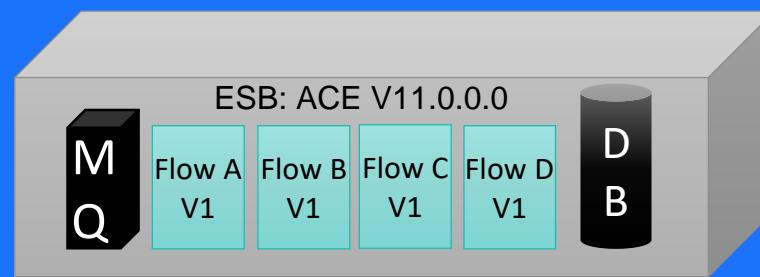
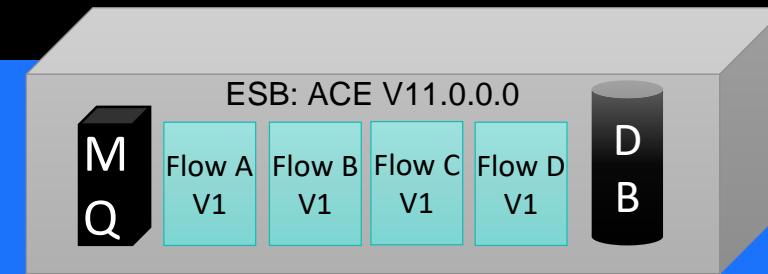
Shift DevOps “Left”: Performance and HA/DR in DEV

IBM



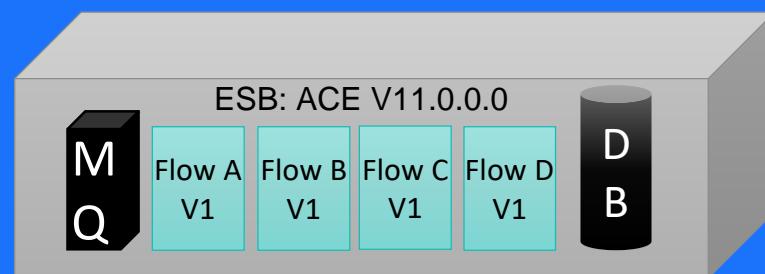
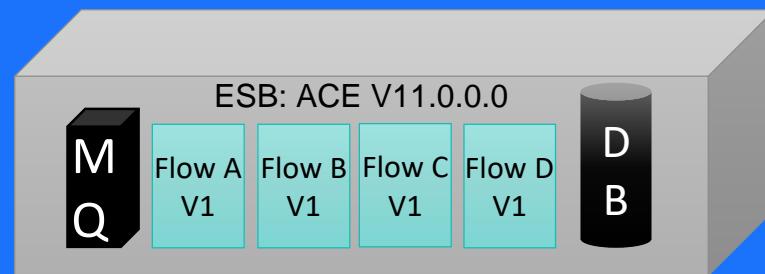
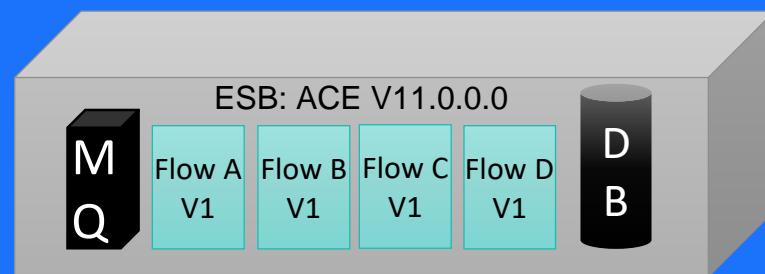
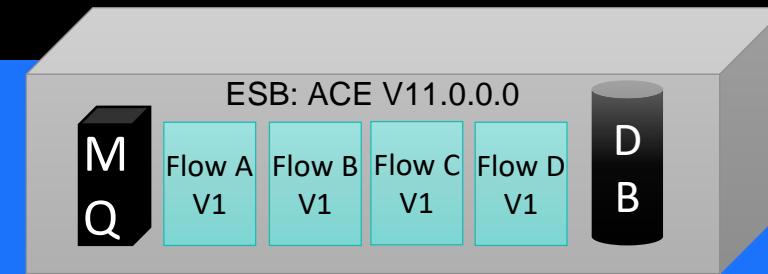
Shift DevOps “Left”: Performance and HA/DR in DEV

IBM



“Think Cloud” – Only spin up when you need it

IBM



“PROD”

“UAT”

“TEST”

“DEV”



Agile Integration Architecture

Modernizing integration
to enable
business agility

Fine grained deployment

Architecture & design

Decentralized Ownership

People & Process

Cloud native infrastructure

Infrastructure & Technology



Improve build independence and production velocity
(deployment agility)



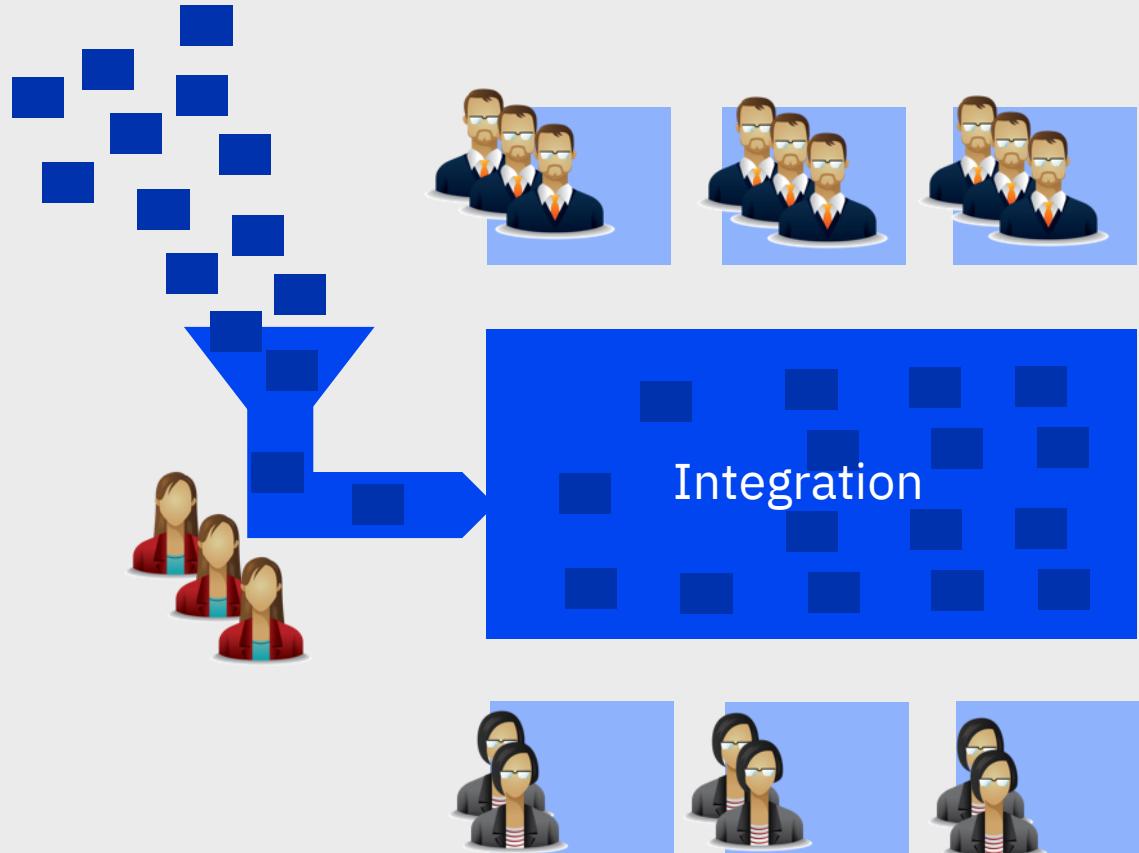
Accelerate agility and innovation
(development agility)



Dynamic scalability and inherent resilience
(operational agility)

Decentralized Ownership

People & Process



Centralized integration facility with plateauing productivity



Increased team agility as a result of greater autonomy and clearer ownership

Independent Deployment & Pipeline

Enables independent Lifecycle Ownership



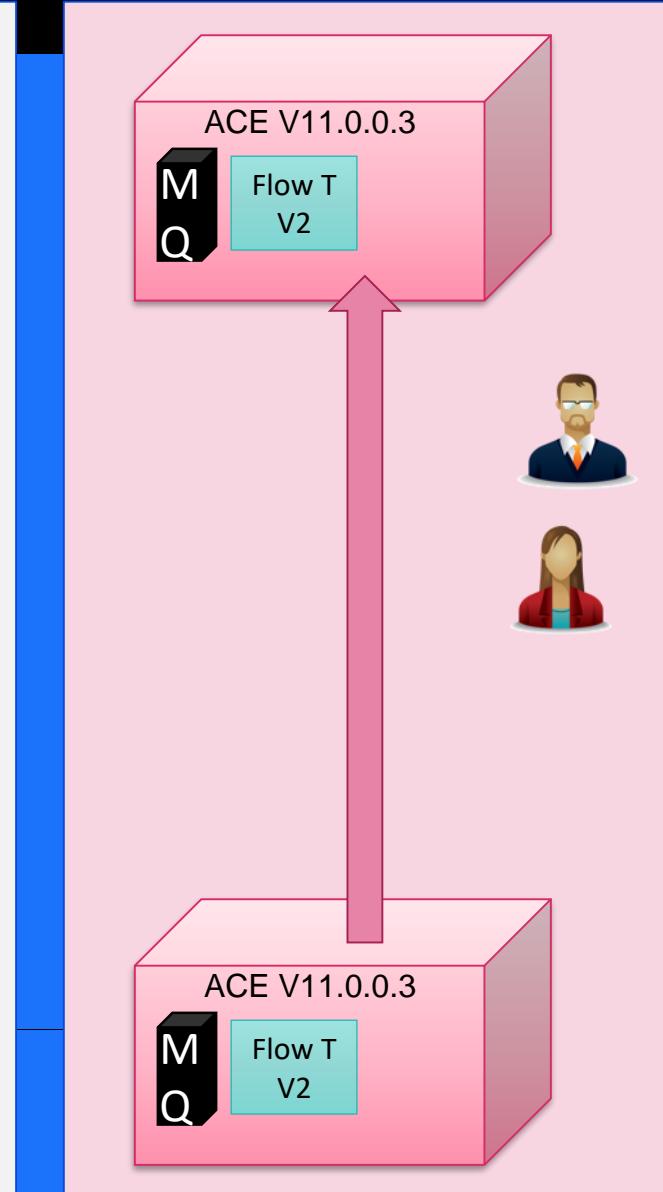
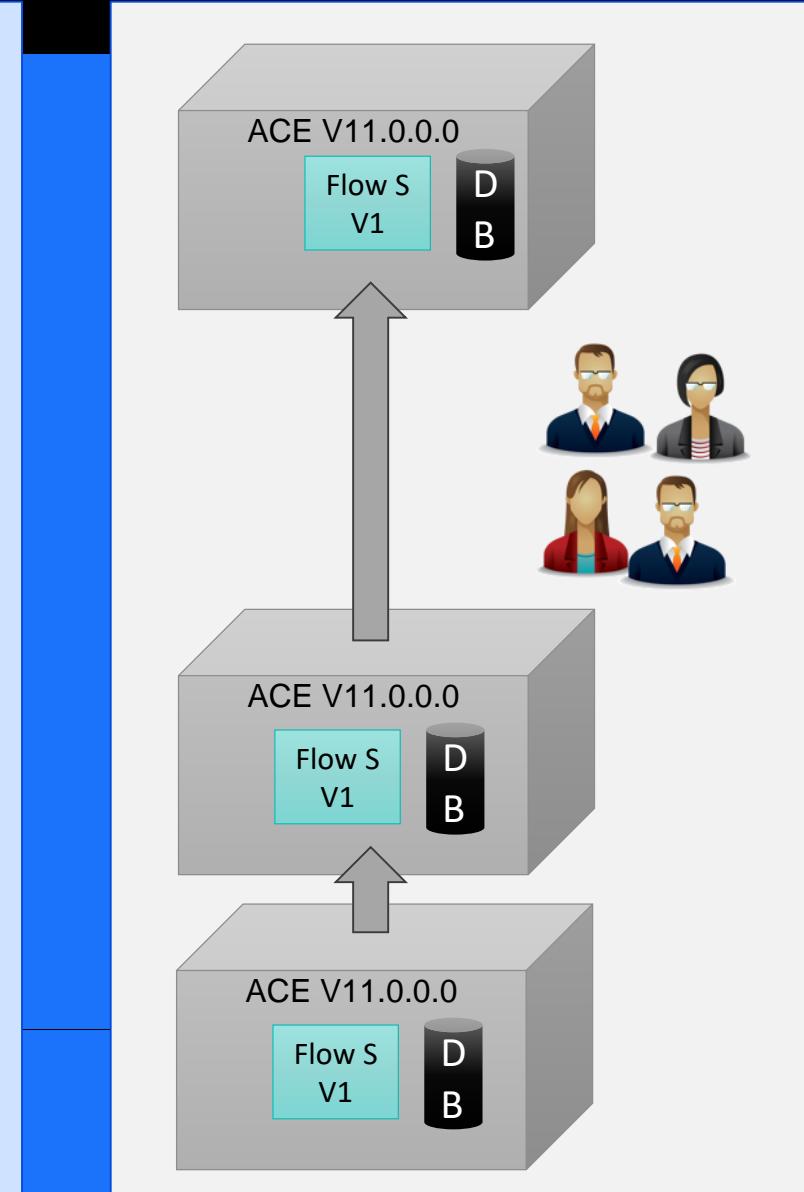
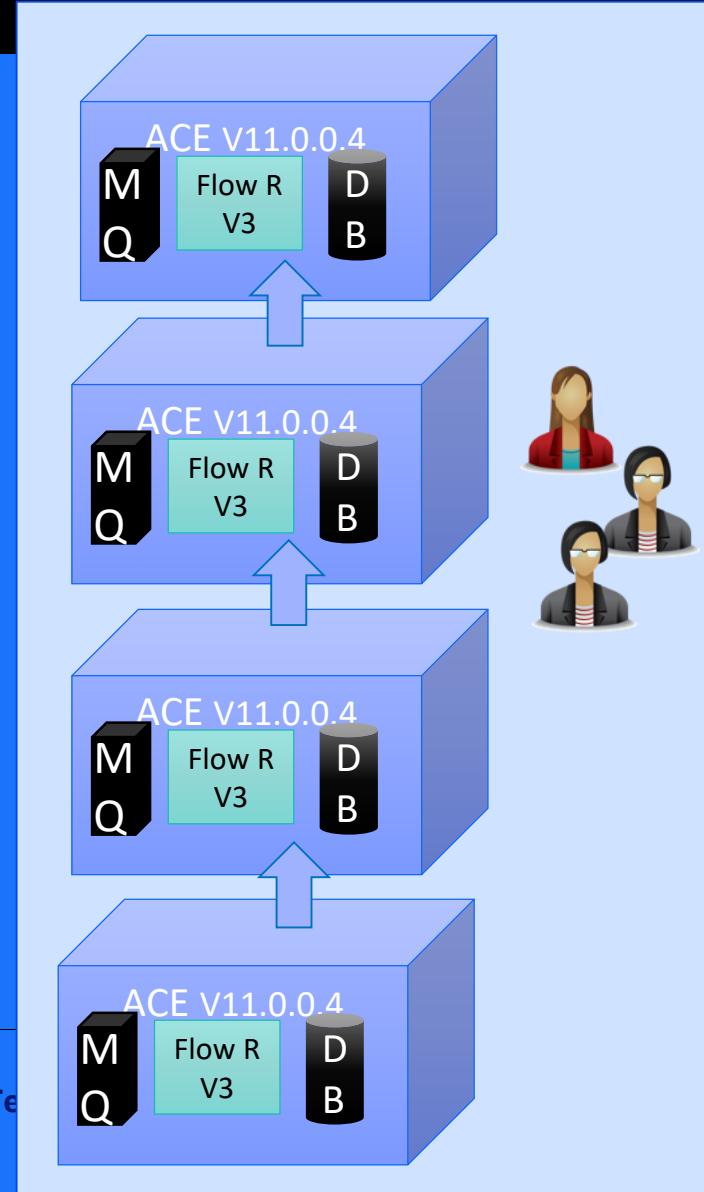
“PROD”

“UAT”

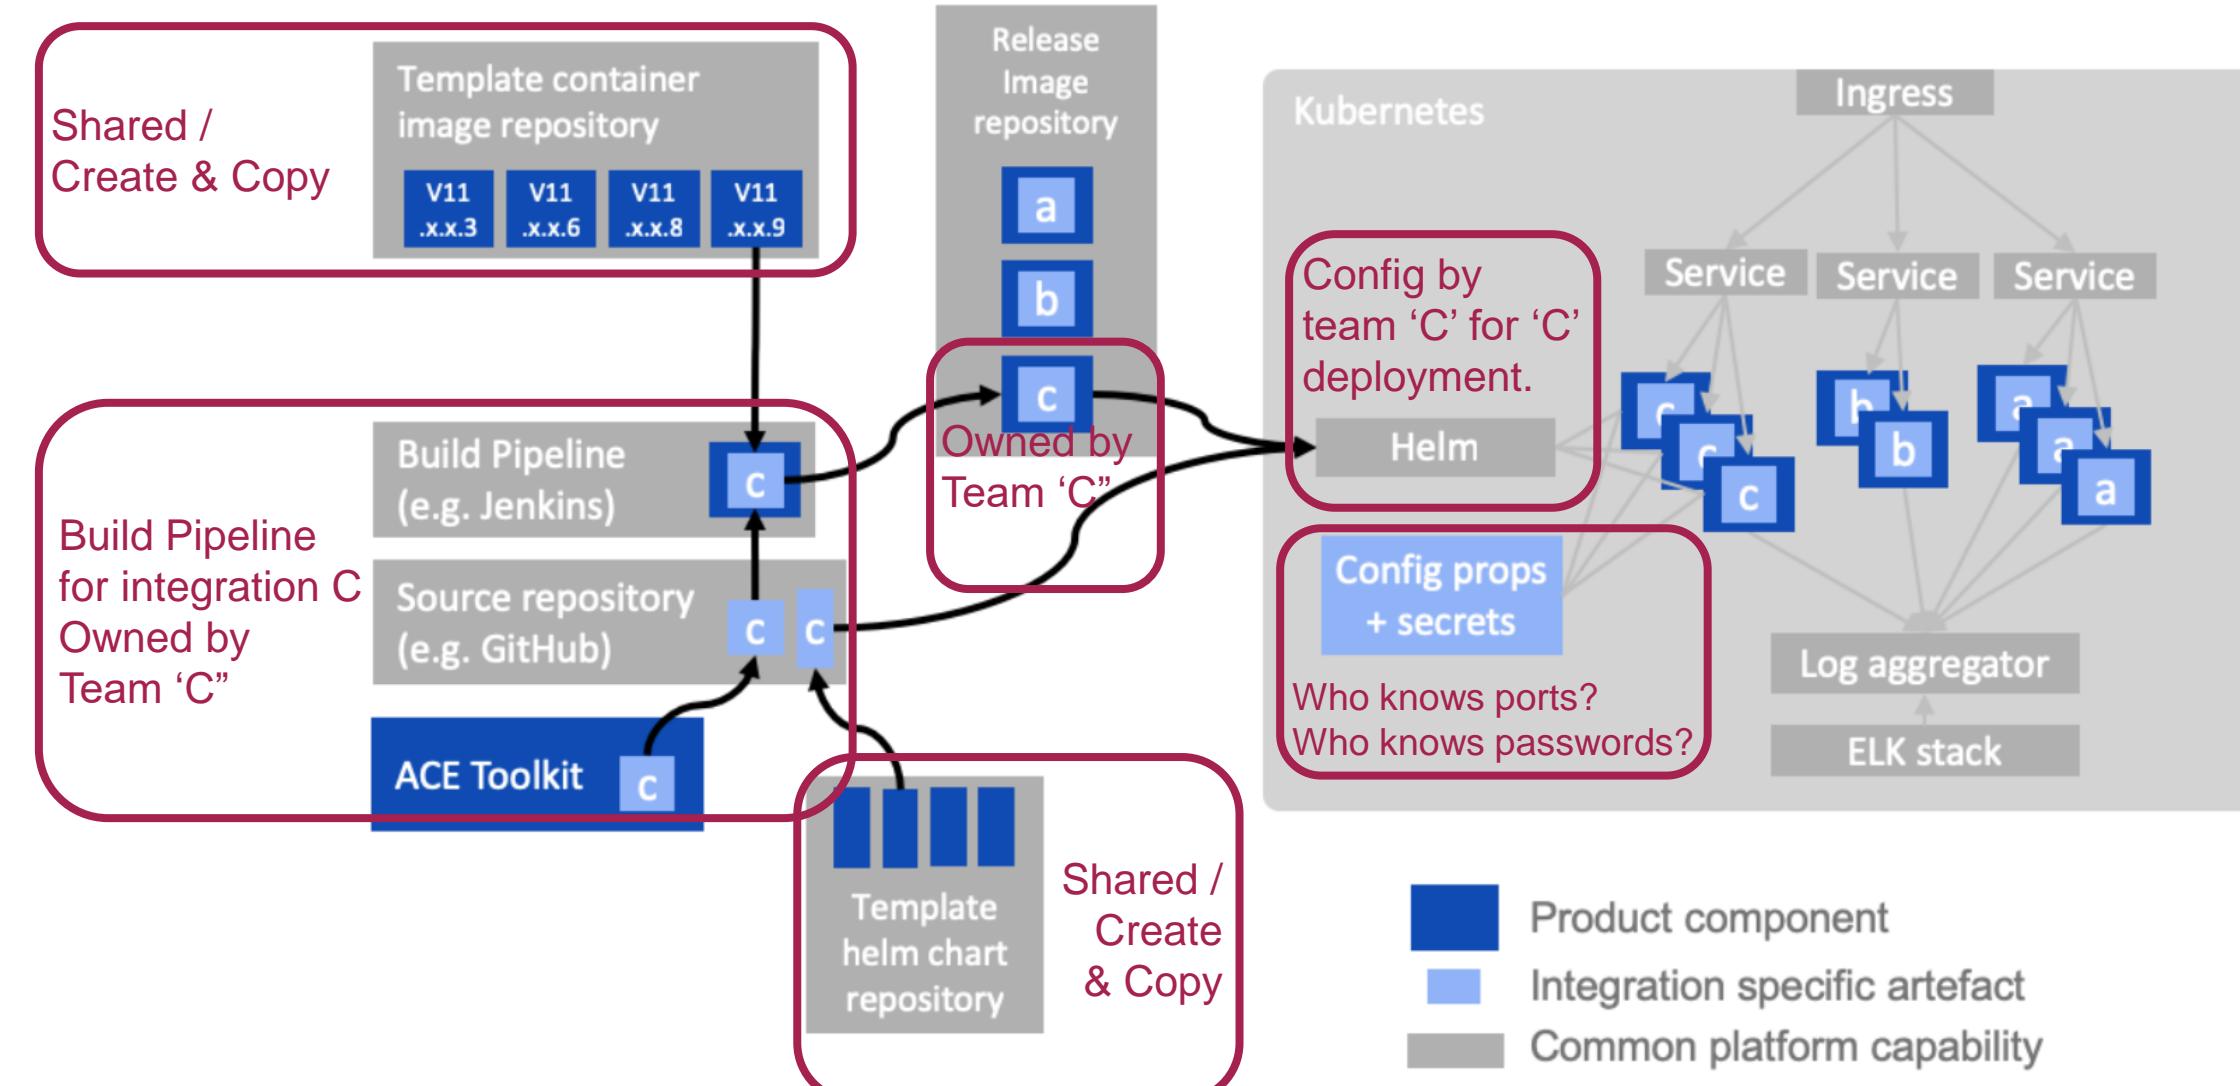
“TEST”

“DEV”

Integration Te



Team owns their pipeline: Code+build+config+deploy



Logging

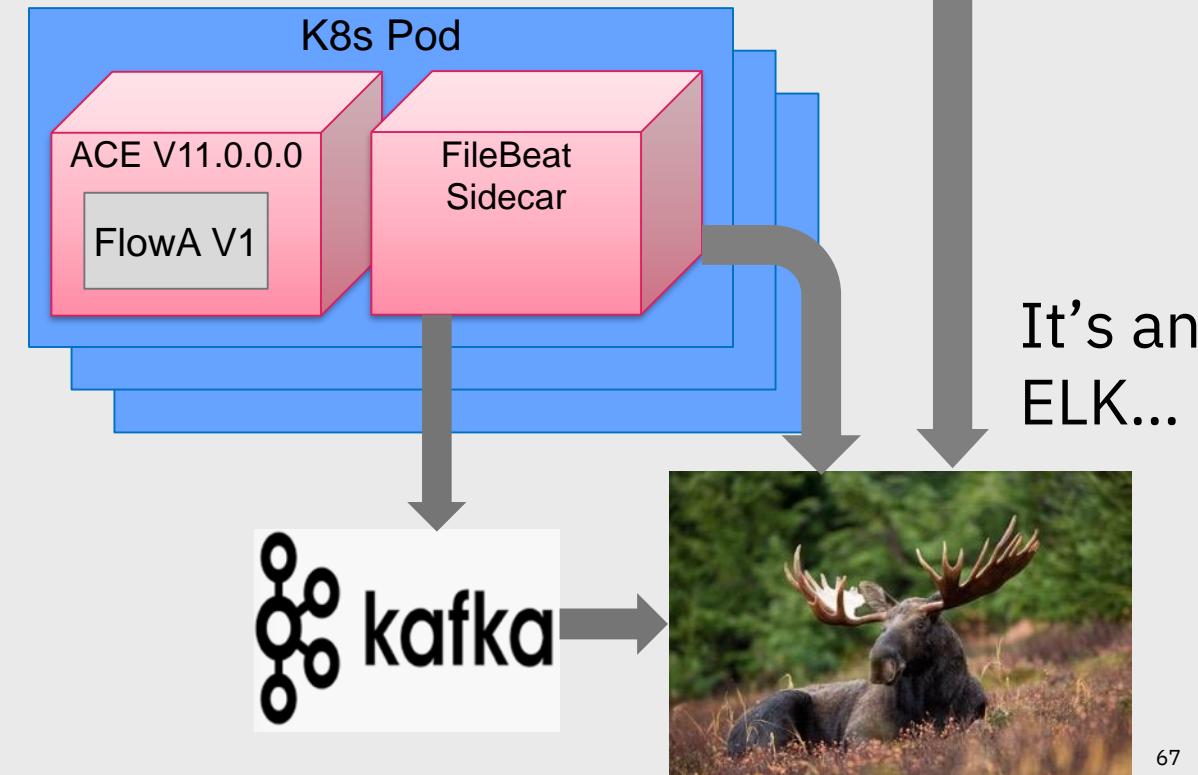
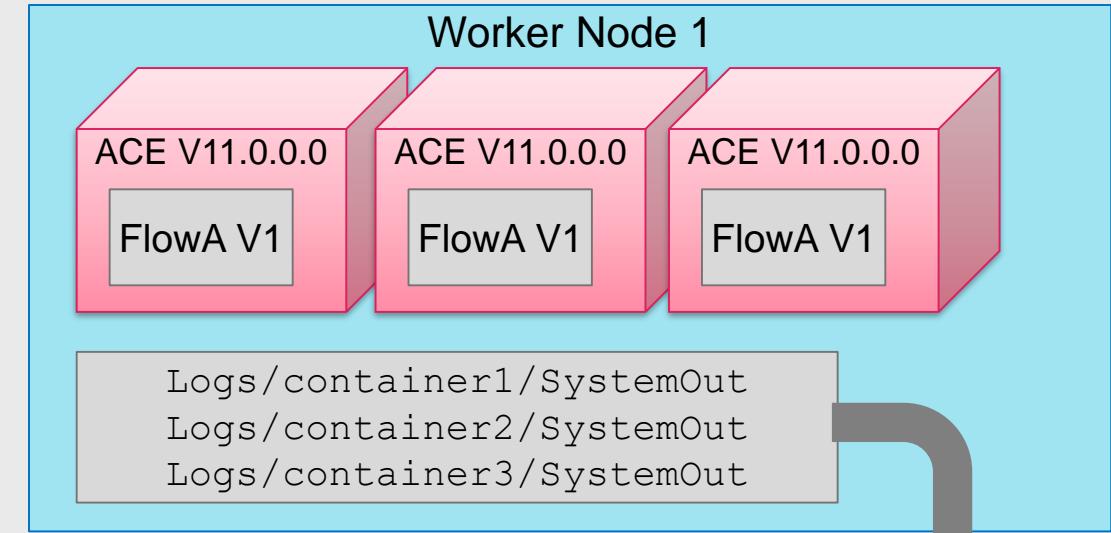
One place to log them all...

One place to find them!

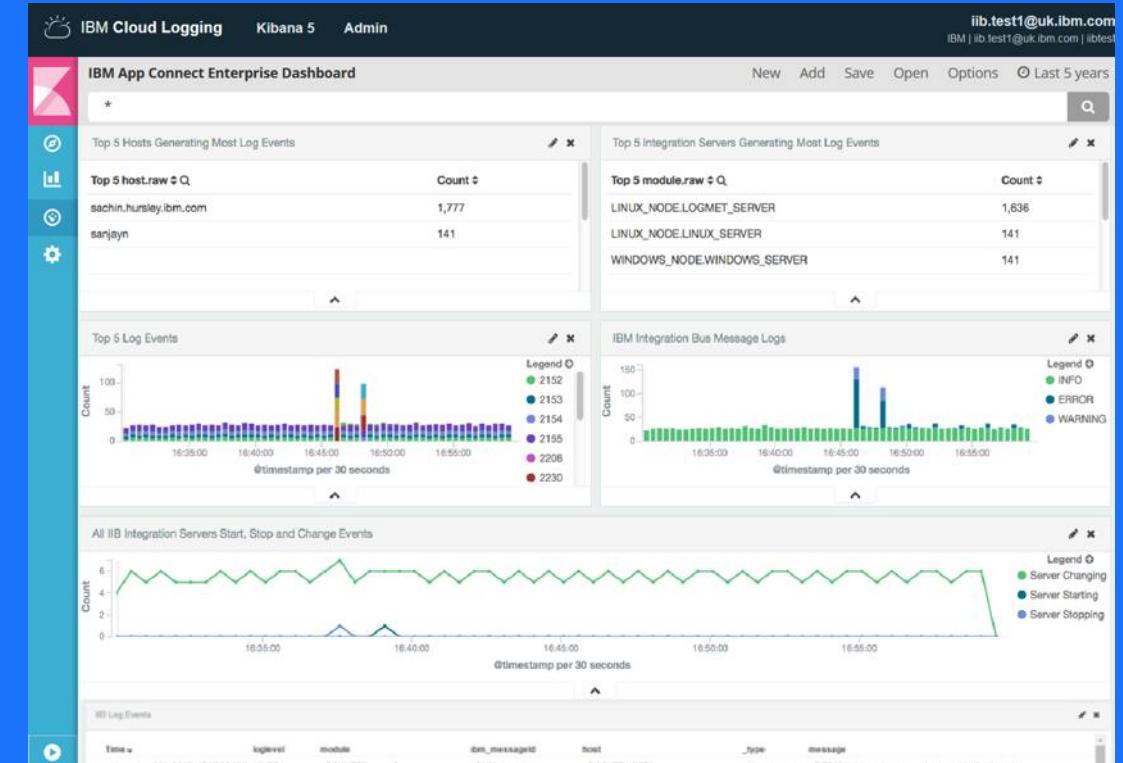
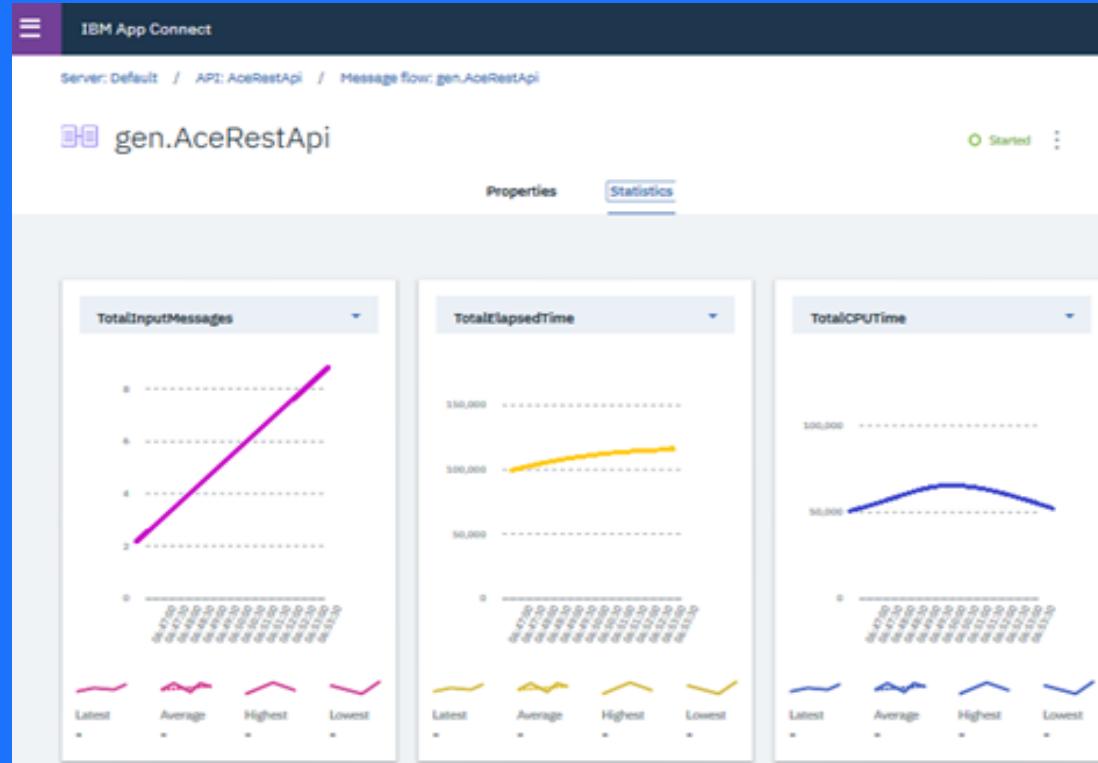
You will create a LOT of logs - the number will change dynamically.

Forward and Manage your logs centrally – then filter as you need.

Standardise format as far as possible.



Monitoring: “Just another container/Pod”?

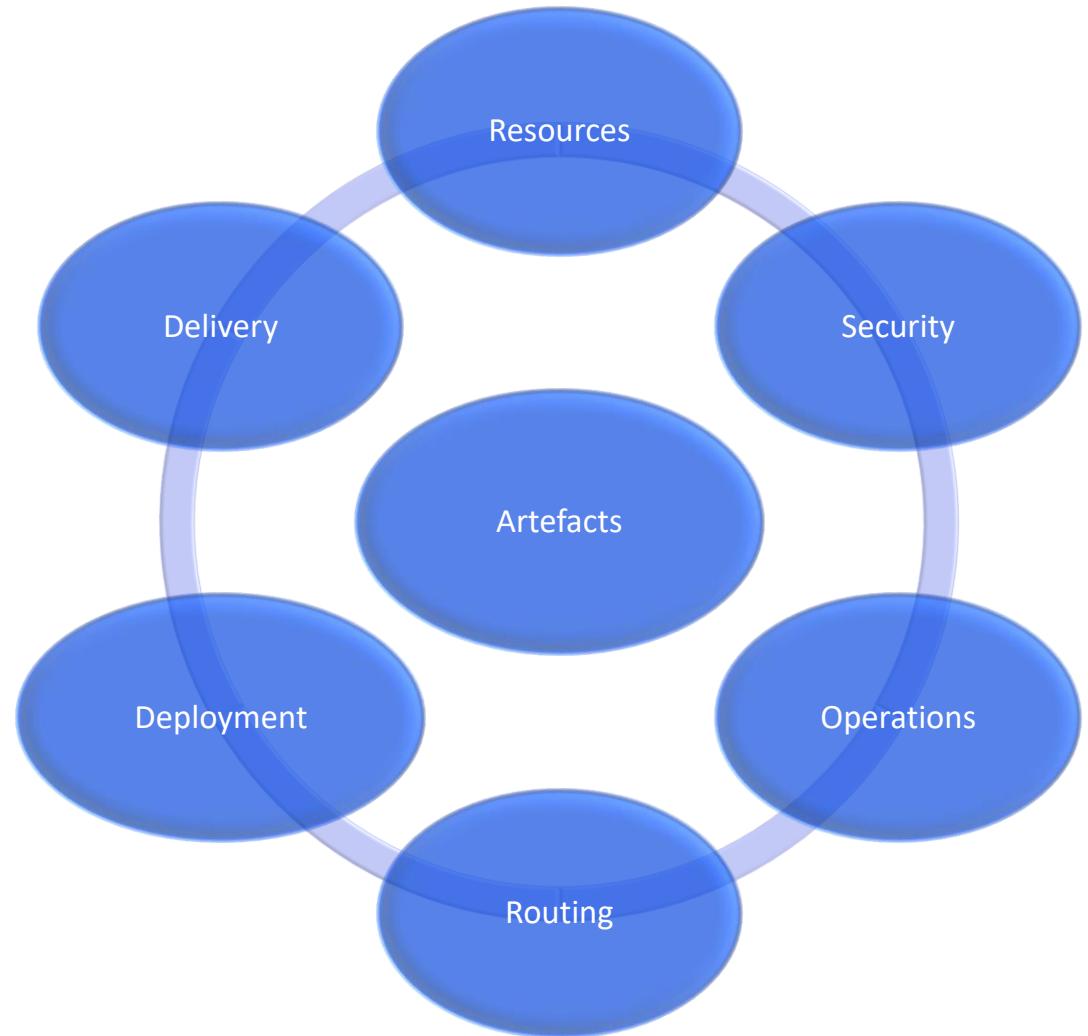


Ops want consistency of infrastructure and operation

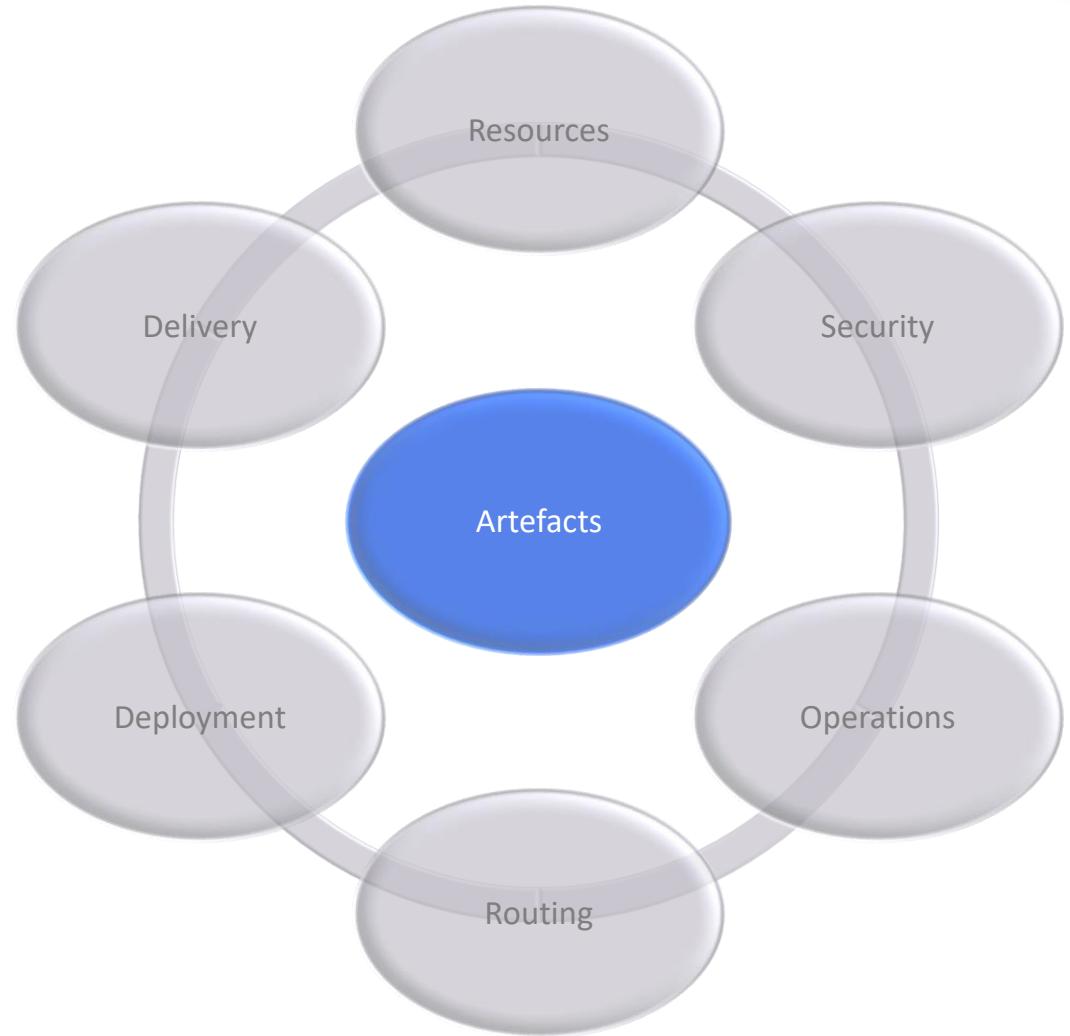
Runtime specific



Provided by platform

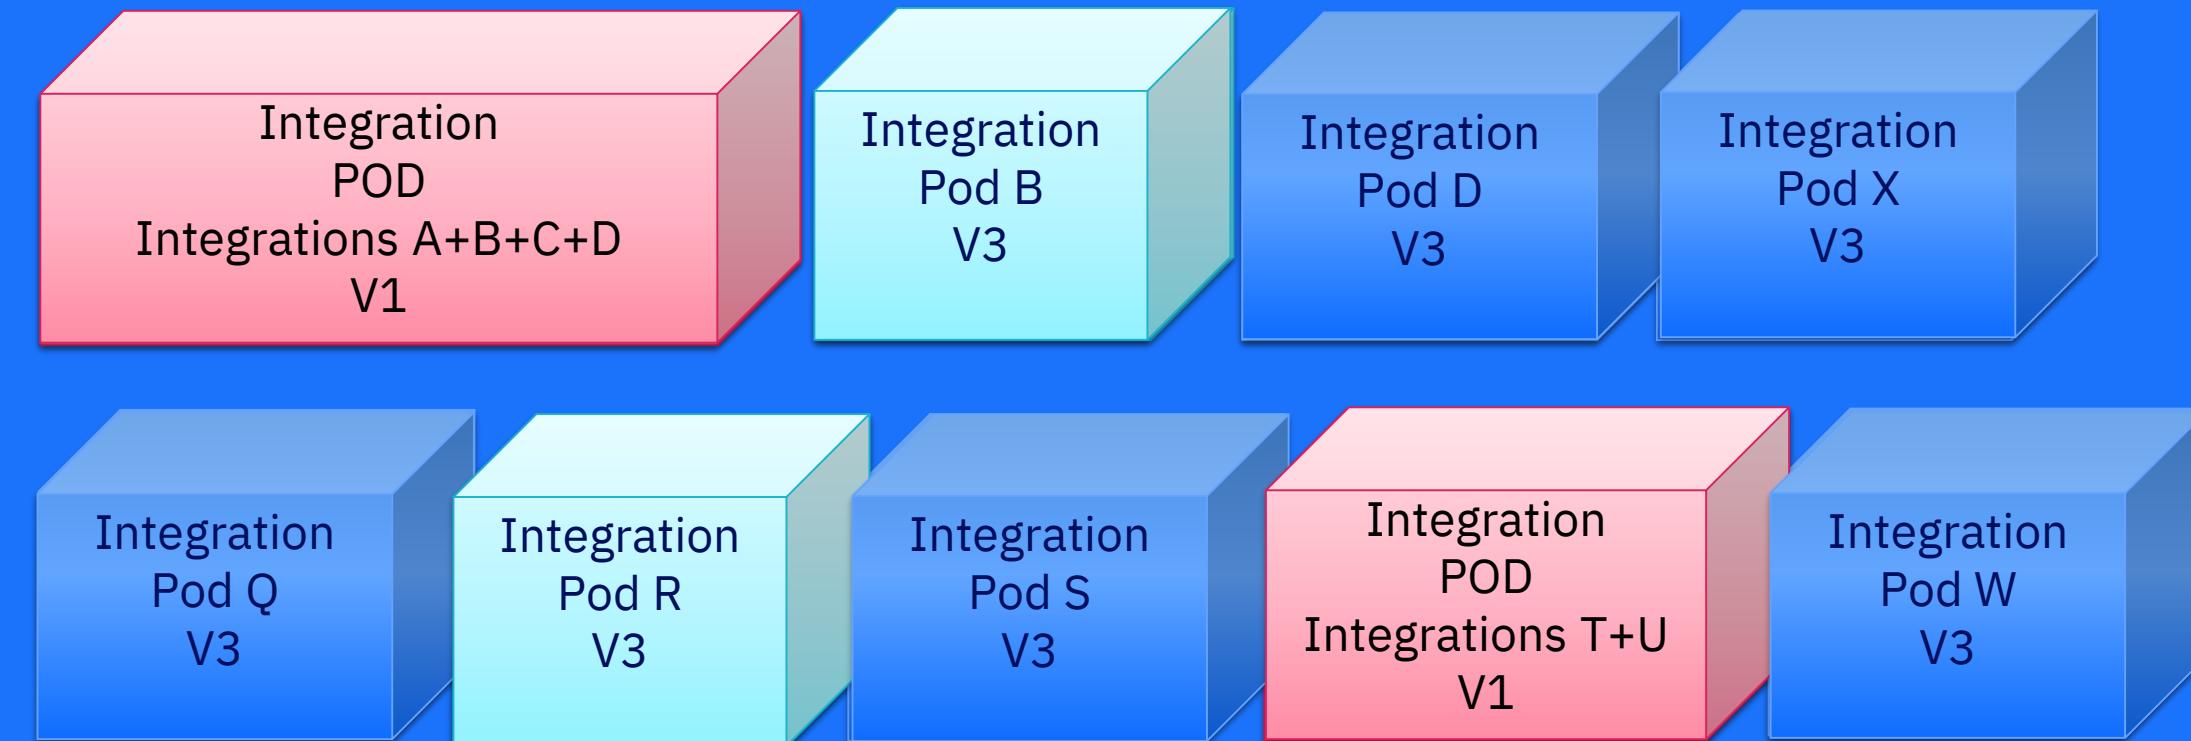


Traditional infrastructure
(Pets)



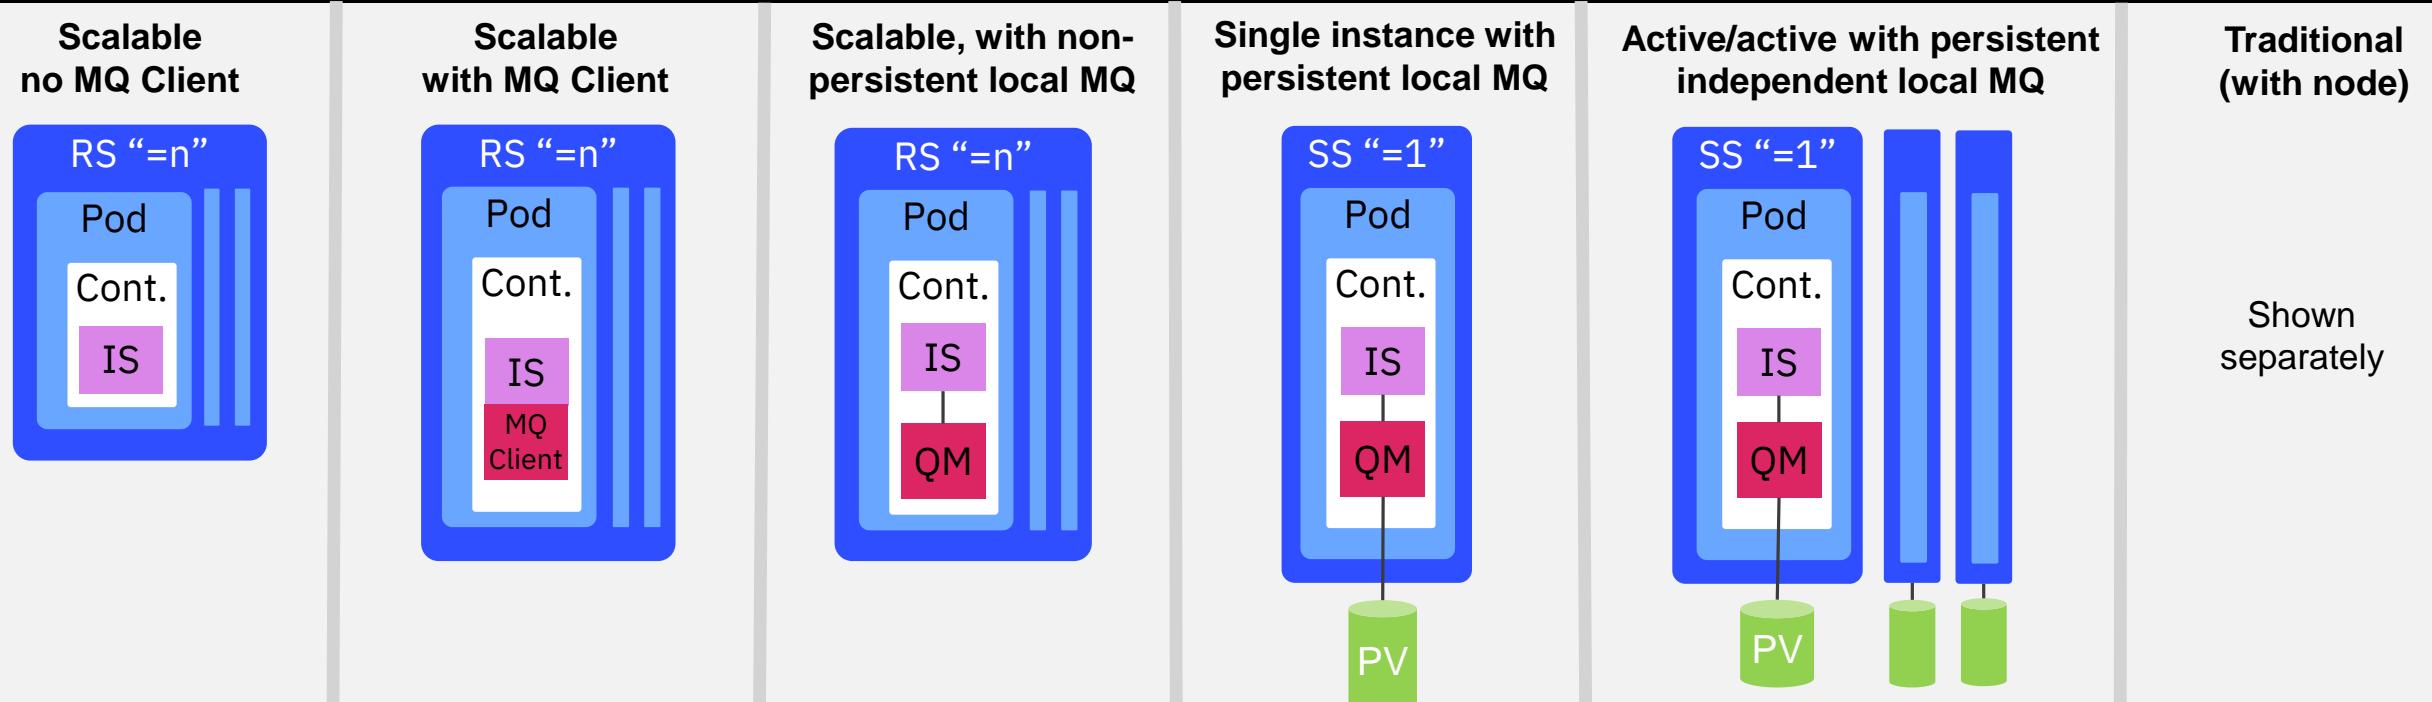
Cloud native infrastructure
(Livestock)

Why does OPS care what's in a pod? “Just another Container” or JAC surely?



OPS still need to know the PATTERN: State can complicate things a lot – know where it is!

IBM



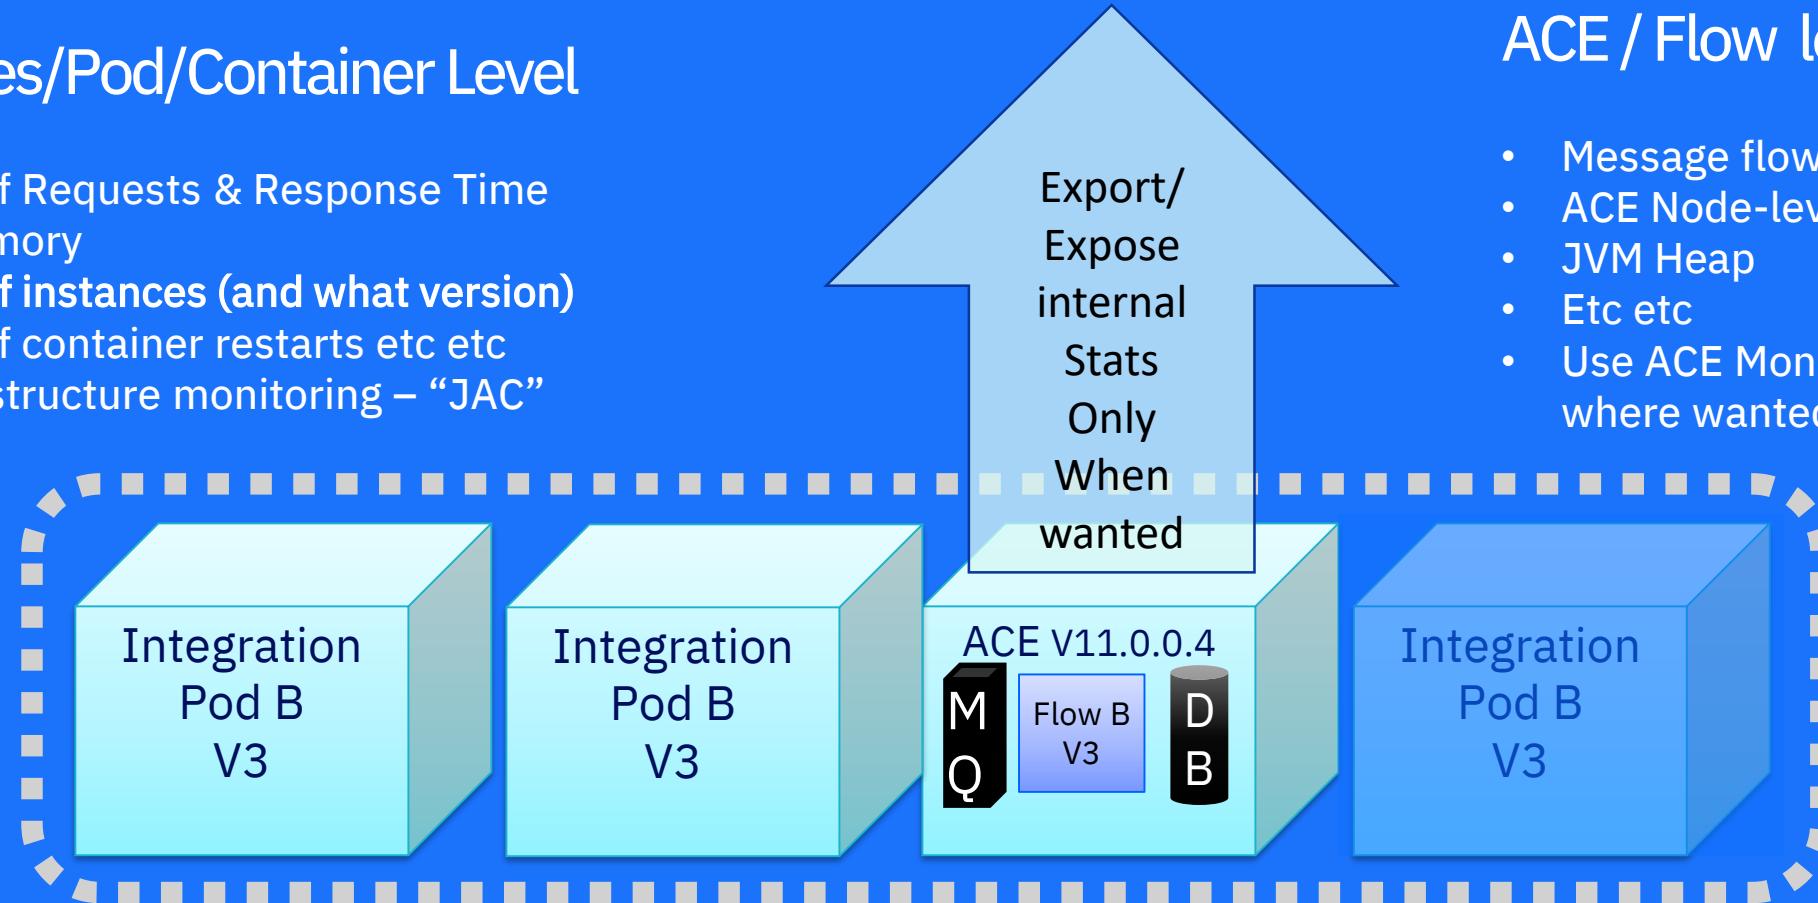
Stateless	Stateless	Stateless	Stateful	Stateful	Stateful
Continuous availability	Continuous availability	Continuous availability	High Availability	Continuous Service Availability High Message Availability	Continuous Availability
Dynamic horizontal scaling	Dynamic horizontal scaling	Dynamic horizontal scaling	Vertical scaling only	Non-dynamic horizontal scaling	Non-dynamic horizontal scaling
No EDA	No EDA	Enables EDA, but note that must be replica=1 for sequencing	Enables EDA and 2PC	Enables EDA and 2PC	Enables EDA and 2PC

Monitoring: Infrastructure vs Application



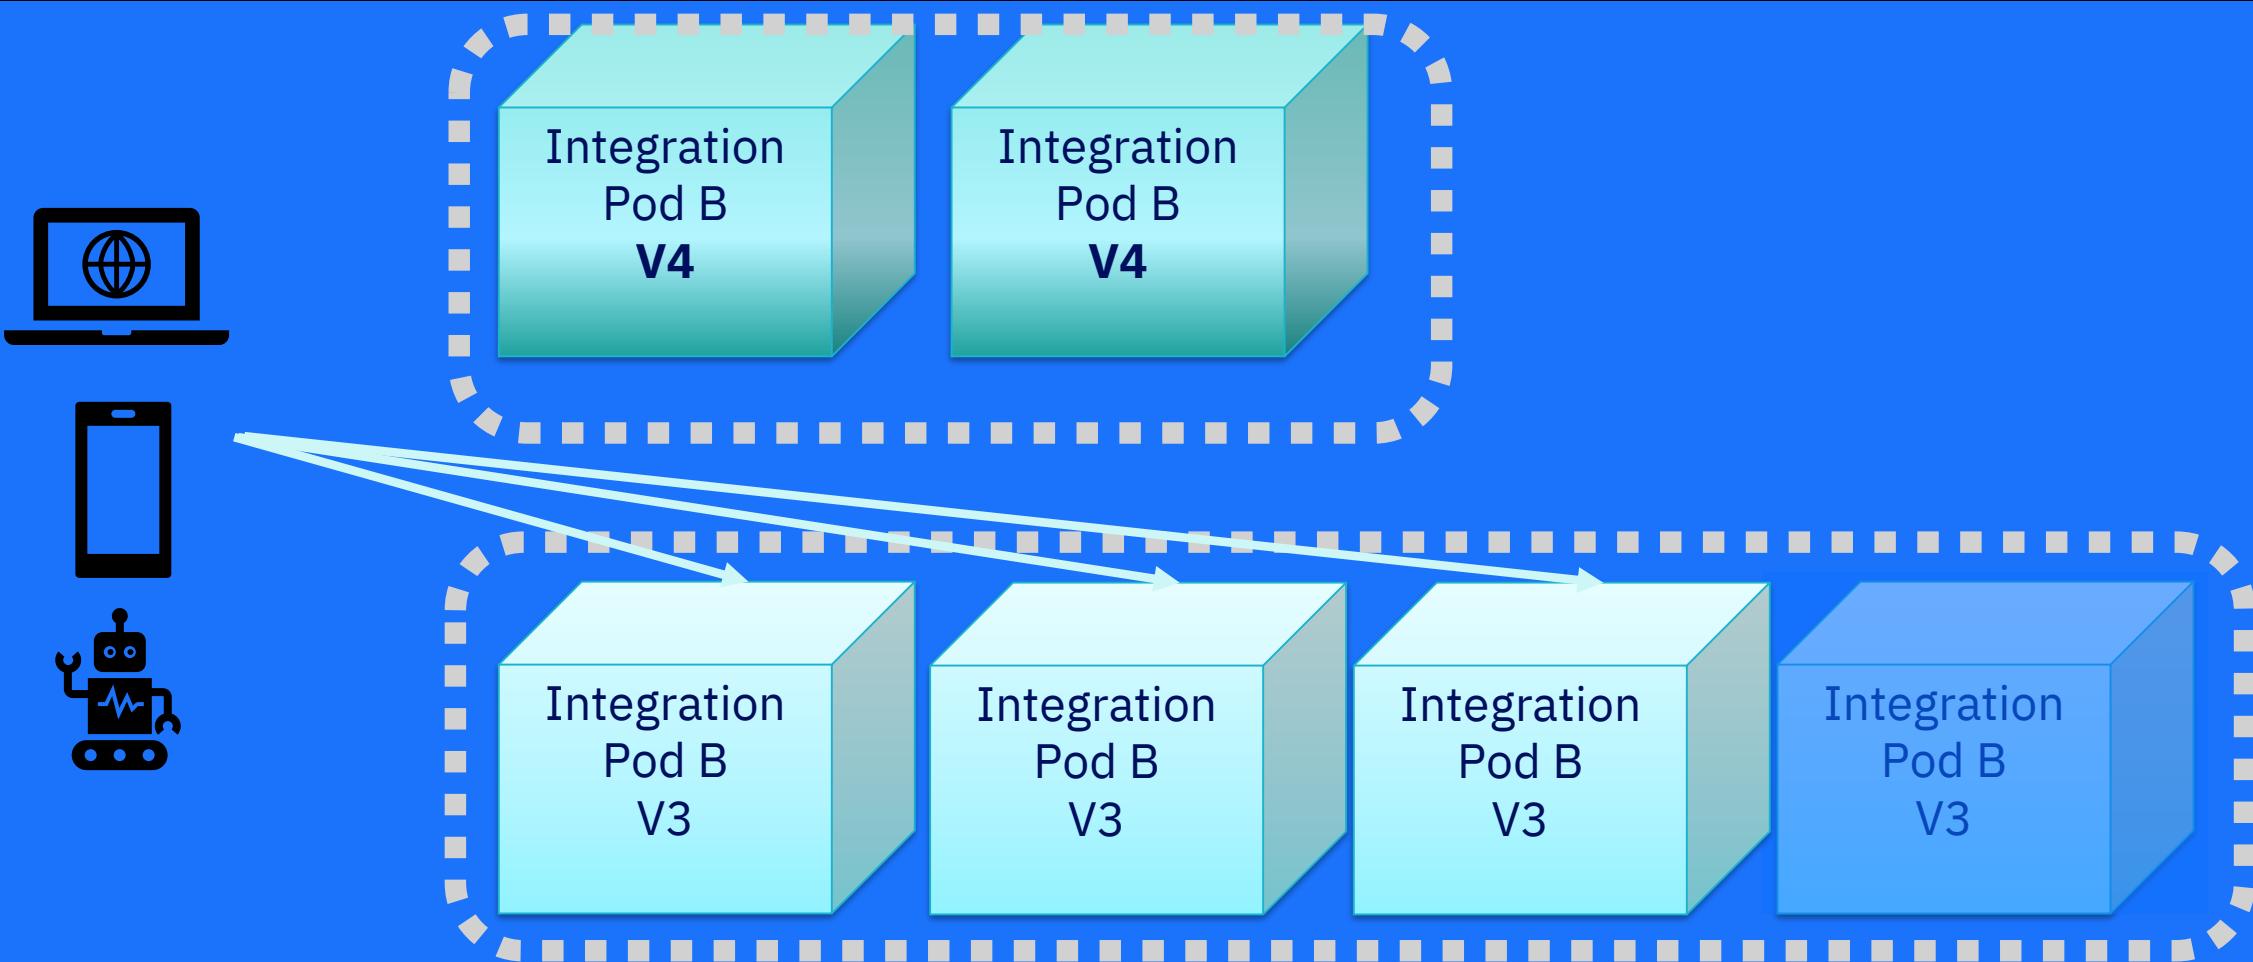
Kubernetes/Pod/Container Level

- Number of Requests & Response Time
- CPU / Memory
- **Number of instances (and what version)**
- Number of container restarts etc etc
- Use Infrastructure monitoring – “JAC”



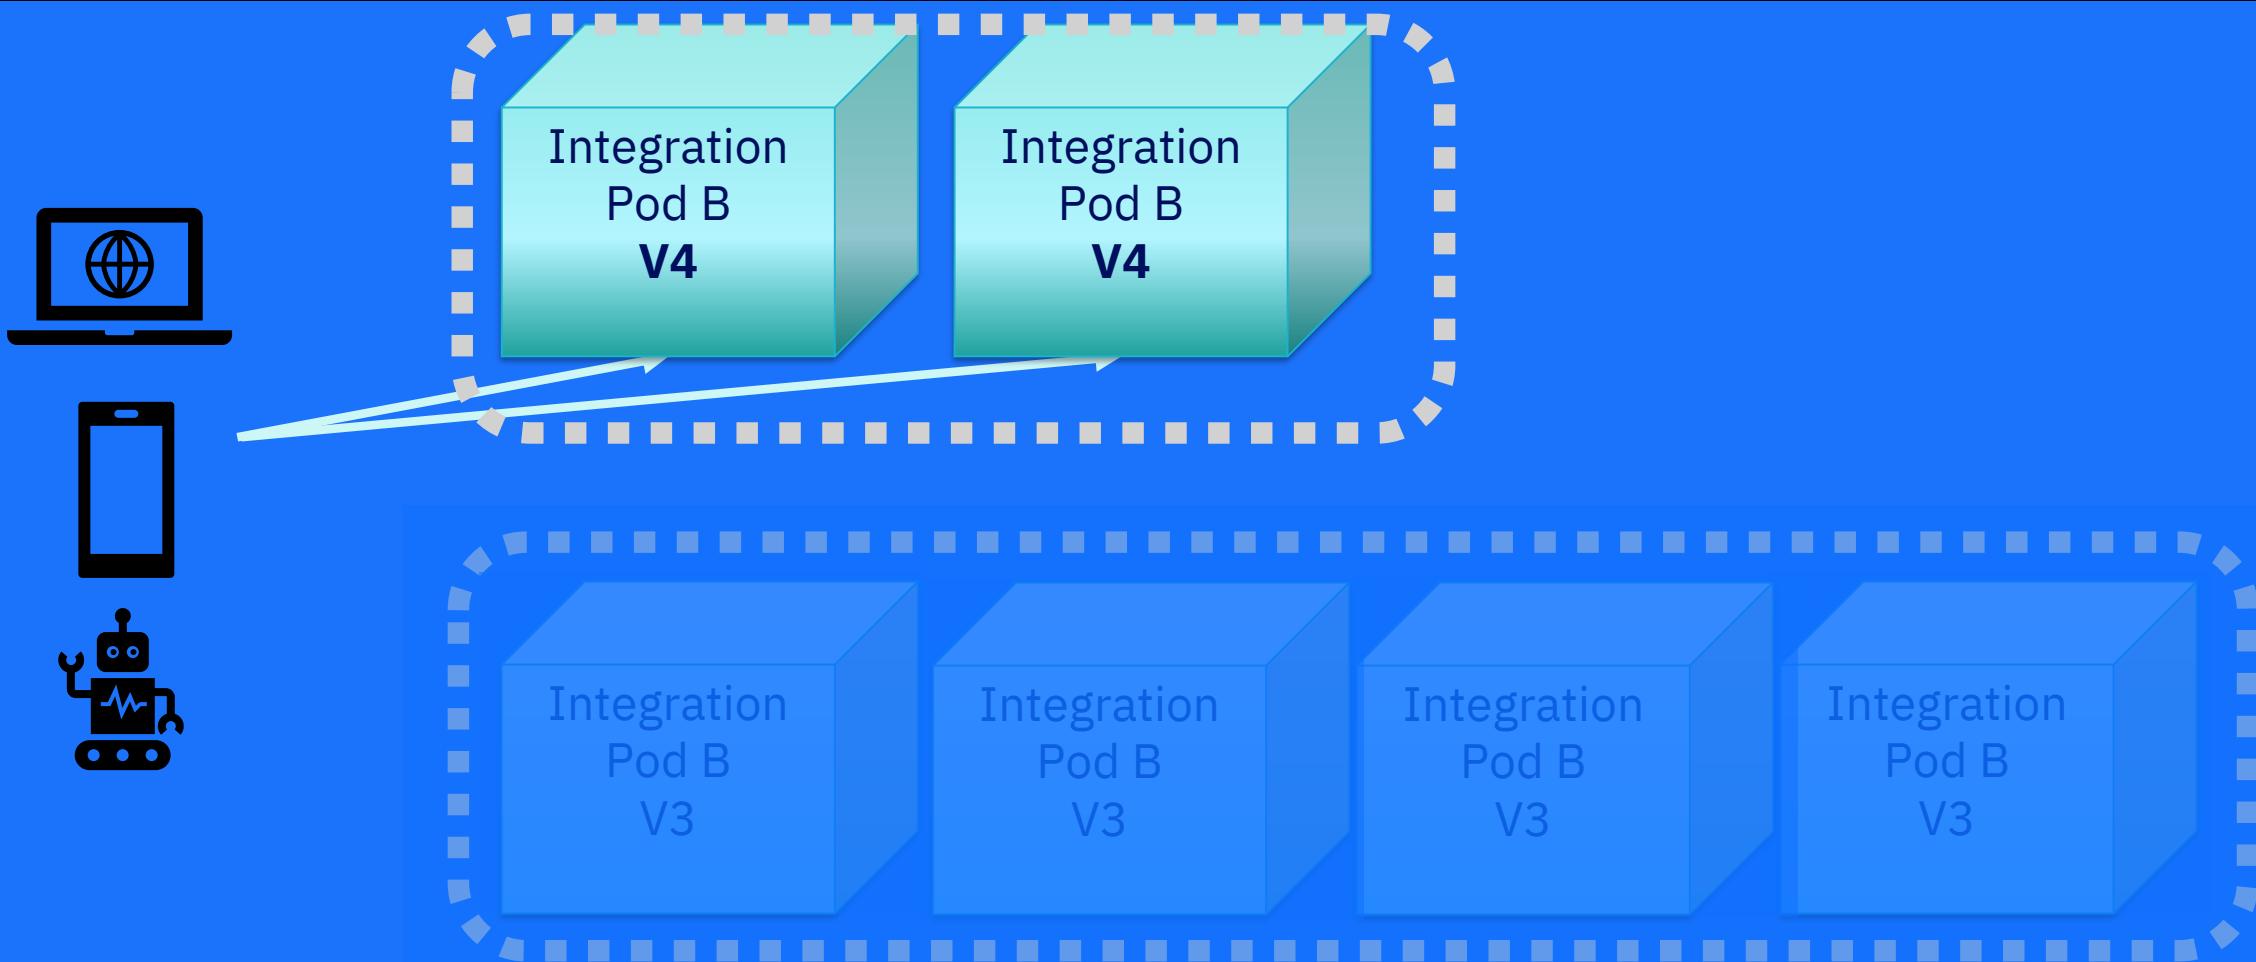
Versioning and upgrading...

IBM



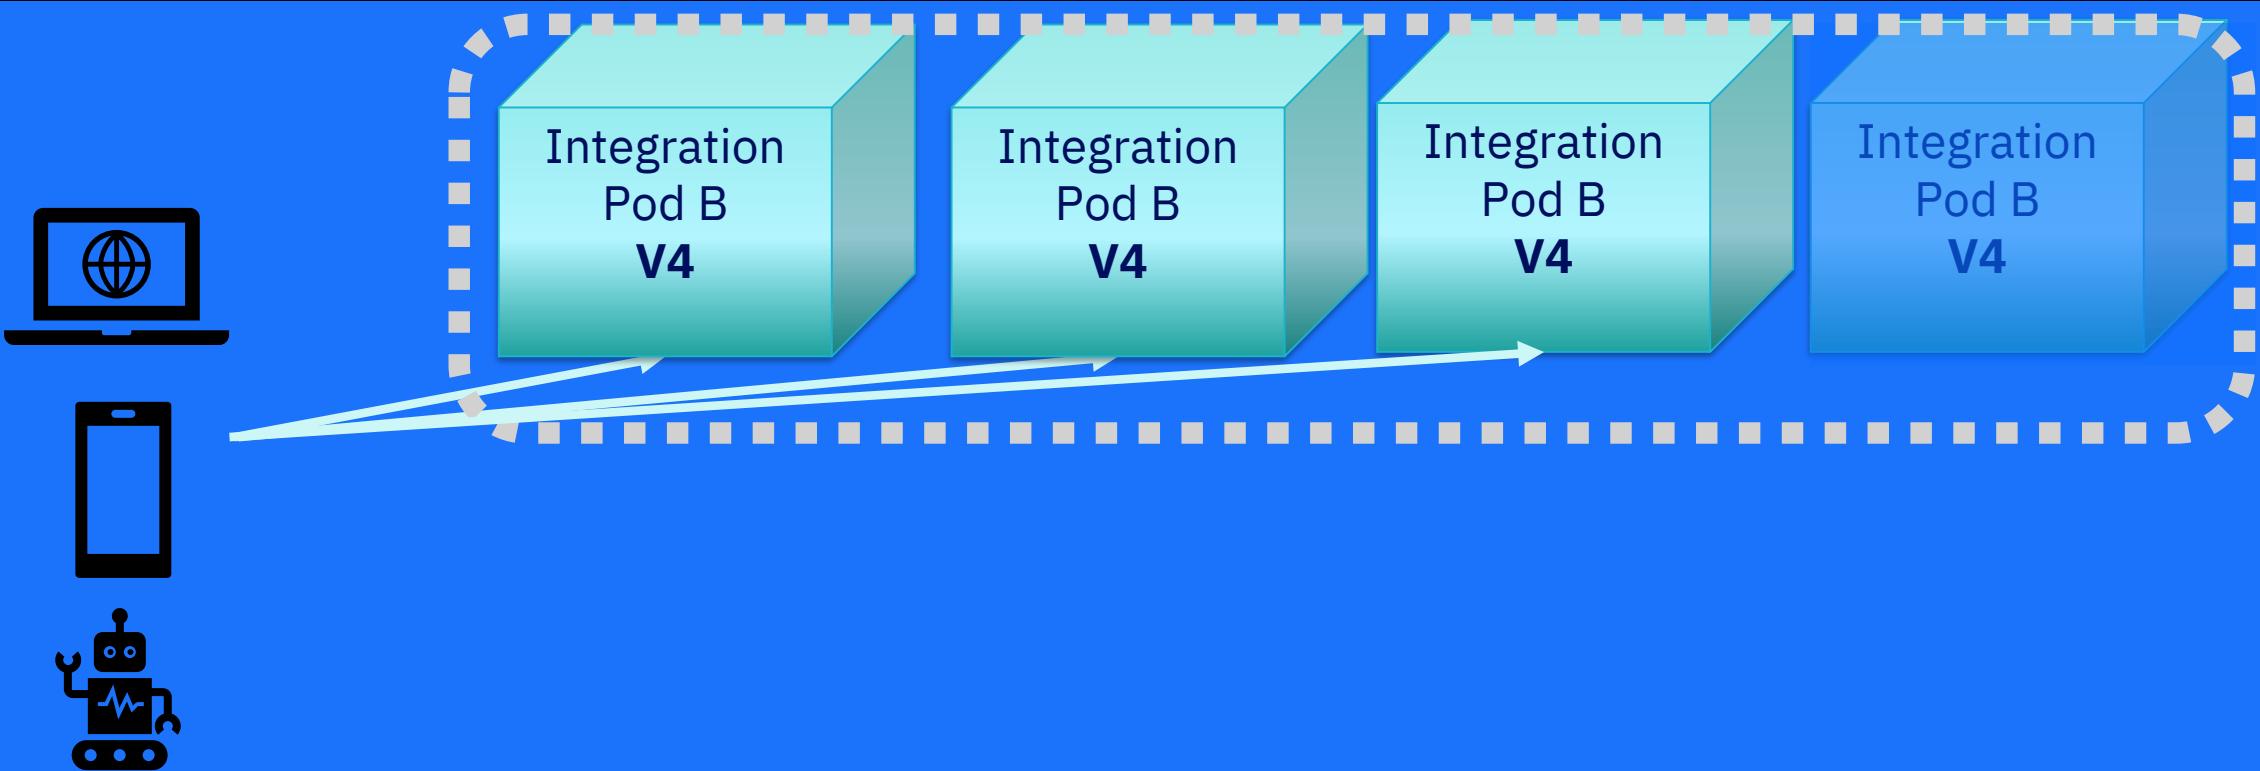
Enough pods? Reroute requests & drop old version

IBM®



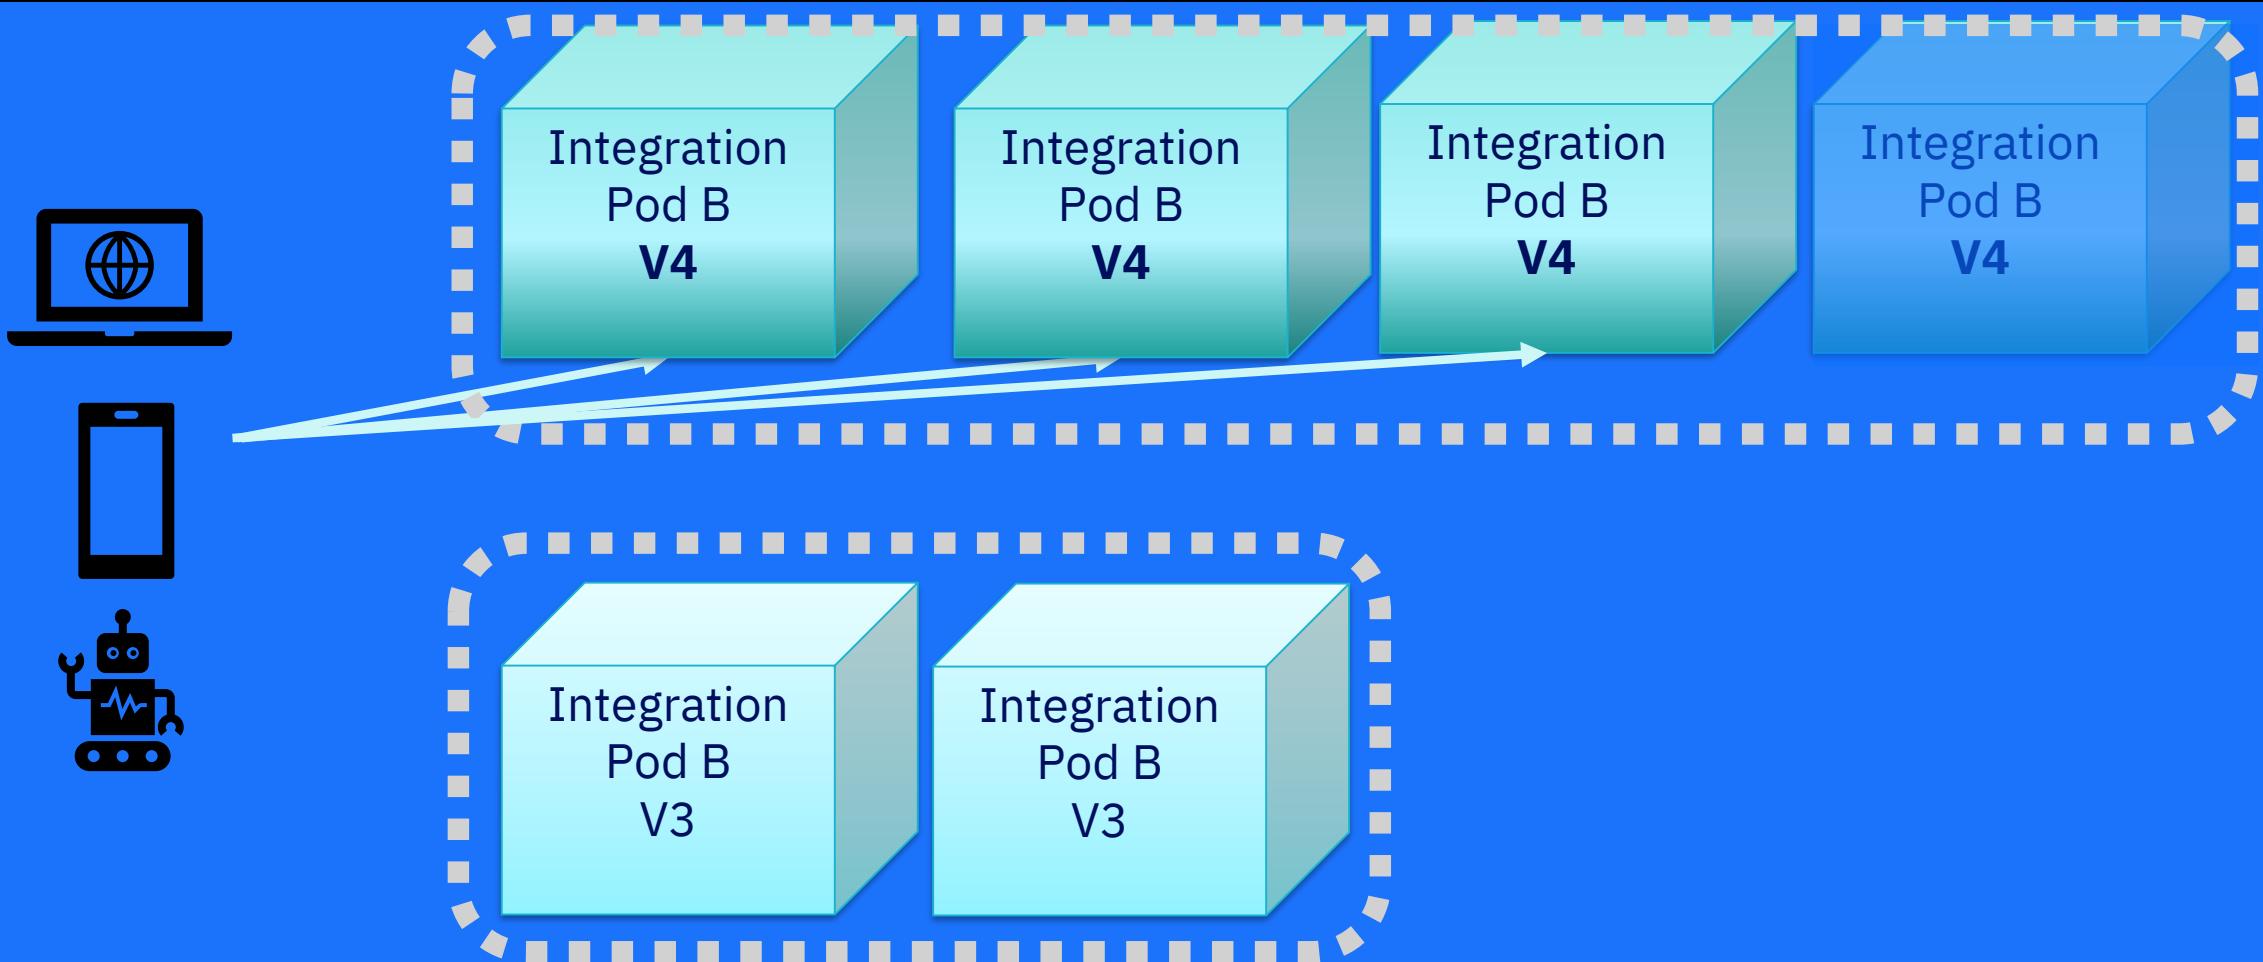
And finally fully roll out new version

IBM®



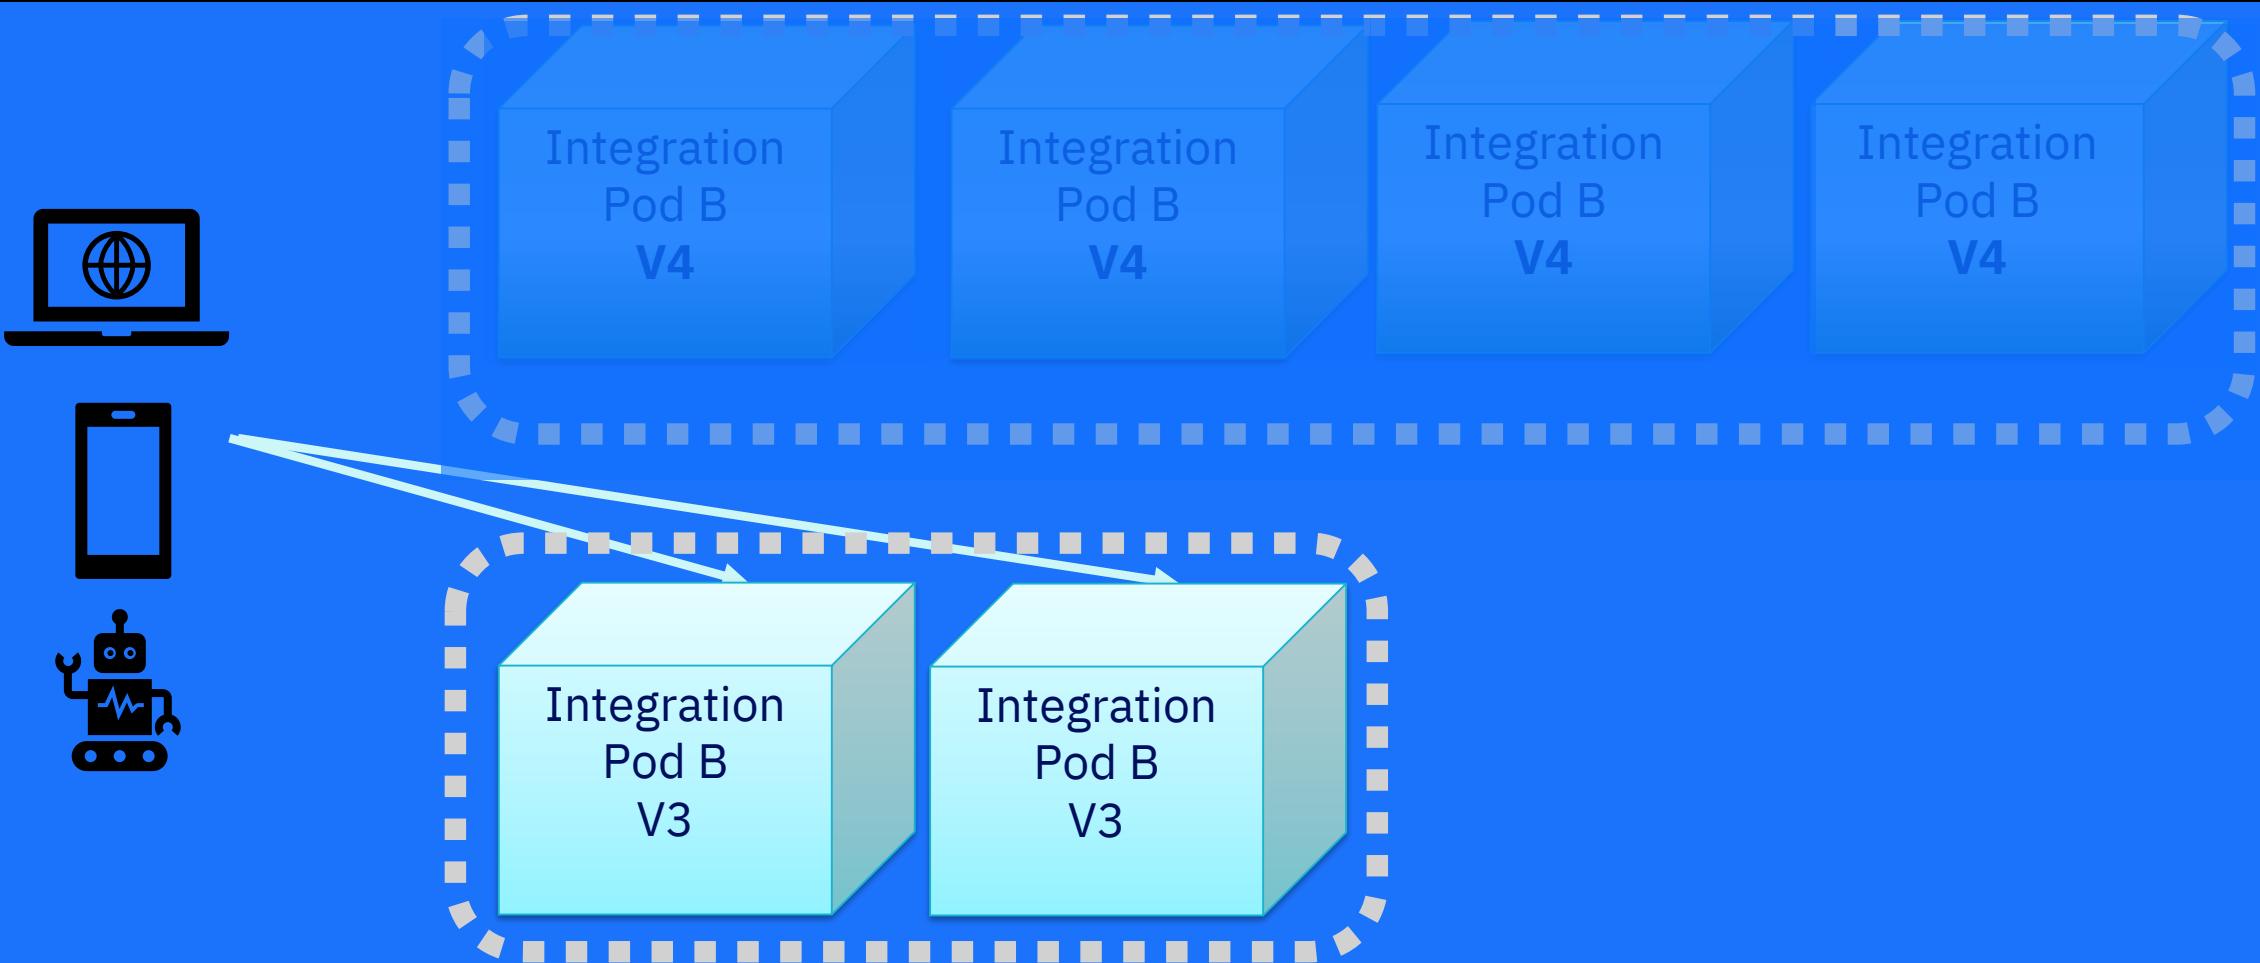
Rollback? Consider same ‘roll forward’ pattern

IBM

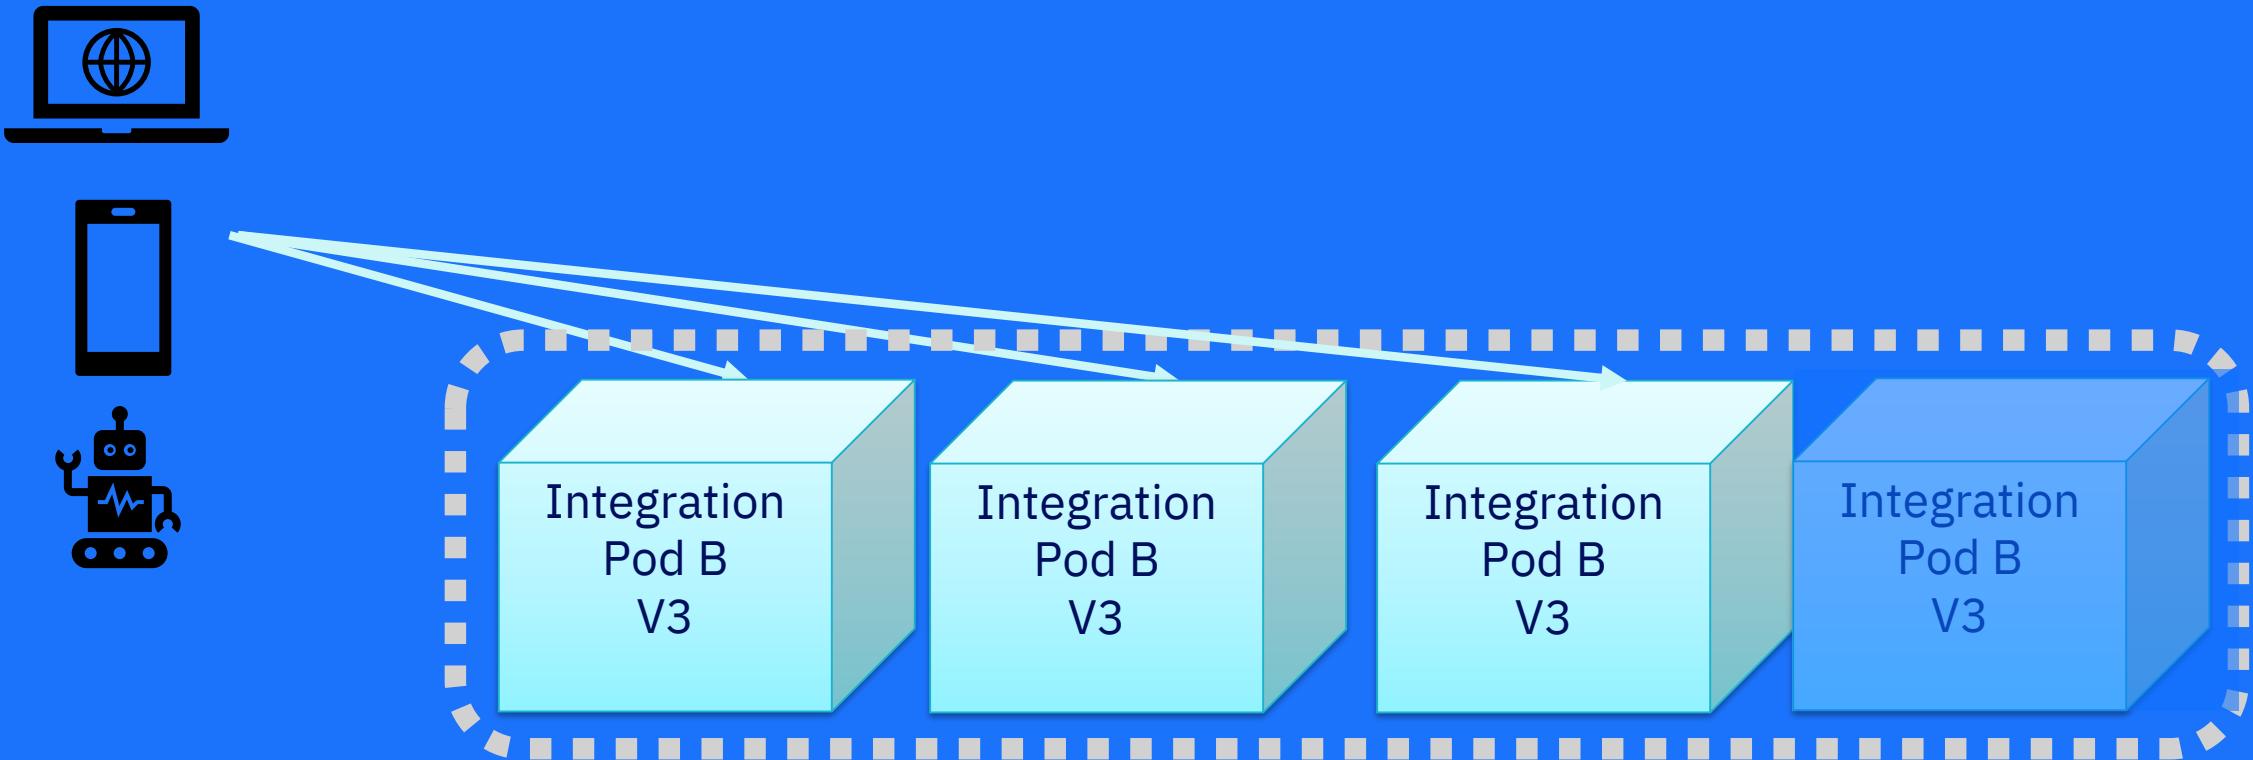


Redeploy immutable V3 containers/pods/chart And remove ‘bad’ version

IBM

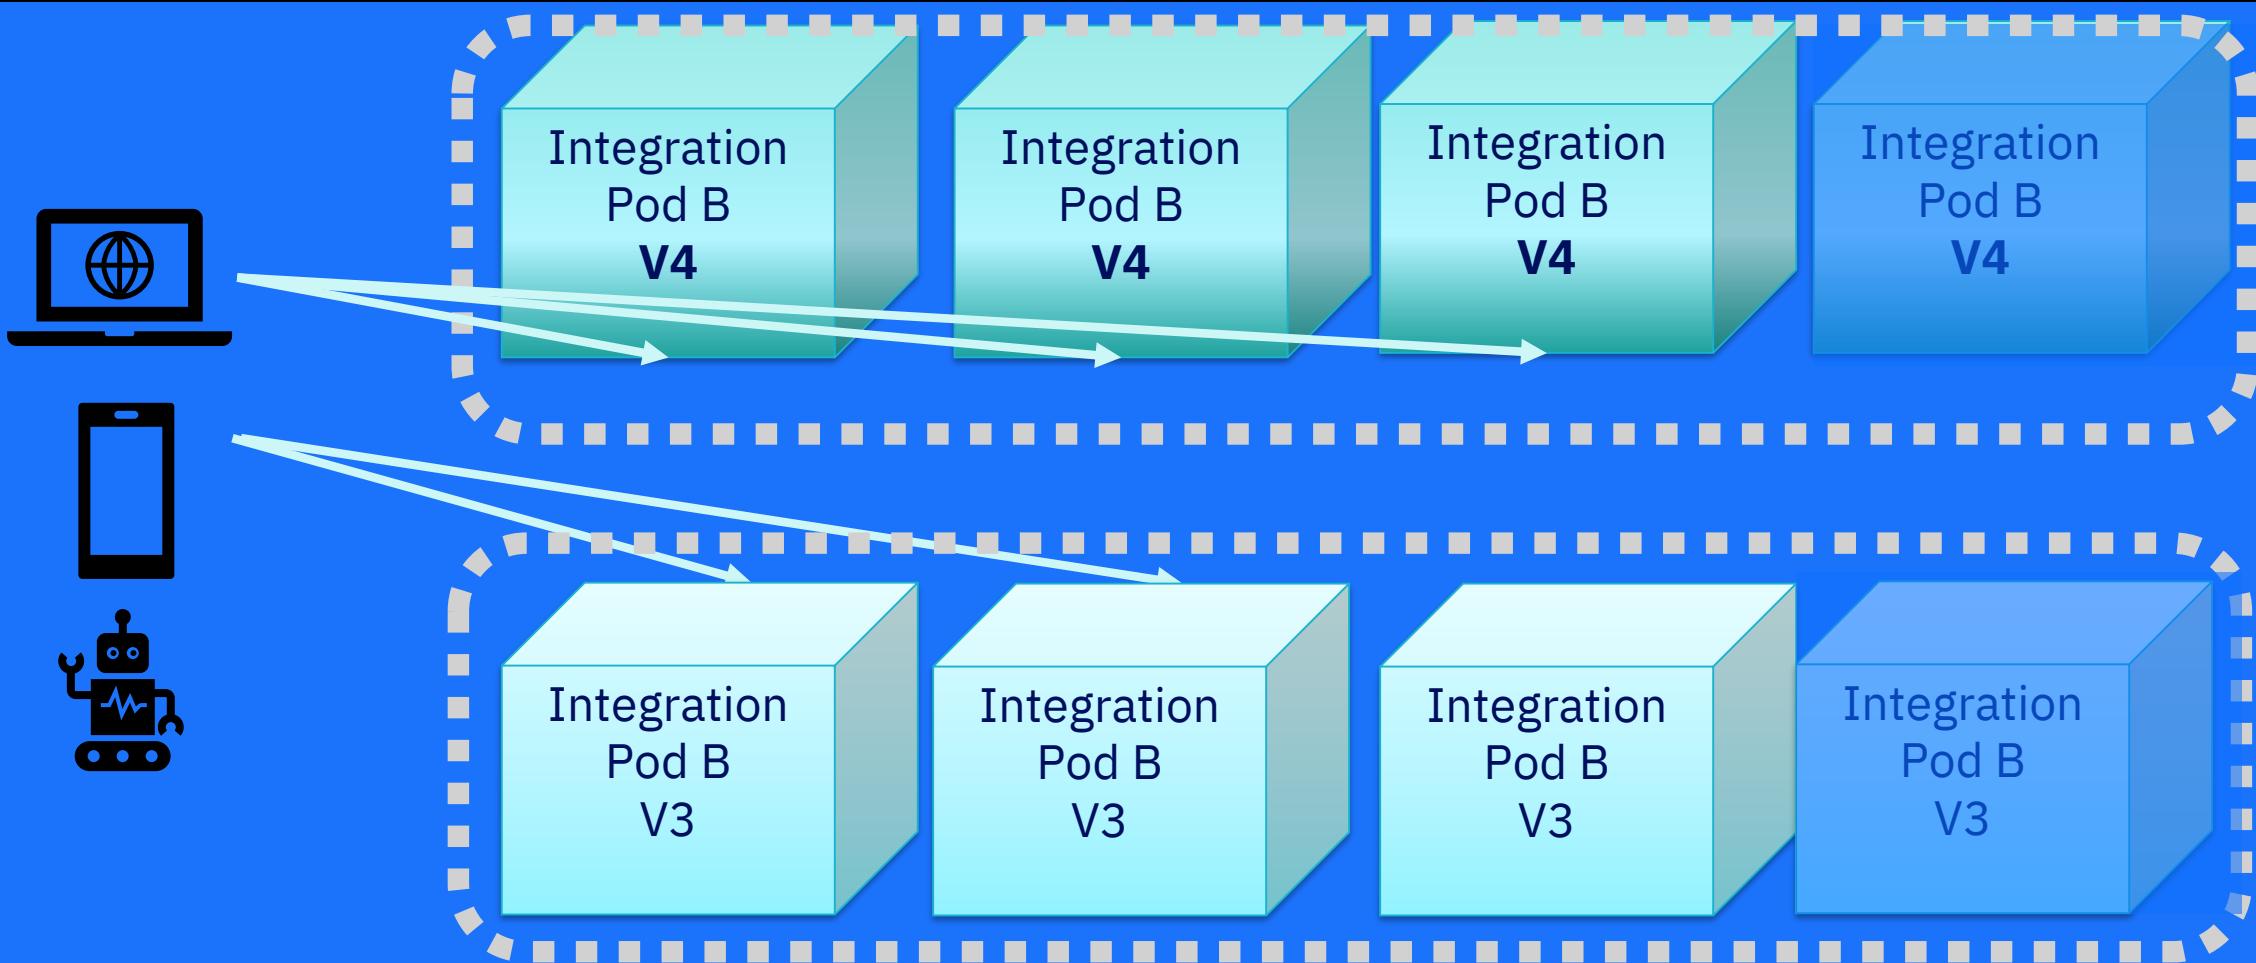


And finally we are back where we were before



Can we run both in parallel? Yes. How to route requests?

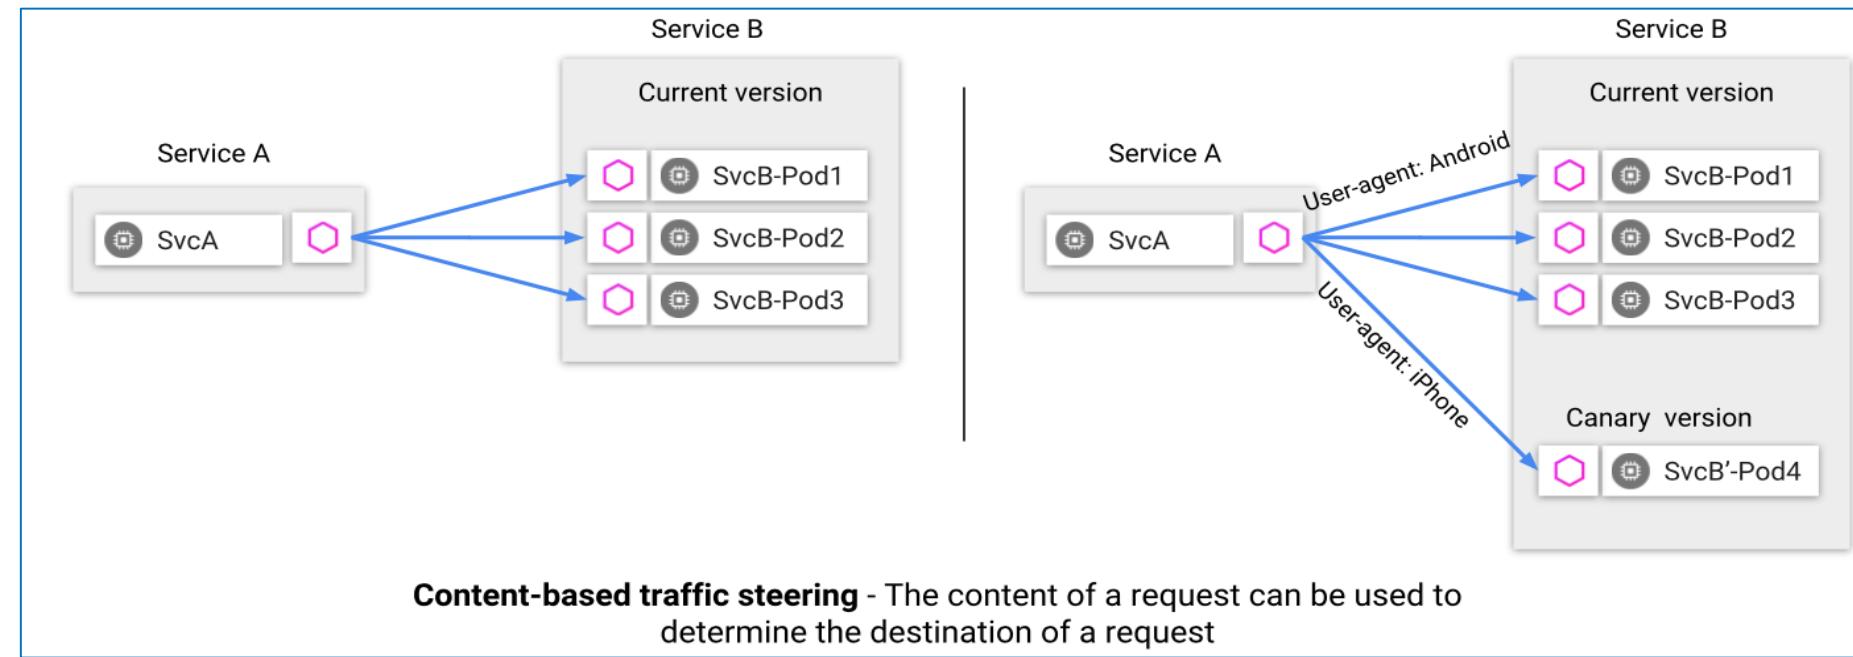
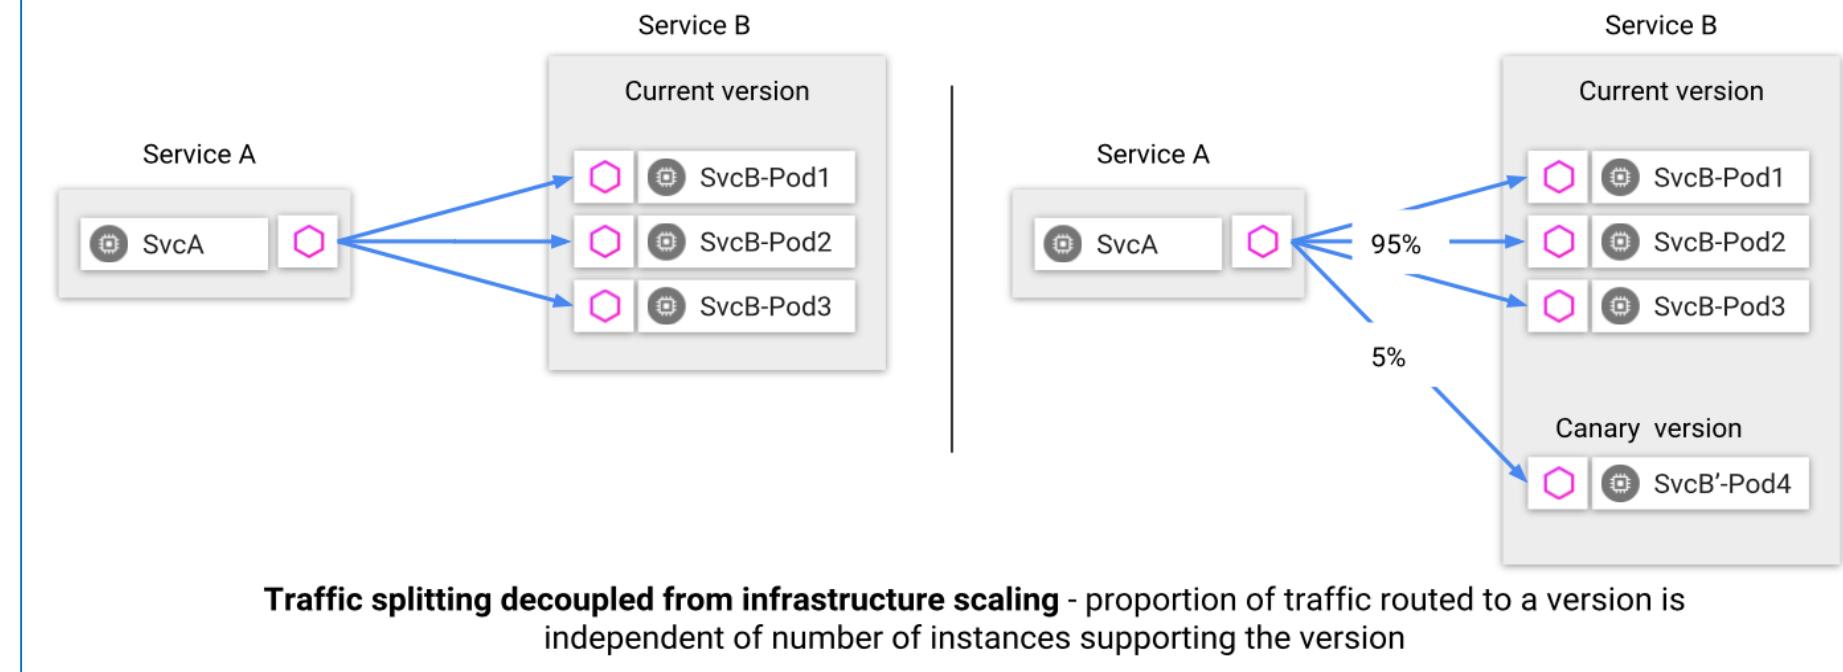
IBM



Use ISTIO Service Meshes and Envoy proxies to route traffic with rules.

Use your infrastructure where you can.

Don't 'Code' it in App Connect if you can delegate to the infrastructure.



Service Mesh Concepts

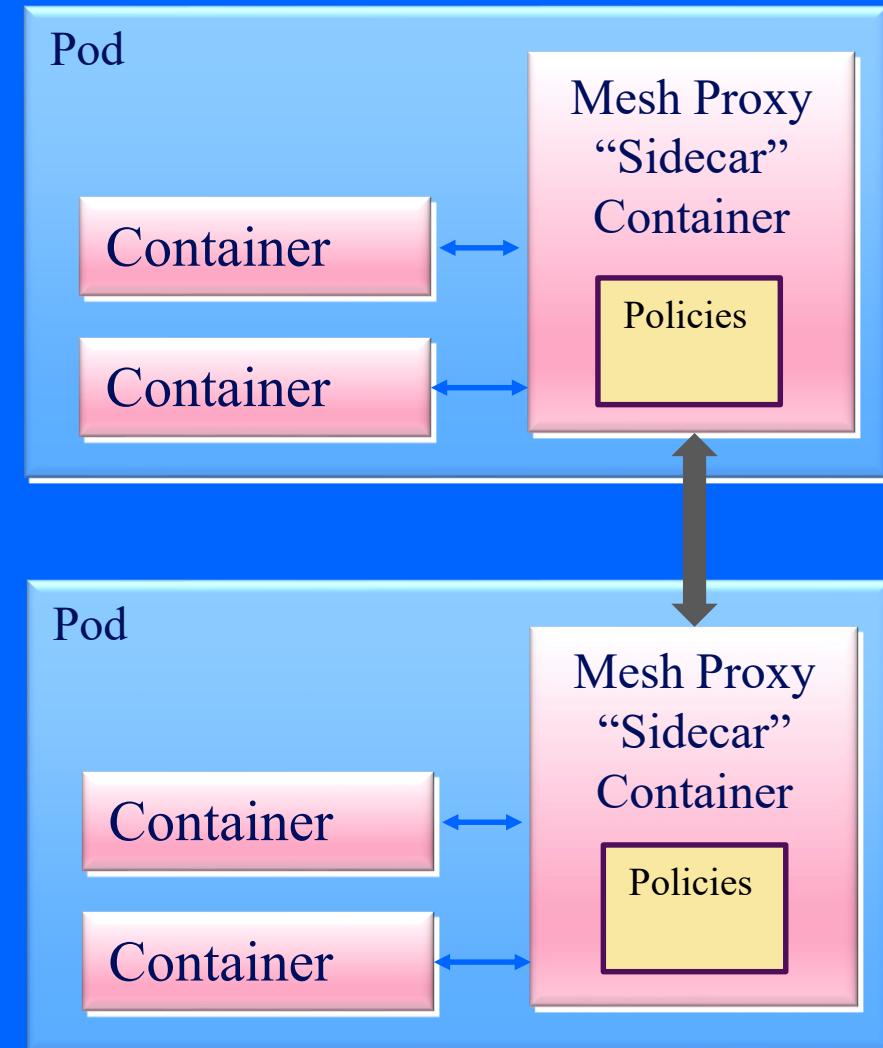
Service Mesh Proxy is deployed as a ‘sidecar’ container in the pod and implements a network proxy.

All traffic in and out MUST go through the proxy (enforced at ‘network’ level – the sidecar updates iptables on init)

One proxy per pod, so scalable

No code change to application /integration containers

You can put logging, MTLS, load balancing, routing etc etc in the mesh Proxy, not App Connect – keeps integration .bars simpler



Agile Integration Architecture

Modernizing integration
to enable
business agility

Fine grained deployment

Architecture & design

Decentralized Ownership

People & Process

Cloud native infrastructure

Infrastructure & Technology



Improve build independence and production velocity
(deployment agility)



Accelerate agility and innovation
(development agility)



Dynamic scalability and inherent resilience
(operational agility)

Key Takeaways

Know WHY you want to ‘Be Agile’

Know what outcomes and value you want Agile to give you.

Remove existing pain points

**Embrace the container concept:
Remove dependencies.
Bind in your flows and App
Connect versions.**

**Same image in Dev, Test, Prod:
Inject configuration**

Let your containers run anywhere

**Consider your granularity:
it will affect your agility.**

Define your deployment and test boundaries

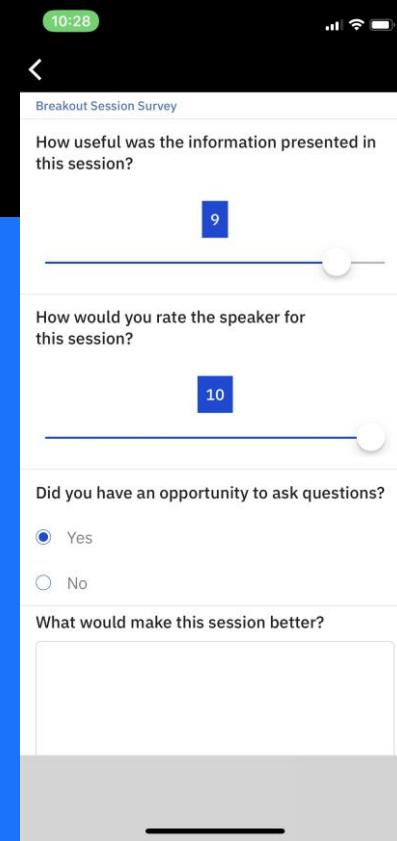
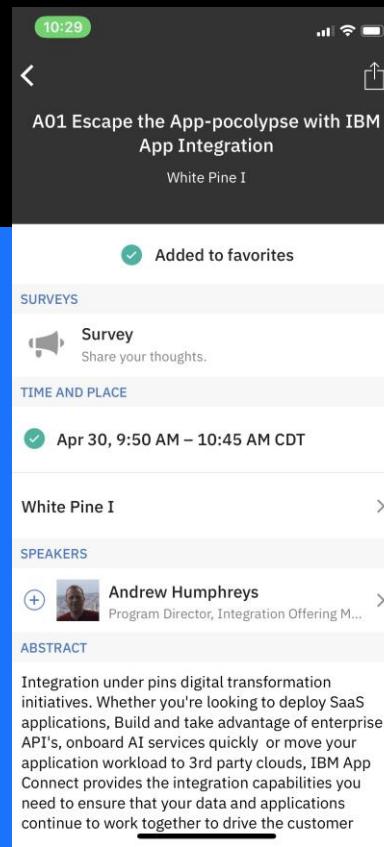
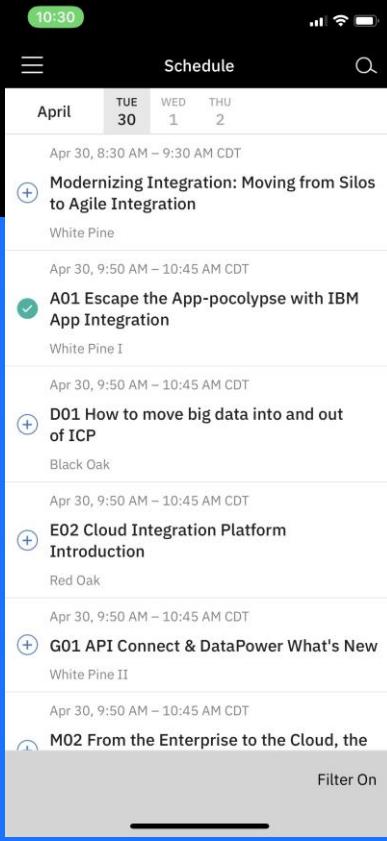
Find out what changes – and how often!

Delegate NFRs to your Agile Architecture

Use K8s and Service Meshes for endpoint routing & config

Use K8s for continuous availability

Use centralized logging: The number of containers will change dynamically.



IBM

Don't forget to fill out the survey!

Select your session, select survey, rate the session (A06) and submit!

Thank You

