

Integration Technical Conference 2019

# M15

## Event-Driven Enterprise using IBM Event Streams

Apache Kafka for the Enterprise

Andrew Schofield

Senior Technical Staff Member

IBM Messaging



IBM Cloud

IBM

# Please note



IBM's statements regarding its plans, directions, and intent are subject to change or withdrawal without notice and at IBM's sole discretion.

Information regarding potential future products is intended to outline our general product direction and it should not be relied on in making a purchasing decision.

The information mentioned regarding potential future products is not a commitment, promise, or legal obligation to deliver any material, code or functionality. Information about potential future products may not be incorporated into any contract.

The development, release, and timing of any future features or functionality described for our products remains at our sole discretion.

Performance is based on measurements and projections using standard IBM benchmarks in a controlled environment. The actual throughput or performance that any user will experience will vary depending upon many factors, including considerations such as the amount of multiprogramming in the user's job stream, the I/O configuration, the storage configuration, and the workload processed. Therefore, no assurance can be given that an individual user will achieve results similar to those stated here.

# What is Event Streaming?

# Event-Driven in Action

*Getting data to where it's needed, before it's needed*



**Respond to events before  
the moment passes**



**Responsive & personalized  
customer experiences**



**Bring real-time intelligence  
to your apps**

# Isn't This “Just Messaging”?

***Message Queuing & Event Streaming focus on different aspects of Messaging***

## Operations

---

Messages that represent some **current or future processing**. For instance: request and response messages.

## Message Queuing

## Events

---

Messages that represent the **state of the system**. For instance: logging, measurements and notifications.

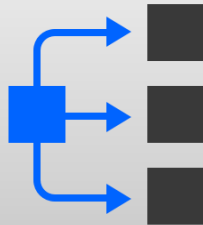
## Event Streaming

# Event Streaming & Message Queuing Need Different Capabilities

## Event Streaming



Stream history



Scalable  
consumption

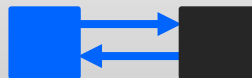


Immutable data

## Message Queuing



Transient data  
persistence



Conversational  
messaging



Assured delivery

# What is Apache Kafka?

# Apache Kafka is an **open-source, distributed streaming platform**



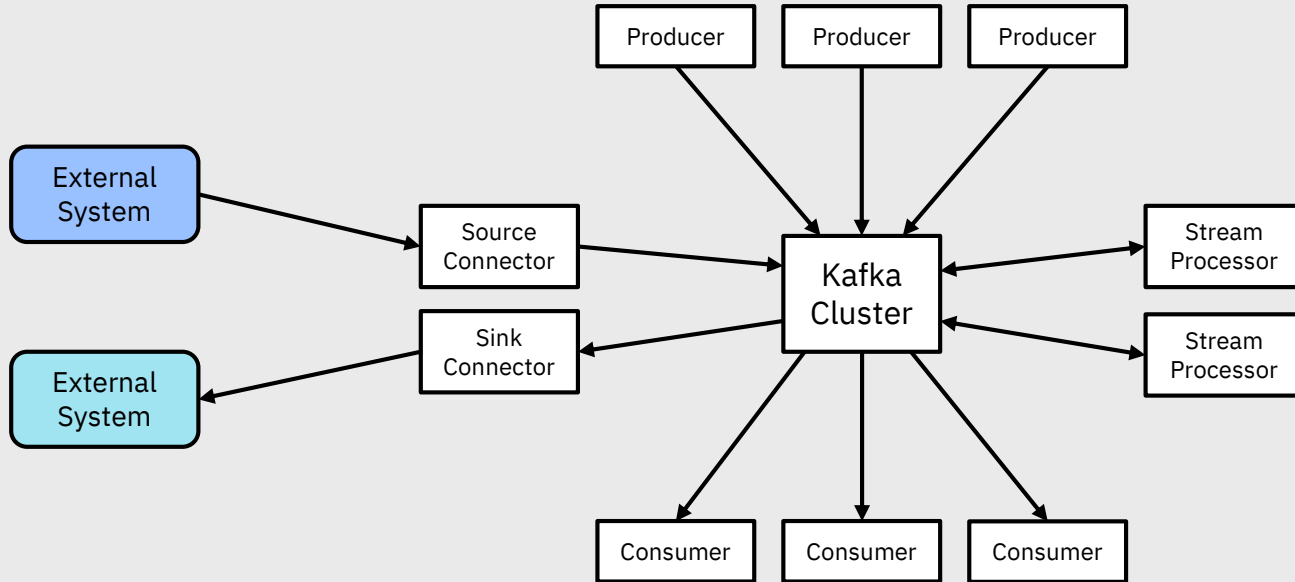
Publish and subscribe to streams of events

Store events in durable way

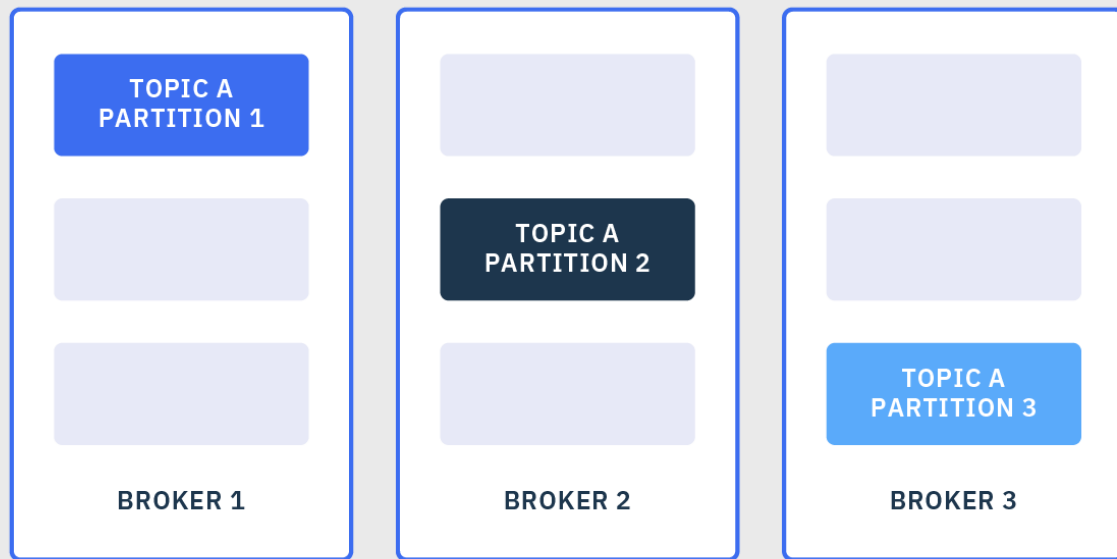
Process streams of events as they occur



# Apache Kafka is an Open-Source Streaming Platform

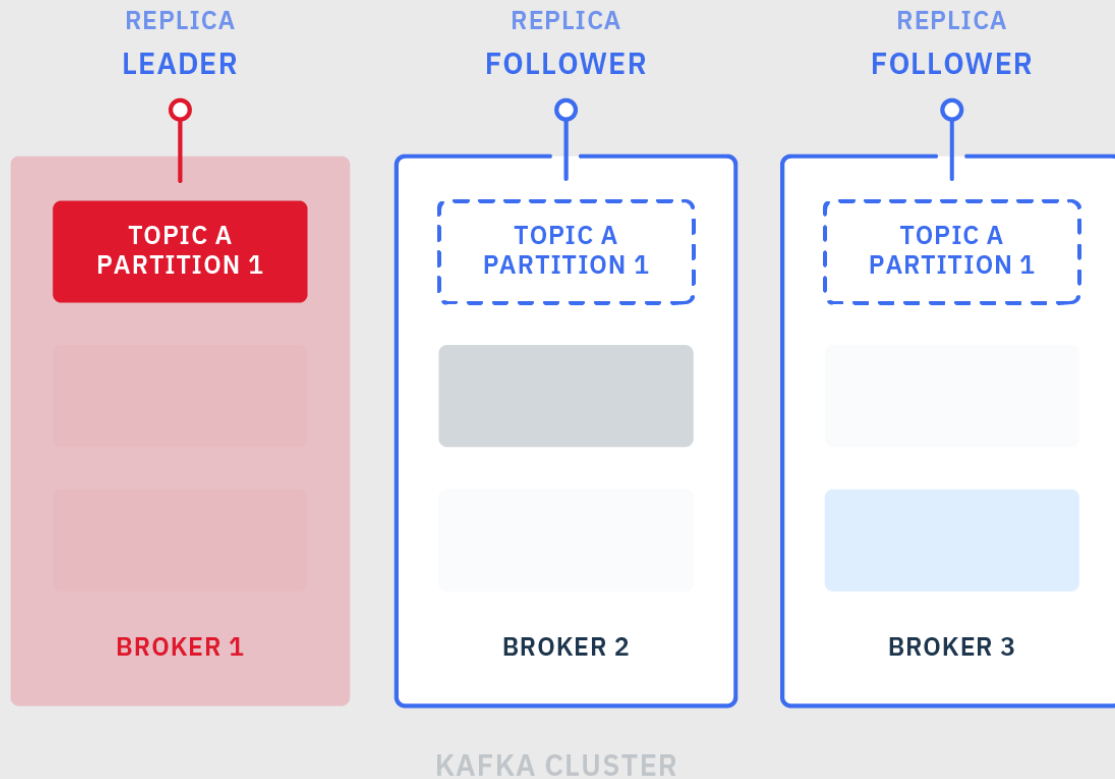


# Kafka cluster

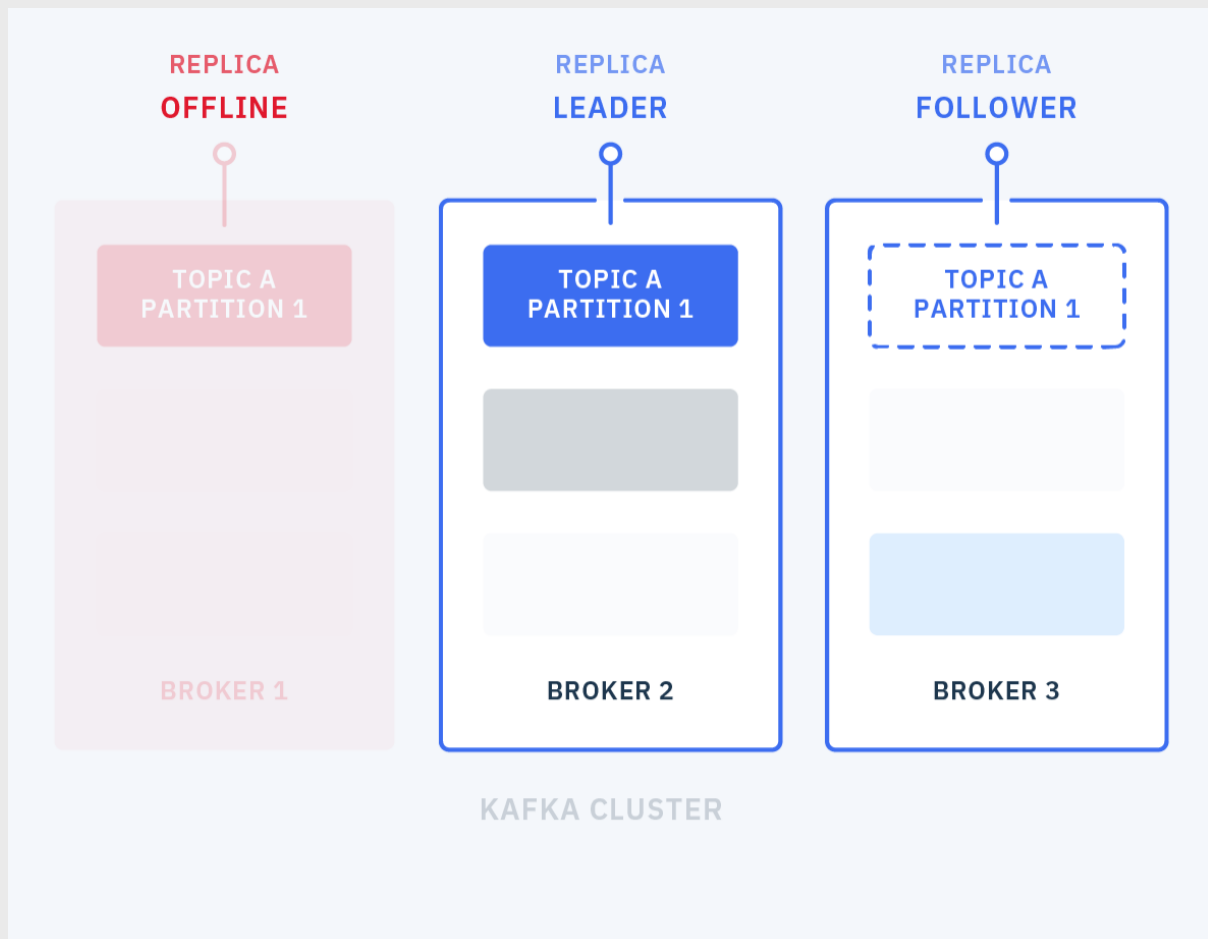


KAFKA CLUSTER

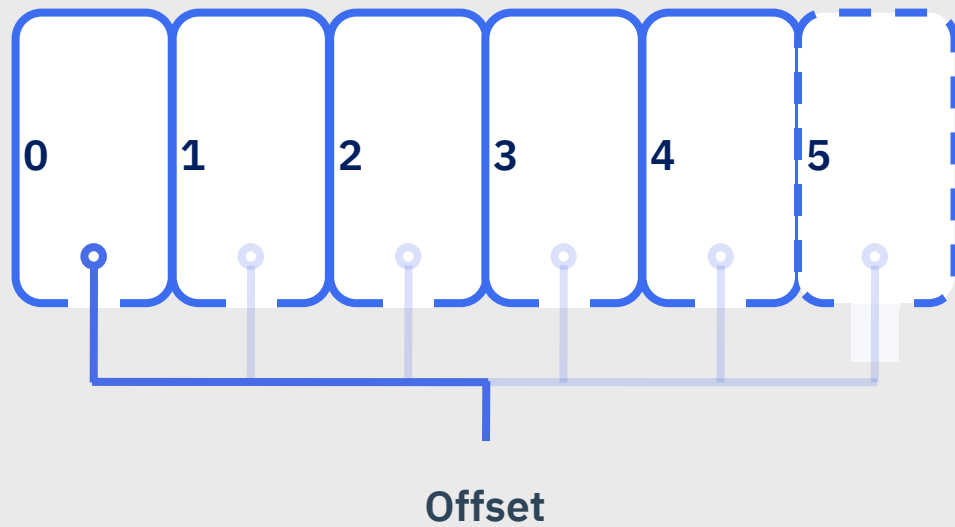
# Replication



# Replication



# Topics



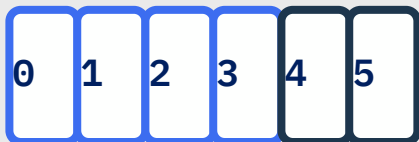
# Producers

TOPIC

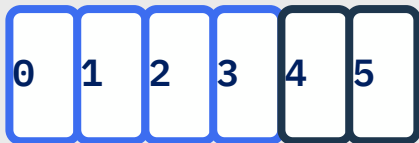
PARTITION 0



PARTITION 1



PARTITION 2



**Producer can choose acknowledgement level:**

0

Fire-and-forget  
*Fast, but risky*

1

Waits for 1 broker to acknowledge

ALL

Waits for all replica brokers to acknowledge

**Producer can choose whether to retry:**

0

Do not retry  
*Loses messages on error*

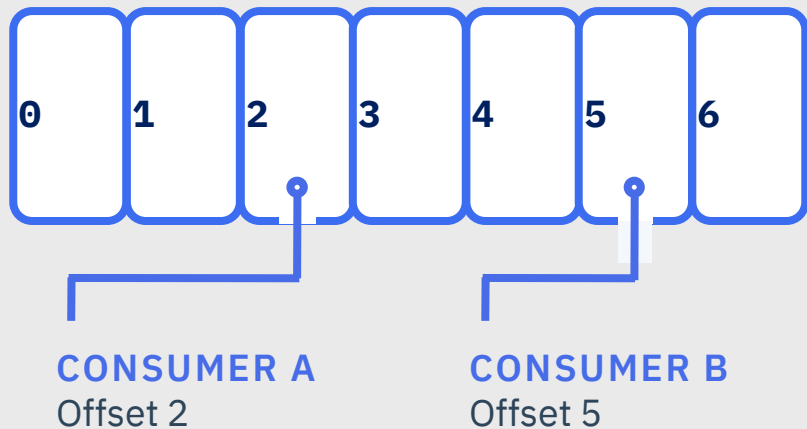
>0

Retry  
*Retry, might result in duplicates on error*

**Producer can also choose idempotence**

Can retry without risking duplicates

# Consumers



## Consumer can choose how to commit offsets:

Automatic

Commits might go faster than processing

Manual,  
asynchronous

Fairly safe, but could re-process messages

Manual,  
synchronous

Safe, but slows down processing

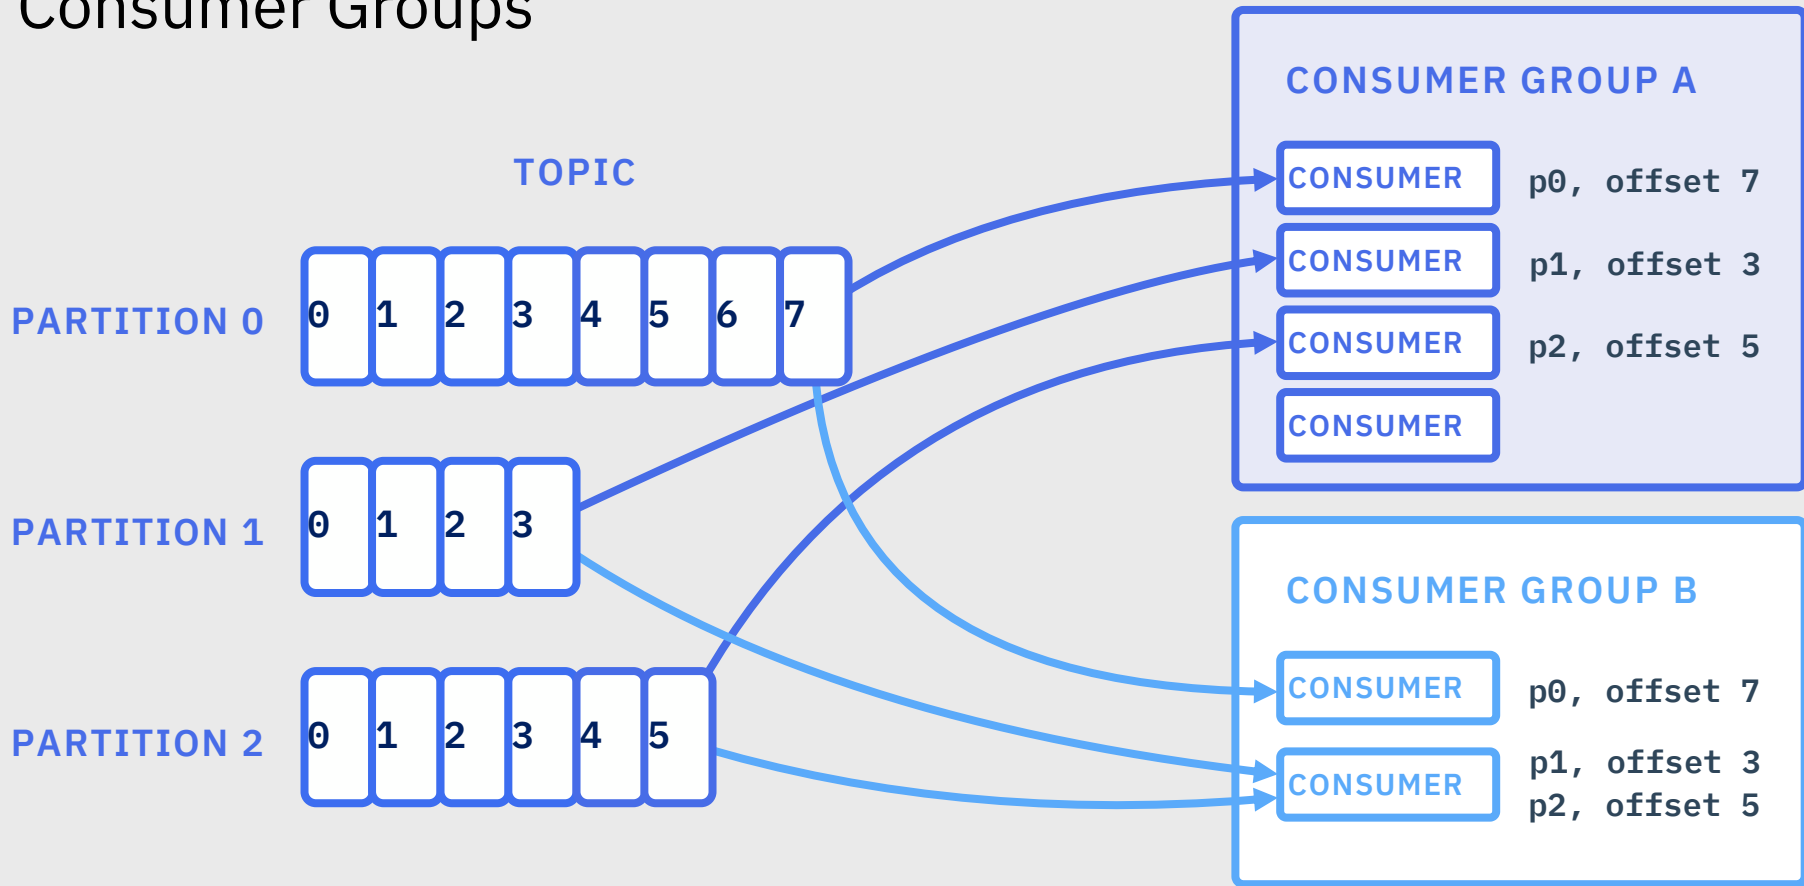
*A common pattern is to commit offsets on a timer*

## Exactly once semantics

Can group sending messages and committing offsets into transactions

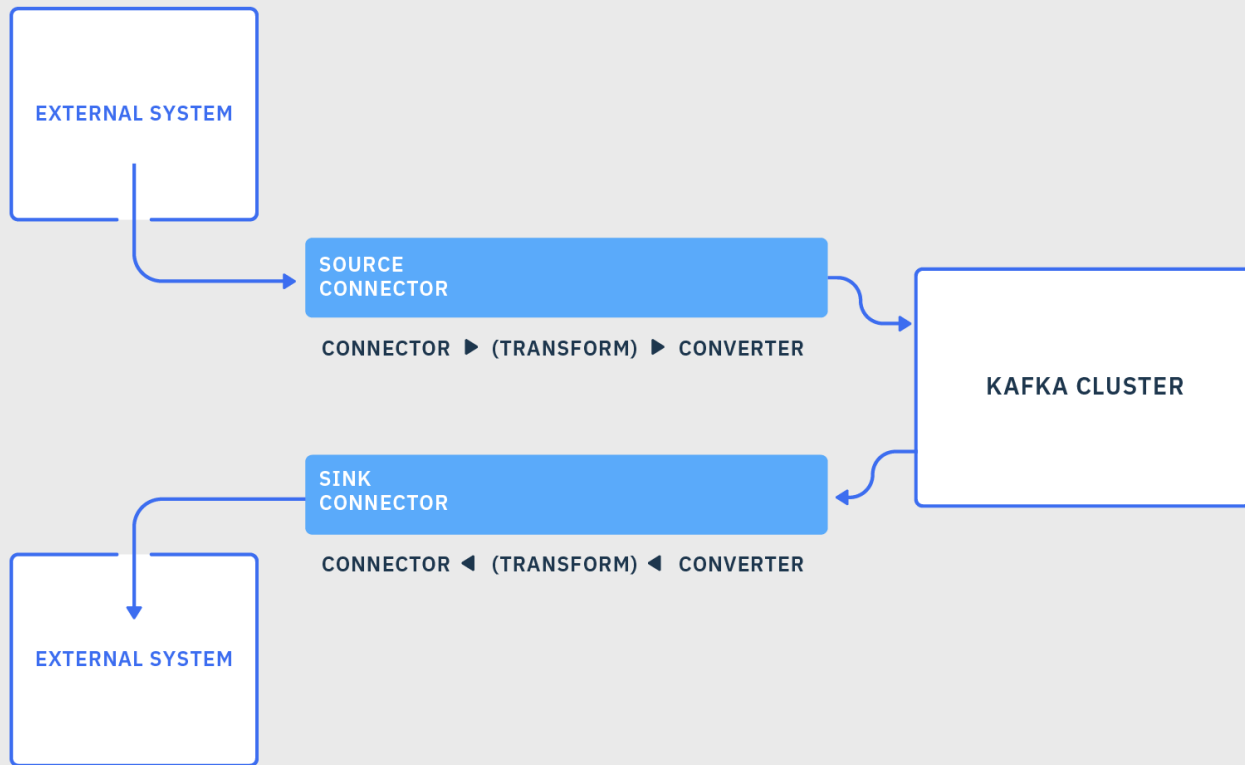
Primarily aimed at stream processing applications

# Consumer Groups





# Kafka Connect



## Over 80 connectors

HDFS

Elasticsearch

MySQL

JDBC

IBM MQ

MQTT

CoAP

*+ many others*

# What is IBM Event Streams?



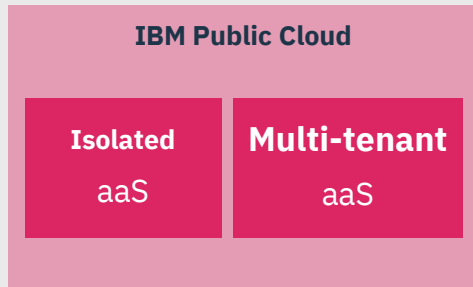
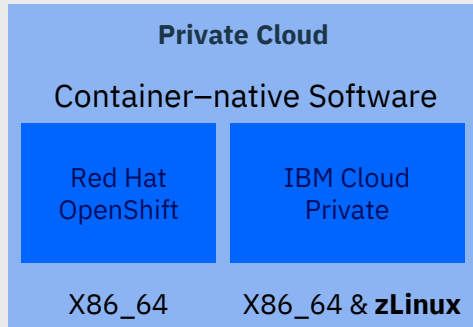
# IBM Event Streams

Fully supported  
Apache Kafka®  
with value-add  
capabilities



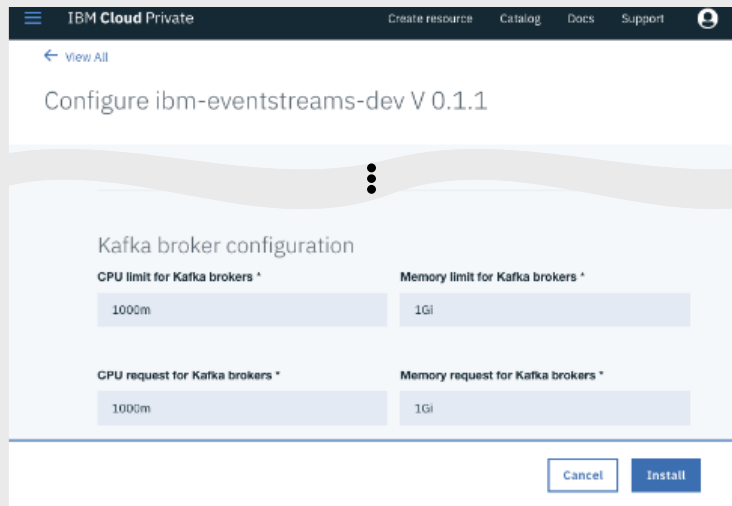
# IBM Event Streams Delivers Differentiated Value

IBM offers **Event Streams** in several form factors:



- **IBM has years of operational expertise** running Apache Kafka for Enterprises
  - This experience has been embedded in the DNA of Event Streams
- Event Streams makes Kafka easy to run, manage & consume, **reducing skill requirements** and increasing speed of deployment for **faster time to value**
- IBM Cloud Private security integration **simplifies Kafka access control** using roles and policies
- **IBM's experience in enterprise-critical software** has shaped features like geo-replication for Disaster Recovery & integration with IBM MQ, to give confidence deploying **mission-critical workloads**
- **Support you can trust** – IBM has decades of experience supporting the World's toughest environments

# Making Apache Kafka intuitive and easy



IBM Cloud Private

Create resource Catalog Docs Support

← View All

Configure ibm-eventstreams-dev V 0.1.1

Kafka broker configuration

CPU limit for Kafka brokers *	Memory limit for Kafka brokers *
1000m	1Gi
CPU request for Kafka brokers *	Memory request for Kafka brokers *
1000m	1Gi

Cancel Install

- Many distinct components to deploy, configure and coordinate secure connectivity
- Container placement critical to ensure production-level availability
- Secured network traffic ingress
- Ensuring consistent and repeatable deployment

think-2019

**Getting started**

Topics

Consumer groups

Monitor

Toolbox

[Connect to this cluster](#) 

# Welcome to IBM Event Streams, let's get you up and running...



## Use a simulated topic

Start exploring what IBM Event Streams has to offer with our simulated topic. You can do this even if your brokers aren't ready



## Generate a starter application

Download and install our starter Kafka application and view data flowing to and from IBM Event Streams in just a few minutes

**Learn more...**[FAQs](#) [GitHub](#) [Documentation](#)

Kafka basics

**Learn the basics of Apache Kafka,  
the heart of Event Streams.** **System is healthy**



Topics

## Airline\_delays

Messages

Consumer groups

Connection information



All partitions

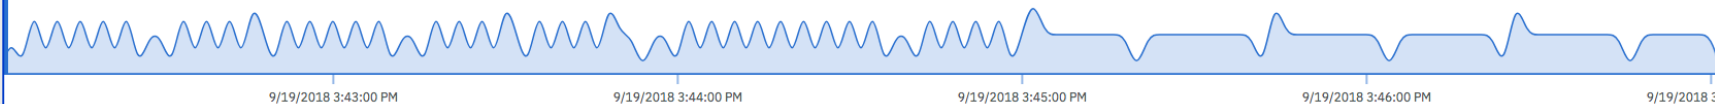


Showing 871 message(s) across all partitions

Find message



19 September 2018, 15:42:18



View live data



Select timeframe of data to display

Hours



Select start date of data

19/09/2018

SEPTEMBER 2018						
S	M	T	W	Th	F	S
26	27	28	29	30	31	1
2	3	4	5	6	7	8
9	10	11	12	13	14	15
16	17	18	19	20	21	22
23	24	25	26	27	28	29

Select start time of data

14:46:42

Timestamp

Partition

Offset

19/09/2018, 15:04:52

0

0

19/09/2018, 15:04:56

0

1

19/09/2018, 15:04:57

0

2

19/09/2018, 15:05:01

0

3

19/09/2018, 15:05:01

0

4

19/09/2018, 15:05:03

0

5

19/09/2018, 15:05:05

0

6

19/09/2018, 15:05:07

0

7

19/09/2018, 15:05:11

0

8

19/09/2018, 15:05:12

0

9



System is healthy

[← Topics / MY.FIRST.TOPIC](#)[Messages](#)[Consumer groups](#)Topic: MY.FIRST.TOPIC [< Previous offset](#)[Next offset >](#)

Partition

0

Offset

2

[View live data](#) 

Select timeframe of data to display

Hours 

Select start date of data

 09/02/2019

&lt;

FEBRUARY 2019

S	M	T	W	Th	F	S
27	28	29	30	31	1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	1	2
3	4	5	6	7	8	9

Indexed timestamp 

Partition

09/02/2019, 22:46:56

0

09/02/2019, 22:46:58

0

09/02/2019, 22:47:00

0

09/02/2019, 22:47:02

0

09/02/2019, 22:47:04

0

09/02/2019, 22:47:06

0

09/02/2019, 22:47:08

0

09/02/2019, 22:47:10

0

09/02/2019, 22:47:12

0

09/02/2019, 22:47:14

0

 Message size  
4 B Kafka timestamp   
09/02/2019, 22:47:00 Key  
-[Raw payload](#)demo 

Select start time of data

21:47:30



[← Topics](#)

## MY.FIRST.TOPIC

[Messages](#)[Consumer groups](#)[All partitions](#) ▾

Showing 34 messages

2 February 2019, 22:48:50

2:00:00 AM


2/5/2019

[View live data](#) ▶

Select timeframe of data to display

[Hours](#) ▾

Select start date of data

 09/02/2019

FEBRUARY 2019

S M T W Th F S

## Topic connection

[Connect a client](#)[Sample code](#)[Geo-replication](#)

## Sample connection code

Use this snippet of code to set the properties in your Kafka client to connect securely. Replace the values in <brackets>.

[Java](#)

```
import java.util.Properties;

import org.apache.kafka.clients.CommonClientConfigs;
import org.apache.kafka.common.config.SaslConfigs;
import org.apache.kafka.common.config.SslConfigs;

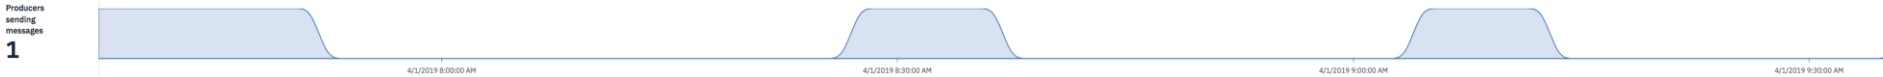
Properties properties = new Properties();
properties.put(CommonClientConfigs.BOOTSTRAP_SERVERS_CONFIG, "9.20.192.113:31934");
properties.put(CommonClientConfigs.SECURITY_PROTOCOL_CONFIG, "SASL_SSL");
properties.put(SslConfigs.SSL_PROTOCOL_CONFIG, "TLSv1.2");
properties.put(SslConfigs.SSL_TRUSTSTORE_LOCATION_CONFIG, "<certs.jks_file_location>");
properties.put(SslConfigs.SSL_TRUSTSTORE_PASSWORD_CONFIG, "<truststore_password>");
```

[Show more ▾](#)

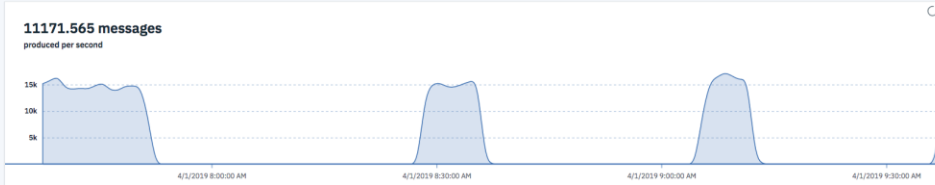
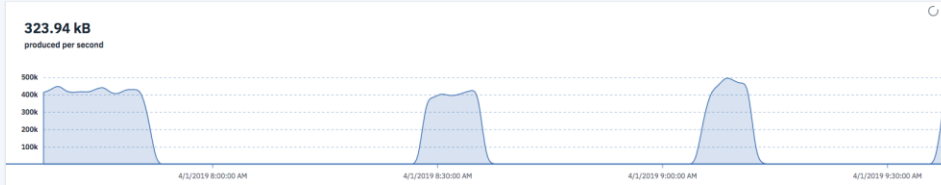
## Sample configuration properties

Use this snippet to create a properties file for use by Kafka tools to connect securely. Replace the values in <brackets>.

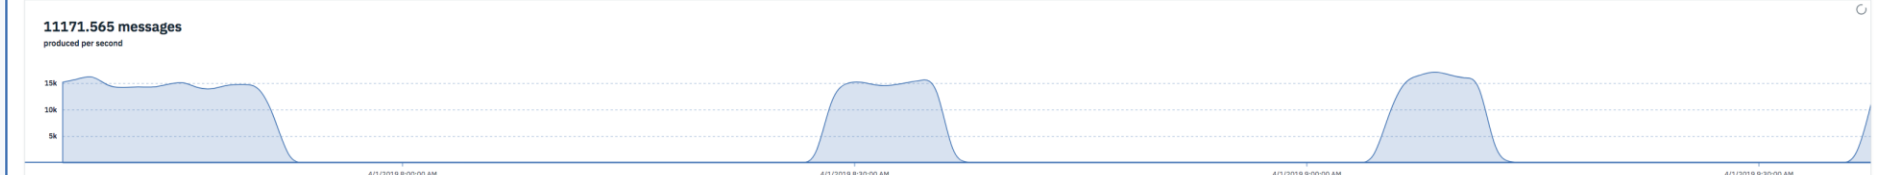
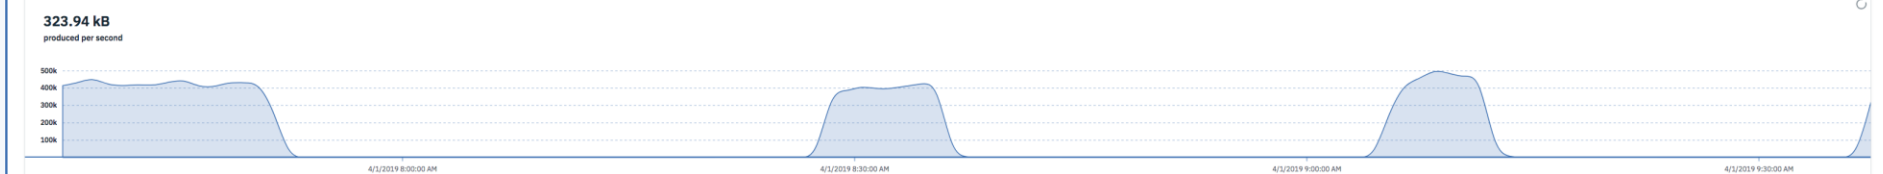
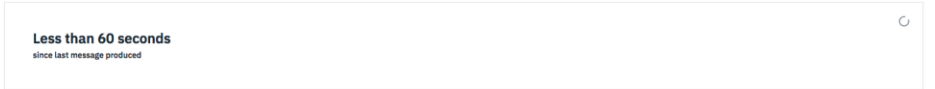
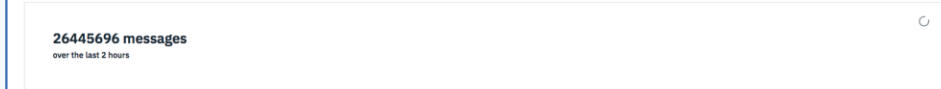
Producers sending messages  
1



Type to filter producers



Kafka producer ID	Min data point	Max data point	Data range	Warnings
<input checked="" type="checkbox"/> producer-1 ServiceId=4f1268369-8e59-4e1c-a23e-08cda193a2d5	0 B	507.26 kB	507.26 kB	None



# Effortless geo-replication

1

## Geo-replication

Geo-replication is a way to duplicate your topics to different instances of Event Streams. These are usually located in different locations to minimize the risk of data loss in the event of a cluster failure.

### Origin locations 1

Want to replicate topics to this cluster?

To add this cluster as a destination location on an origin cluster, you will need to generate connection information.

Generate connection information for this cluster

### Destination locations 1

Add destination cluster +

rc6-repl	TOPICS 1	WORKERS 2	→
----------	----------	-----------	---

2

## Back to geo-replication

rc6-repl	TOPICS 0	WORKERS 2	⋮
----------	----------	-----------	---

### Geo-replicated topics

Name	Replication health
timtopic	Awaiting creation

You are about to create a replicator for 1 topic.

☐ Add prefix to destination topic names ☒ Include message history

Cancel Create

3

## Back to geo-replication

rc6-repl	TOPICS 0	WORKERS 2	⋮
----------	----------	-----------	---

### Geo-replicated topics

Name	Replication health
timtopic	CREATING

## Topics

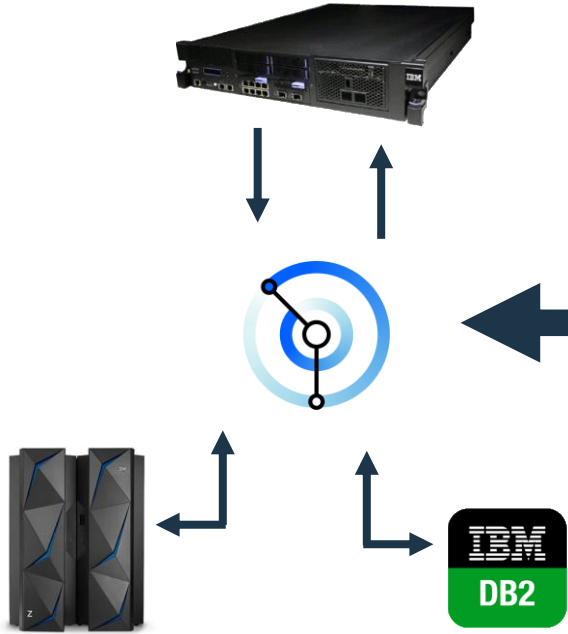
🔍 Type to search topics

Geo-replication 📍

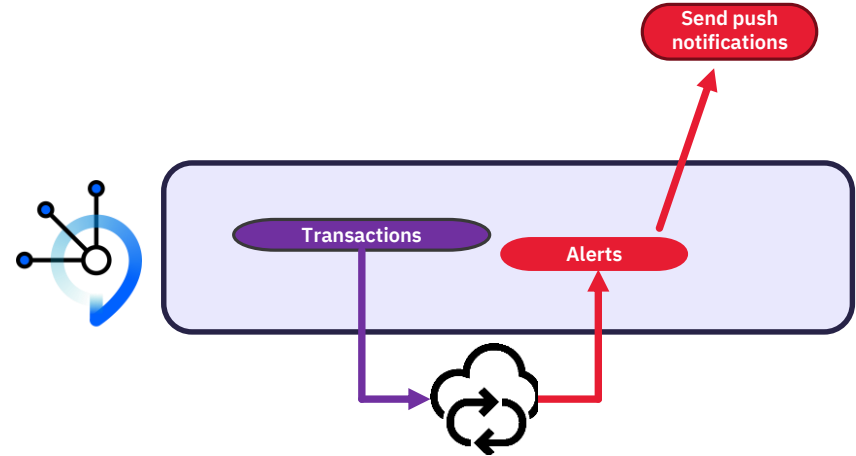
Create topic +

Name	Replicas	Partitions	
testfest	3	3	⋮
timtopic	2	1	📍 (1) ⋮

# Integrates with IBM MQ



**IBM MQ** connects mission-critical Systems of Record, requiring **transactional, once-only delivery**



**IBM Event Streams** distributes and processes streams of events in real-time to intelligently engage with customers

# It's Easy to Connect IBM MQ to Apache Kafka

IBM has created a pair of connectors, available as source code or as part of IBM Event Streams

## **Source Connector**

From MQ queue to Kafka topic

<https://github.com/ibm-messaging/kafka-connect-mq-source>

## **Sink Connector**

From Kafka topic to MQ queue

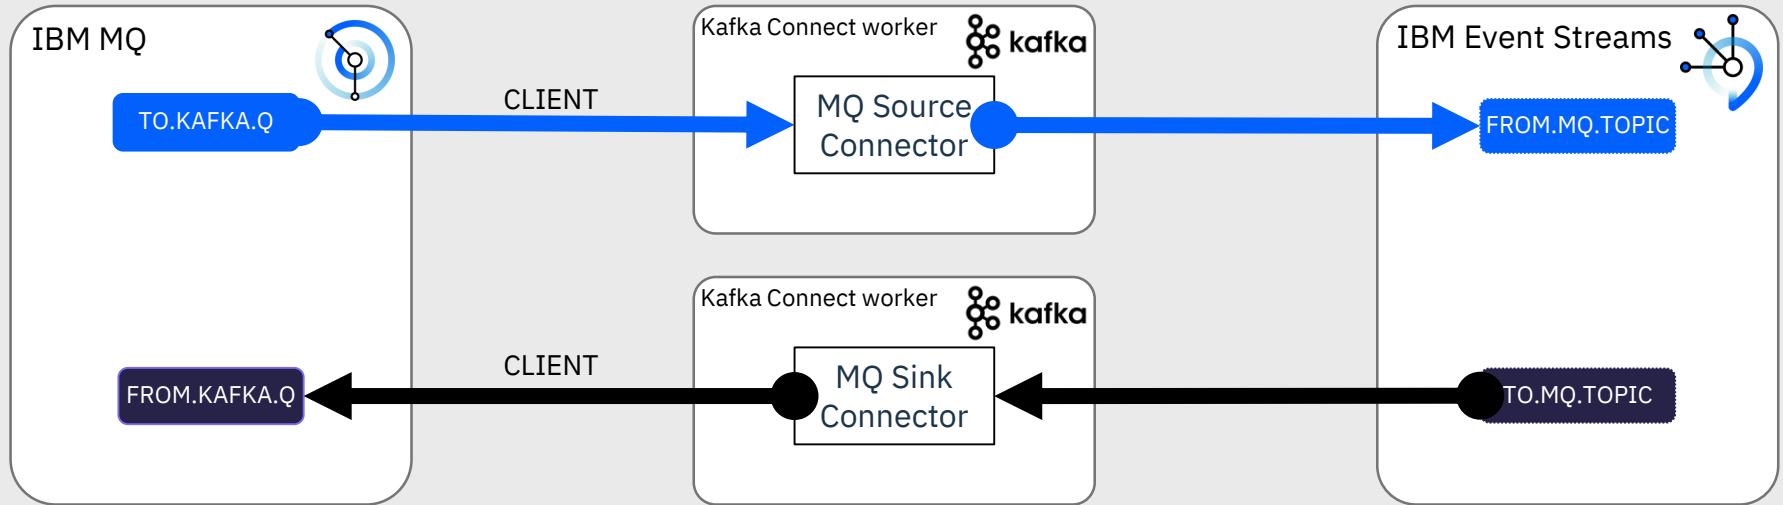
<https://github.com/ibm-messaging/kafka-connect-mq-sink>

- Copies messages from MQ queues to Event Streams topics and vice versa
- Supports all current MQ versions (MQ v8 or later, all platforms)
- Extend the connector to support any business-specific message format
- Fully supported by IBM for customers with support entitlement for IBM Event Streams

# Running the Connectors for IBM MQ

The connectors are deployed into a component of Apache Kafka called a Kafka Connect worker

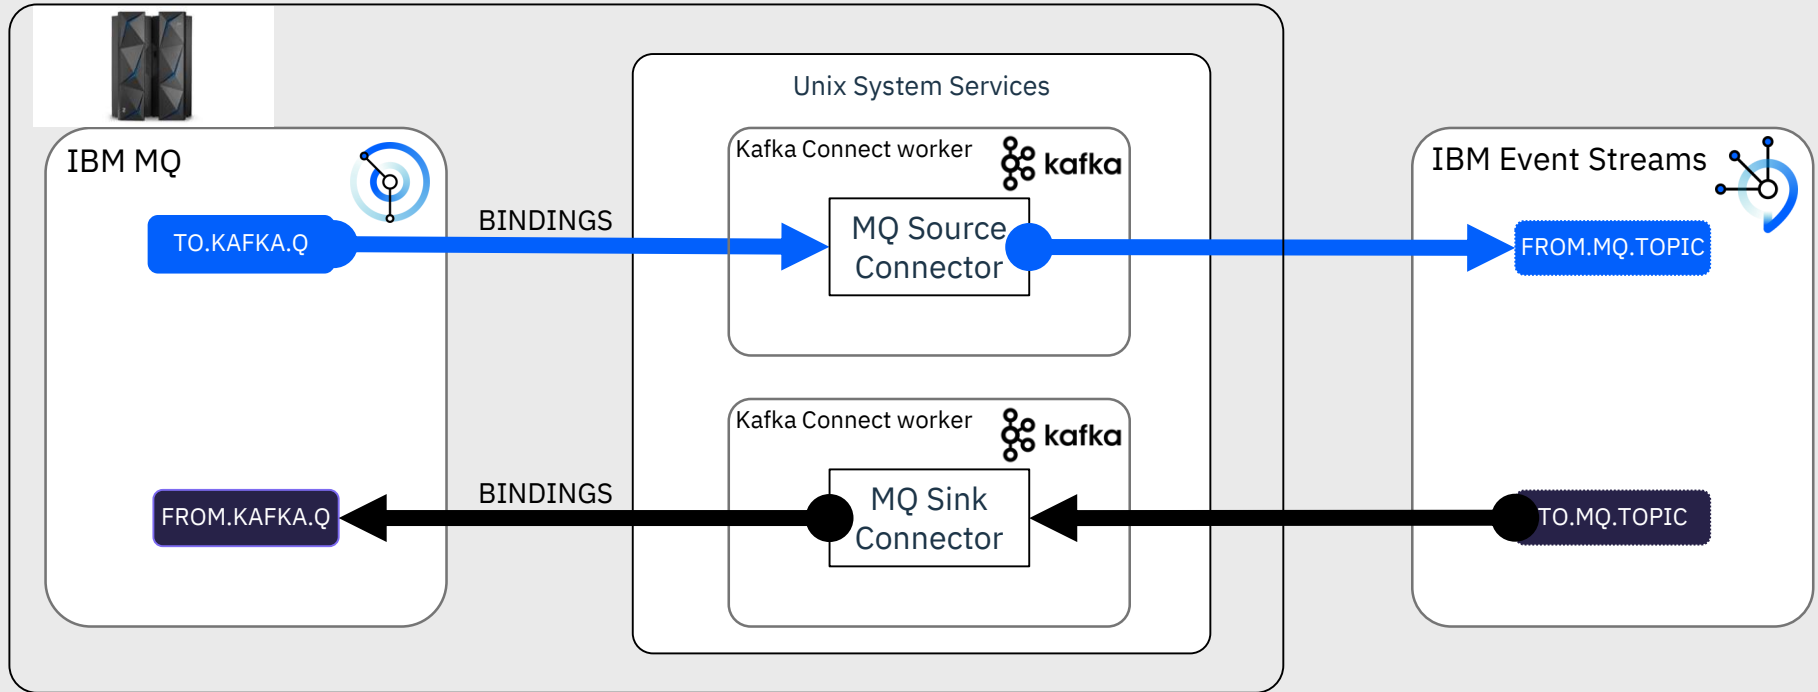
This runs between IBM MQ and IBM Event Streams (or open-source Apache Kafka)



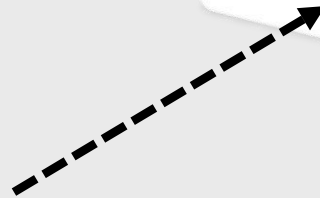
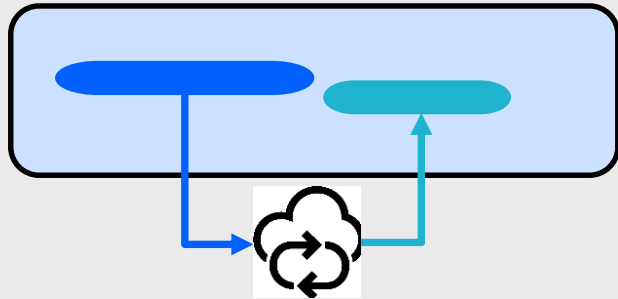
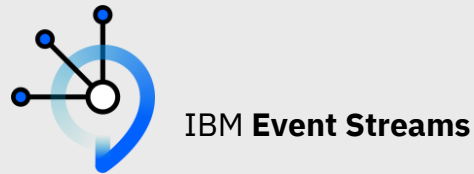
# Running the Connectors for IBM MQ on z/OS

The Kafka Connect workers can be deployed onto z/OS Unix System Services

Then, the connection to MQ can be a bindings connection



# IBM Event Streams | Integrated with Key Monitoring Tools



**External monitoring tools**  
Datadog, Splunk, etc



# Additional capabilities in 1Q19

**IBM Event Streams 2019.1.1** | **GA:** 29th March

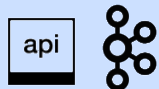


**OpenShift  
support**

## **What:**

Deploy IBM Event Streams into existing OpenShift private cloud environments

**Clients that use Red Hat OpenShift can now deploy Event Streams into the same platform as their applications**



**Kafka  
REST API**

Scalable REST interface for inbound event data

**Unlock events from systems that cannot connect easily to Kafka**

☞ REST connectivity helps Mainframe and DataPower users especially



**Monitoring**

Integration points for 3<sup>rd</sup> party monitoring tools

**Connect existing monitoring tools to Event Streams to monitor it alongside the other components of their application**

☞ Get a single dashboard view of the environment

**COMING SOON!** (Statement of Direction: 1Q19)



**Schema Registry**

Associate data schemas with topics to ensure messages are well-formed

**Speeds up the rate of new application development as expected data formats are recorded and understood**

# Unlocking events in existing systems

ENTERPRISE IT

DB

IBM MQ

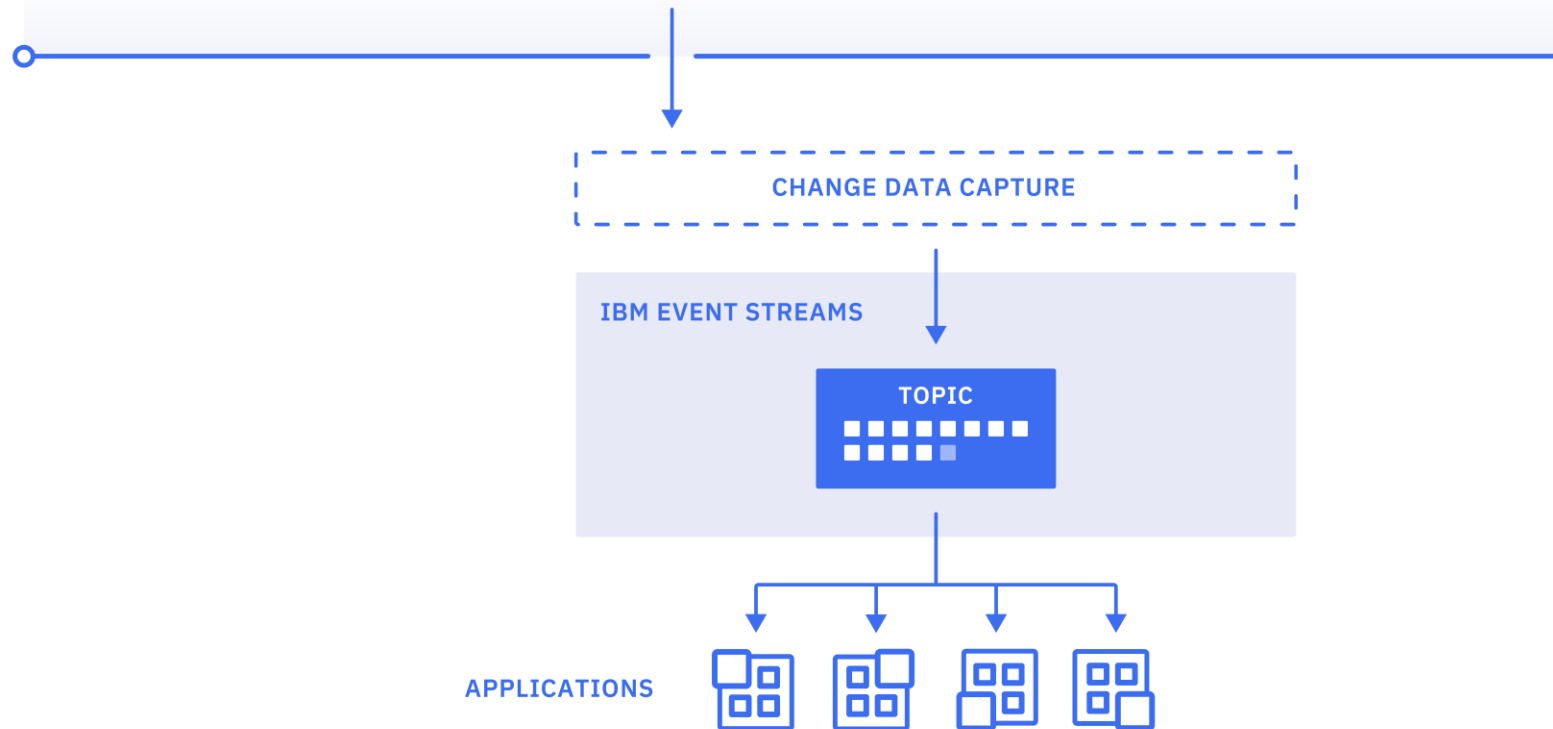


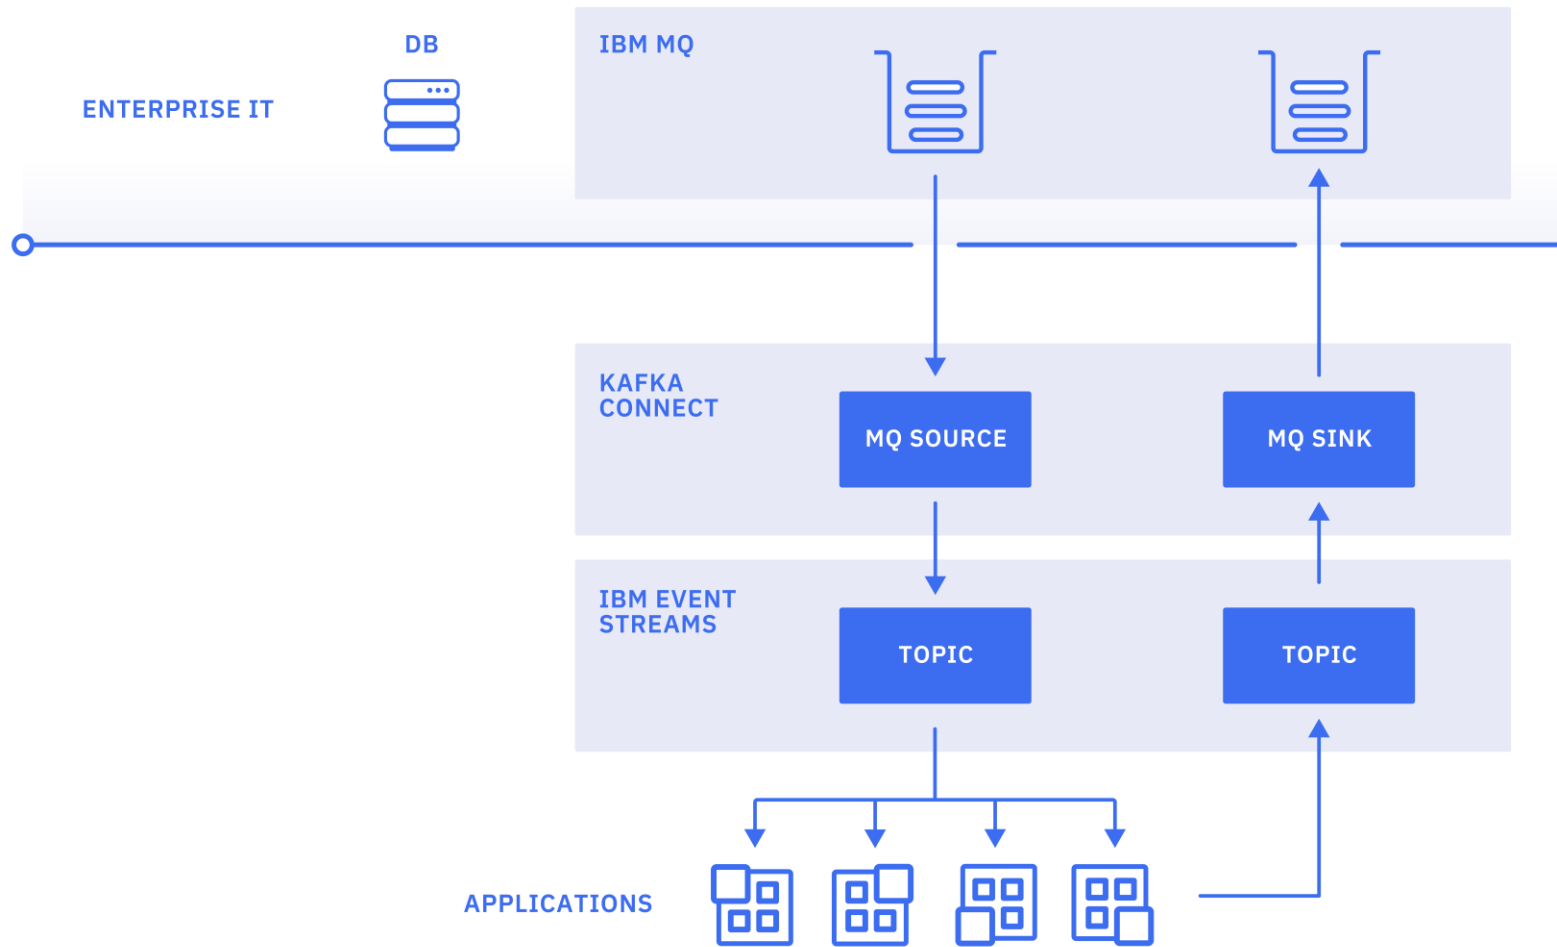
CHANGE DATA CAPTURE

IBM EVENT STREAMS

TOPIC

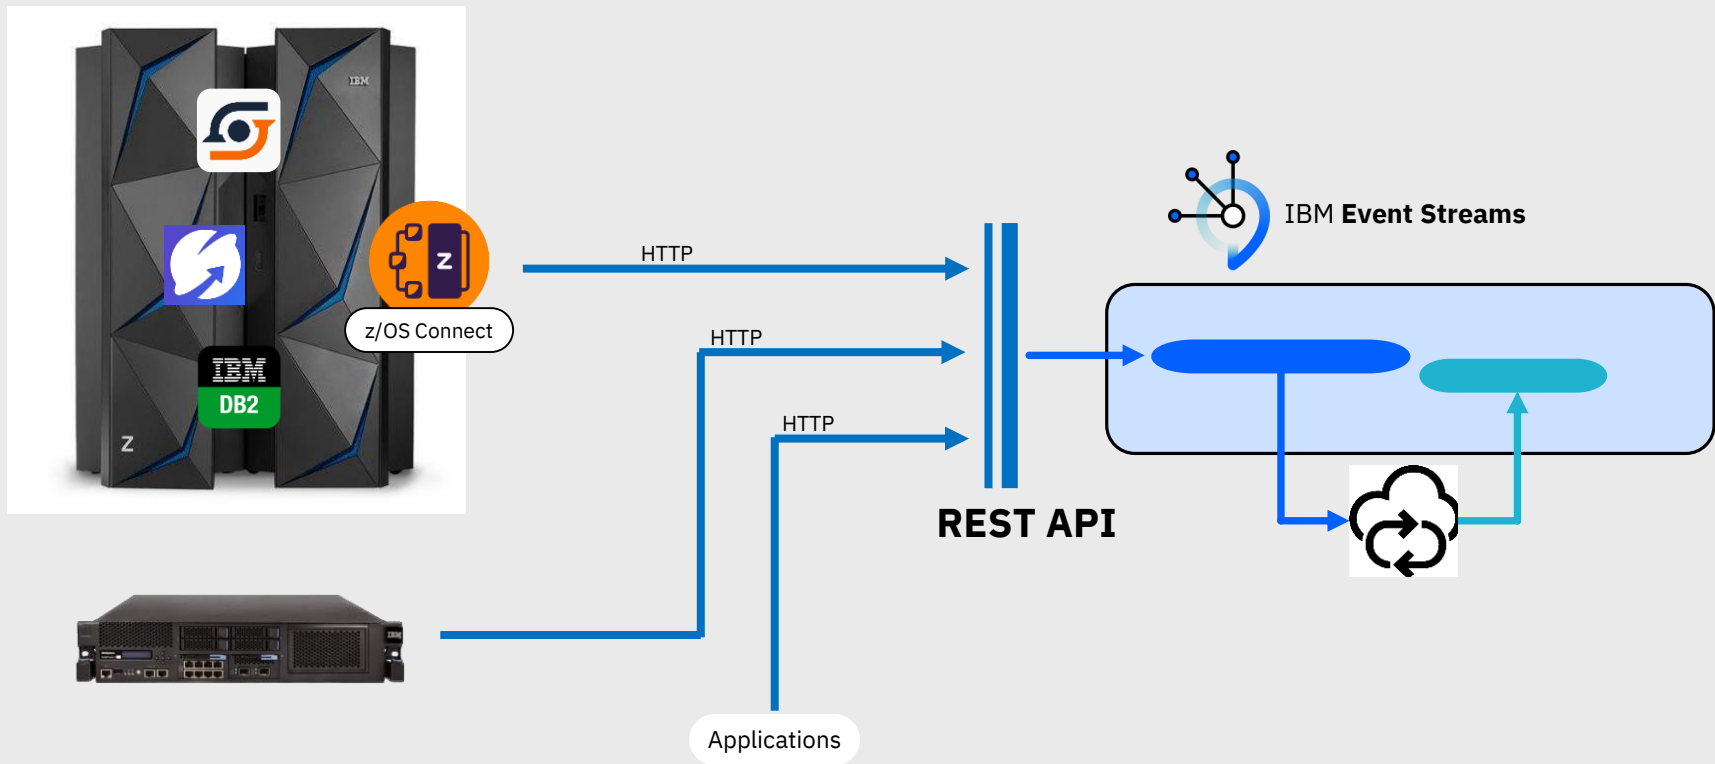
APPLICATIONS





# Unlock Events from Systems where Kafka clients are not available

## *REST API for Inbound Data*



# Publish Events from Anywhere with the REST Producer API

IBM has created a new easy-to-use REST Producer API

## **POST /topics/{topic\_name}/records**

Content-Type: text/plain

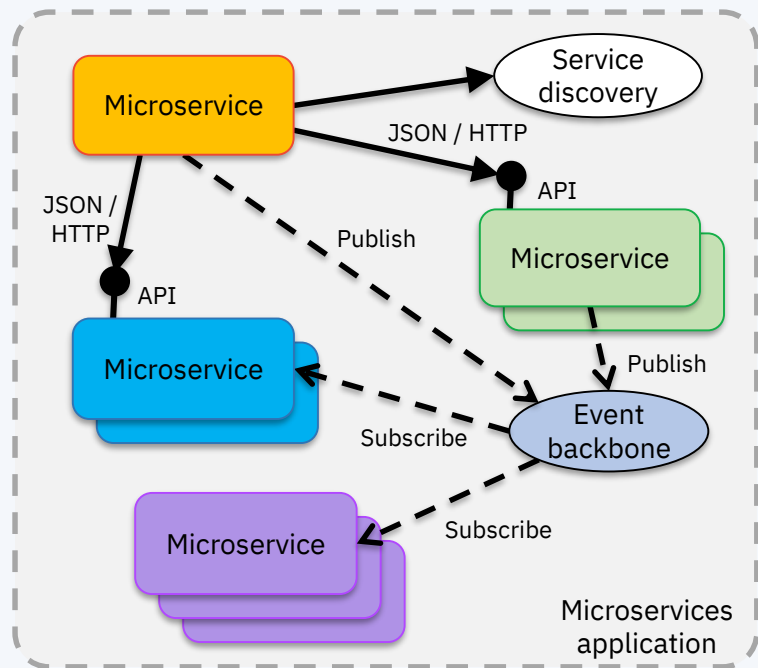
Authorization: Bearer {bearer\_token}

Hello Event Streams

- Use it wherever it's difficult to use a real Kafka client e.g. DataPower, z/OS
- Straightforward design makes it easy to use from the command line and developer tools
- Supports partitioning keys and headers
- So easy you can use it from the command line with cURL

# Event-driven microservices

# Event-driven microservices



Microservices communicate primarily using events, with APIs where required

Microservices can produce and consume events using the publish/subscribe pattern

Events are delivered using an event backbone

Data is eventually consistent

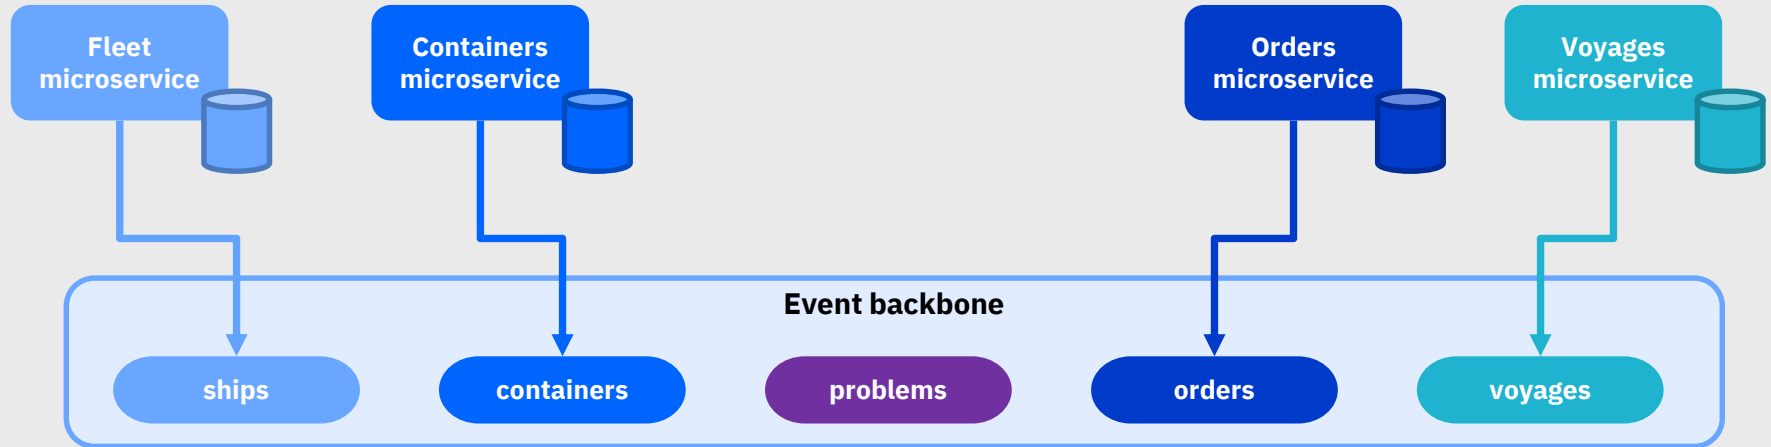


# Pattern – database per microservice

Each microservice persists its own data

Protects independence of the microservice against external change

An event stream (topic) is associated with each microservice for a log of events



# Summary

# Event-Driven Enterprise using IBM Event Streams

- Event streaming is a powerful new paradigm for messaging
- It's ideal for building responsive applications, particularly using microservices principles
- IBM Event Streams offers an easy, supported way to use Apache Kafka in your enterprise

# Find out more

- Try out Event Streams
  - <https://ibm.github.io/event-streams/installing/trying-out/> (On premise)
  - <https://console.bluemix.net/catalog/services/event-streams> (IBM hosted)
- Get in touch
  - <https://ibm.github.io/event-streams/support/>
- Apache Kafka
  - <https://kafka.apache.org/>
- Find out more
  - <https://www.ibm.com/cloud/event-streams>

Thank You

